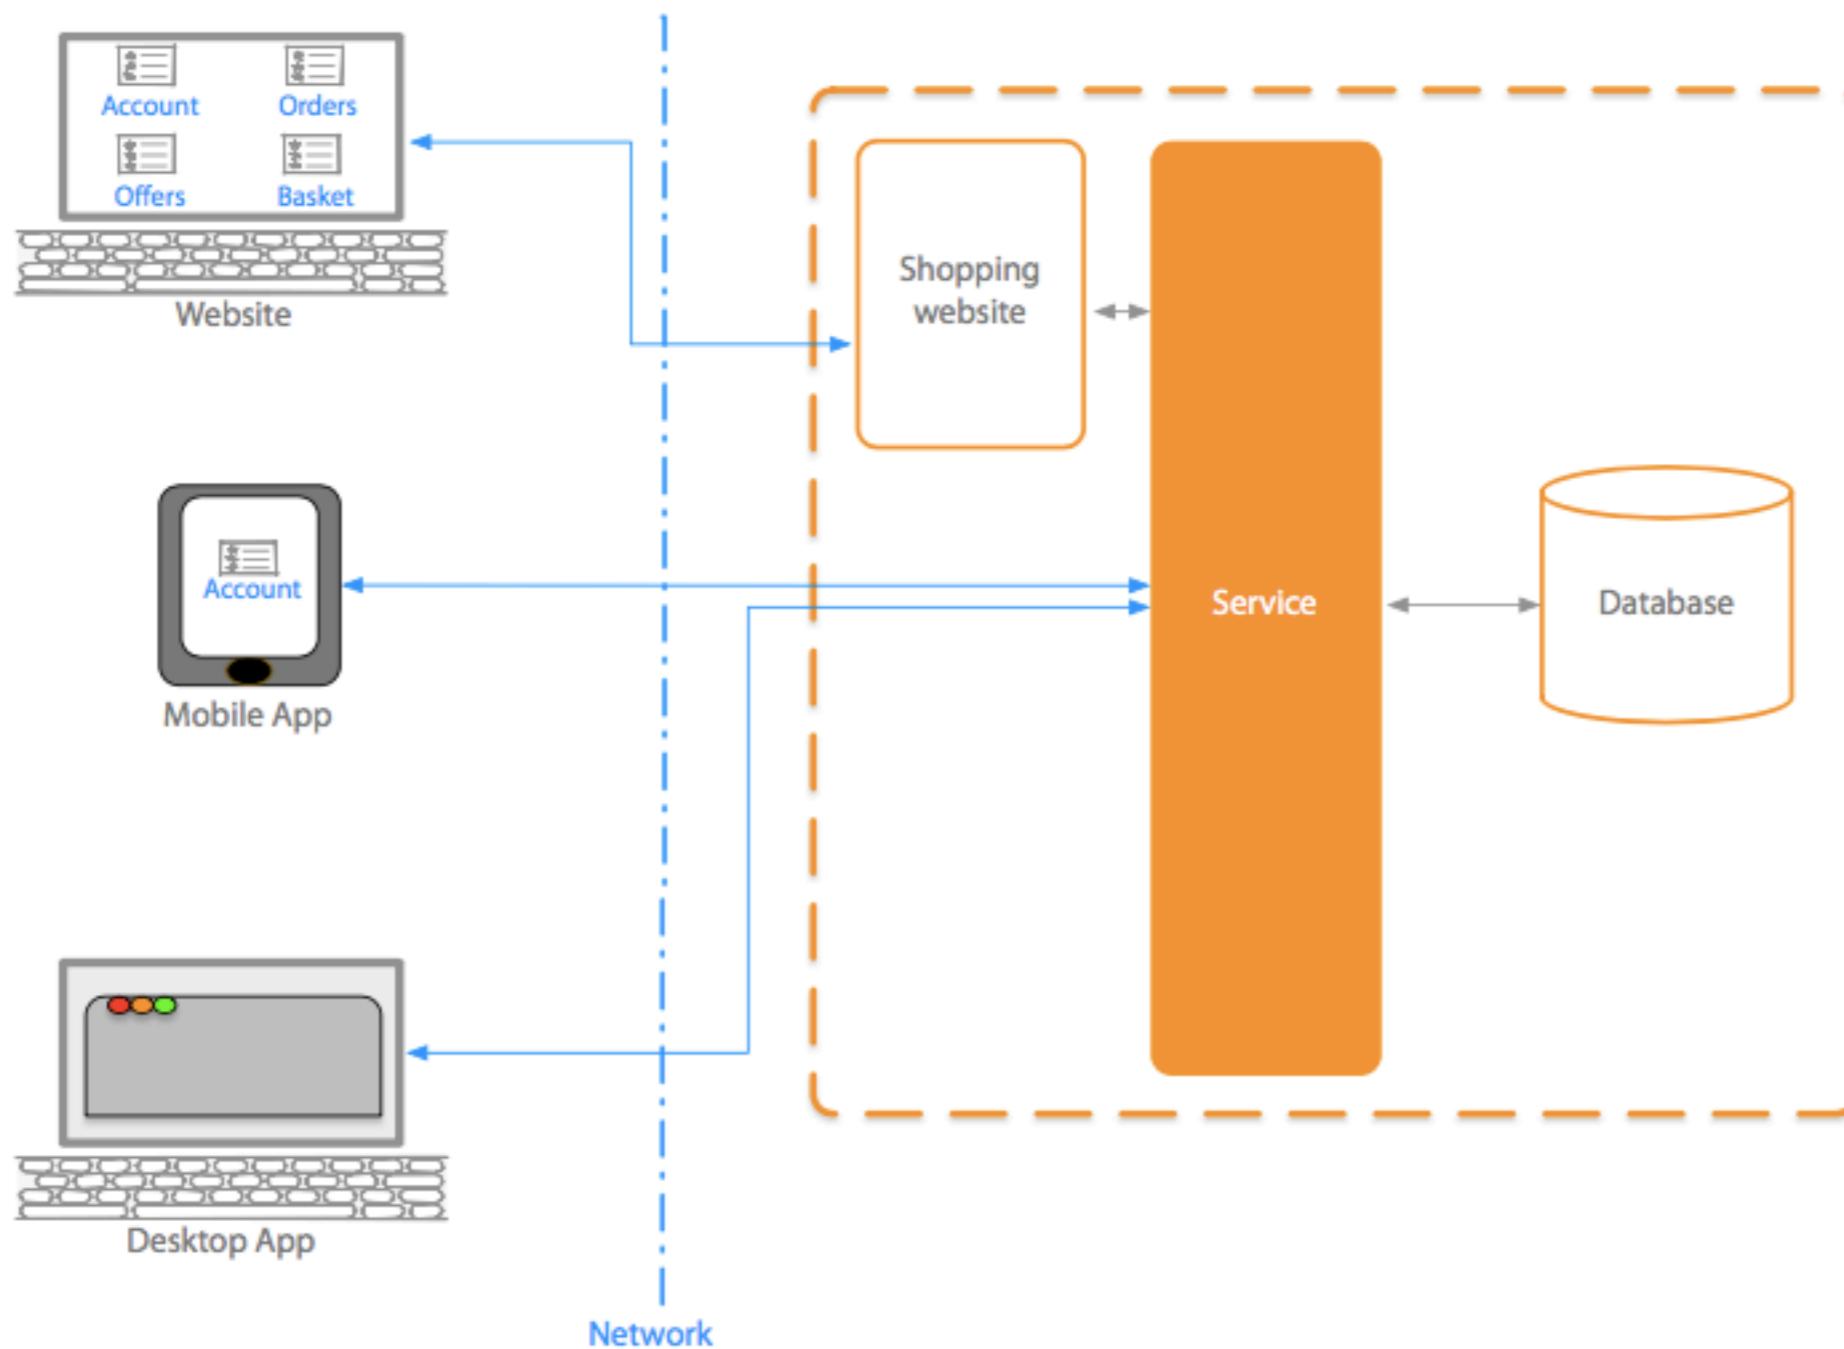


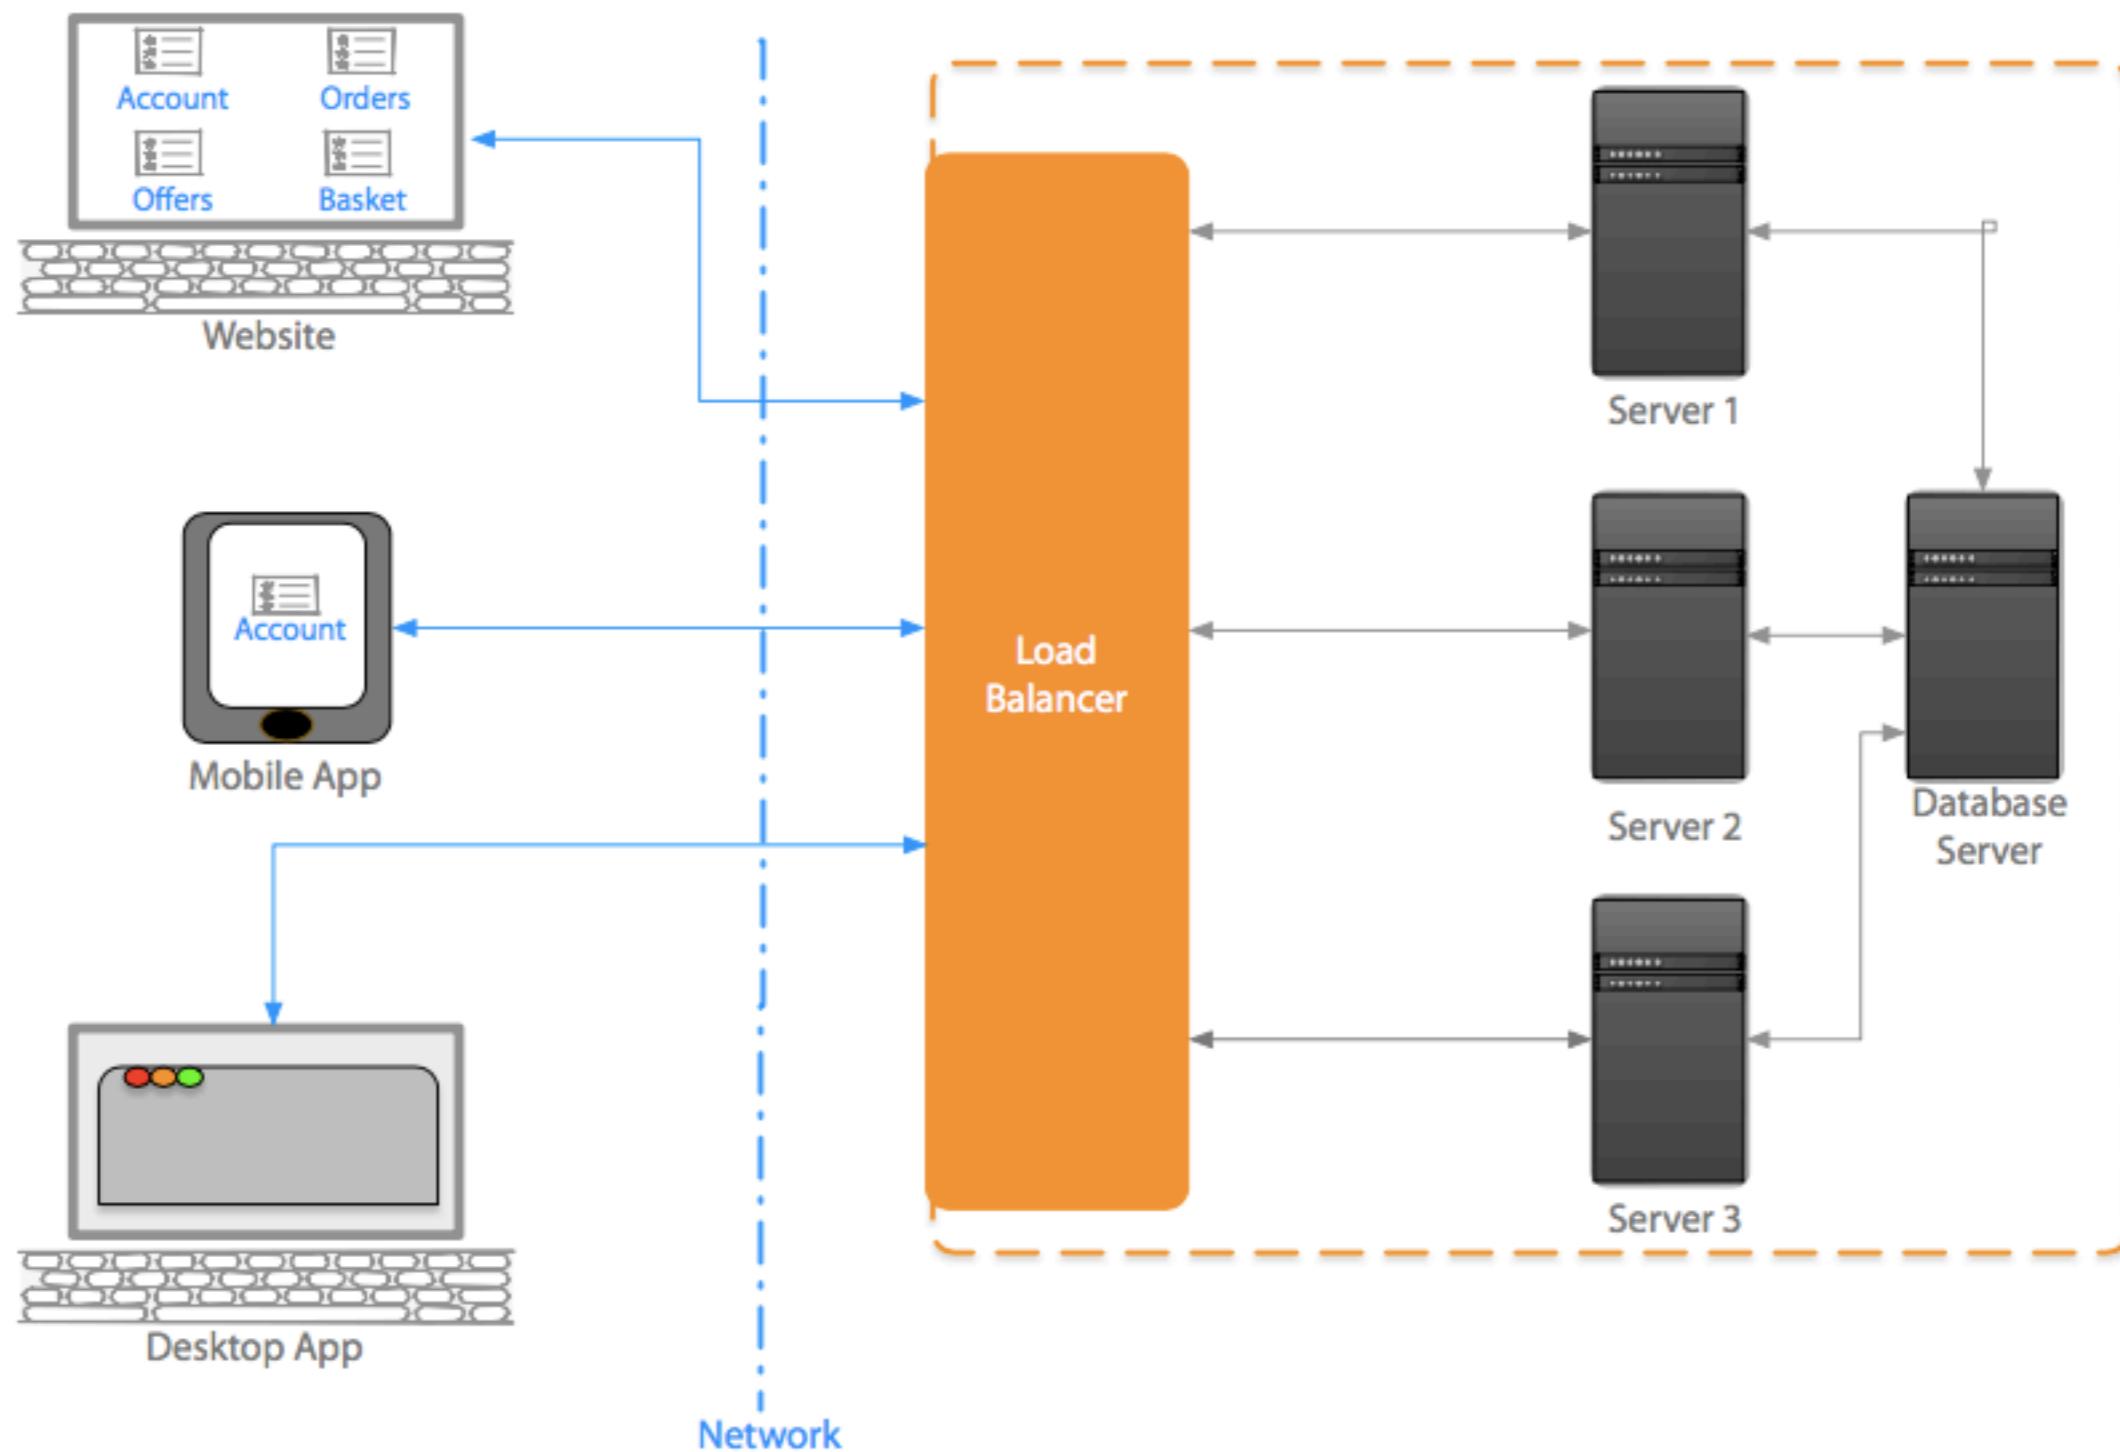
Microservices Architecture

What is a Service?

Service



Service

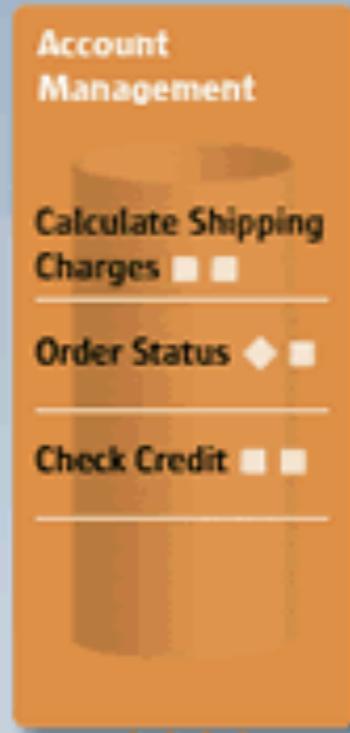
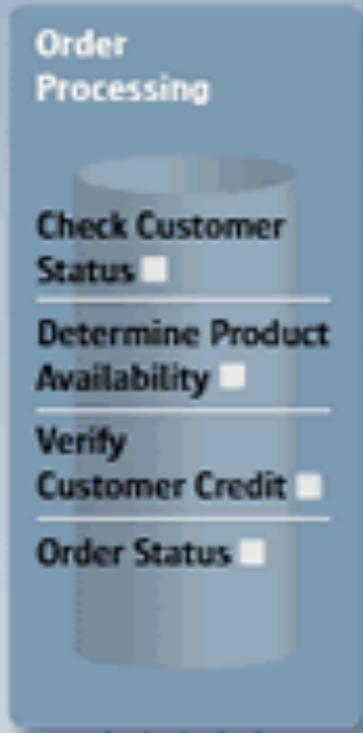


SOA

Before SOA

Siloed • Closed • Monolithic • Brittle

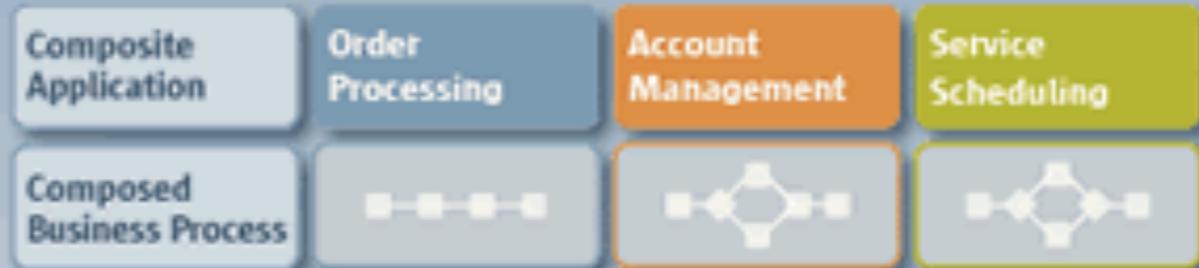
Application Dependent Business Functions



After SOA

Shared services • Collaborative • Interoperable • Integrated

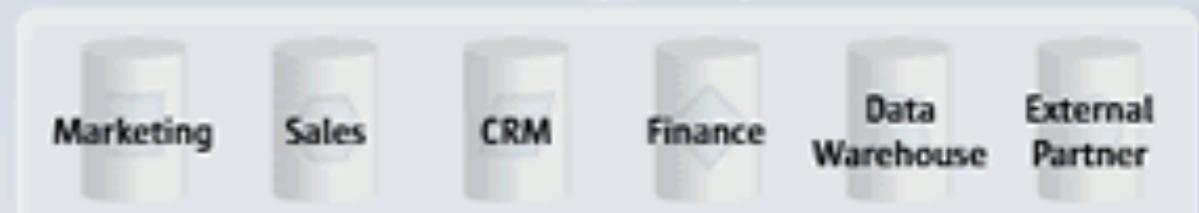
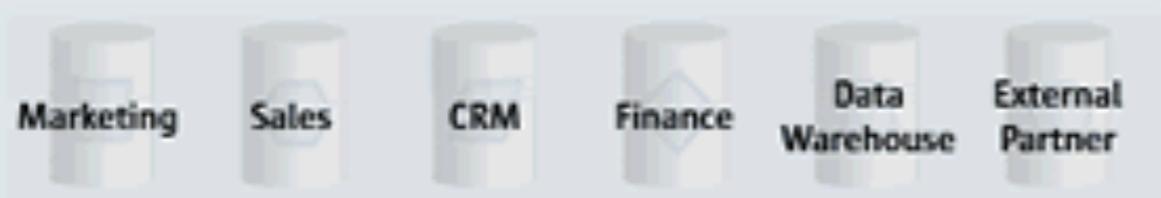
Composite Applications



Reusable Business Services



Data Repository

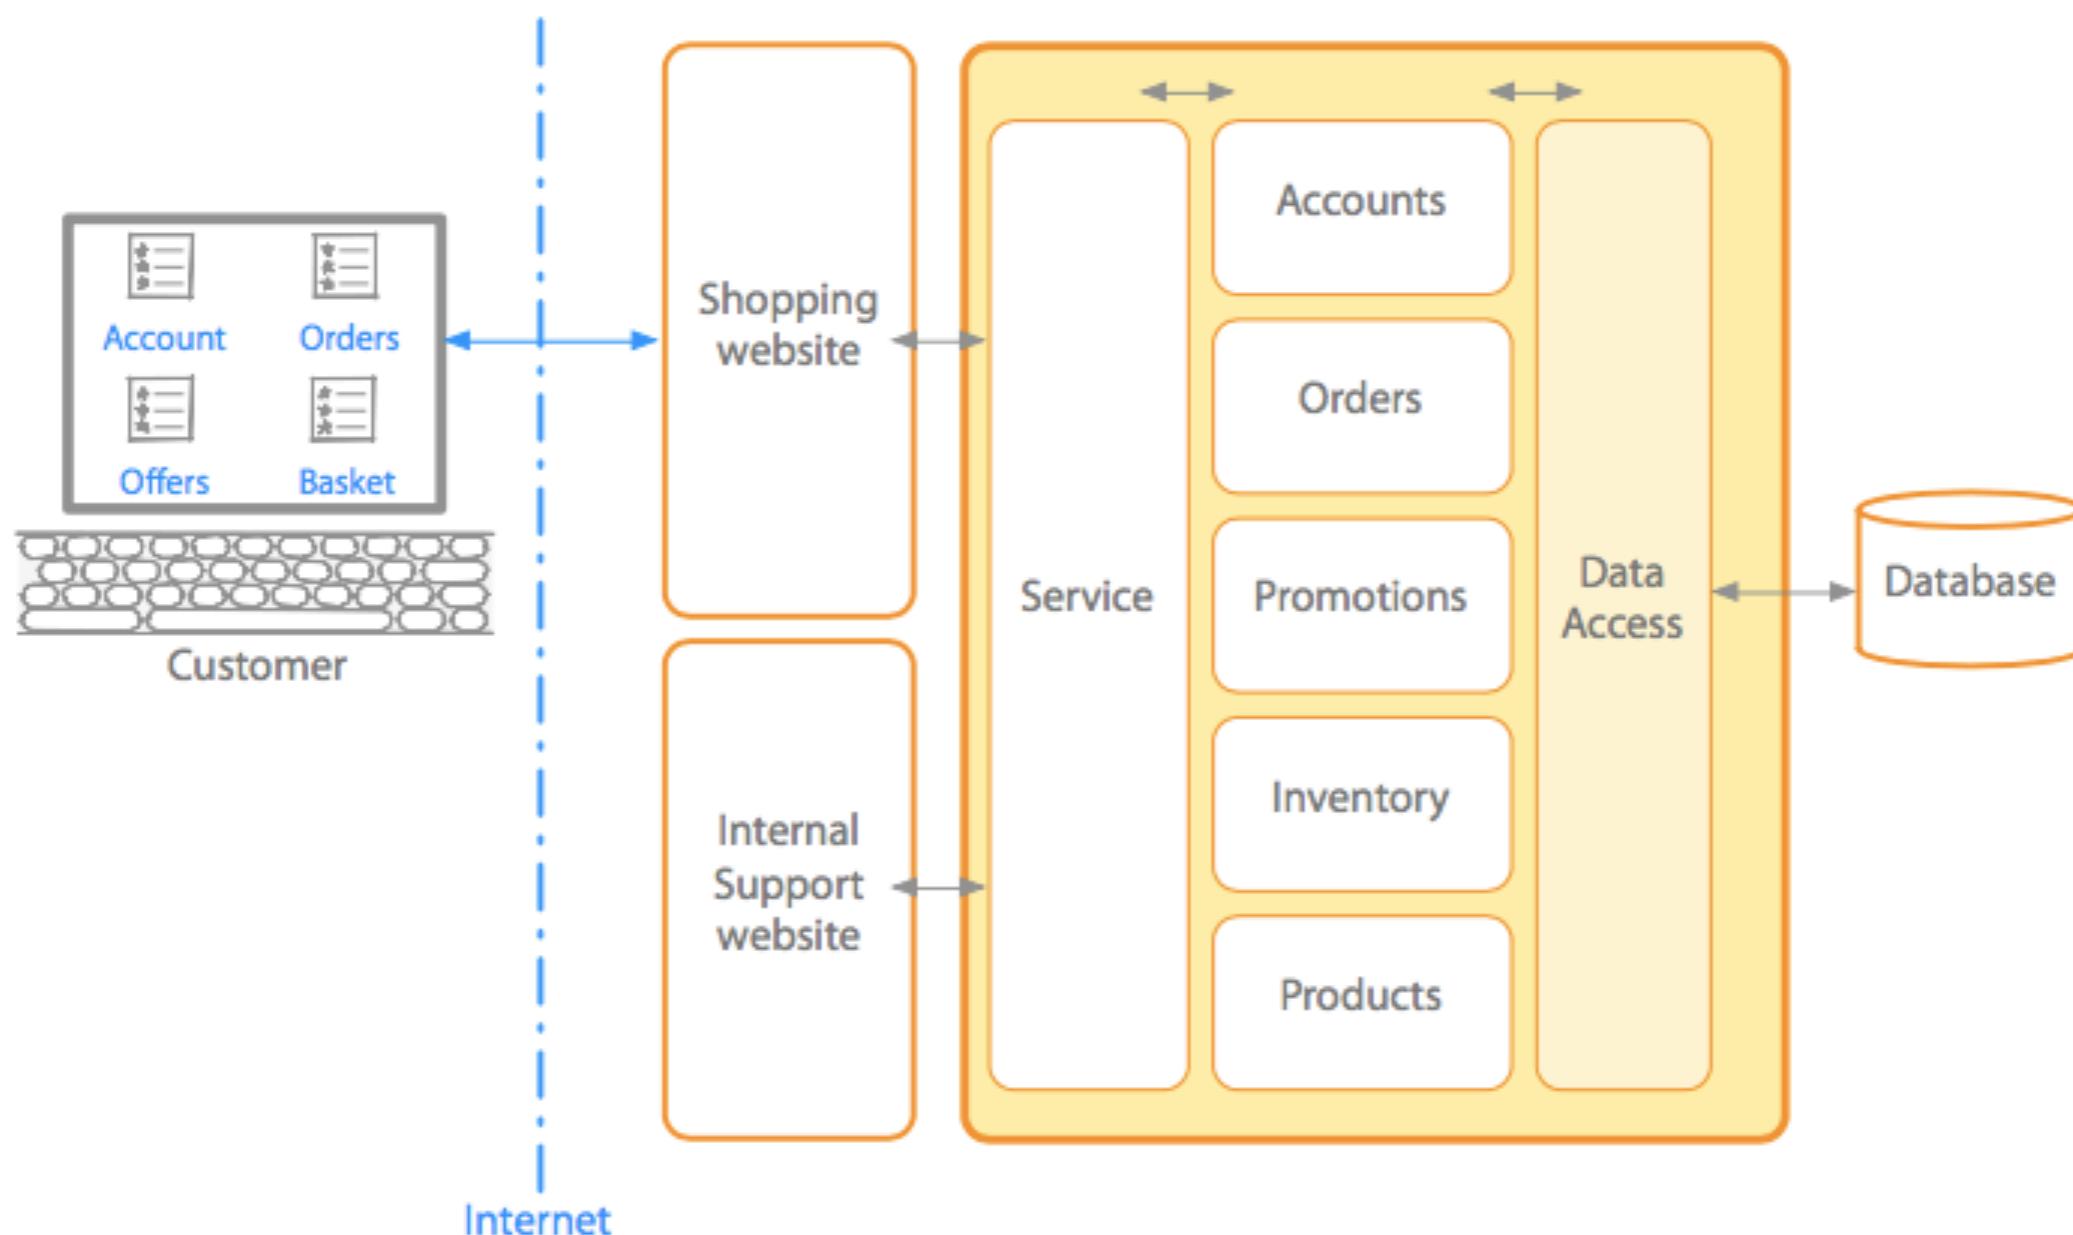


SOA

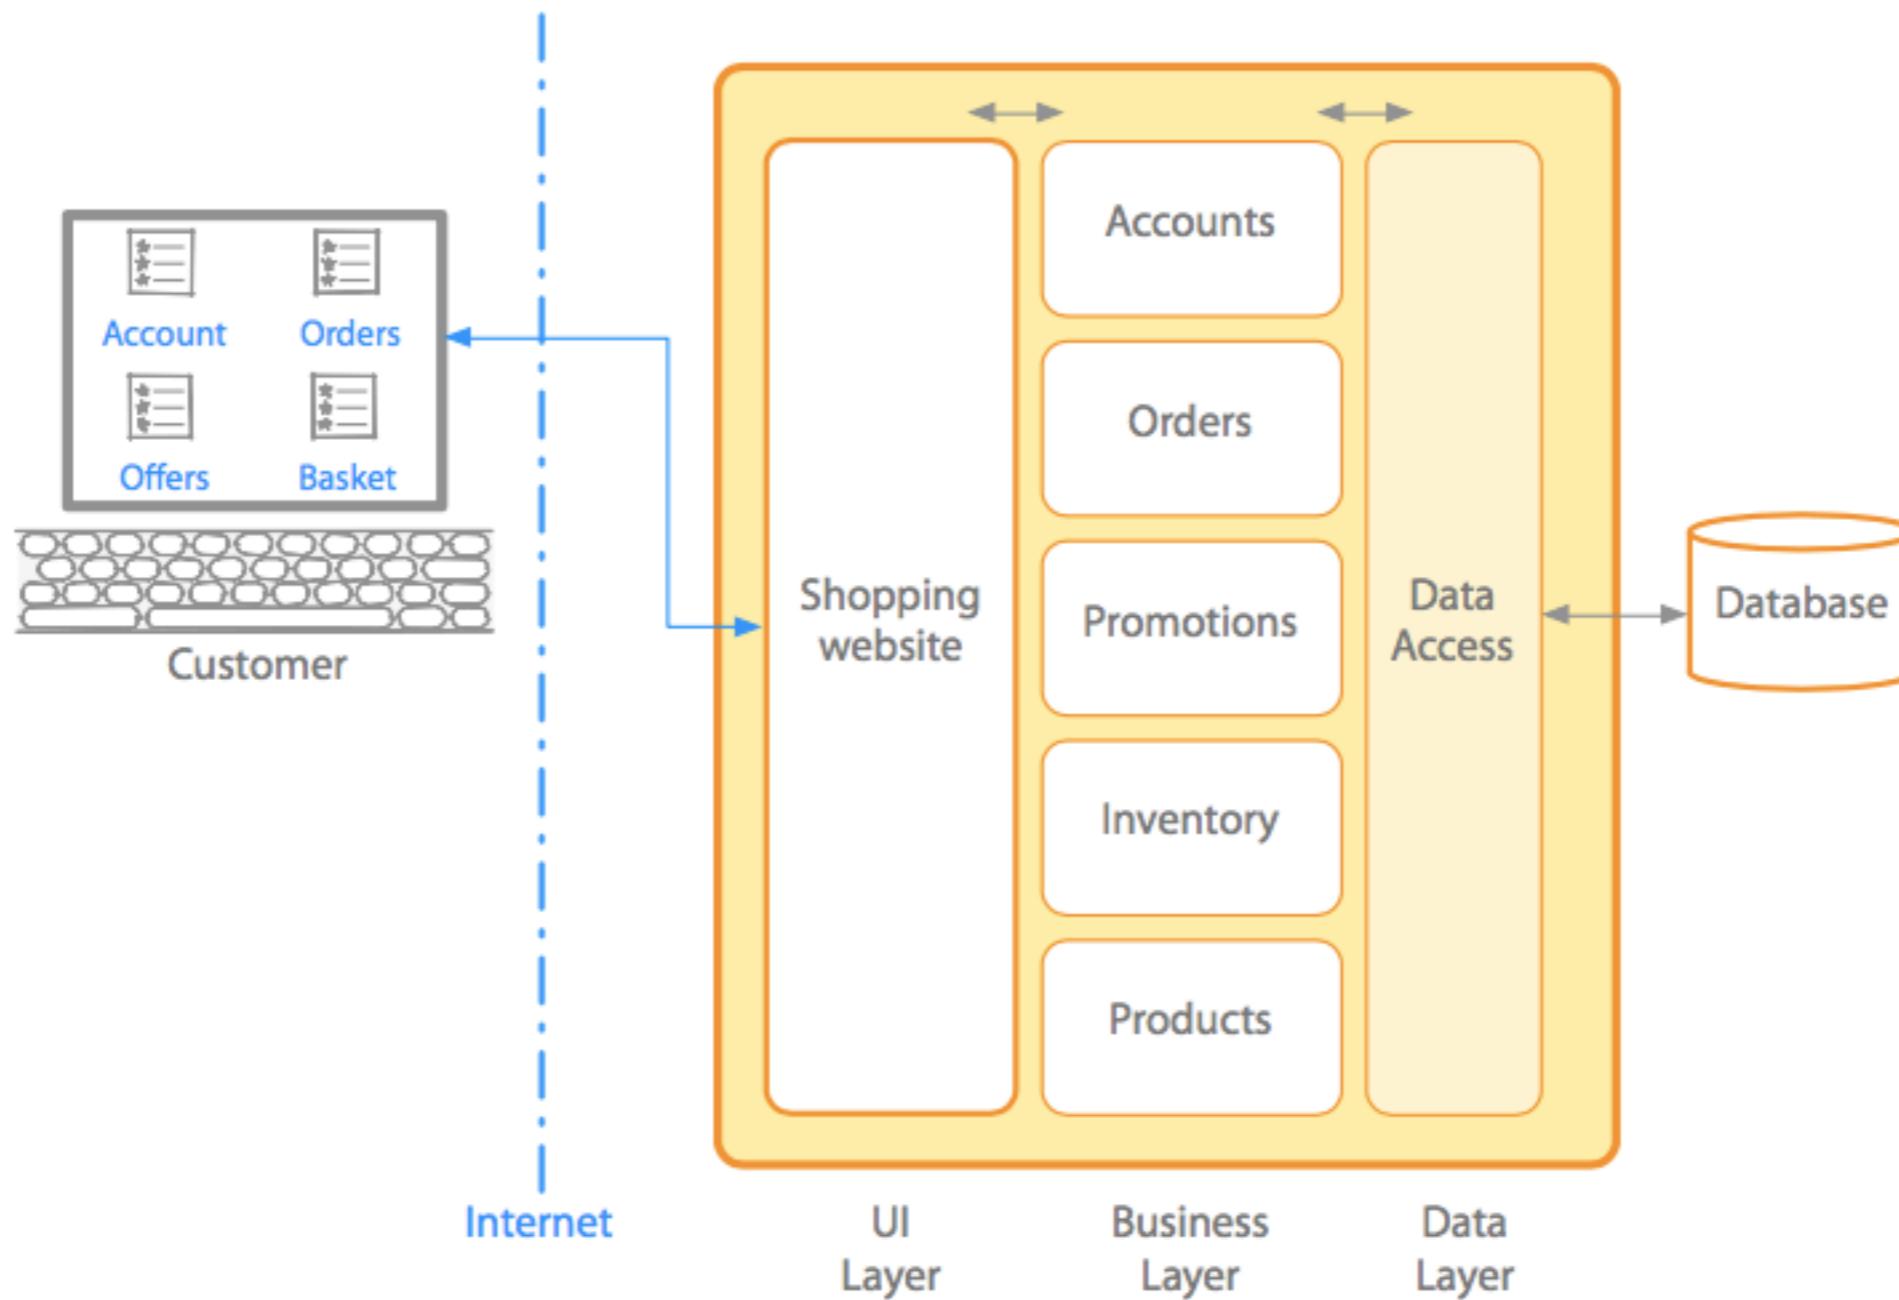
SOA done well

- Knowing how to size a service
- Traditional SOA resulted in a monolithic services

Monolithic



Monolithic



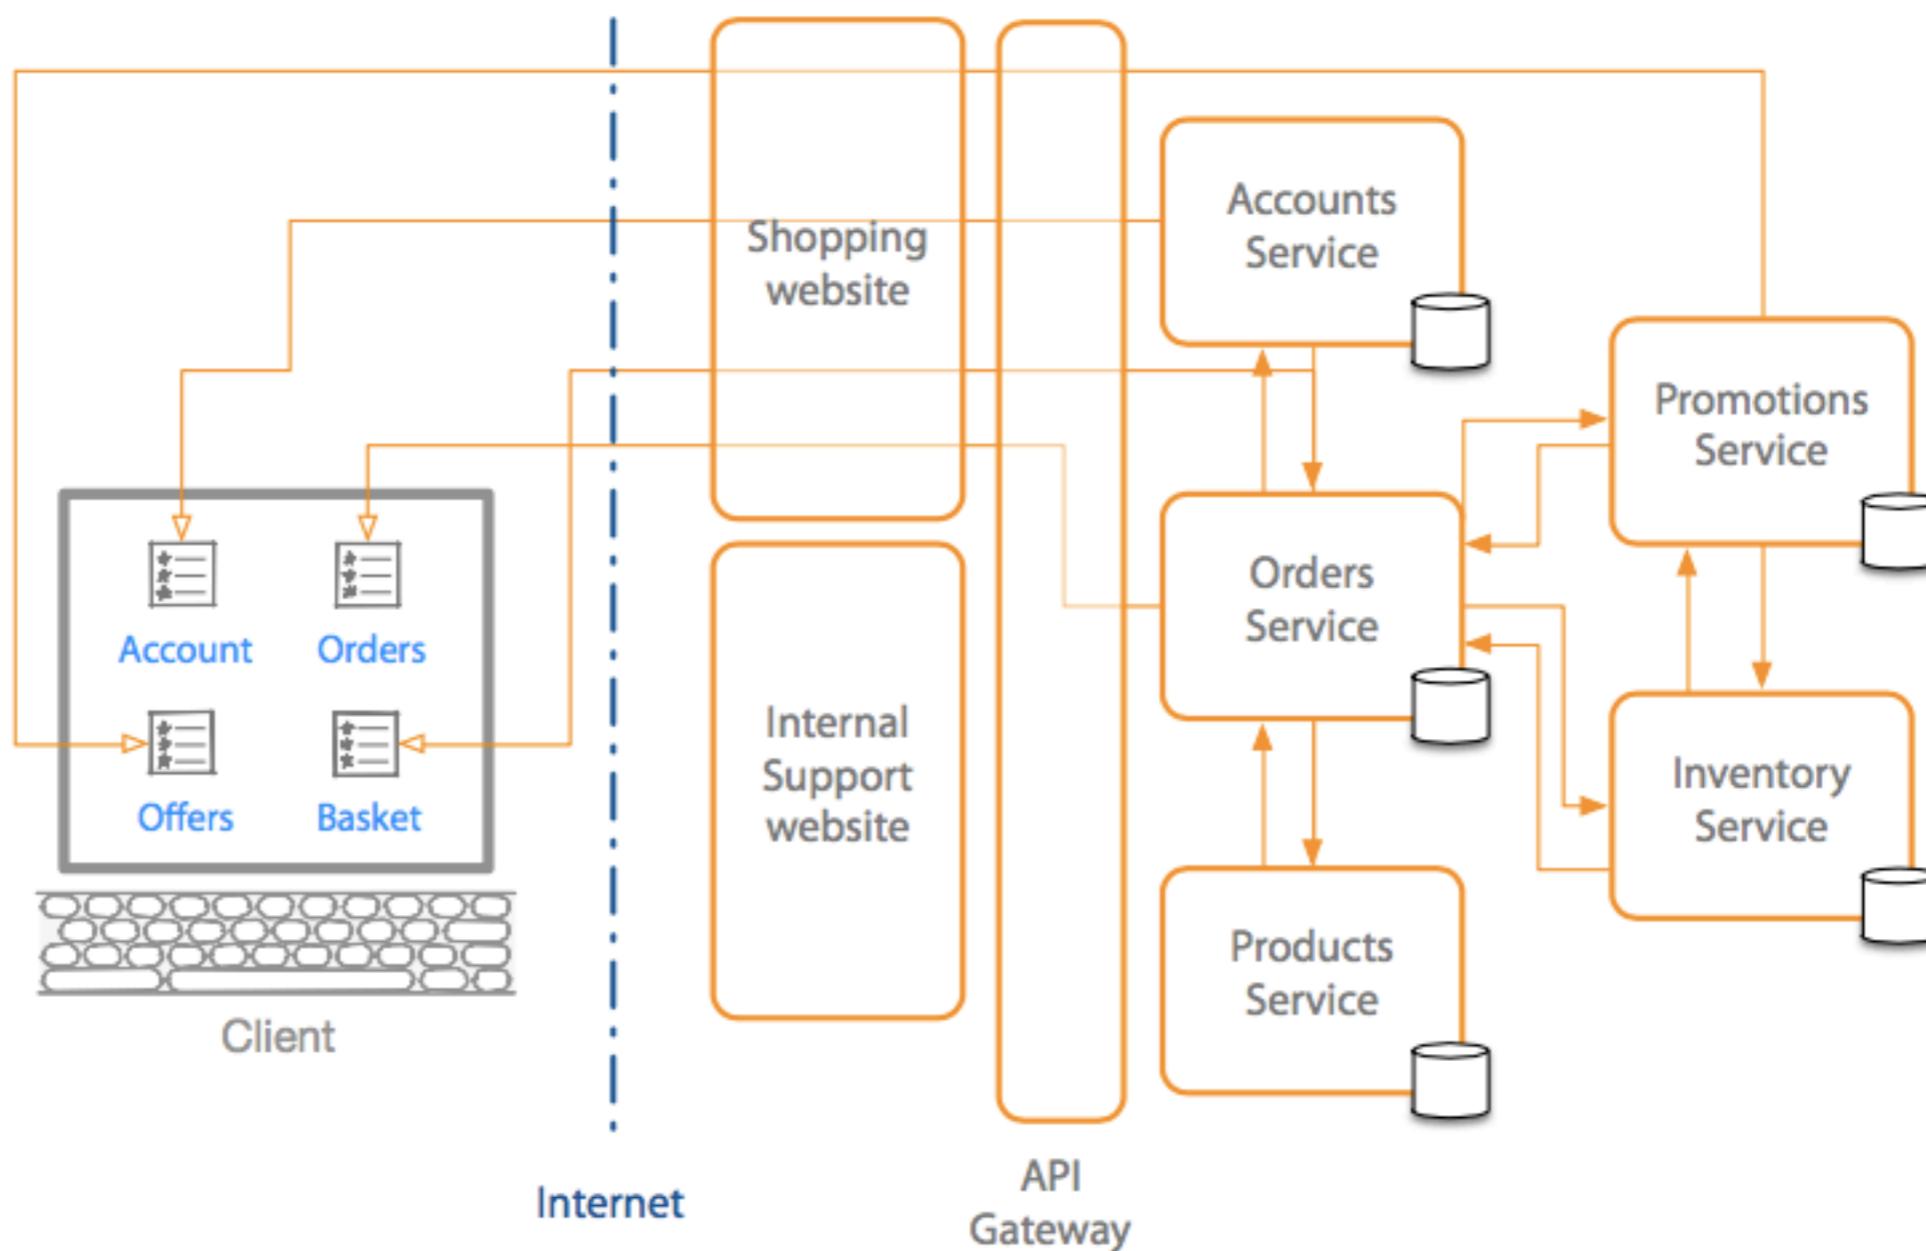
Monolithic

- Typical enterprise application (No restriction on size)
- Large codebase
- Longer development times
- Challenging deployment
- Inaccessible features
- Fixed technology stack

Monolithic

- High levels of coupling
 - Between modules
 - Between services
- Failure could affect whole system
- Scaling requires duplication of the whole
- Single service on server
- Minor change could result in complete rebuild
- Easy to replicate environment

Microservices



Resources

<https://microservices.io>

[Microservice Architecture](#)

Supported by Kong

Posts

22 Mar 2018 » [microXchg 2018 - Managing data consistency in a microservice architecture using Sagas](#)

20 Feb 2018 » [CodeFreeze 2018 - There is no such thing as a microservice!](#)

04 Dec 2017 » [QCONSF 2017 presentation - ACID Is So Yesterday, Maintaining Data Consistency with Sagas](#)

04 Dec 2017 » [Upcoming public training - Microservices at UMN, Minneapolis in January](#)

01 Aug 2017 » [Presentation - Solving distributed data management problems in a microservice architecture](#)

24 Jul 2017 » [Revised data patterns](#)

26 Mar 2017 » [There is no such thing as a microservice!](#)

24 Feb 2017 » [New book - Microservice patterns](#)

12 Feb 2017 » [How to apply the pattern language](#)

06 Jun 2016 » [Fantastic presentation by Eric Evans on DDD and microservices](#)

29 Feb 2016 » [One day microservices class in Oakland, CA](#)

21 Feb 2016 » [Microservice chassis pattern](#)

28 Jul 2015 » [What's new with #microservices - July 28, 2015: integration platforms, new article on microservices and IPC](#)

23 Jun 2015 » [What's new with #microservices - June 23, 2015: @crichton and microservices training](#)

06 Jun 2015 » [What's new with #microservices - June 6, 2015: @martinfowler, @crichton, @adrianco](#)

[Patterns](#) [Articles](#) [Presentations](#) [Resources](#) [Assessment Platform new](#) [Other Languages](#) [About](#)

What are microservices?

Microservices - also known as the microservice architecture - is an architectural style that structures an application as a collection of loosely coupled services, which implement business capabilities. The microservice architecture enables the continuous delivery/deployment of large, complex applications. It also enables an organization to evolve its technology stack.

Microservices are not a silver bullet

The microservice architecture is not a silver bullet. It has several drawbacks. Moreover, when using this architecture there are numerous issues that you must address. The microservice architecture pattern language is a collection of patterns for applying the microservice architecture. It has two goals:

1. The pattern language enables you to decide whether microservices are a good fit for your application.
2. The pattern language enables you to use the microservice architecture successfully.

Where to start?

A good starting point is the [Monolithic Architecture pattern](#), which is the traditional architectural style that is still a good choice for many applications. It does, however, have numerous limitations and issues and so a better choice for large/complex applications is the [Microservice architecture pattern](#).

Example microservices applications

Want to see an example? Check out Chris Richardson's Money Transfer and Kanban board examples.

[See code](#)

How to apply the pattern language

An article that describes how to develop a microservice architecture by applying the patterns

[Learn more](#)

Microservices adoption: Who is using microservices?

Many companies are either using microservices or considering using them. Read the case studies...



[About Microservices.io](#)

Microservices.io is brought to you by Chris Richardson. Experienced software architect, author of POJOs in Action and the creator of the original CloudFoundry.com. His latest startup is [eventuate.io](#), a microservices application platform.

Learn more about microservices

Chris offers a comprehensive set of resources for learning about microservices including articles, an O'Reilly training video, and example code.

[Learn more](#)

[Microservices](#)

Resources

[PRODUCTS](#)[SOLUTIONS](#)[RESOURCES](#)[SUPPORT](#)[PRICING](#)[BLOG](#)[FREE TRIAL](#)[CONTACT US](#)[Home](#) > [Blog](#) > [Tech](#) > [Introduction to Microservices](#)[BLOG](#) [TECH](#)

Chris Richardson of Eventuate, Inc. May 19, 2015

Introduction to Microservices

[microservices](#), [monolithic application](#), [Docker](#), [cloud](#) [TWITTER](#) [LINKEDIN](#) [Y](#) [F](#) [@](#)

Editor – This seven-part series of articles is now complete:

1. [Introduction to Microservices \(this article\)](#)
2. [Building Microservices: Using an API Gateway](#)
3. [Building Microservices: Inter-Process Communication in a Microservices Architecture](#)
4. [Service Discovery in a Microservices Architecture](#)
5. [Event-Driven Data Management for Microservices](#)
6. [Choosing a Microservices Deployment Strategy](#)
7. [Refactoring a Monolith into Microservices](#)

You can also download the complete set of

<https://www.nginx.com/blog/introduction-to-microservices>

ABOUT NGINX

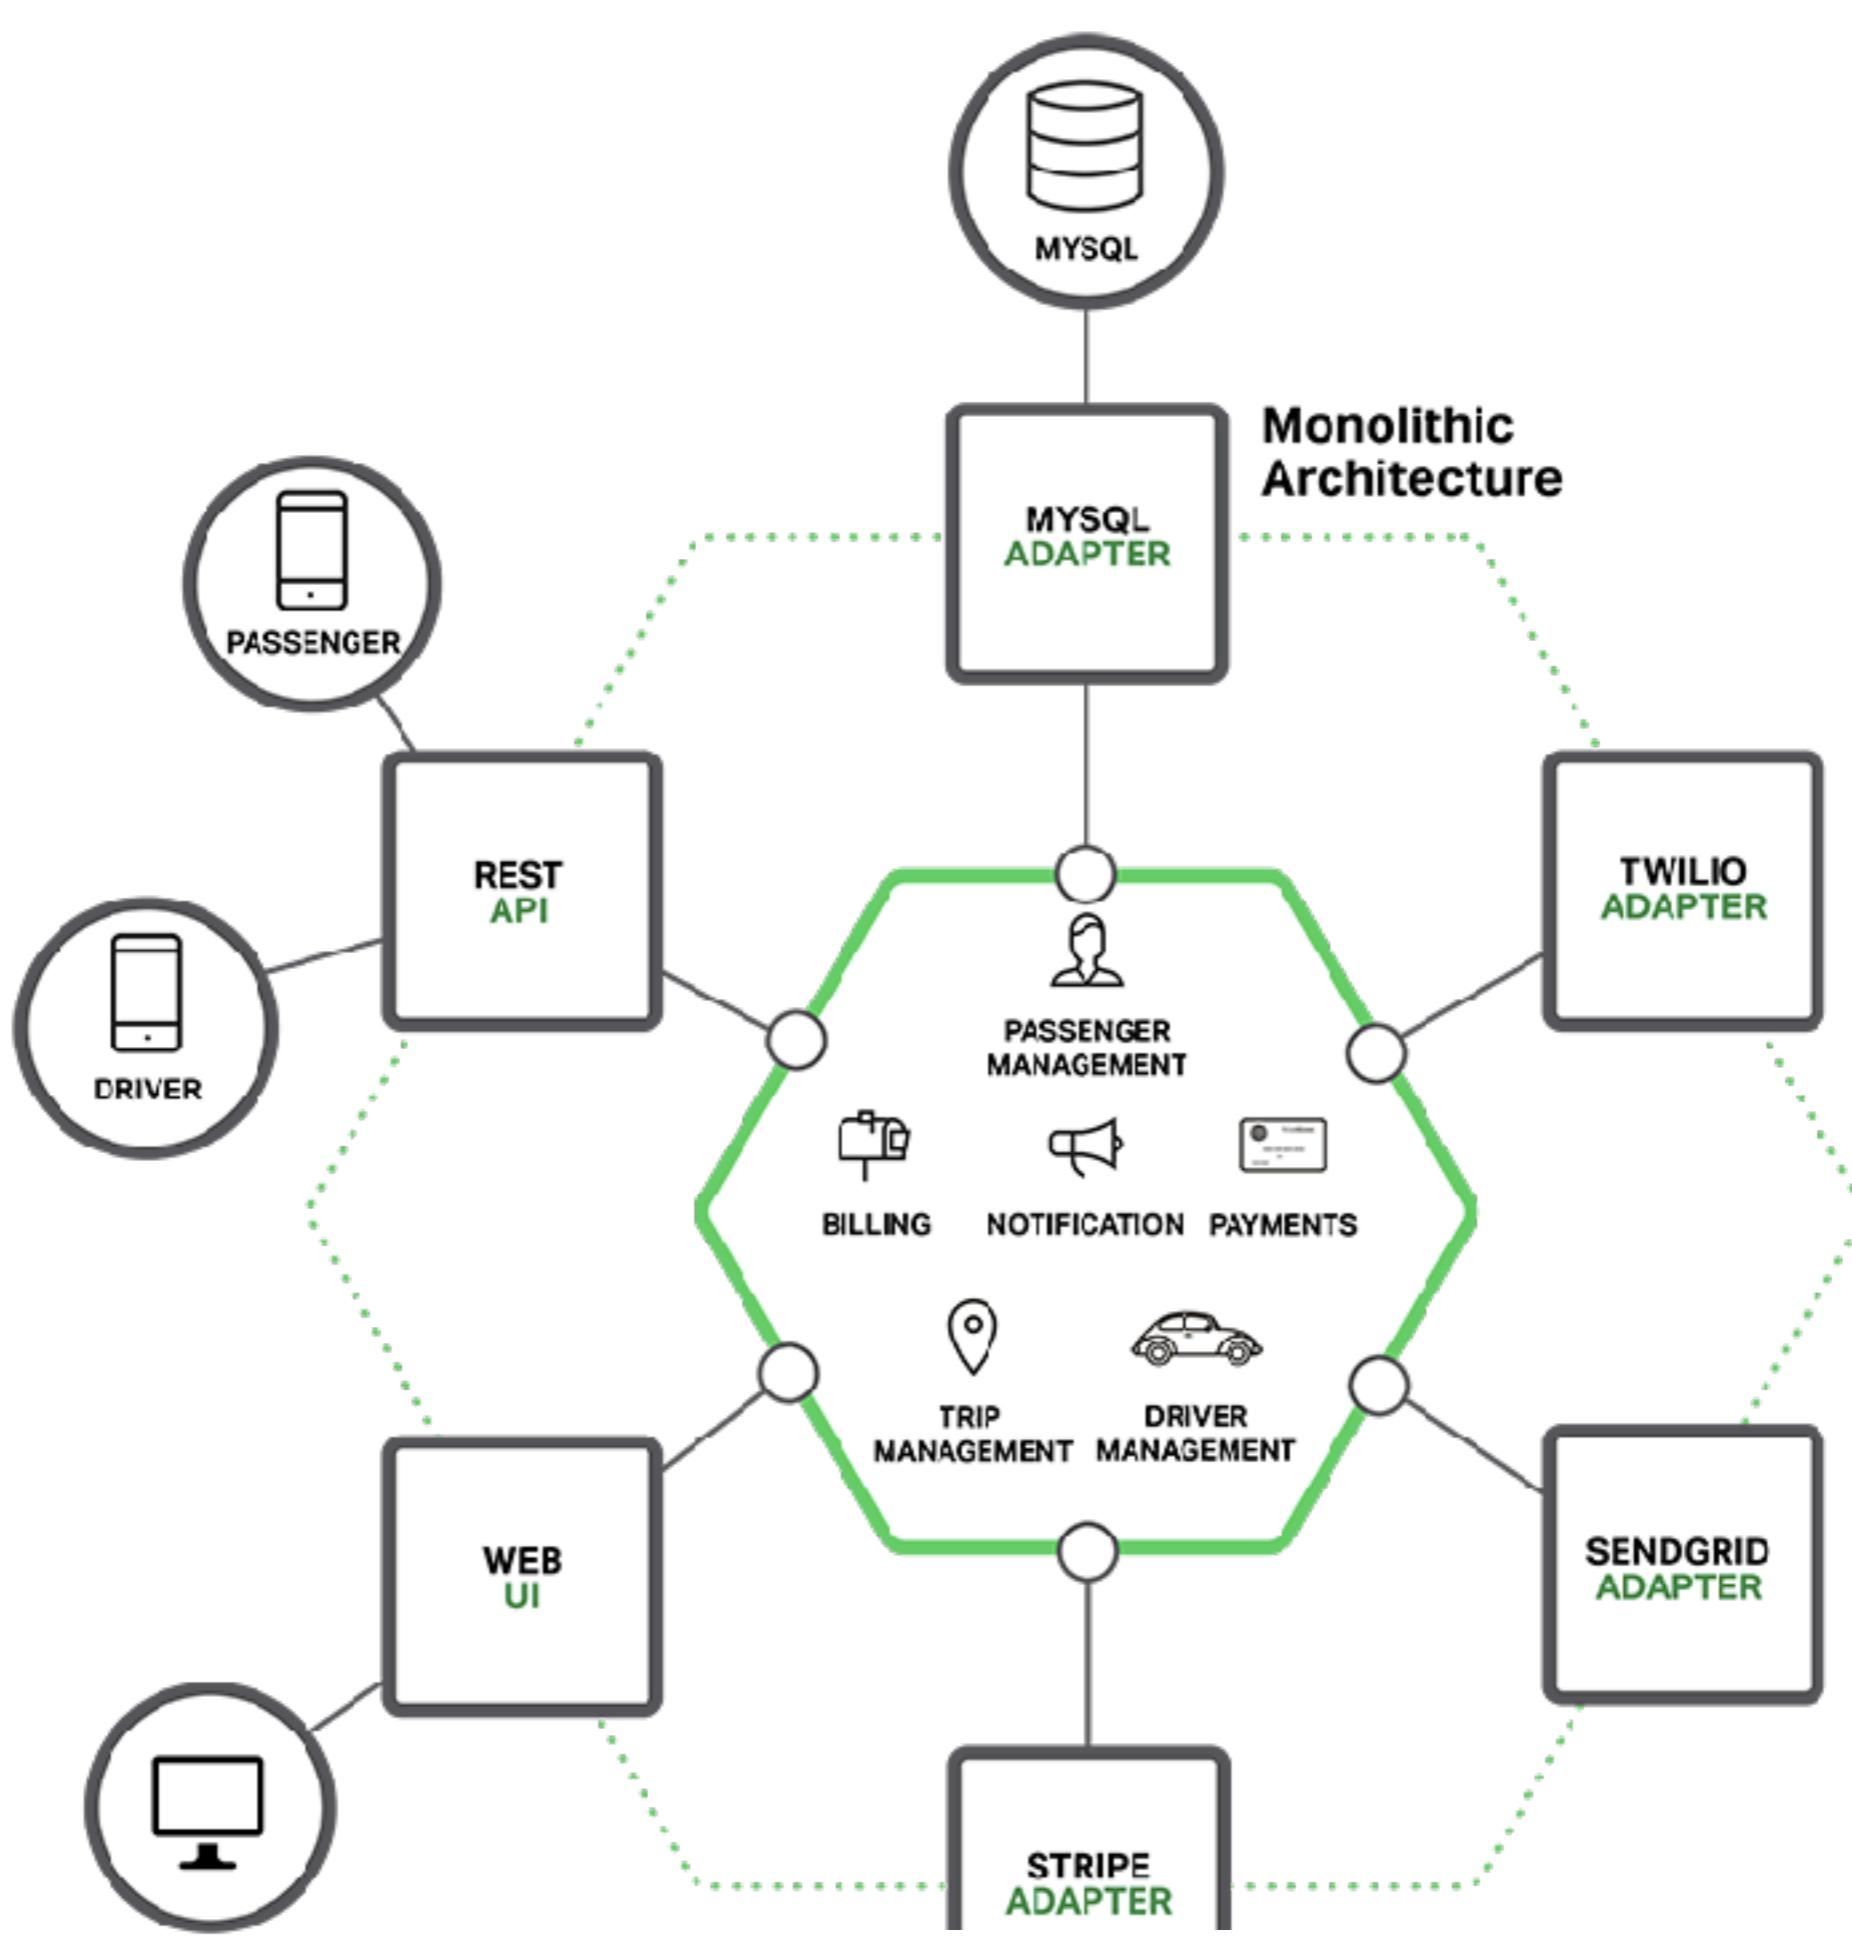
NGINX, Inc., now part of F5, is the company behind the popular open source project. We offer a suite of technologies for developing and delivering modern applications. Together with F5, our combined solution bridges the gap between NetOps and DevOps, with multi-cloud application services that span from code to customer.

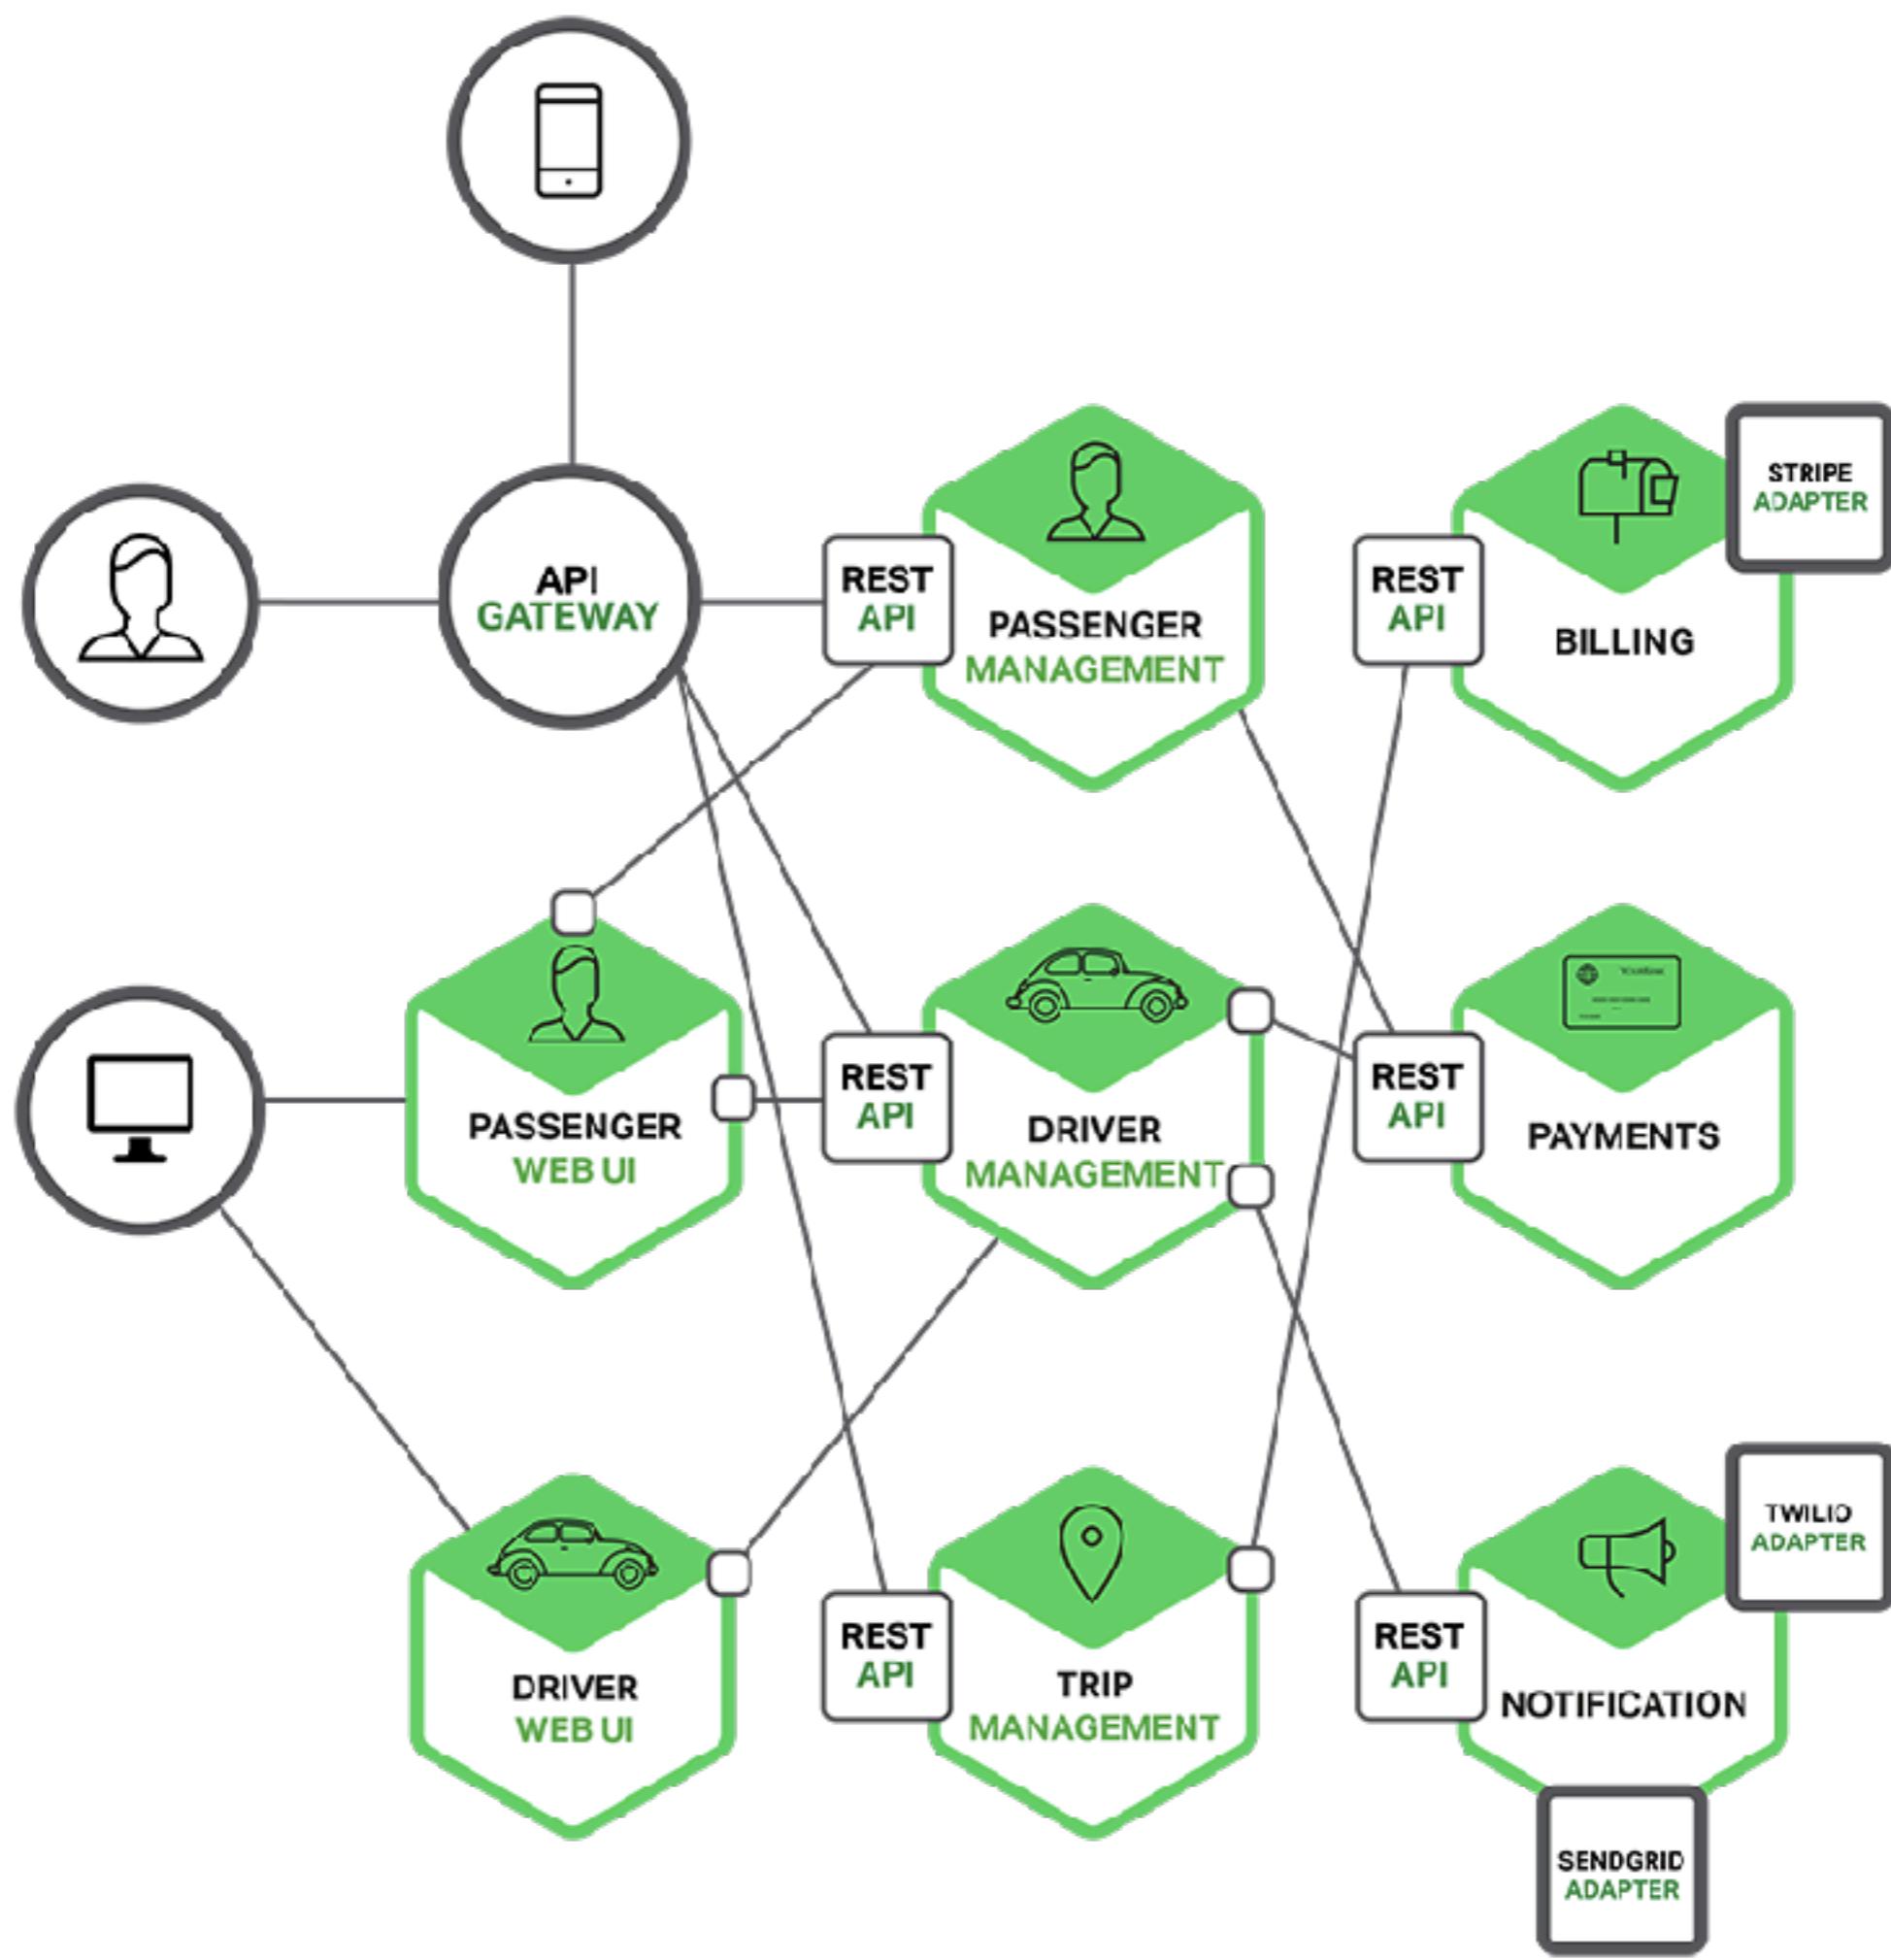
Learn more at [nginx.com](#) or join the conversation by following [@nginx](#) on Twitter.

 SEARCH

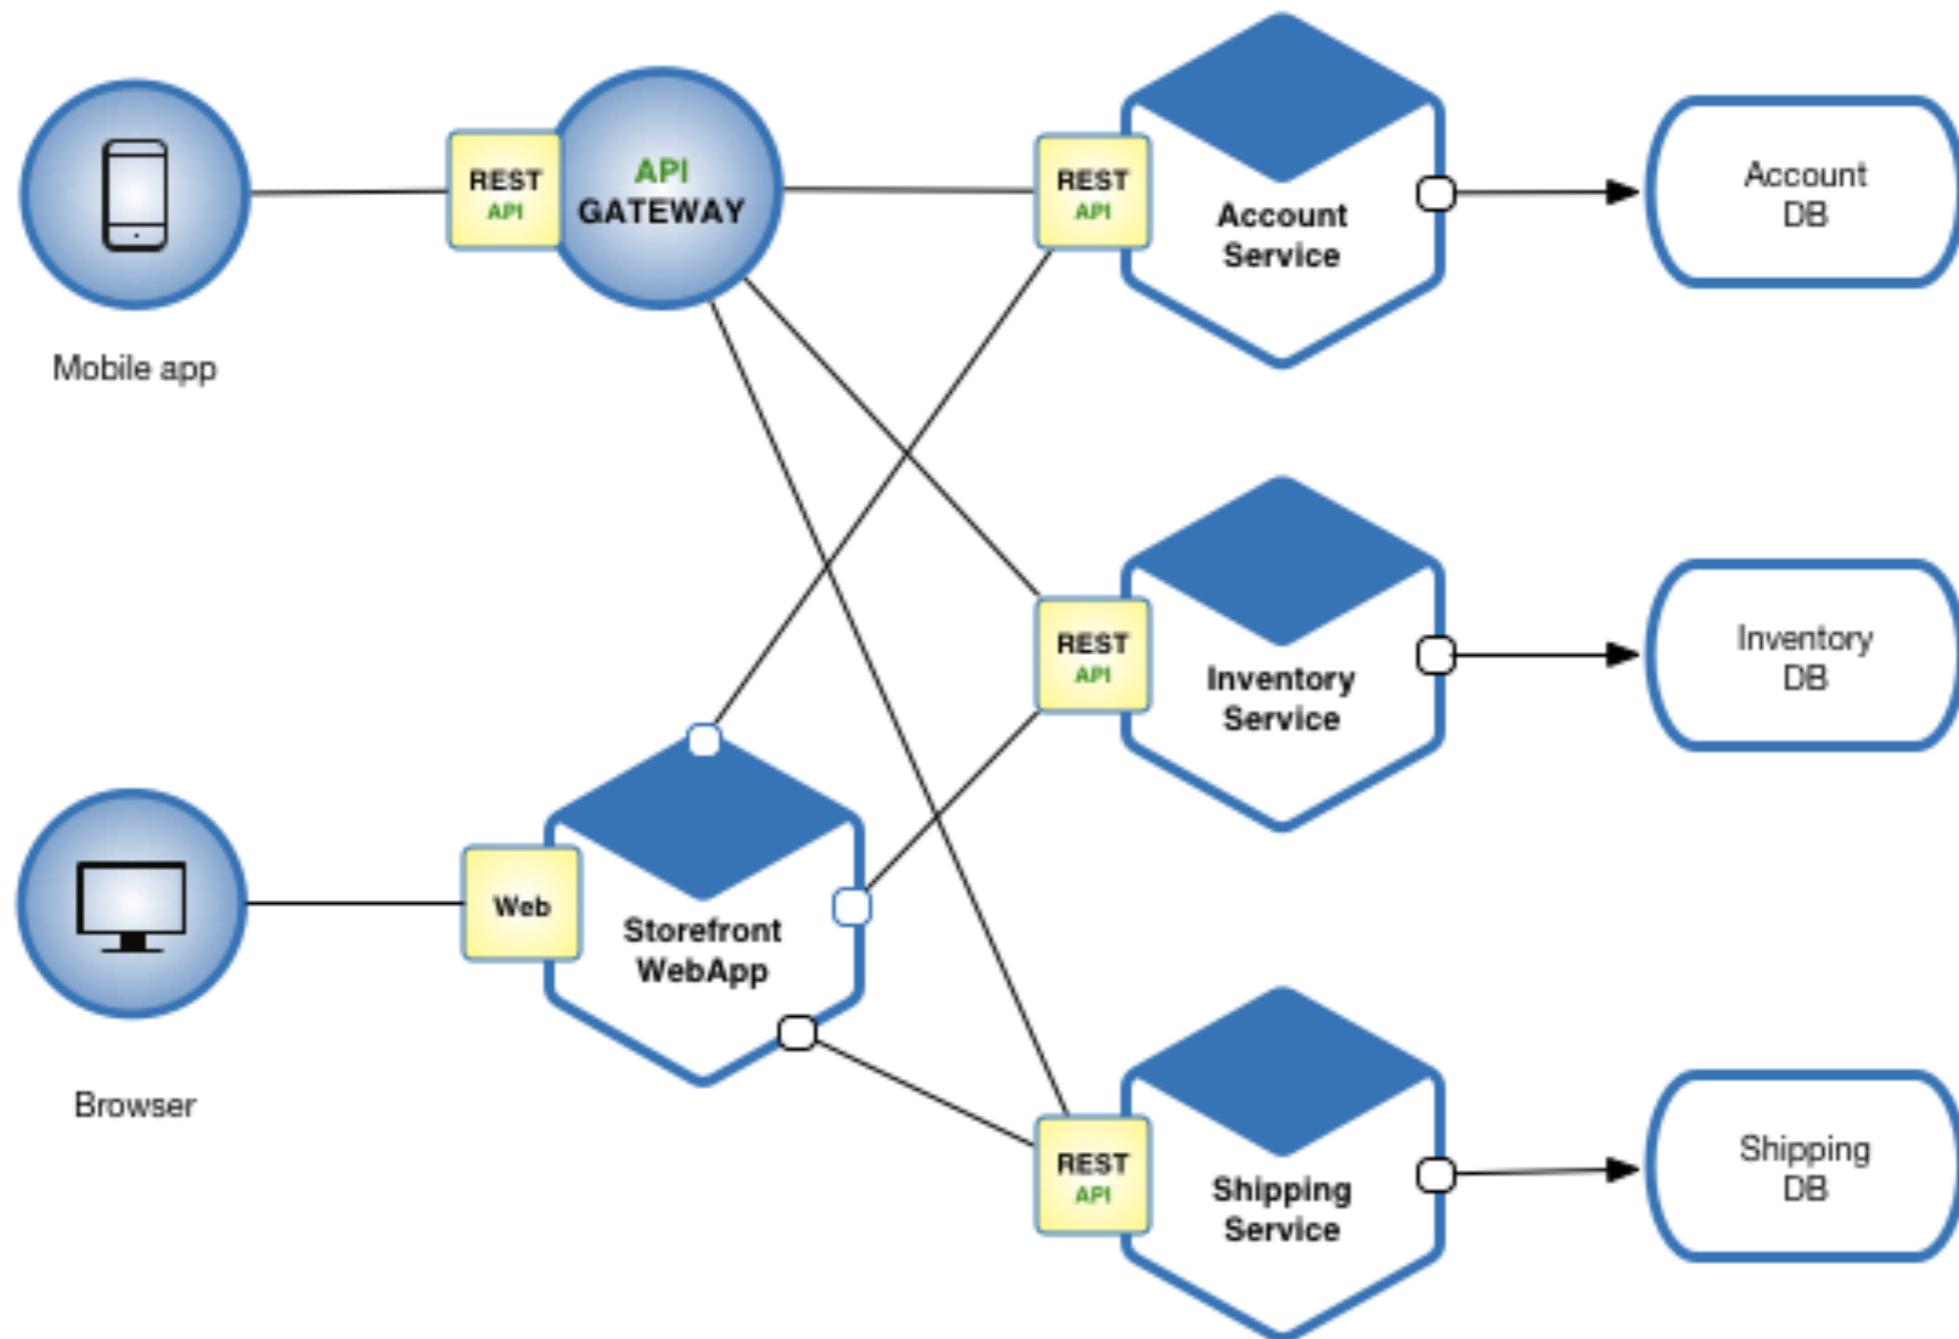
CATEGORIES

[Events](#)
[News](#)
[Tech](#)
[Opinion](#)
[More »](#)





E-commerce Application



Netflix OSS



Getting Started

How can you get started quickly?

For the simple approach, try out our [ZeroToDocker](#) container images. After downloading the images, you can be up and running NetflixOSS in just a few minutes.

After you've tackled that, check out the [IBM ACME Air](#) and [Flux Capacitor](#) apps, and the [Zero-to-Cloud workshop](#).

See these [CloudFormation templates](#) on Answers For AWS for use of NetflixOSS through CloudFormation.



[zerotodocker](#)

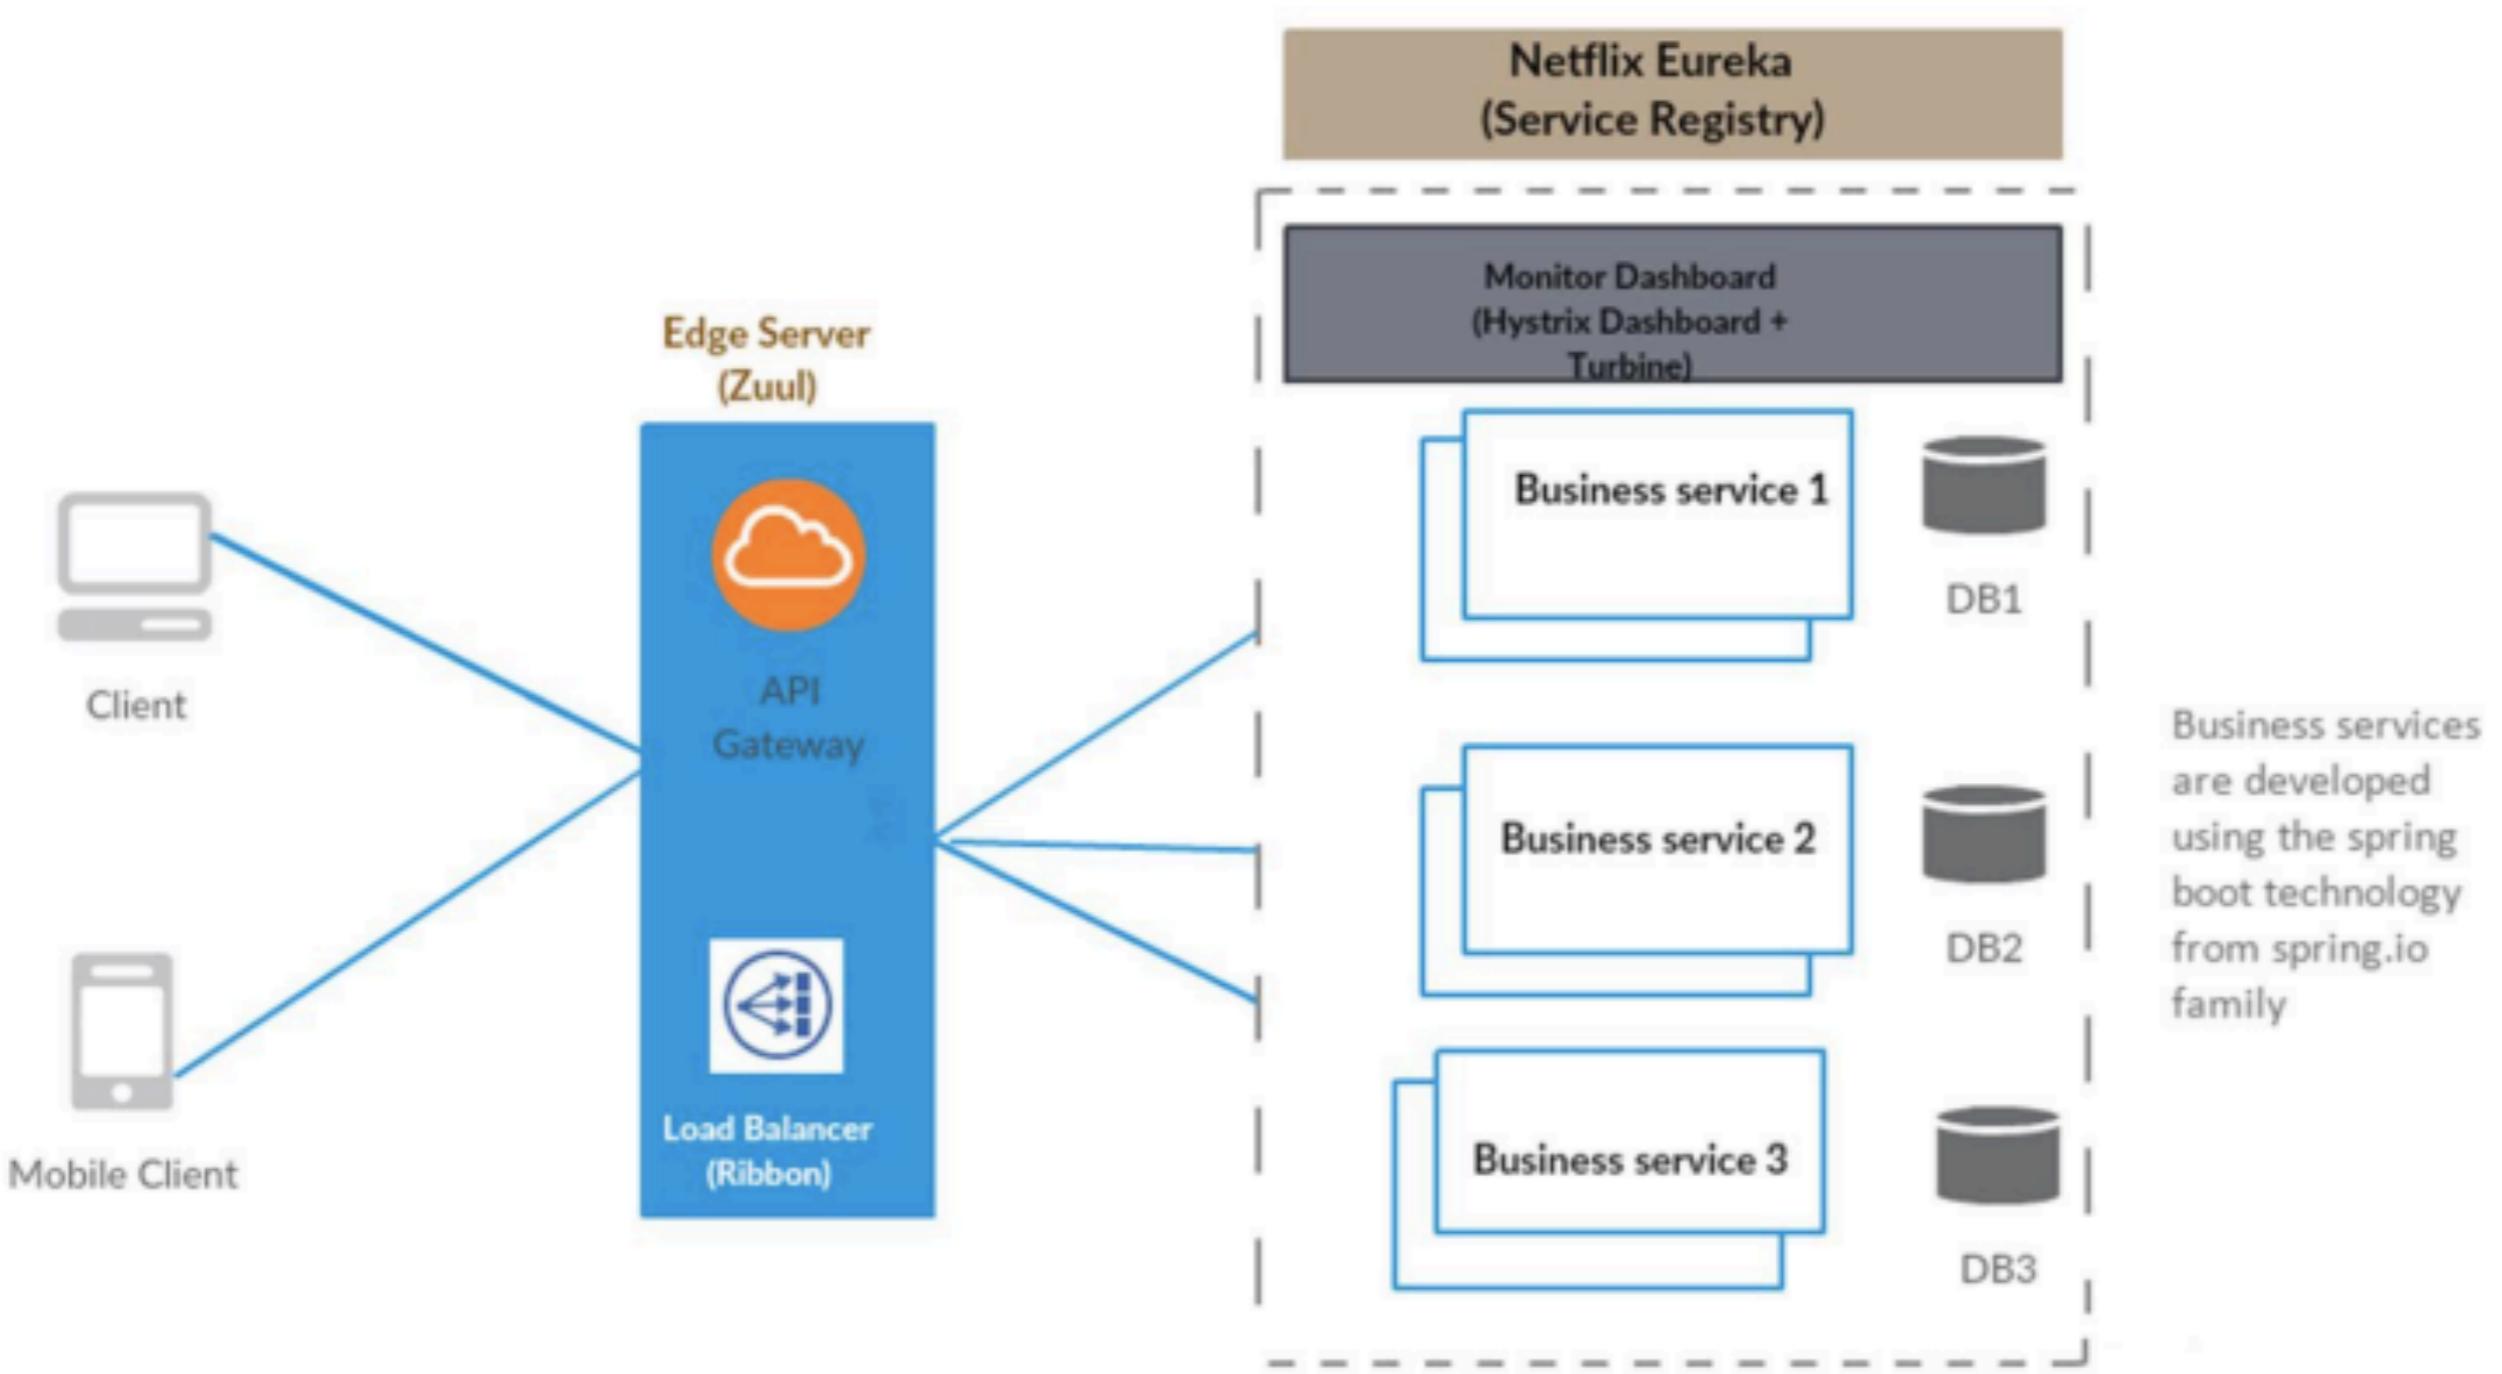
0 ● 0

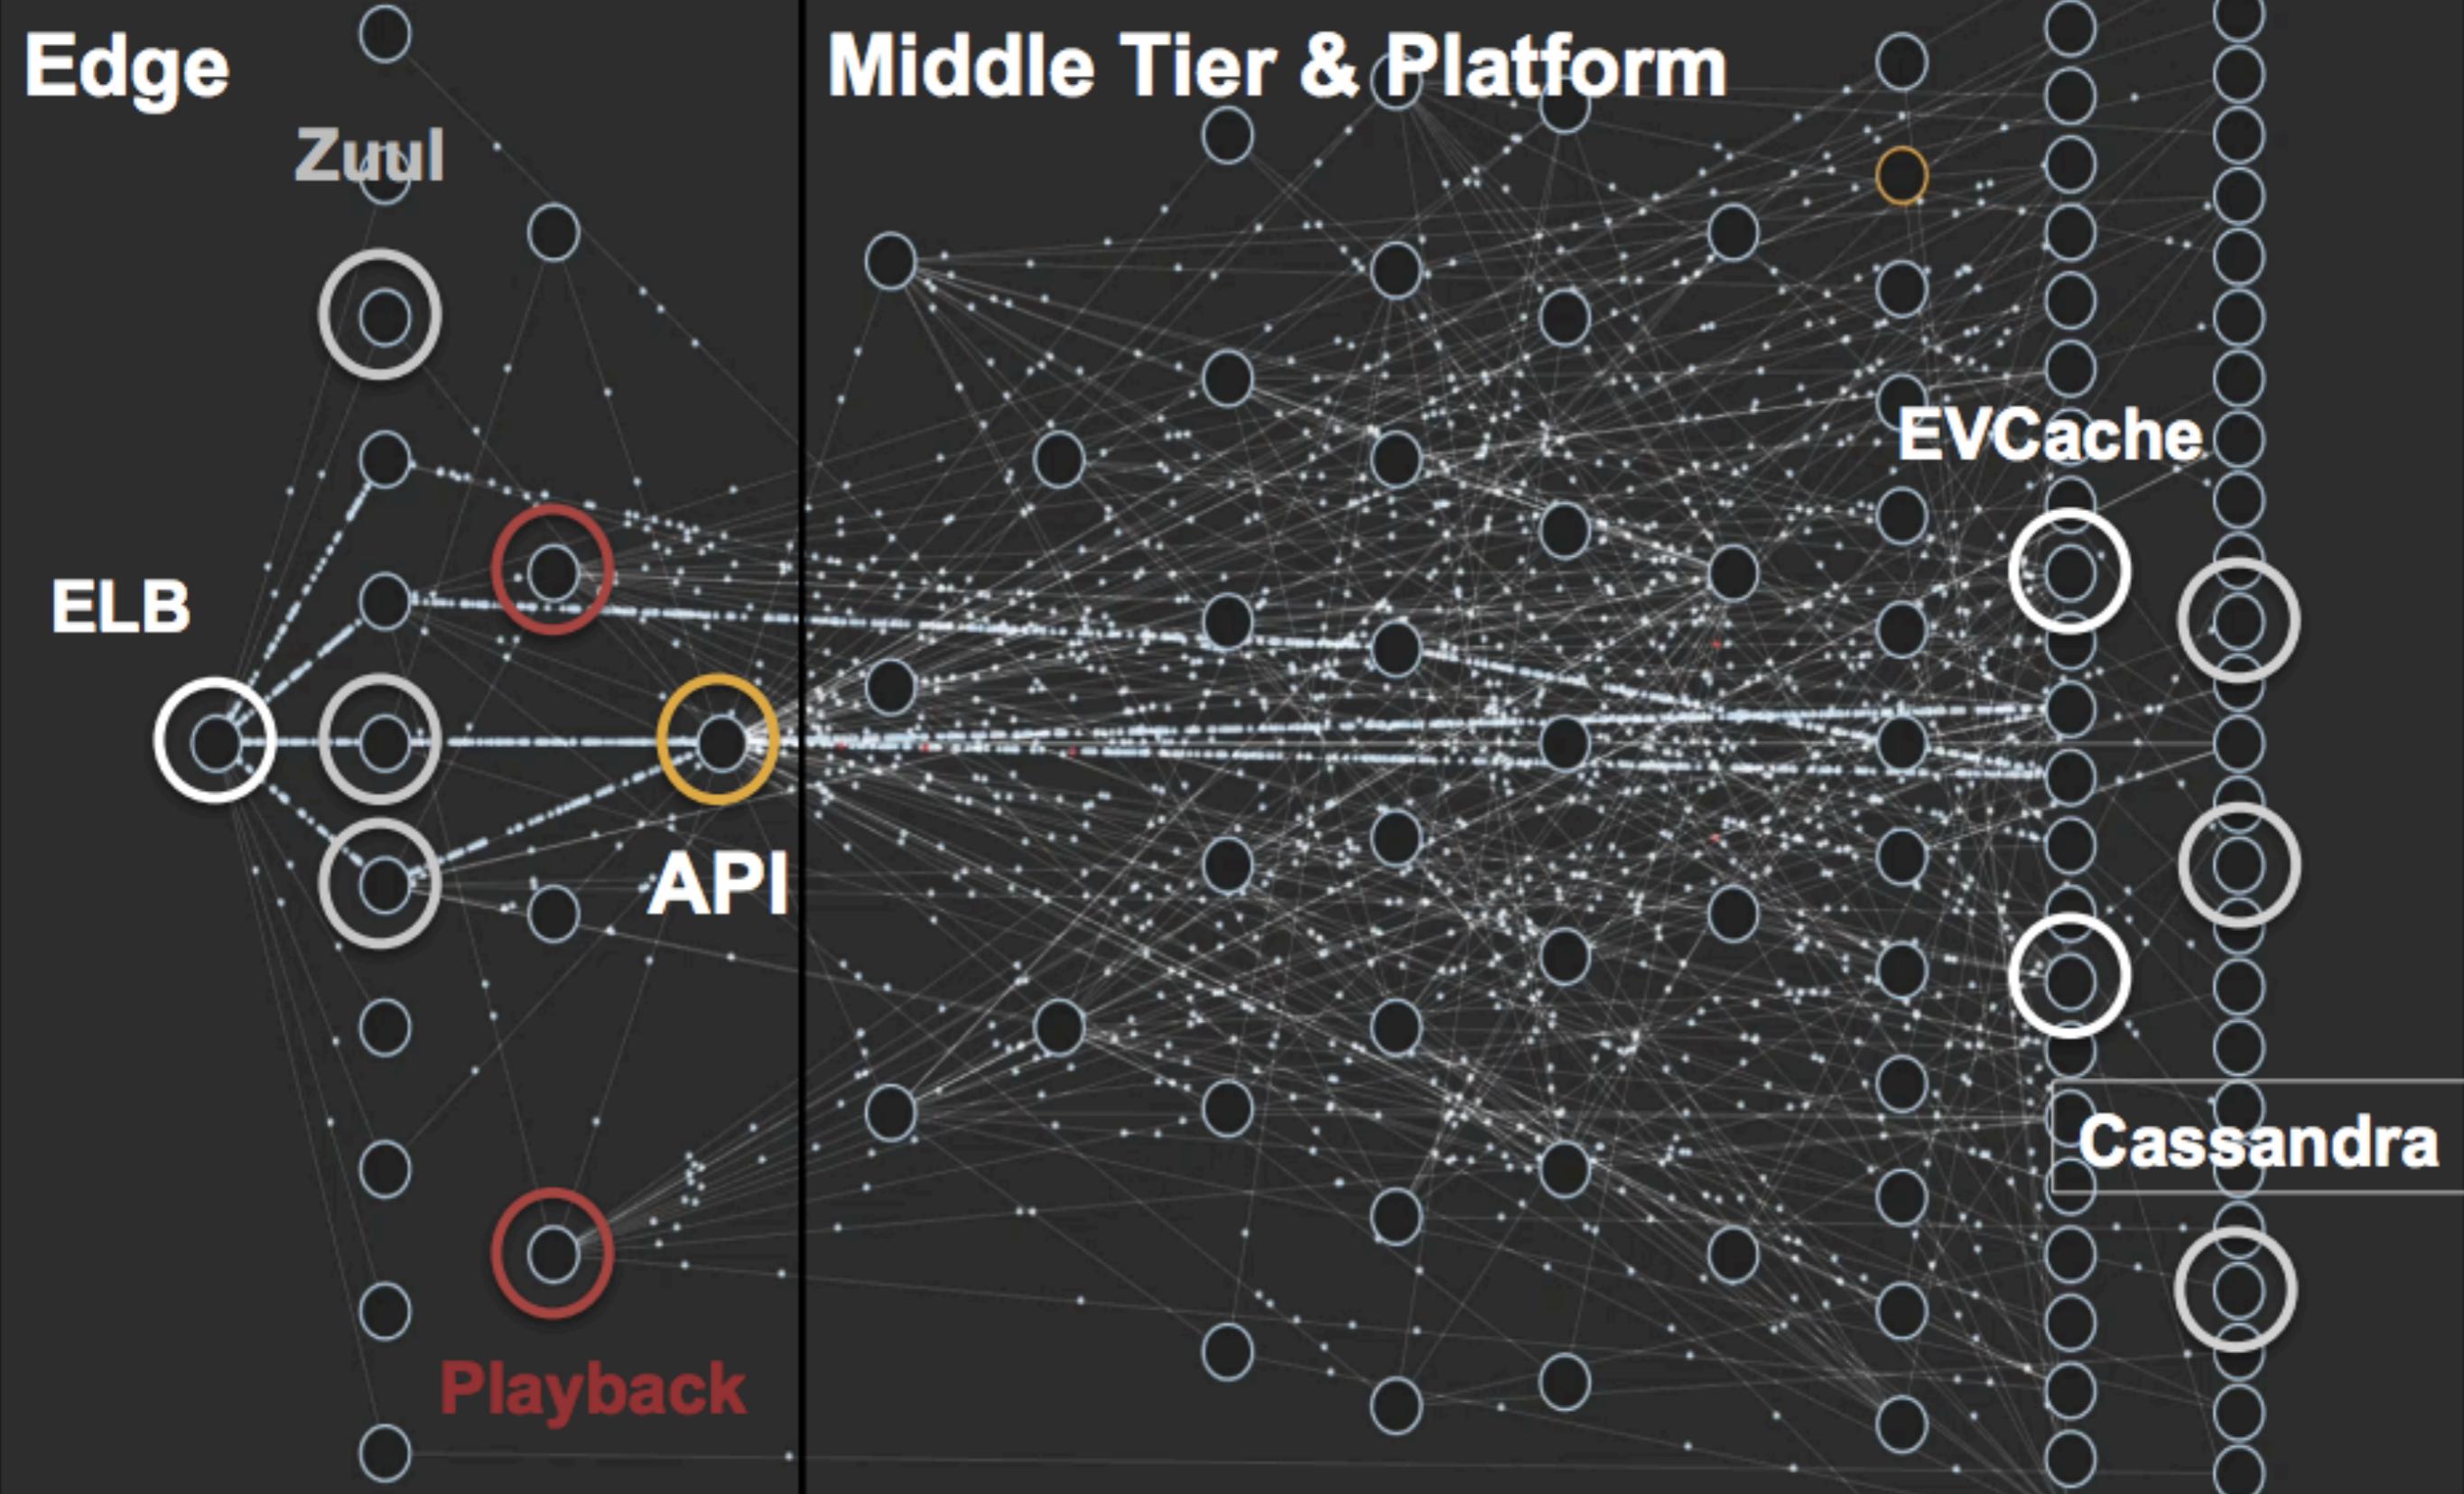
Netflix Open Source Software Center

Netflix is committed to open source. Netflix both leverages and provides open source technology focused on providing the leading Internet television network. Our technology focuses on providing immersive experiences across all Internet-connected screens. Netflix's deployment technology allows for continuous build and integration into our worldwide deployments serving members in over 50 countries. Our focus on reliability defined the bar for cloud based elastic deployments with several layers of failover. Netflix also provides the technology to operate services responsibly with operational insight, peak performance, and security. We provide technologies for data (persistent & semi-persistent) that serve the real-time load to our 62 million members, as well as power the big data analytics that allow us to make informed decisions on how to improve our service. If you want to learn more, jump into any of the functional areas below to learn more.

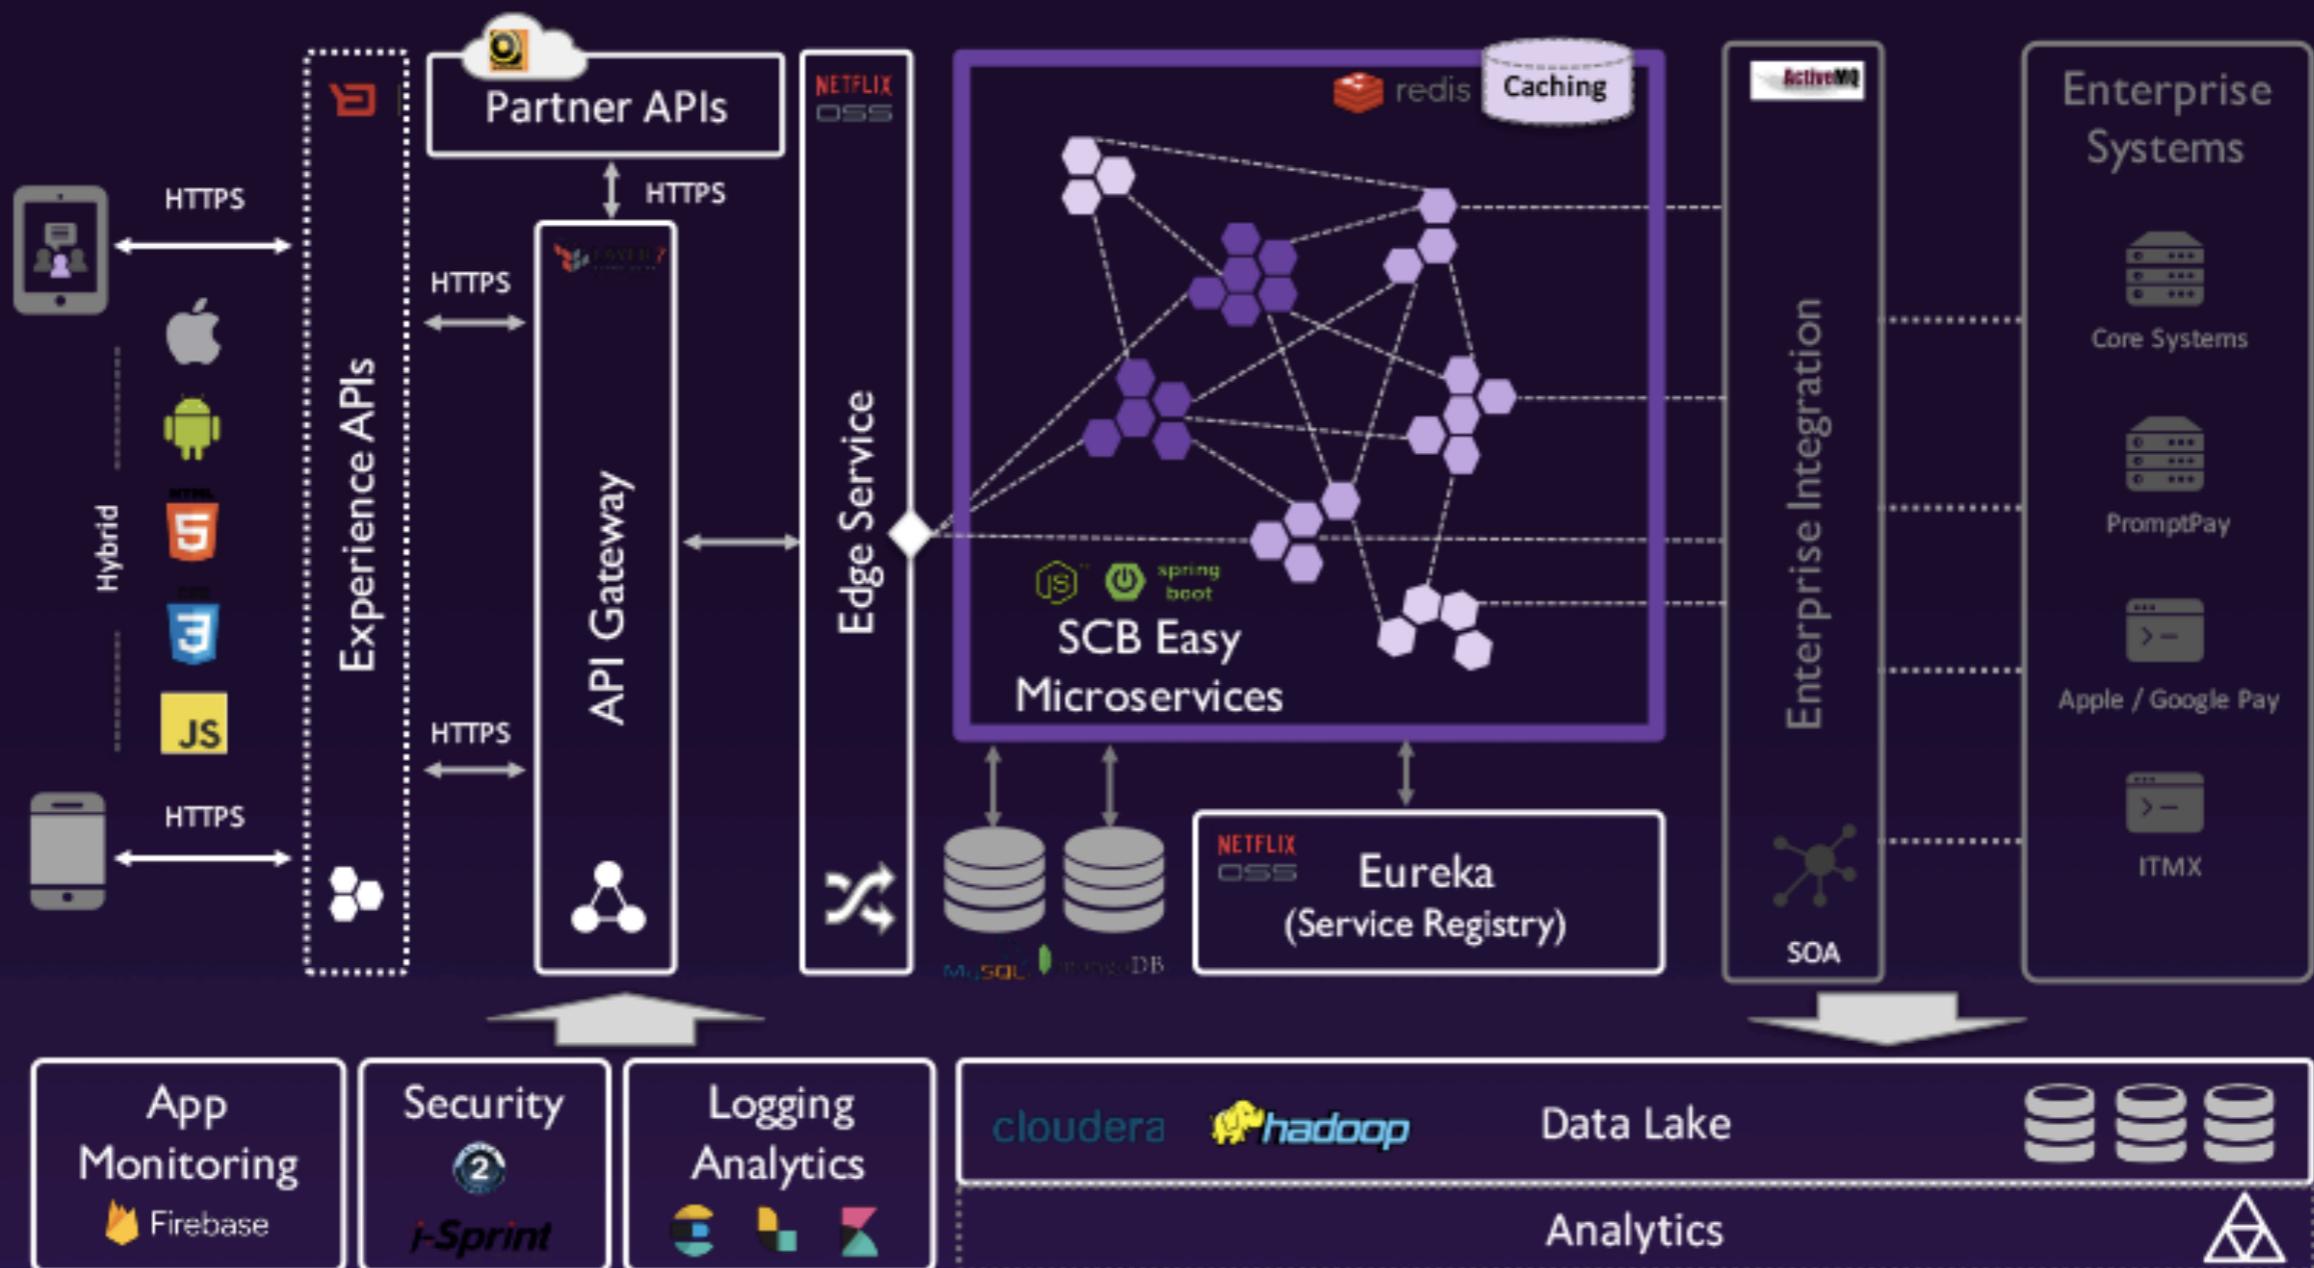
<https://netflix.github.io>

NetFlix Architecture

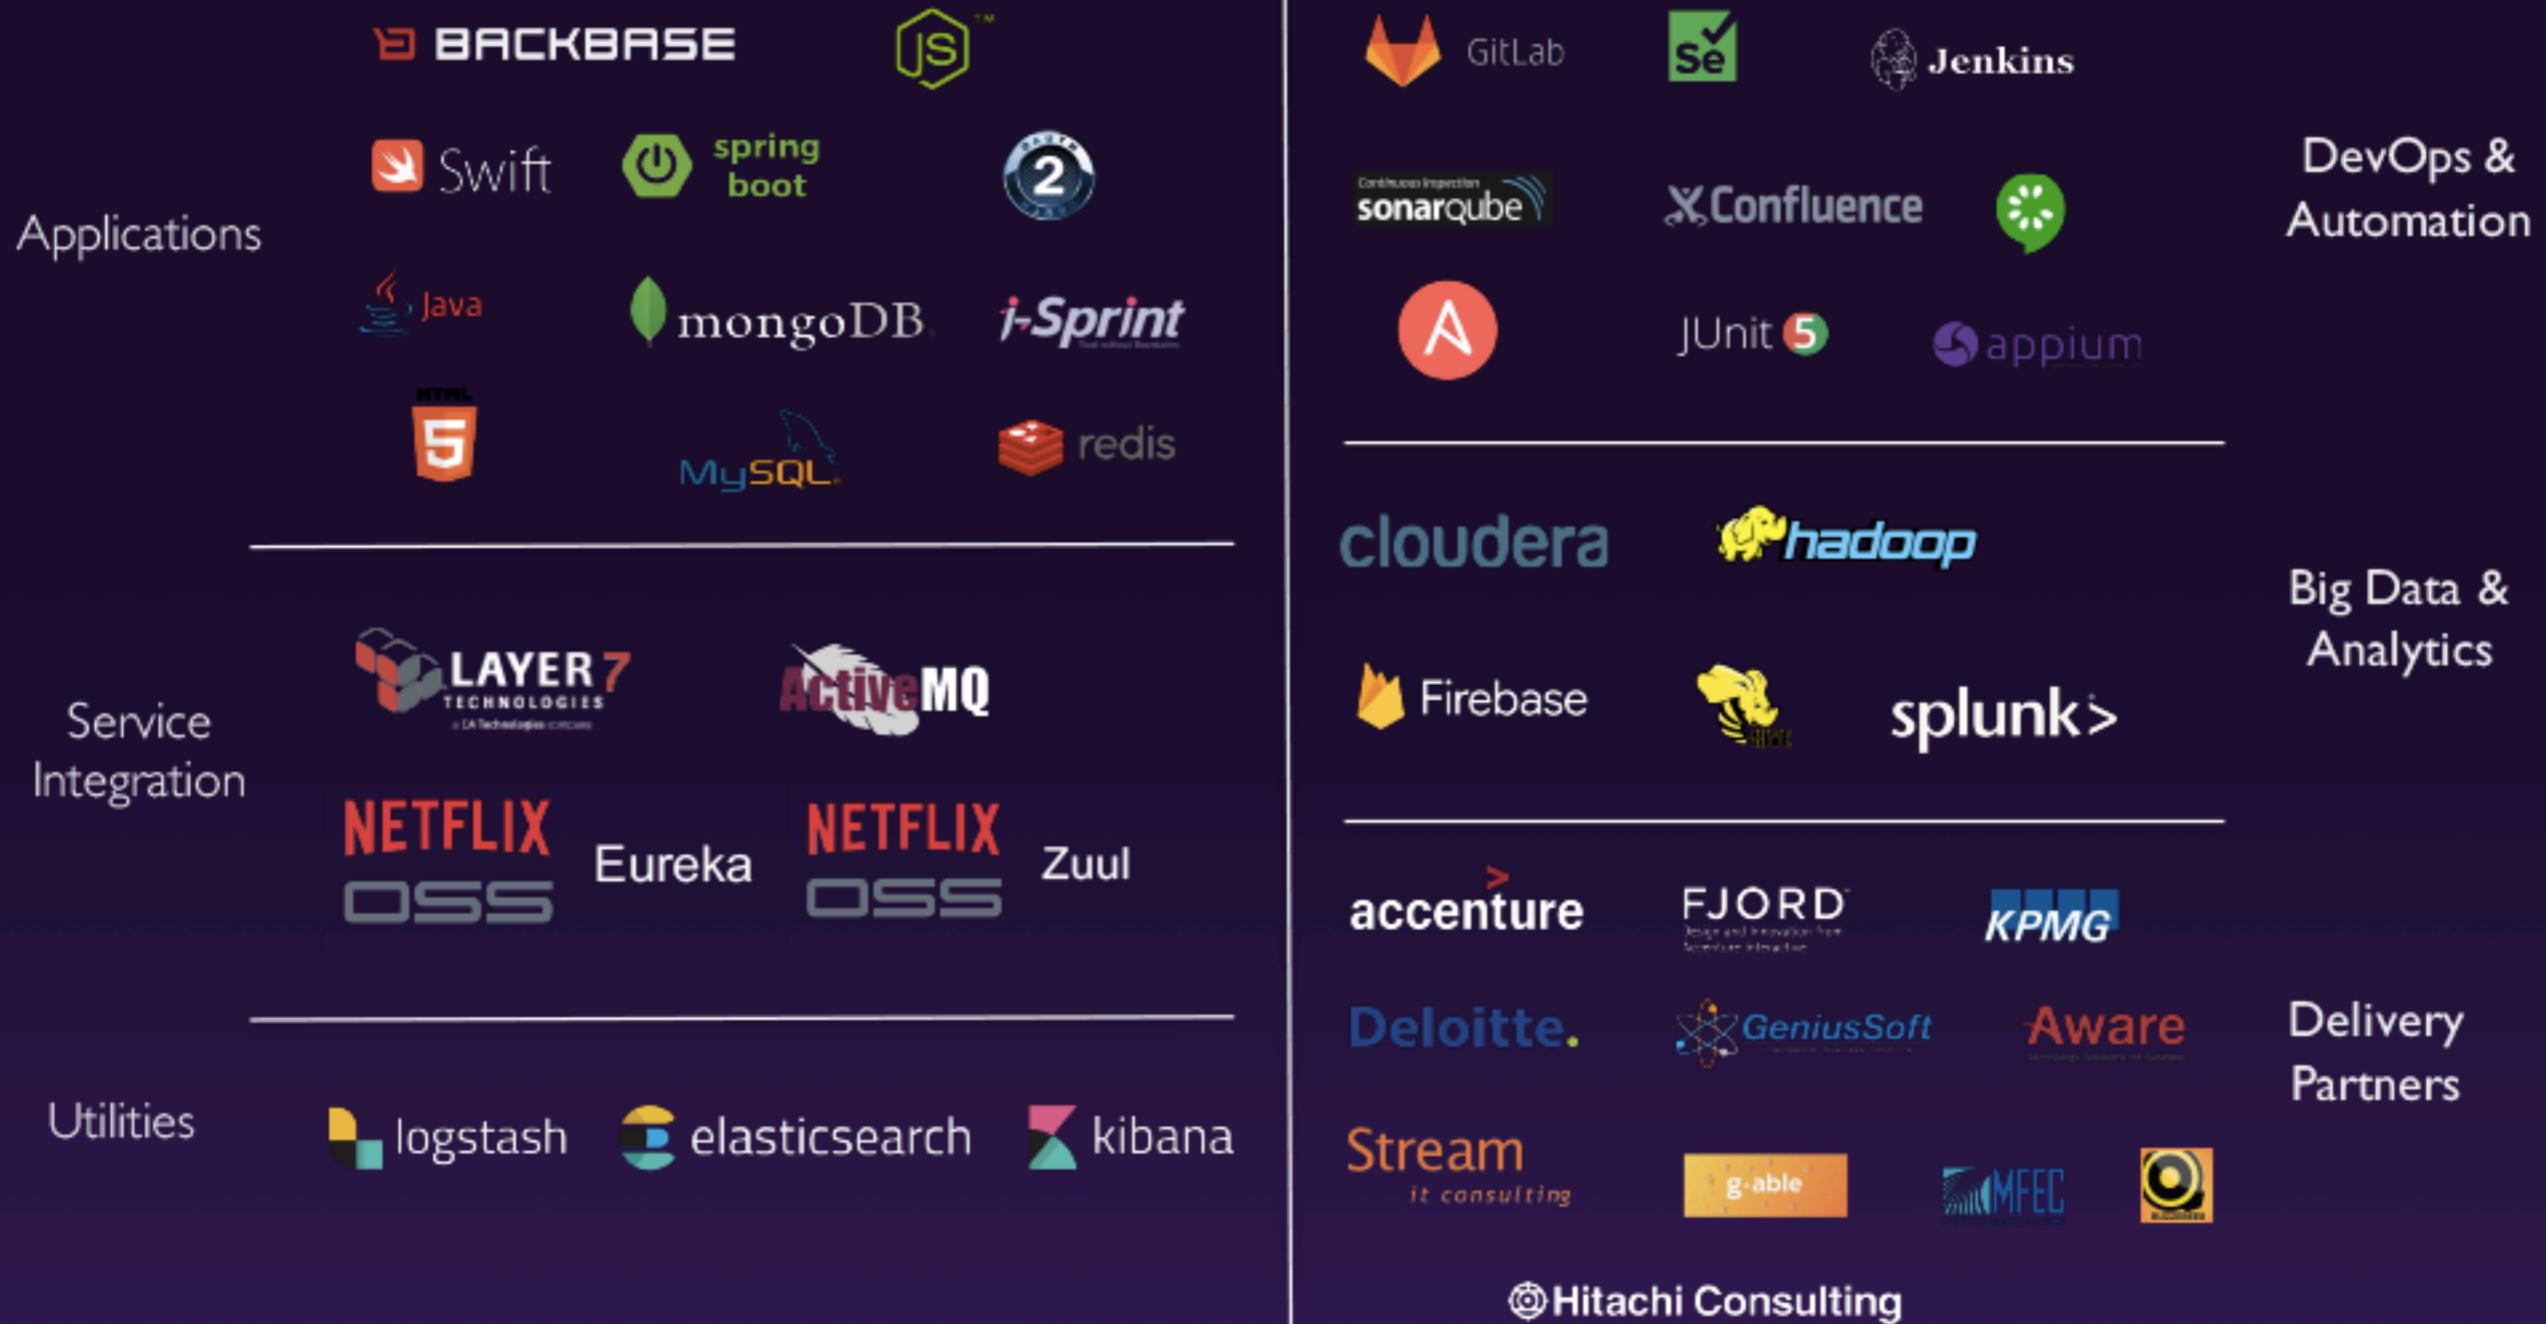




SCB EASY MICROSERVICES ARCHITECTURE



SCB EASY TECHNOLOGY



Microservices

- Application(s) powered by multiple services
- Small service with a single focus
- Lightweight communication mechanism
 - Both client to service and service to service
- Technology agnostic API
- Independent data storage

Microservices

Micro sized services provide

- Efficiently scalable applications
- Flexible applications
- High performance applications

Microservices

- Independently changeable
- Independently deployable
- Distributed transactions
- Centralized tooling for management

Why Now?

- Need to respond to change quickly
- Need for reliability
- Business domain-driven design
- Automated test tools
- Release and deployment tools
- On-demand hosting technology
- On-line cloud services
- Need to embrace new technology
- Asynchronous communication technology
- Simpler server side and client side technology

Microservices Pros

- Shorter development times
- Reliable and faster deployment
- Enables frequent updates
- Decouple the changeable parts
- Security
- Increased uptime
- Fast issue resolution
- Highly scalable and better performance
- Better ownership and knowledge
- Right technology
- Enables distributed teams

Polyglot Persistence



Martin Fowler

16 November 2011

In 2006, my colleague Neal Ford coined the term [Polyglot Programming](#), to express the idea that applications should be written in a mix of languages to take advantage of the fact that different languages are suitable for tackling different problems. Complex applications combine different types of problems, so picking the right language for the job may be more productive than trying to fit all aspects into a single language.

Over the last few years there's been an explosion of interest in new languages, particularly functional languages, and I'm often tempted to spend some time delving into Clojure, Scala, Erlang, or the like. But my time is limited and I'm giving a higher priority to another, more significant shift, that of the [DatabaseThaw](#). The first drips have been coming through from clients and other contacts and the prospects are enticing. I'm confident to say that if you starting a new strategic enterprise application you should no longer be assuming that your persistence should be relational. The relational option might be the right one - but you should seriously look at other alternatives.

Polyglot Persistence



Twelve Factor

Twelve Factor App

- Source: <http://12factor.net>
- Initial proposed to build SaaS apps for Heroku
- Principles translate well to cloud and container native application
- Translate to microservices

Principle 1: Codebase

One codebase tracked in revision control, many deploys

Principle 2: Dependencies

Explicitly declare and isolate dependencies

Principle 3: Config

Store config in the environment

Principle 4: Backing Services

Treat backing services as attached resources

Principle 5: Build Release Run

Strictly separate build and run stages

Principle 6: Processes

Execute the app as one or more stateless processes

Principle 7: Port Binding

Export services via port binding

Principle 8: Concurrency

Scale out via the process model

Principle 9: Disposability

Maximize robustness with fast startup and graceful shutdown

Principle 10: Dev/Prod Parity

Keep development, staging, and production as similar as possible

Principle 11: Logs

Treat logs as event streams

Principle 12: Admin Processes

Run admin/management tasks as one-off processes

Design Principles

Design Principles

High Cohesion

Autonomous

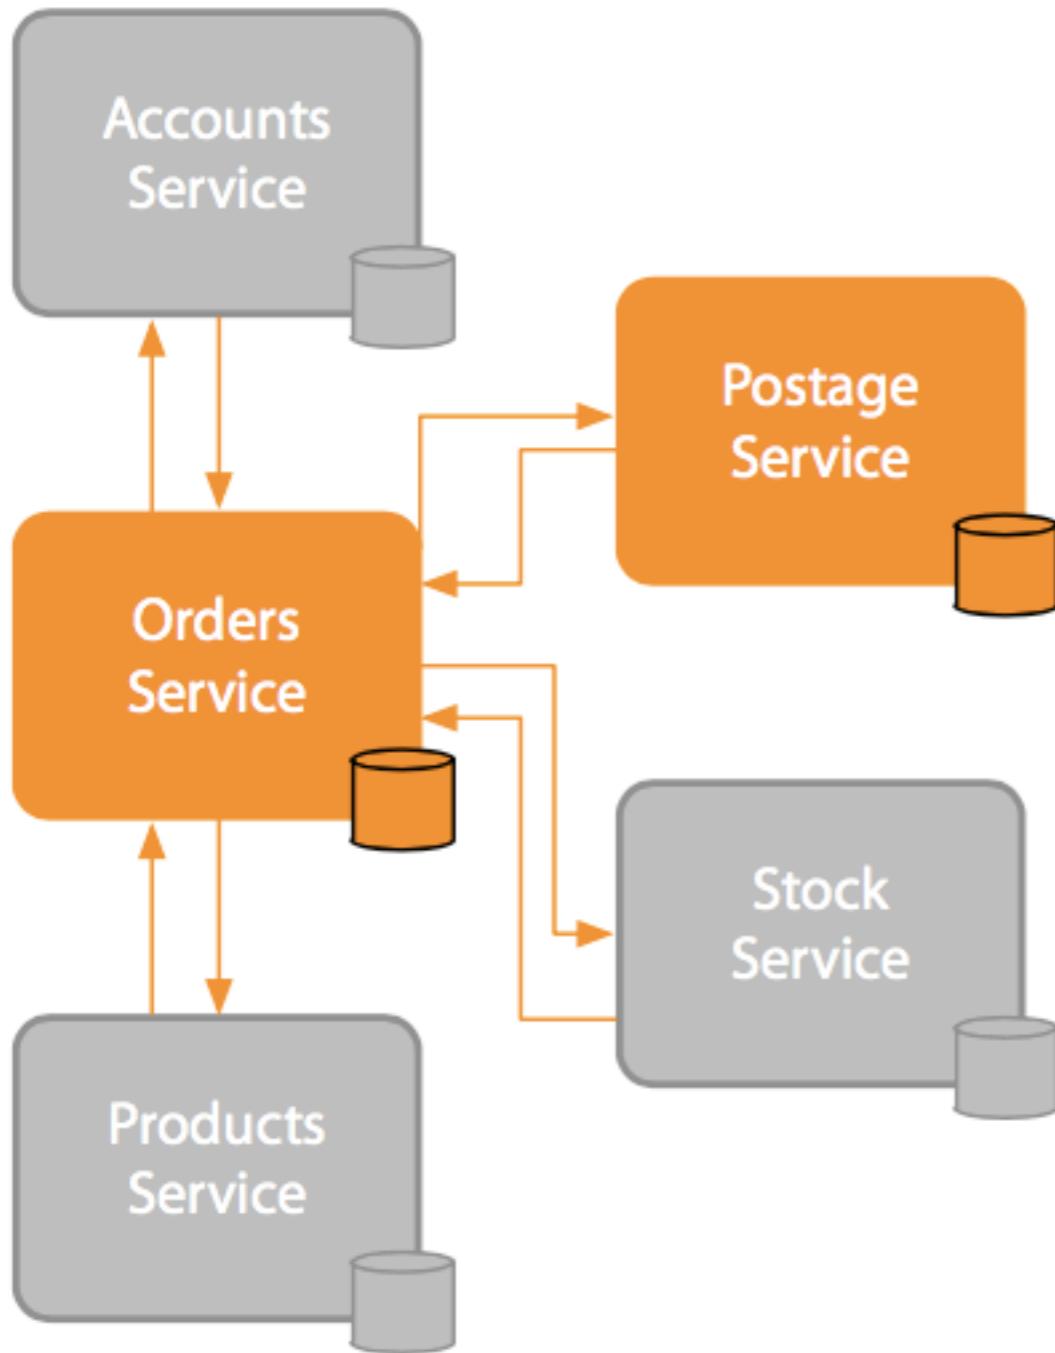
Business Domain
Centric

Resillience

Observable

Automation

High Cohesion

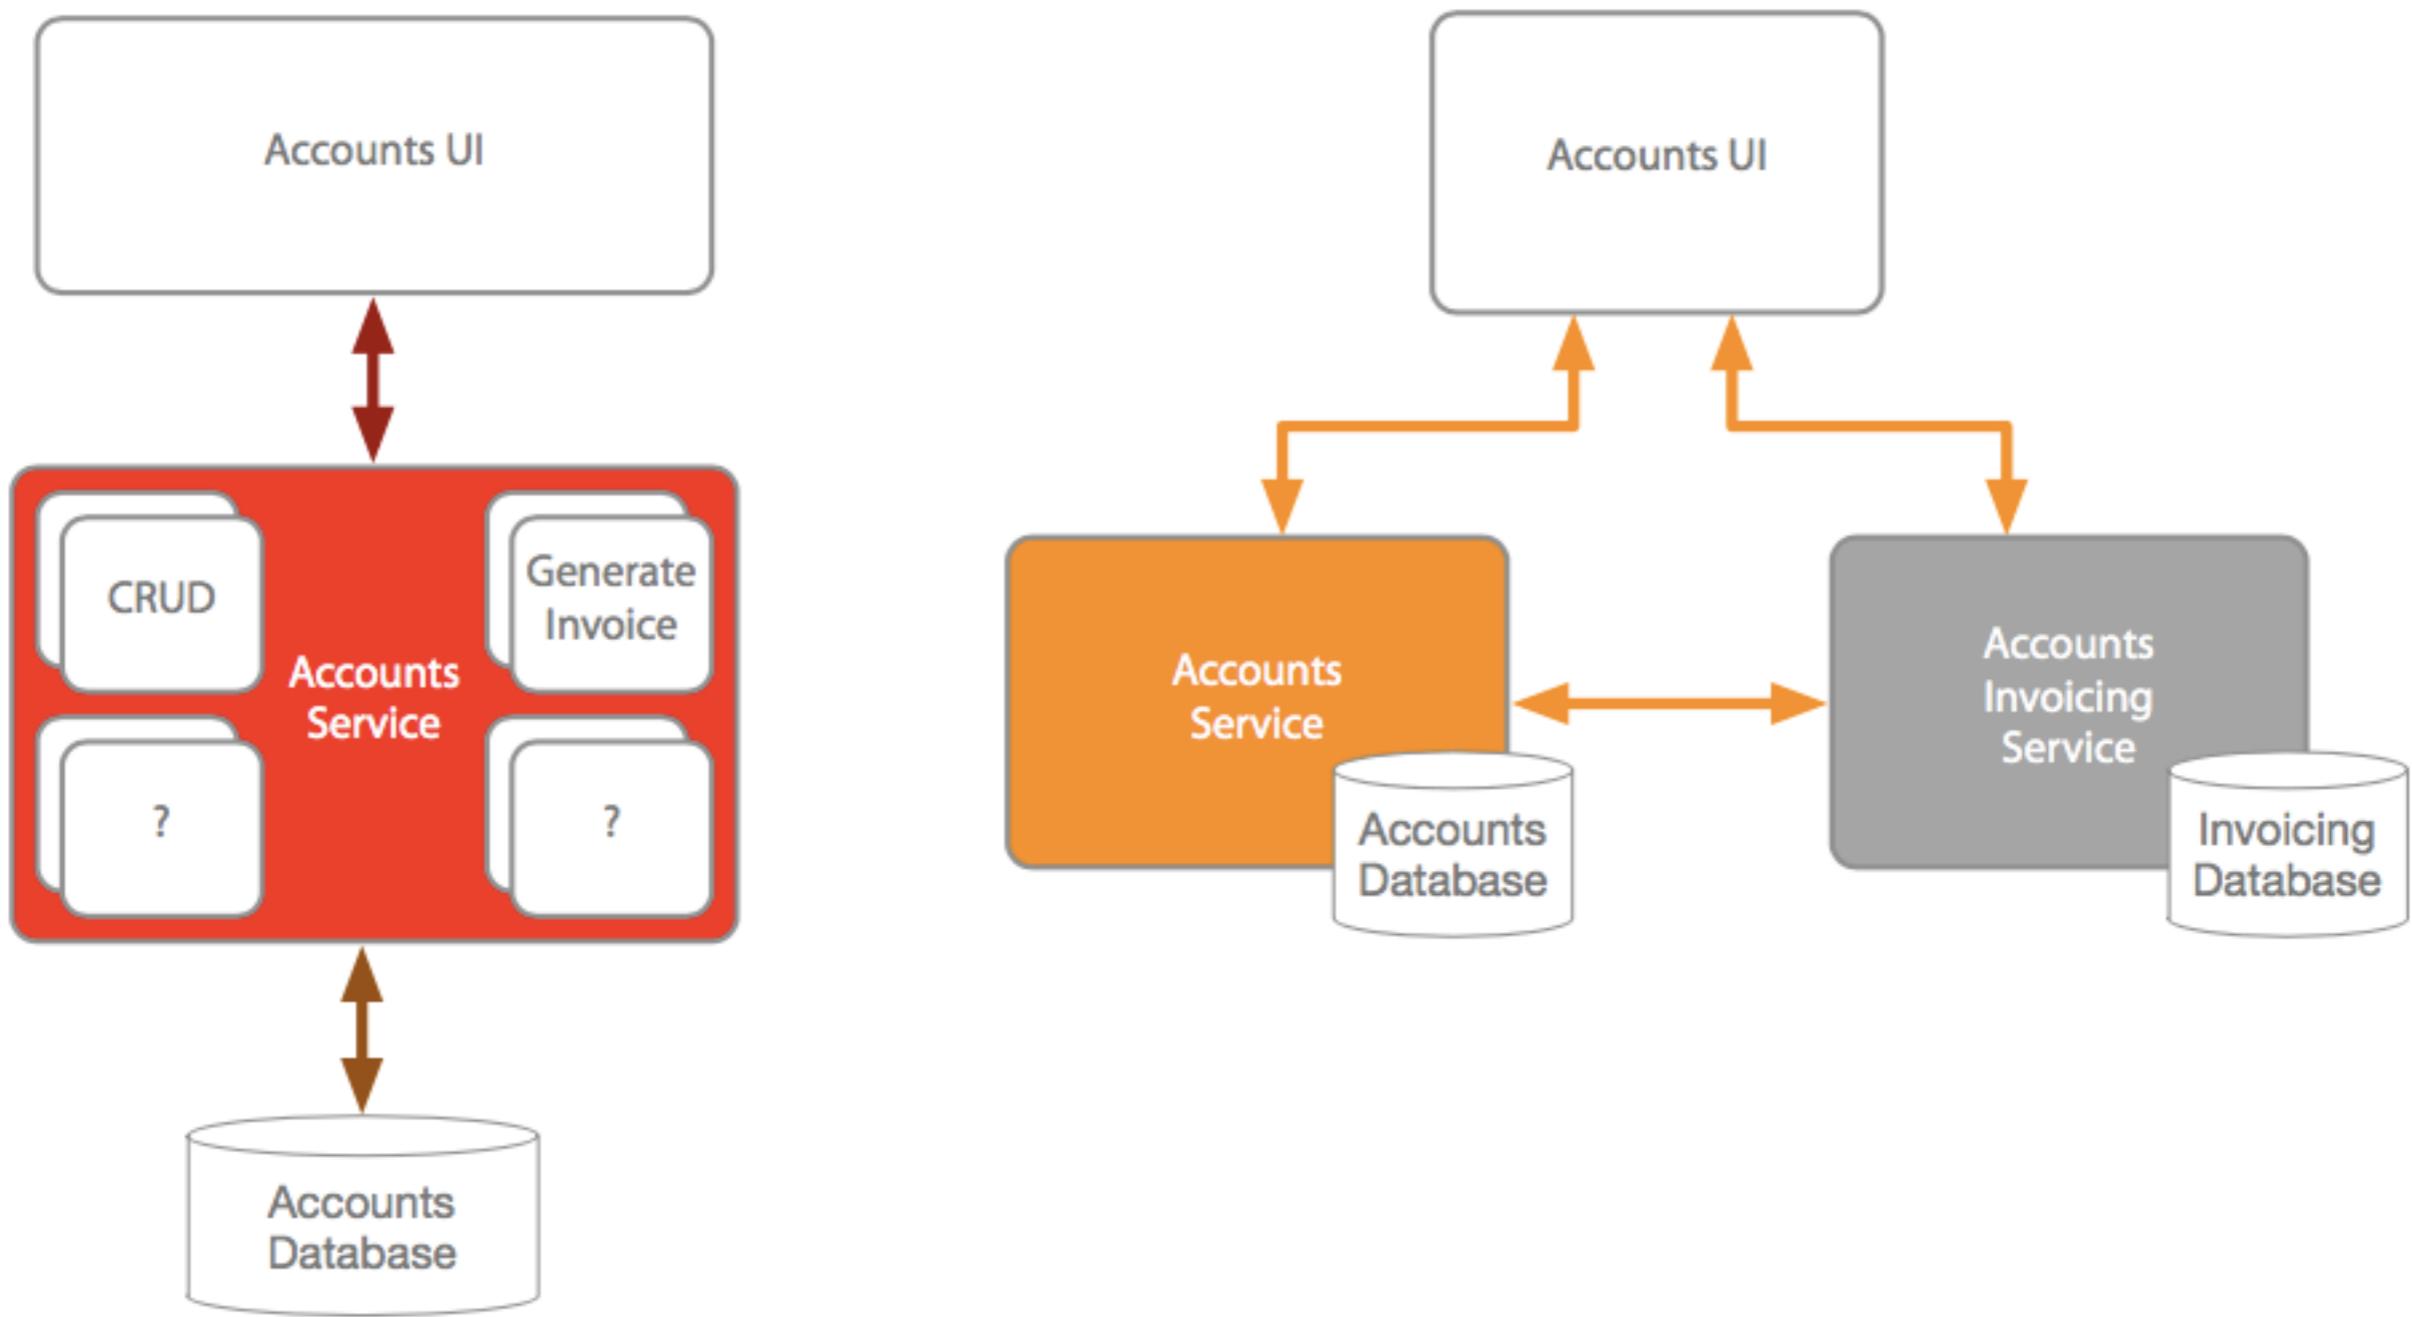


- Single focus
- Single responsibility
 - SOLID principle
 - Only change for one reason
- Reason represents
 - A business function
 - A business domain
- Encapsulation principle
- OOP principle
- Easily rewritable code

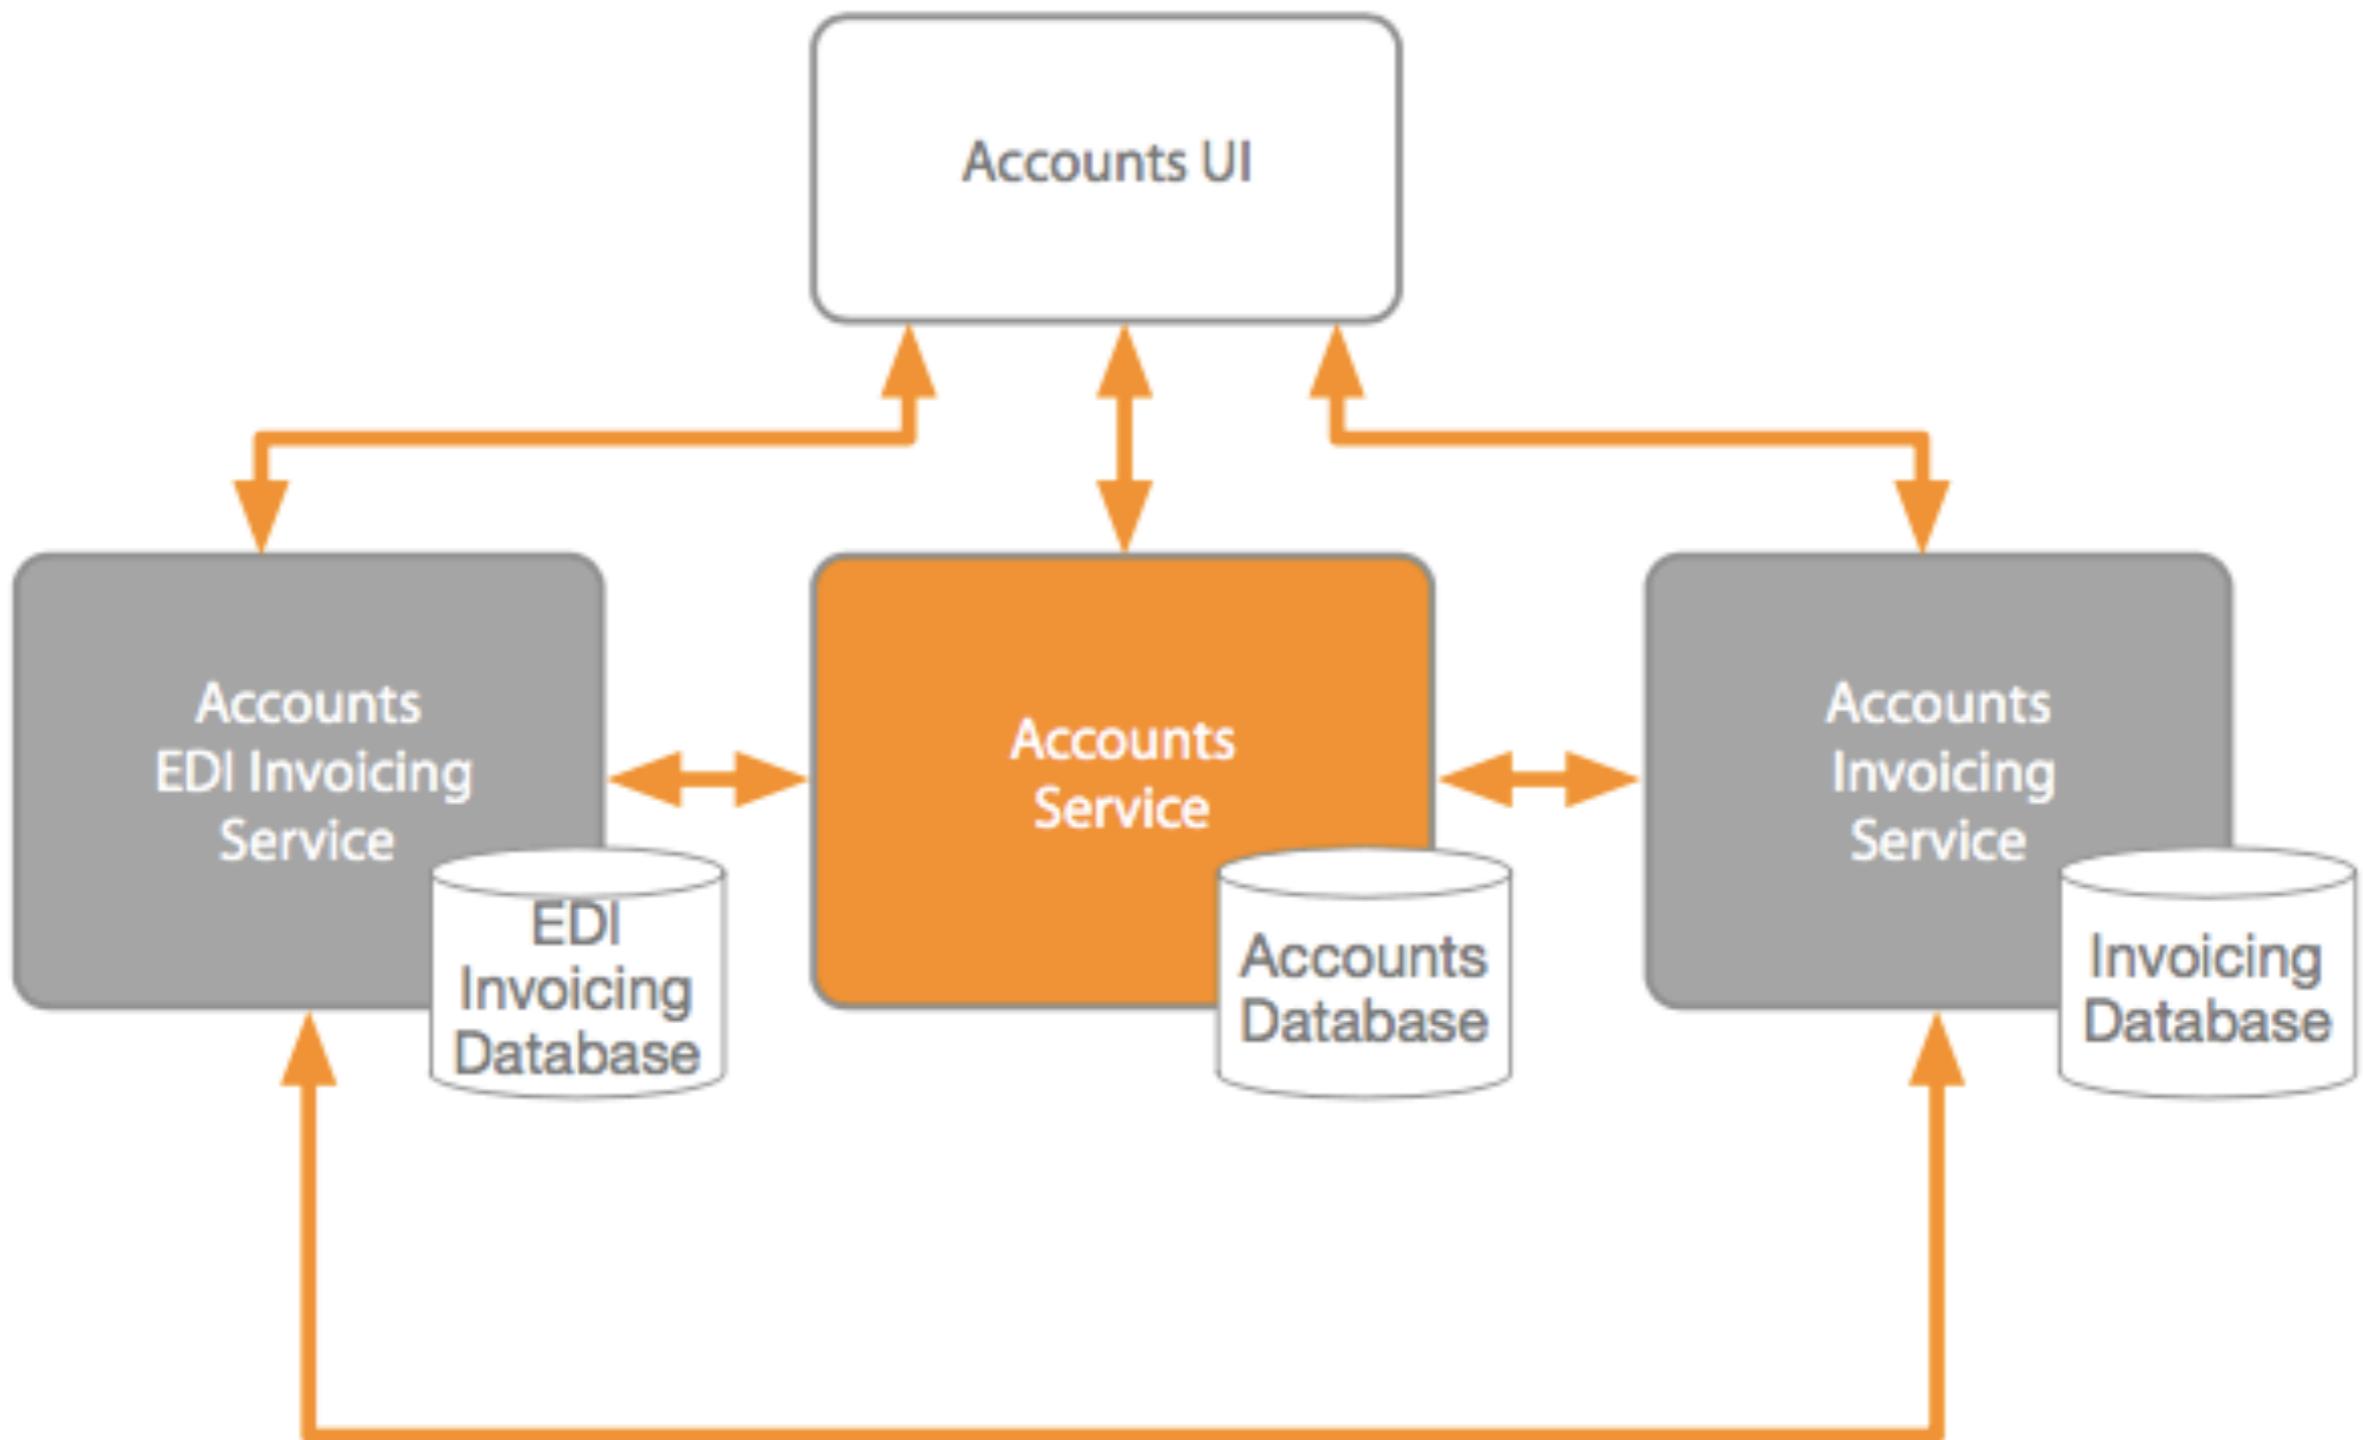
High Cohesion

- Identify a single focus
 - Business function
 - Business domain
- Split into finer grained services
- Avoid “Is kind of the same”
- Don't get lazy!
- Don't be afraid to create many services

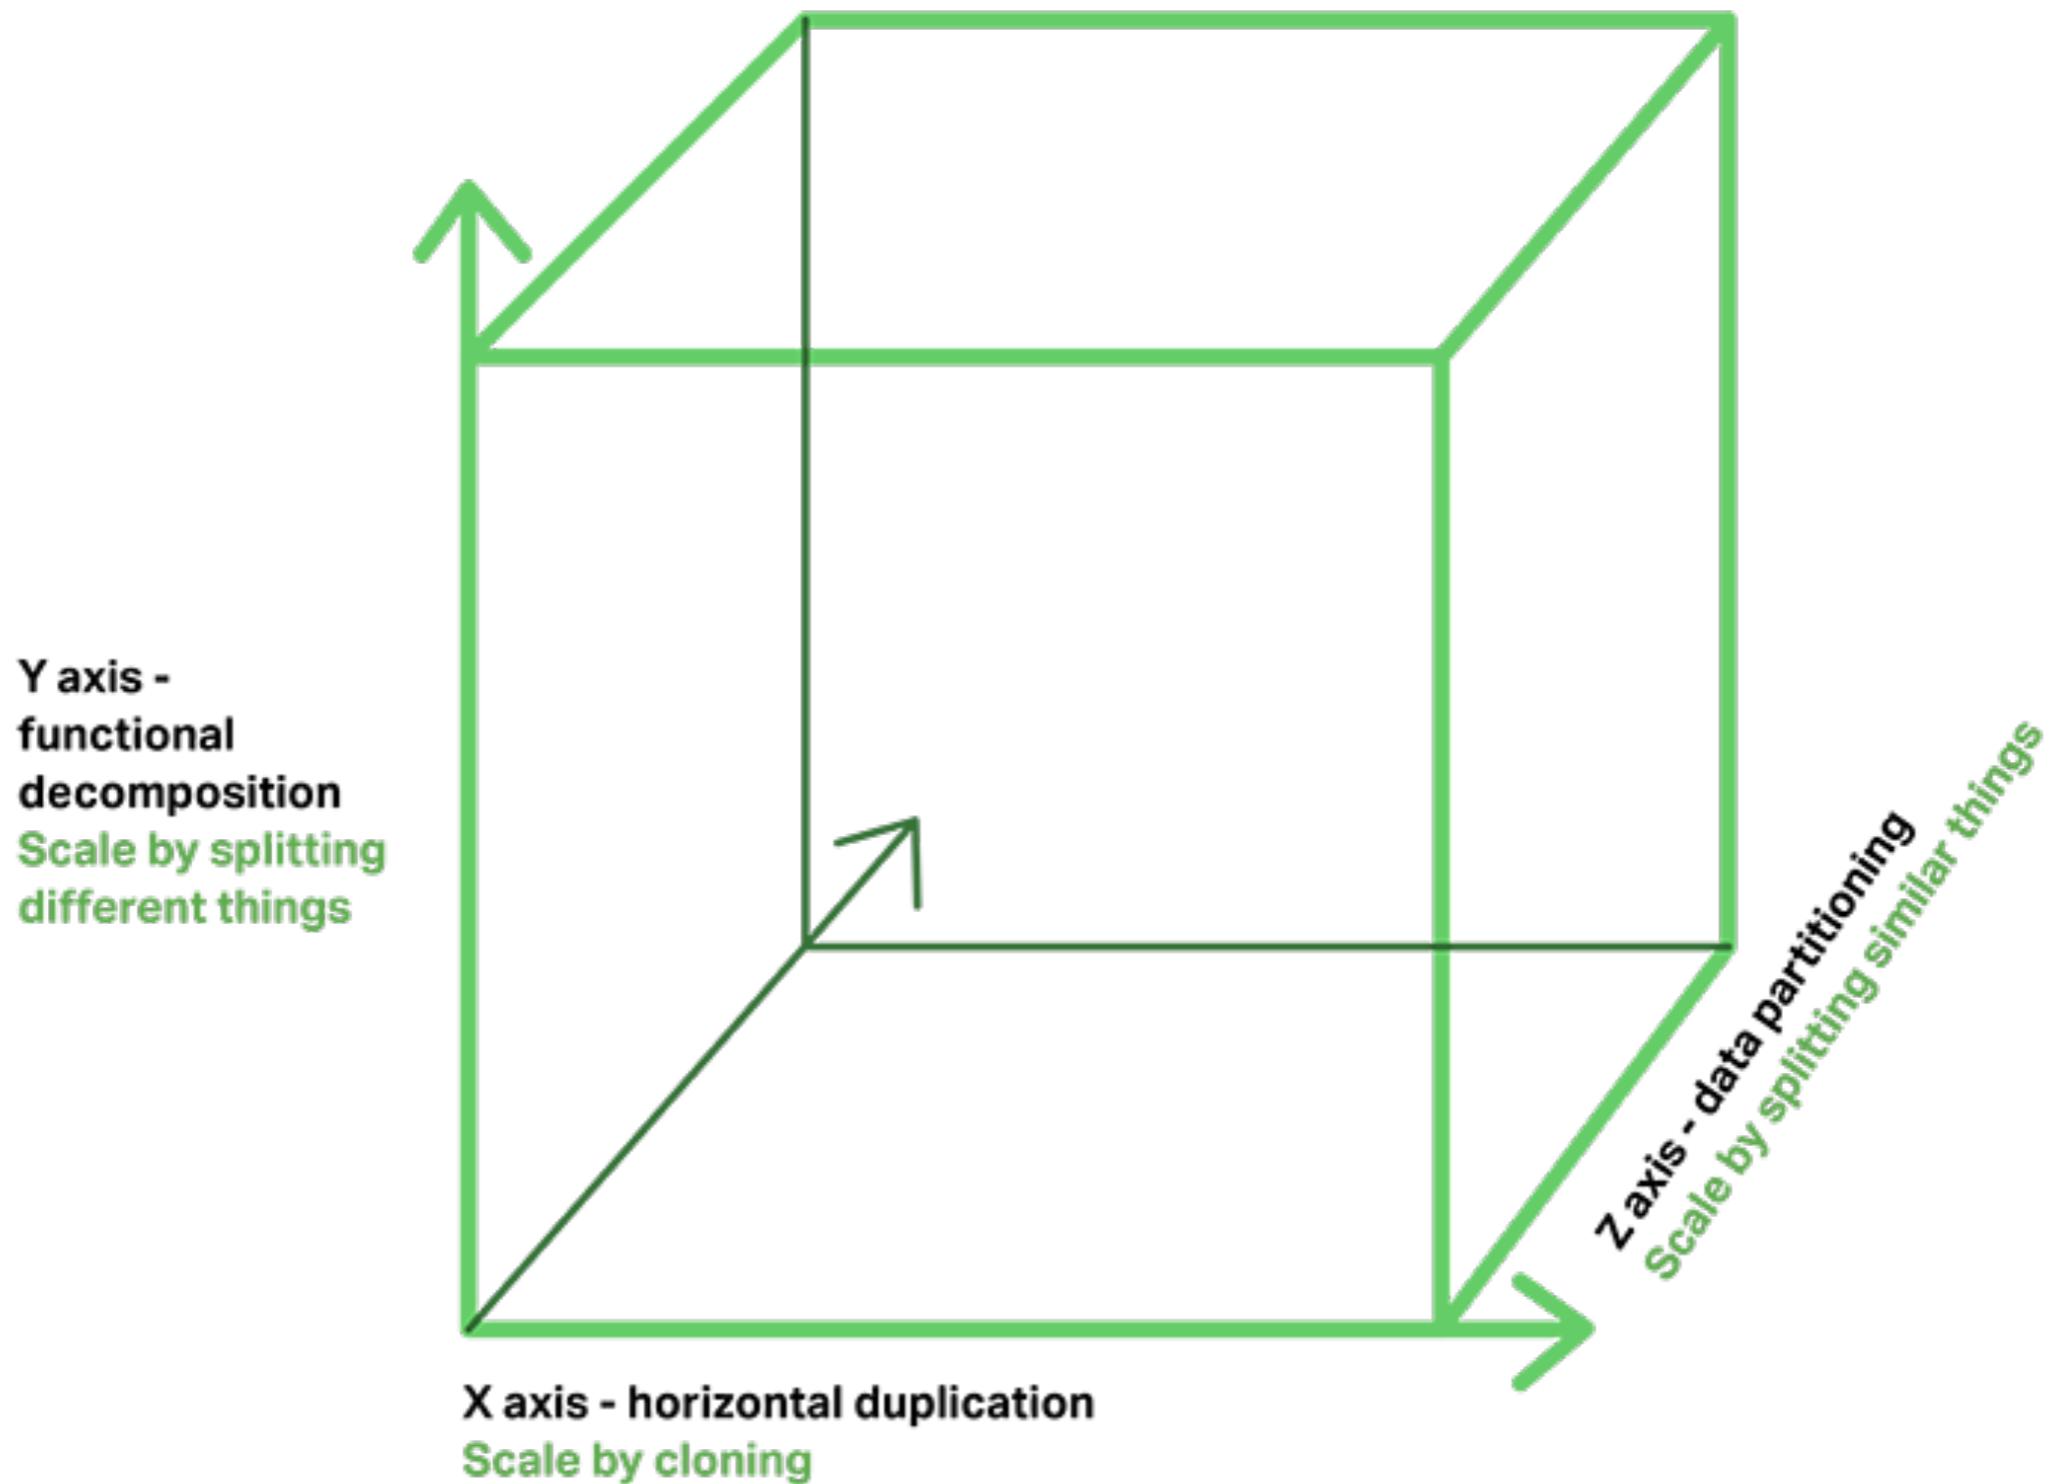
High Cohesion



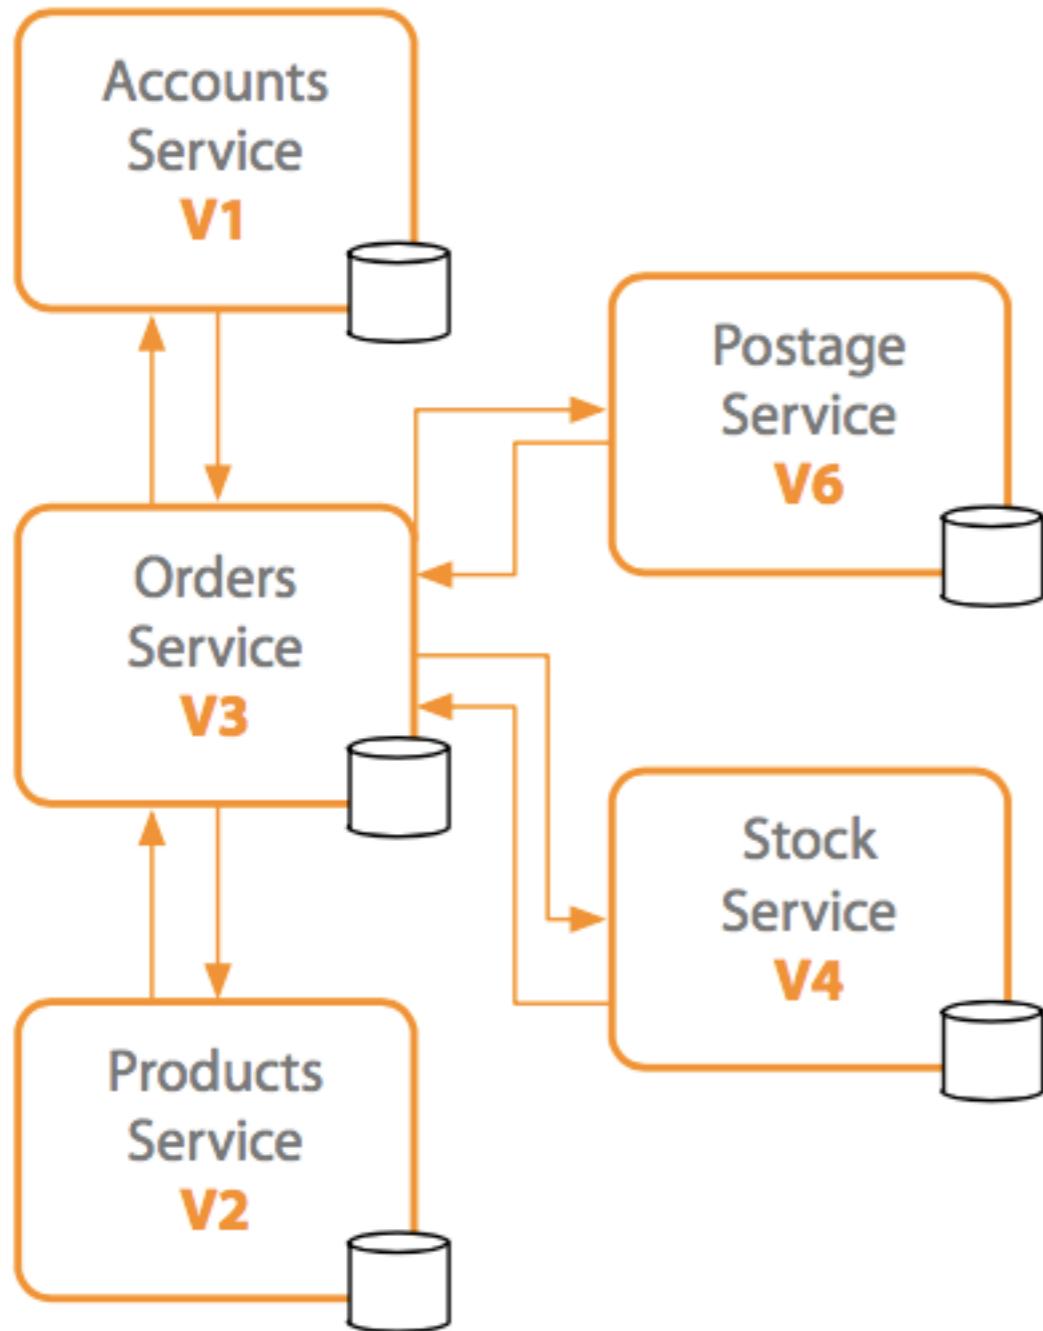
High Cohesion



The Art of Scalability

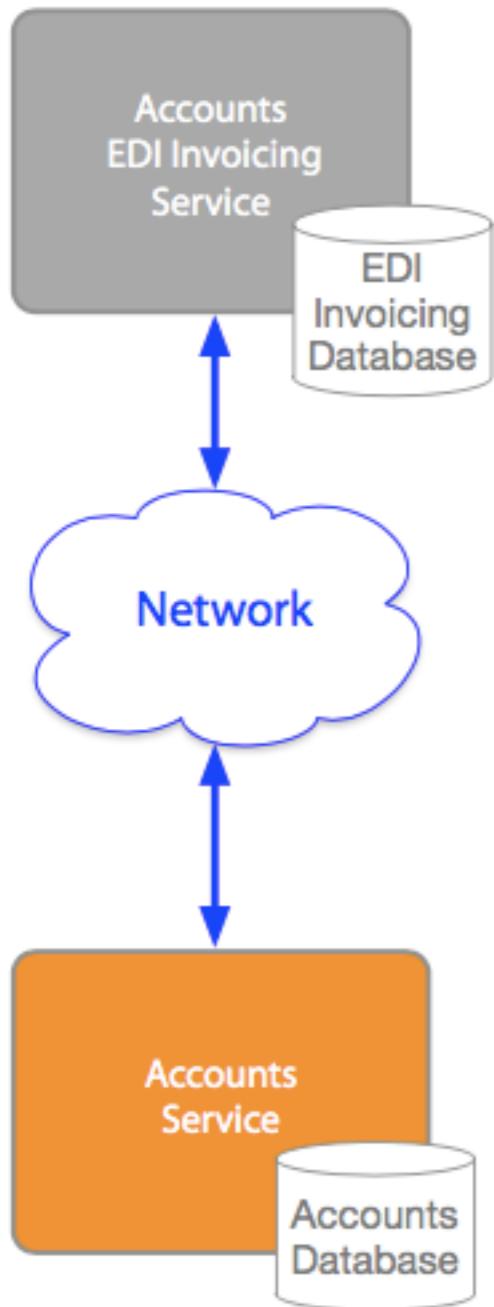


Autonomous



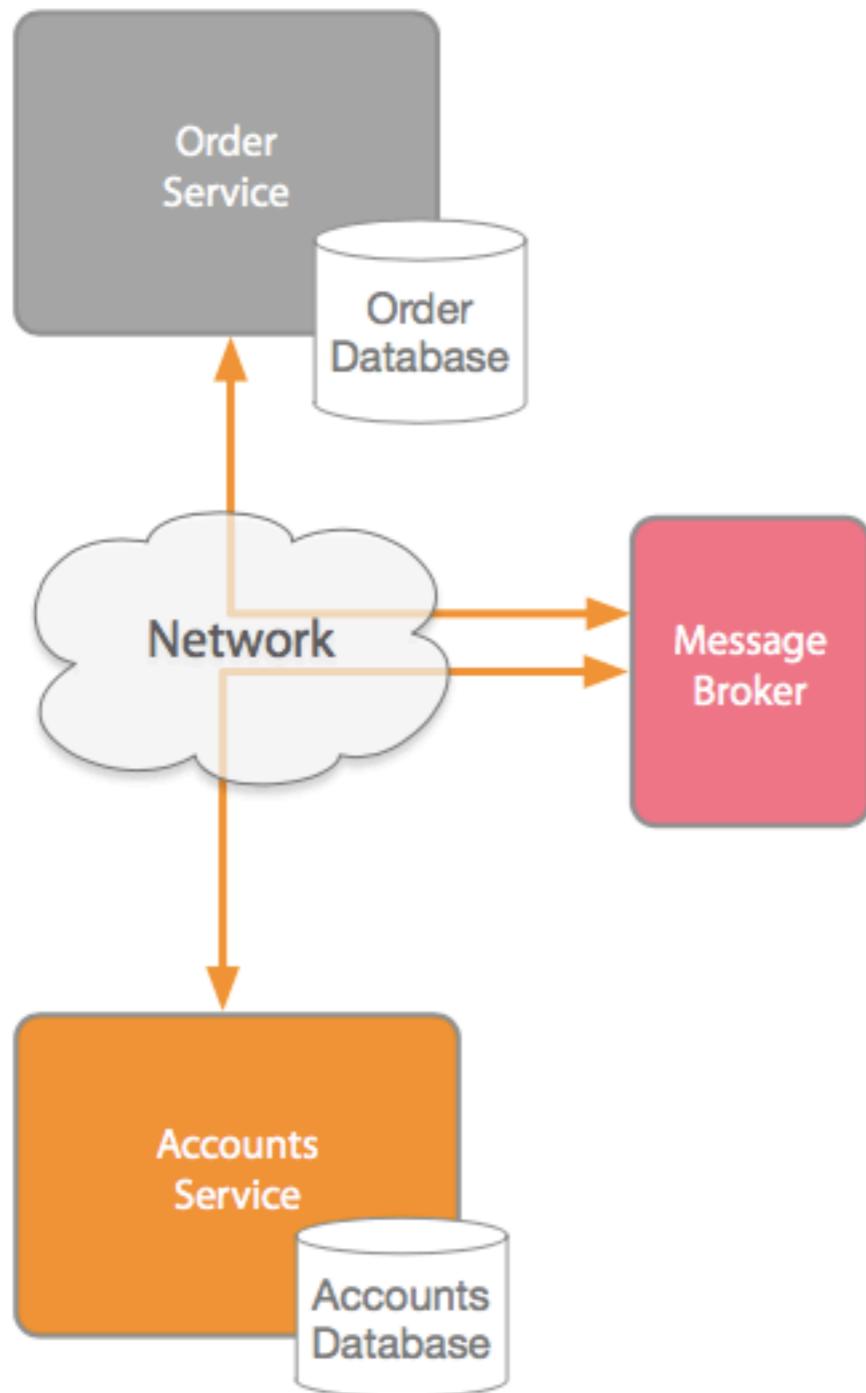
- Loose coupling
- Honor contracts and interfaces
- Stateless
- Independently changeable
- Independently deployable
- Backwards compatible
- Concurrent development

Autonomous



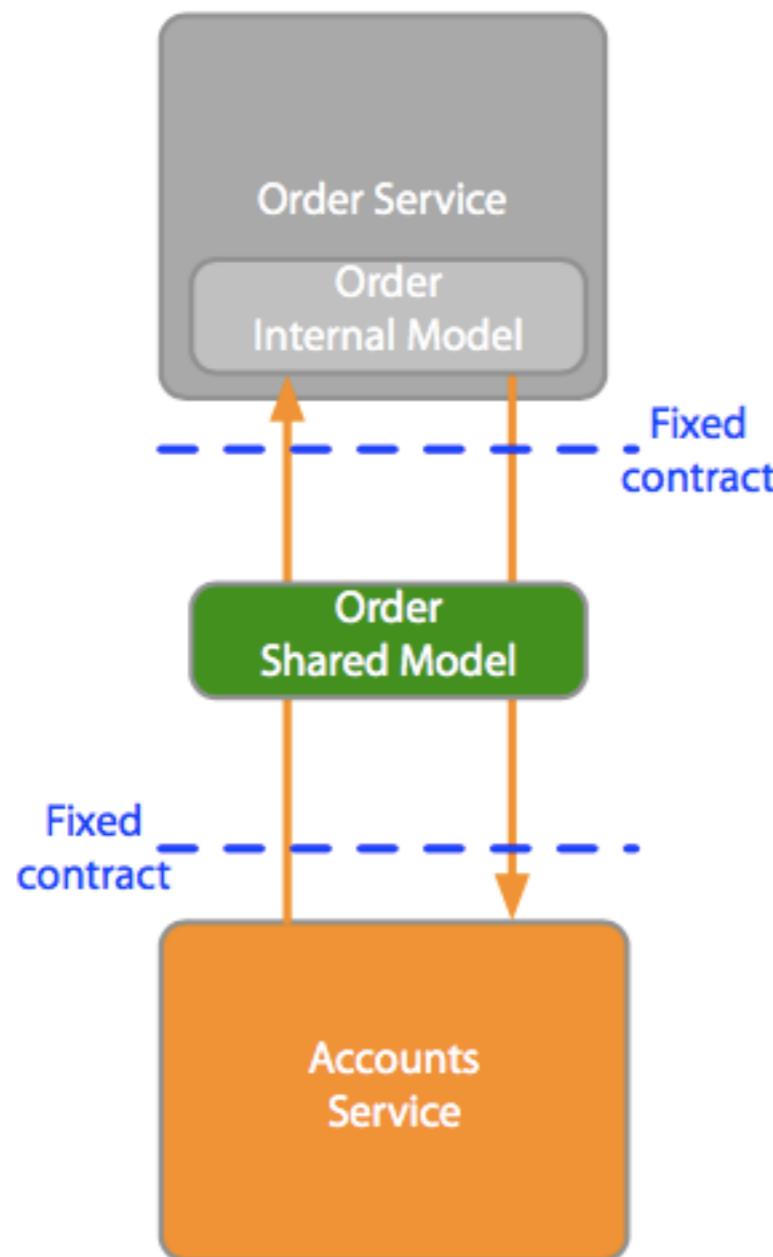
- Communication by network
- Technology agnostic API
- Avoid client libraries
- Contracts between services
- Avoid chatty exchanges between services
- Avoid sharing between services

Autonomous



- Communication by network
 - Synchronous
 - Asynchronous
 - ▶ Publish events
 - ▶ Subscribe to events

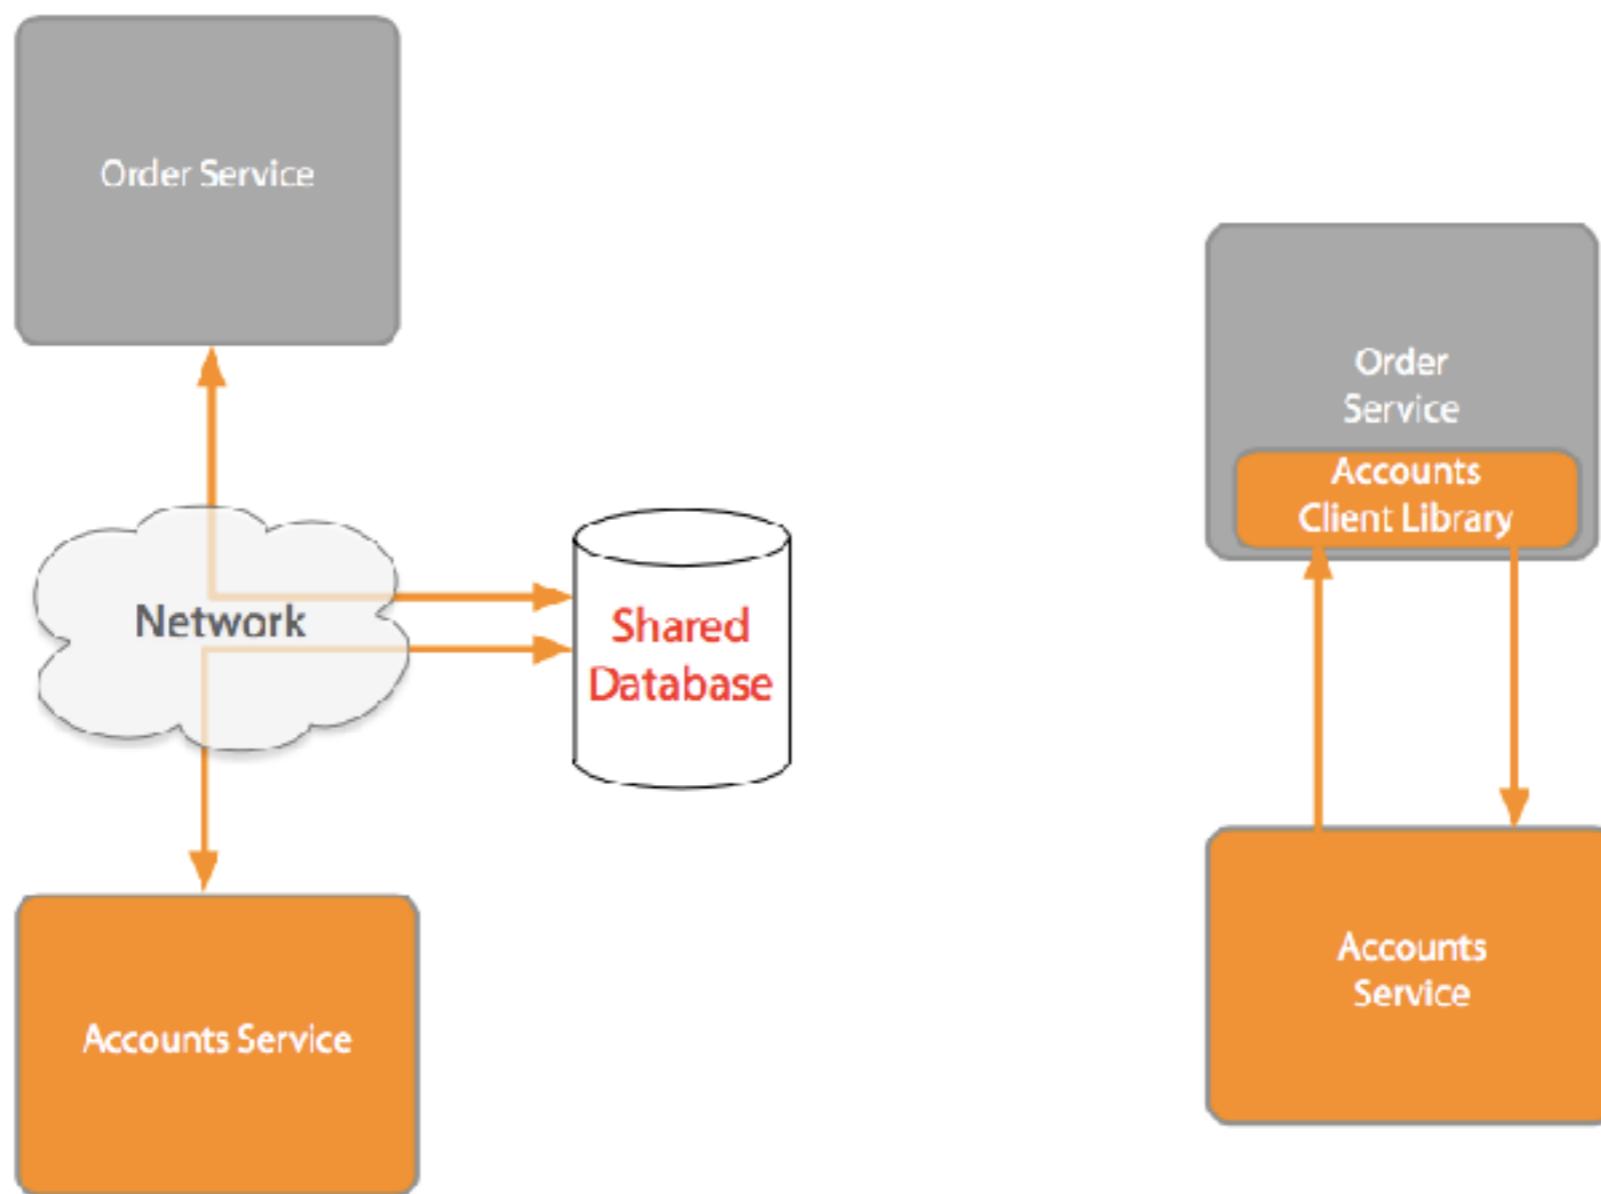
Autonomous



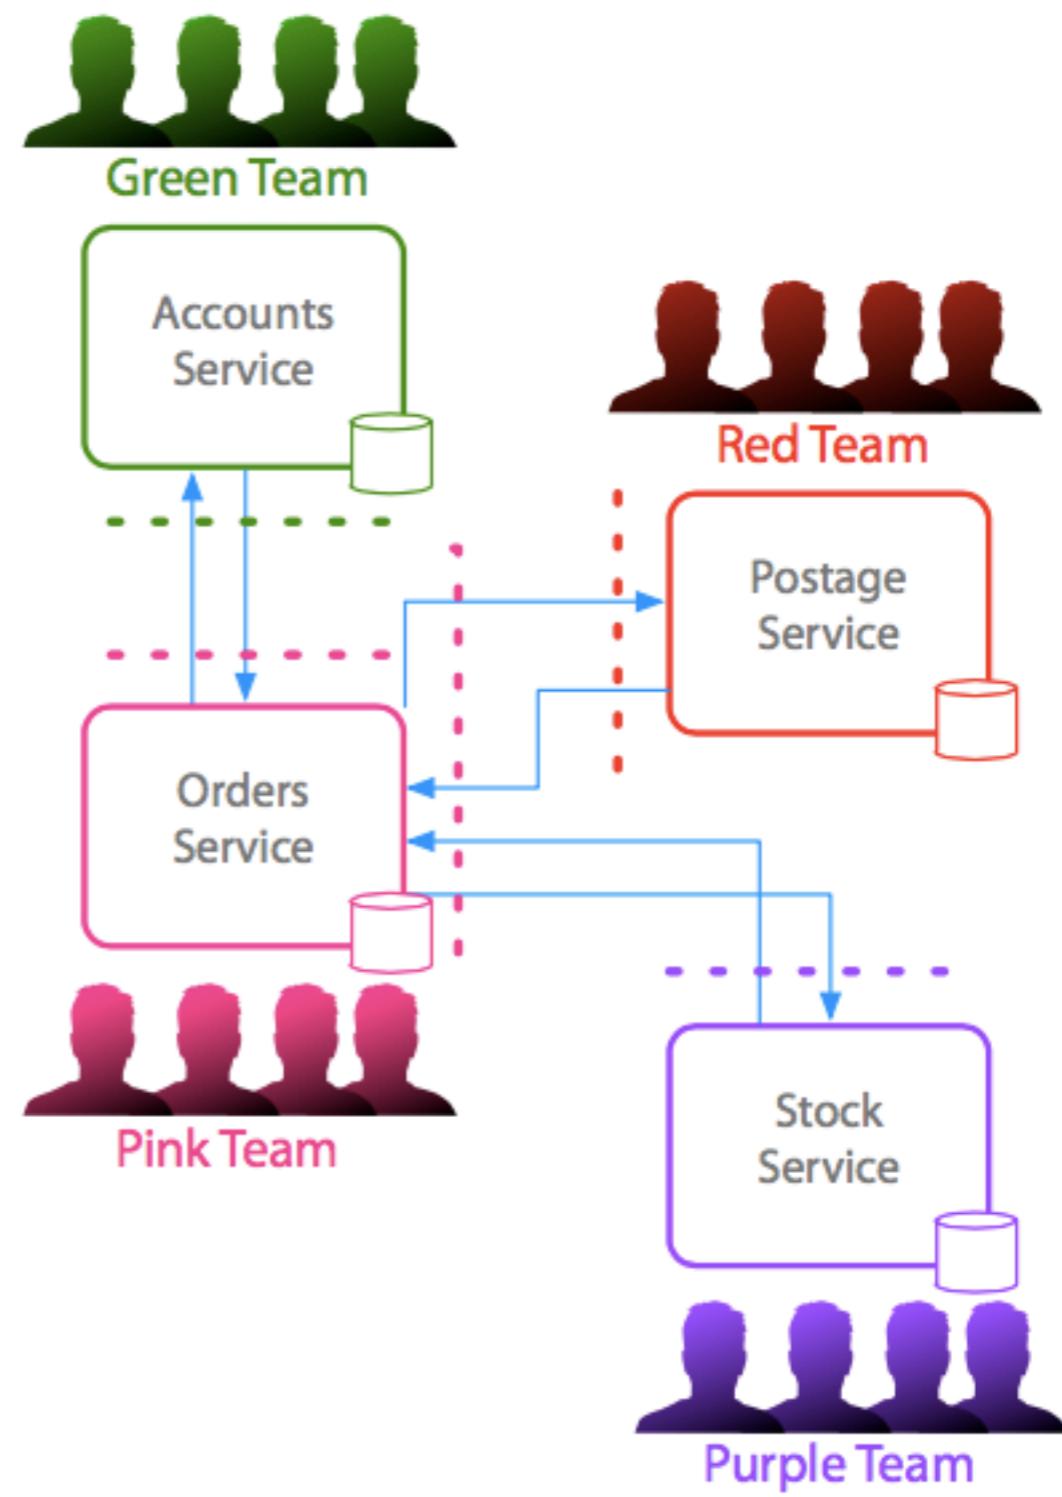
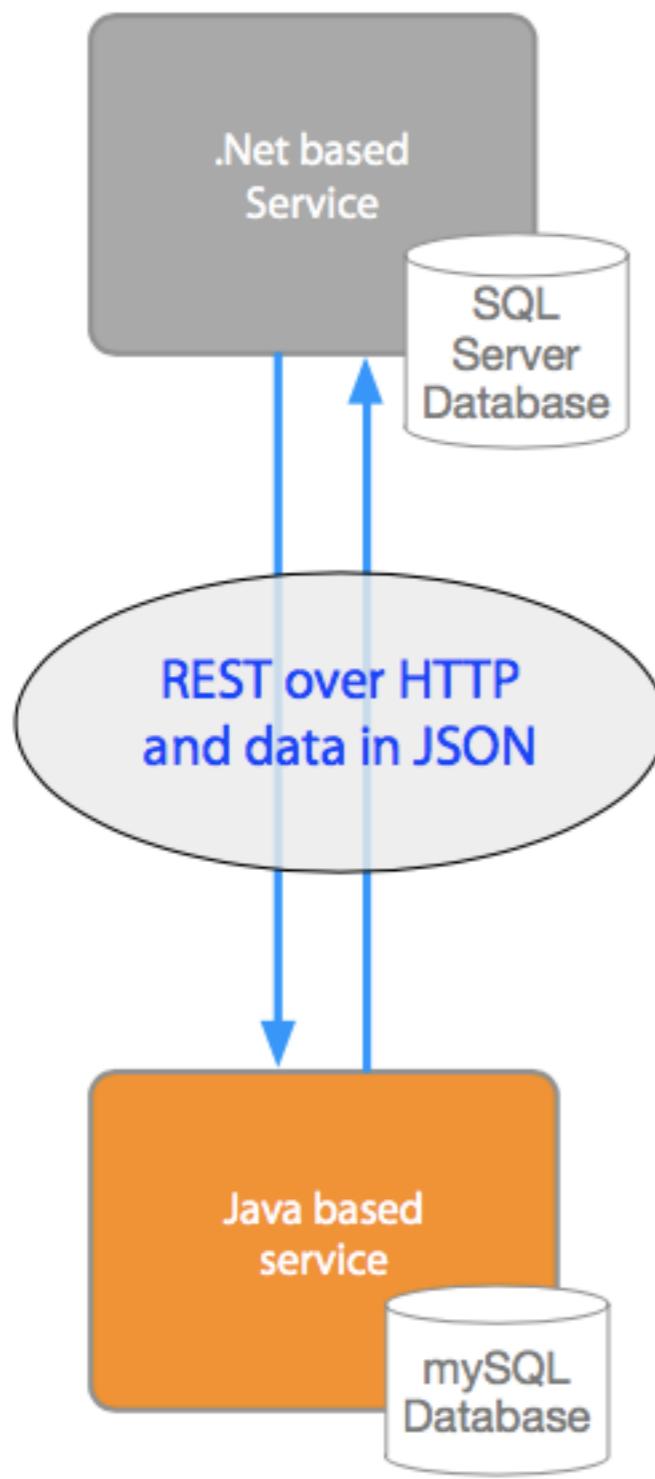
- Contracts between services
 - Fixed and agreed interfaces
 - Shared models
 - Clear input and output

Avoid

Avoid sharing between services

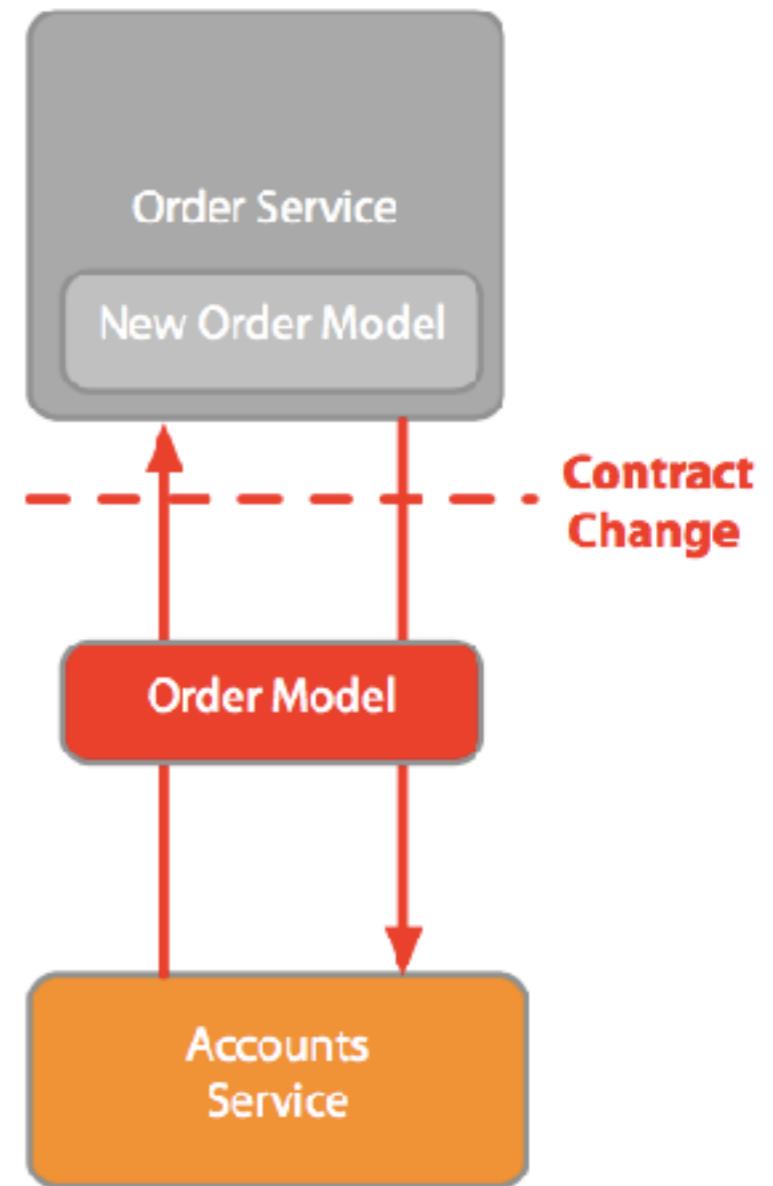


Autonomous



Ownership by team

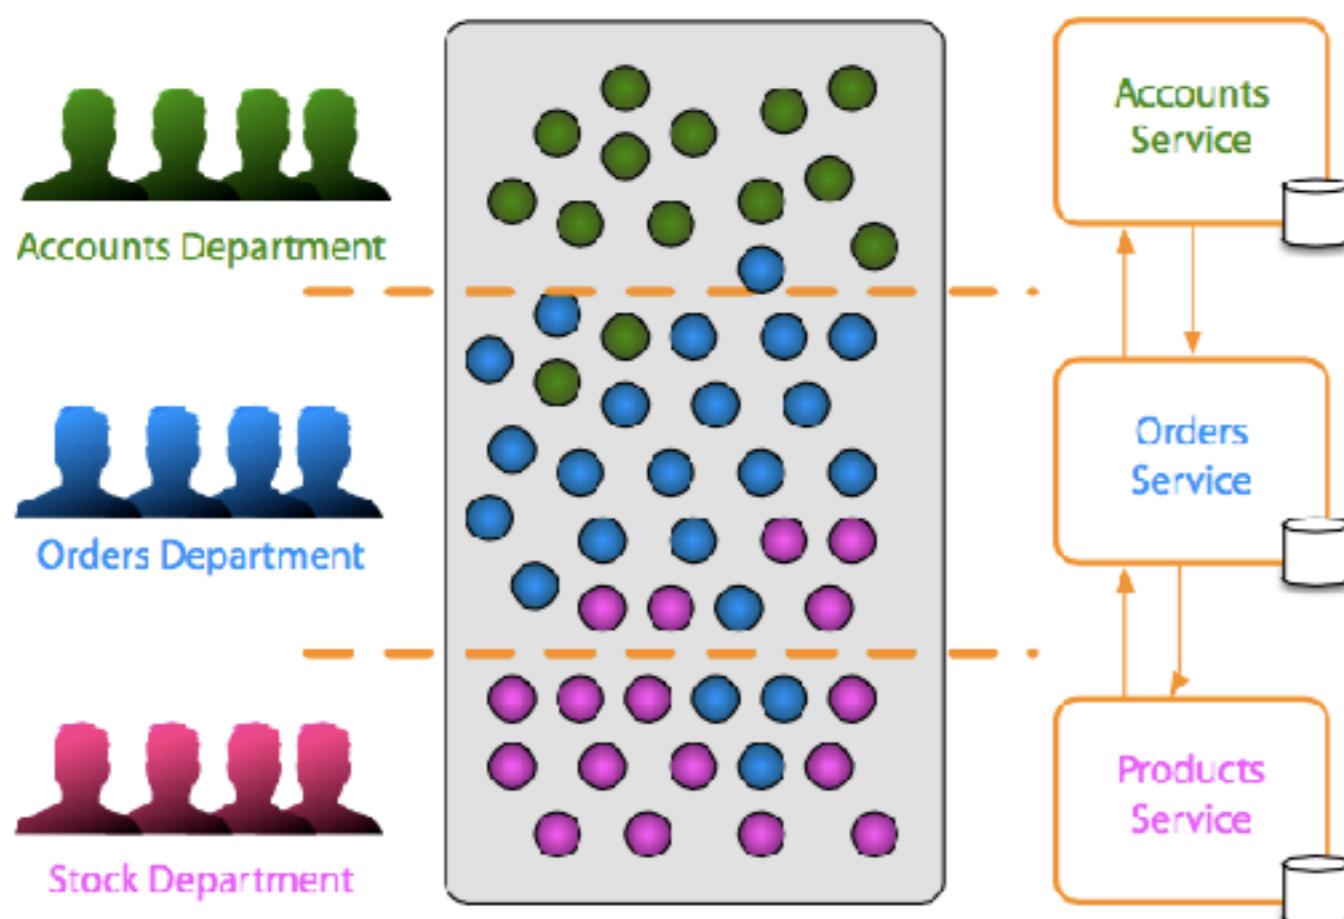
- Responsibility to make autonomous
- Agreeing contracts between teams
- Responsible for long-term maintenance
- Collaborative development
 - Communicate contract requirements
 - Communicate data requirements
- Concurrent development



Versioning

- Avoid breaking changes
- Backwards compatibility
- Integration tests
- Have a versioning strategy
 - Concurrent versions (Old and new)
 - Semantic versioning (Major.Minor.Patch , e.g. 15.1.2)
 - Coexisting endpoints (/V2/customer/)

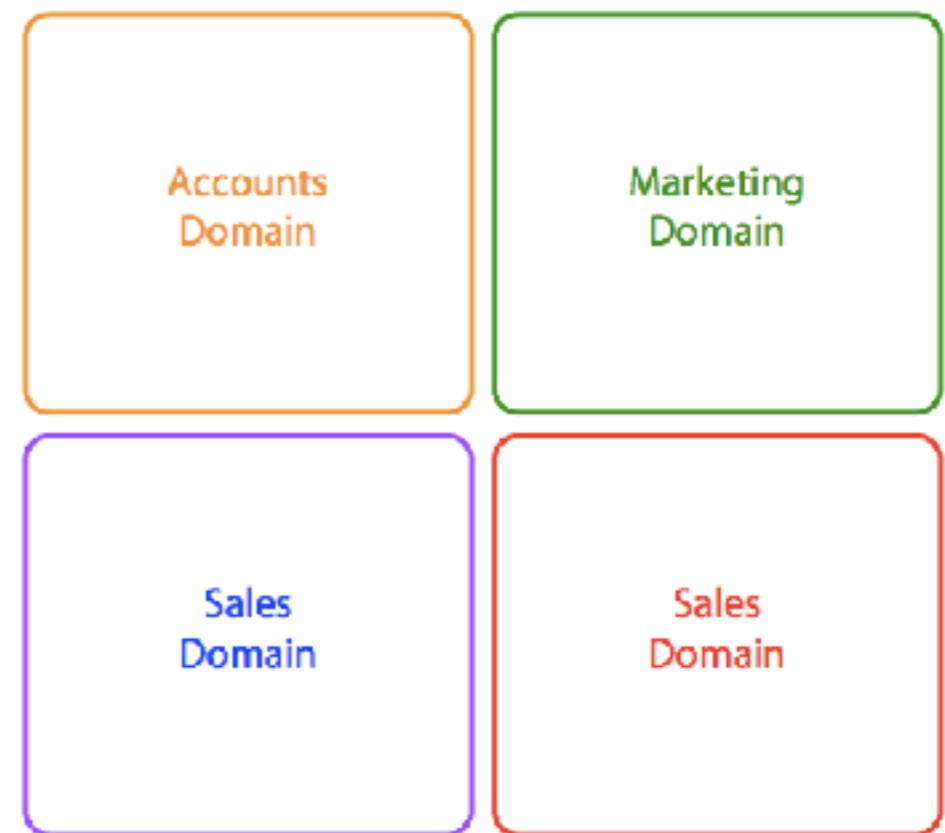
Business Domain Centric



- Service represents business function
- Scope of service
- Bounded context from DDD
- Identify boundary
- Shuffle code if required

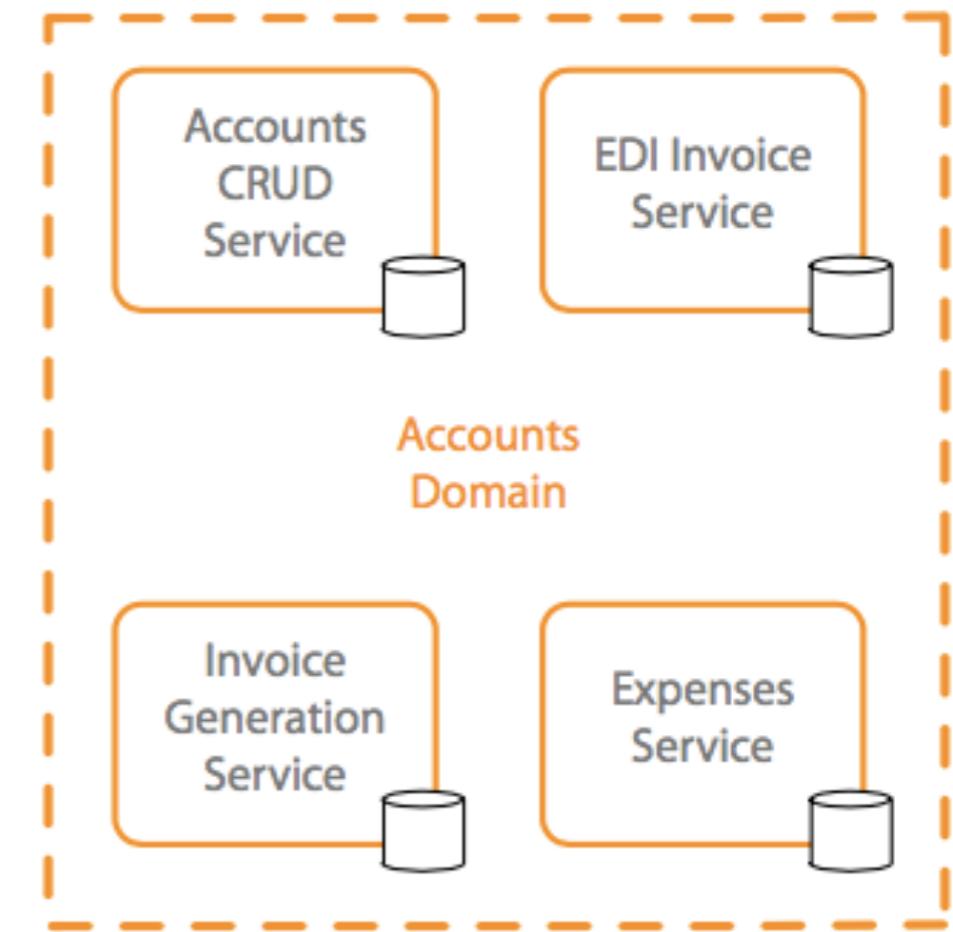
Business Domain Centric

- Identify business domains in a coarse manner
- Review sub groups of business functions or areas
- Review benefits of splitting further
- Agree a common language

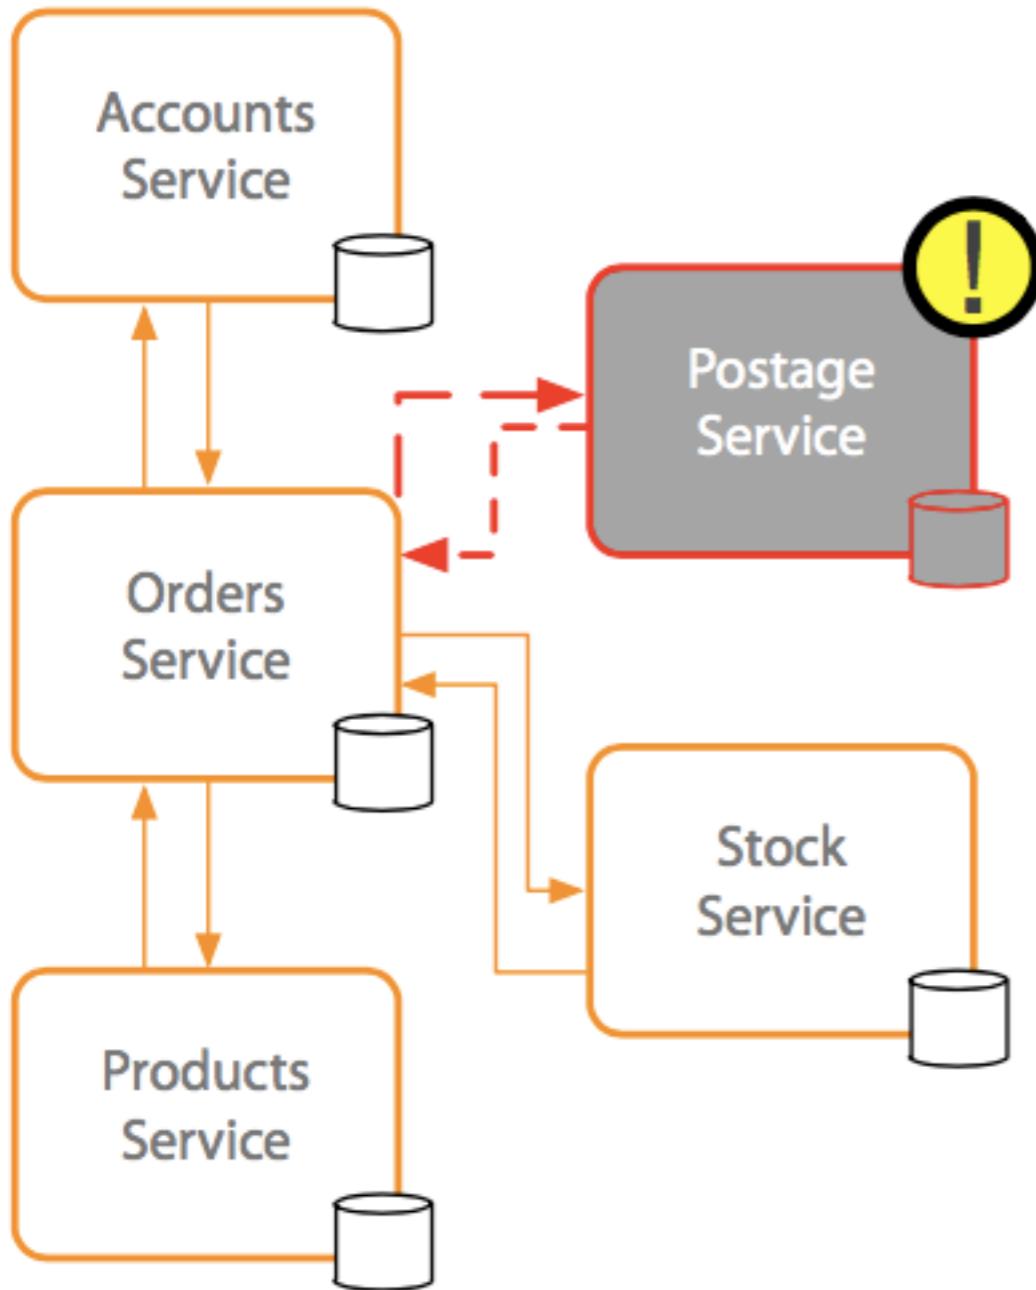


Business Domain Centric

- Fix incorrect boundaries
 - Merge or split
- Explicit interfaces for outside world
- Splitting using technical boundaries
 - Service to access archive data
 - For performance tuning

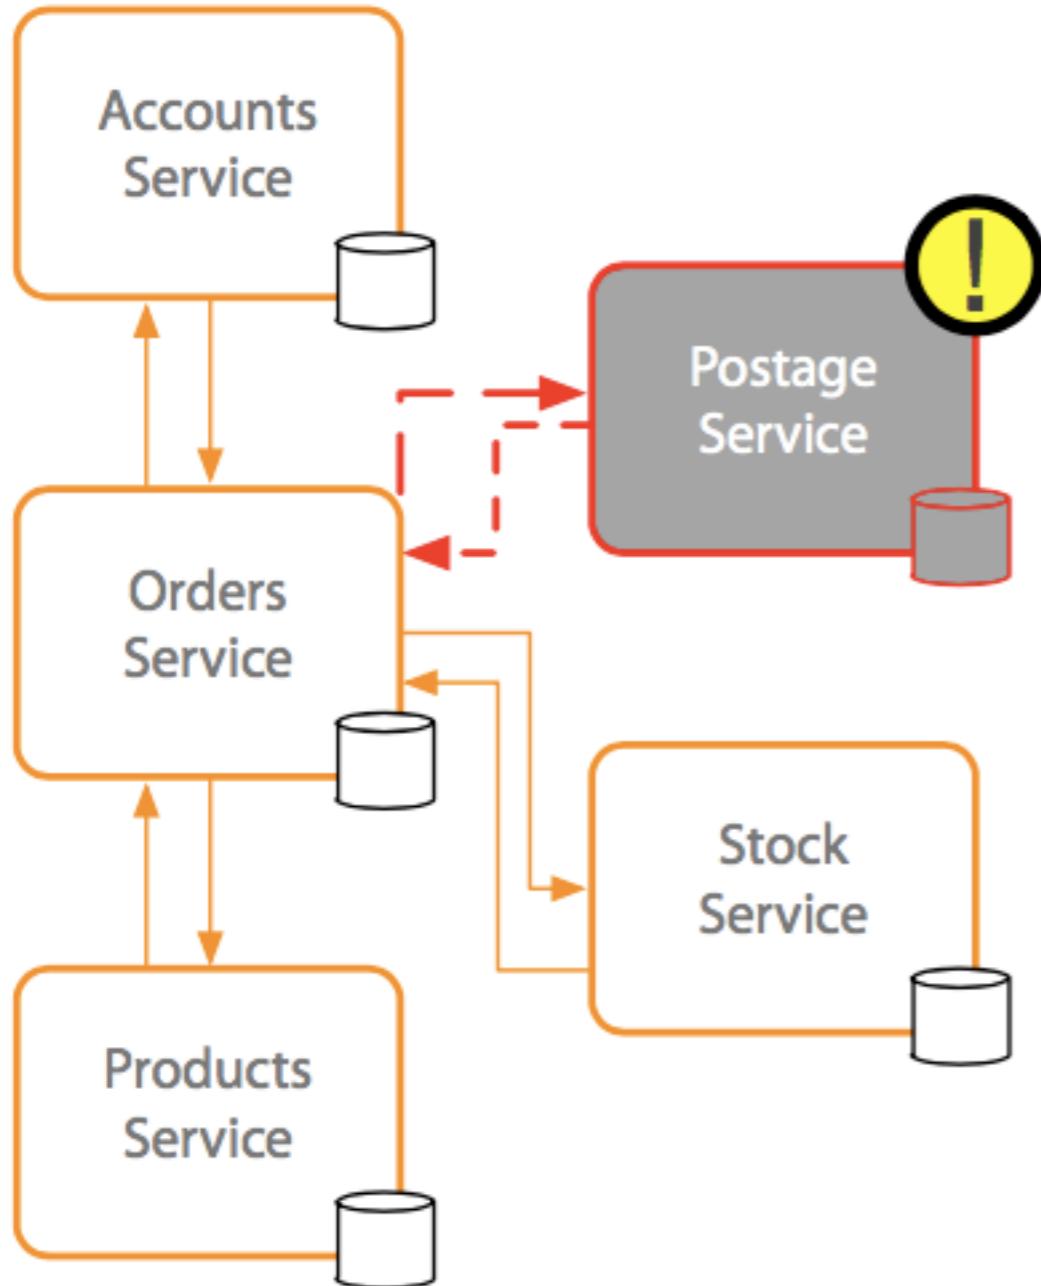


Resilience



- Embrace failure
 - Another service
 - Specific connection
 - Third-party system
- Degrade functionality
- Default functionality
- Multiple instances
 - Register on startup
 - Deregister on failure

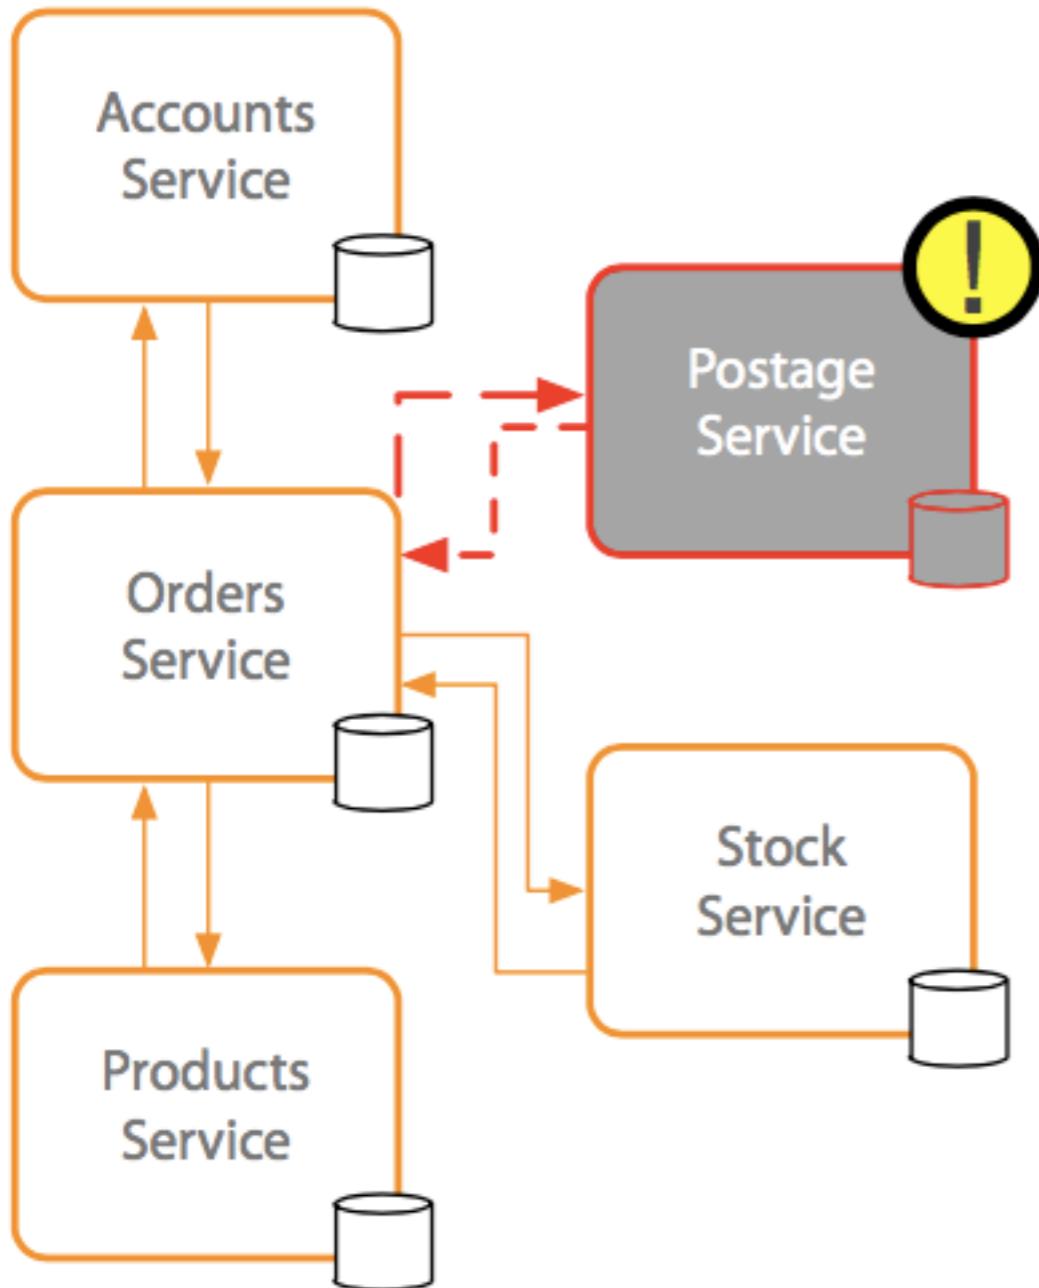
Resilience



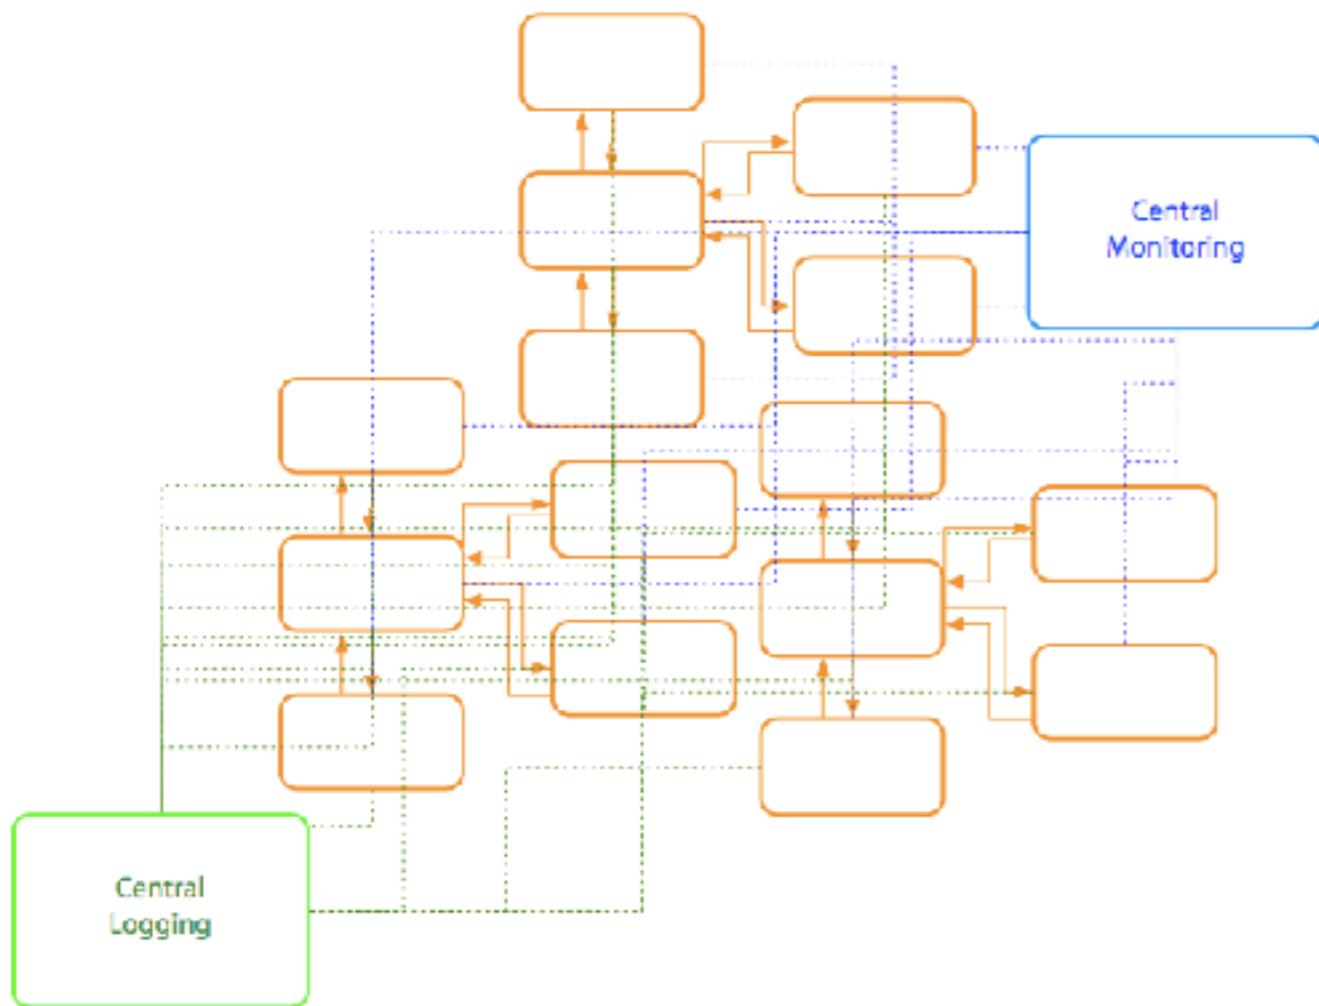
- Types of failure
 - Exceptions\Errors
 - Delays
 - Unavailability
- Network issues
 - Delay
 - Unavailability
- Validate input
 - Service to service
 - Client to service

Resilience

- Design system to fail fast
- Design System for known failures
- Use timeouts
 - Use for connected systems
 - Timeout our requests after a threshold
 - Service to service
 - Service to other systems
 - Standard timeout length
 - Adjust length on a case by case basis
- Network outages and latency
- Monitor timeout
- Log timeouts



Observable



- System Health
 - Status
 - Logs
 - Errors
- Centralized monitoring
- Centralized logging

Centralize Monitoring

- Real-time monitoring
- Monitor the host (CPU, memory, disk usage, etc.)
- Expose metrics within the services
 - Response times
 - Timeouts
 - Exceptions and errors
- Business data related metrics
 - Number of orders
 - Average time from basket to checkout

Centralize Monitoring

- Collect and aggregate monitoring data
 - Monitoring tools that provide aggregation
 - Monitoring tools that provide drill down options
- Monitoring tool that can help visualize trends
- Monitoring tool that can compare data across servers
- Monitoring tool that can trigger alerts

Centralize Logging

- When to log
 - Startup or shutdown
 - Code path milestones
 - Requests, responses and decisions
 - Timeouts, exceptions and errors
- Traceable distributed transactions
 - Correlation ID (Passed service to service)

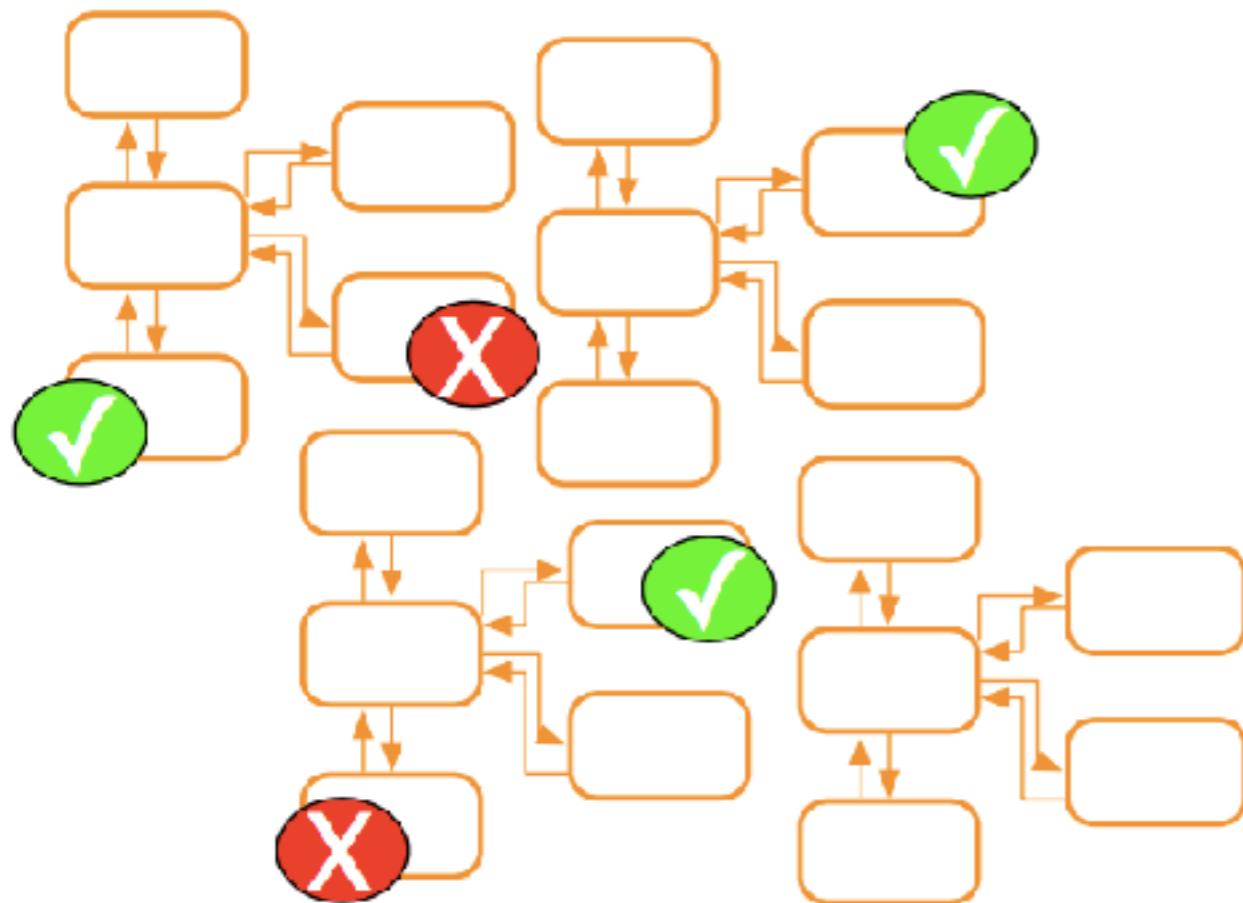
Structure Logging

- Level
 - Information
 - Error
 - Debug
 - Statistic
- Date and time
- Correlation ID
- Host name
- Service name and service instance
- Message

Why Observable

- Distributed transactions
- Quick problem solving
- Quick deployment requires feedback
- Data used for capacity planning
- Data used for scaling
- What's actually used
- Monitor business data

Automation



- Tools to reduce testing
 - Manual regression testing
 - Time taken on testing integration
 - Environment setup for testing
- Tools to provide quick feedback
 - Integration feedback on check in
 - Continuous Integration
- Tools to provide quick deployment
 - Pipeline to deployment
 - Deployment ready status
 - Automated deployment
 - Reliable deployment
 - Continuous Deployment

Why Automation

- Distributed system
- Multiple instances of services
- Manual integration testing too time consuming
- Manual deployment time consuming and unreliable

Automation

Continuous
Intregation

Continuous
Deployment

Continuous Integration

- Work with source control system
- Automatic after check-in
- Unit tests and Integration tests required
- Ensure to quality of check-in
 - Code compiles
 - Test pass
 - Changes integrate
 - Quick feedback

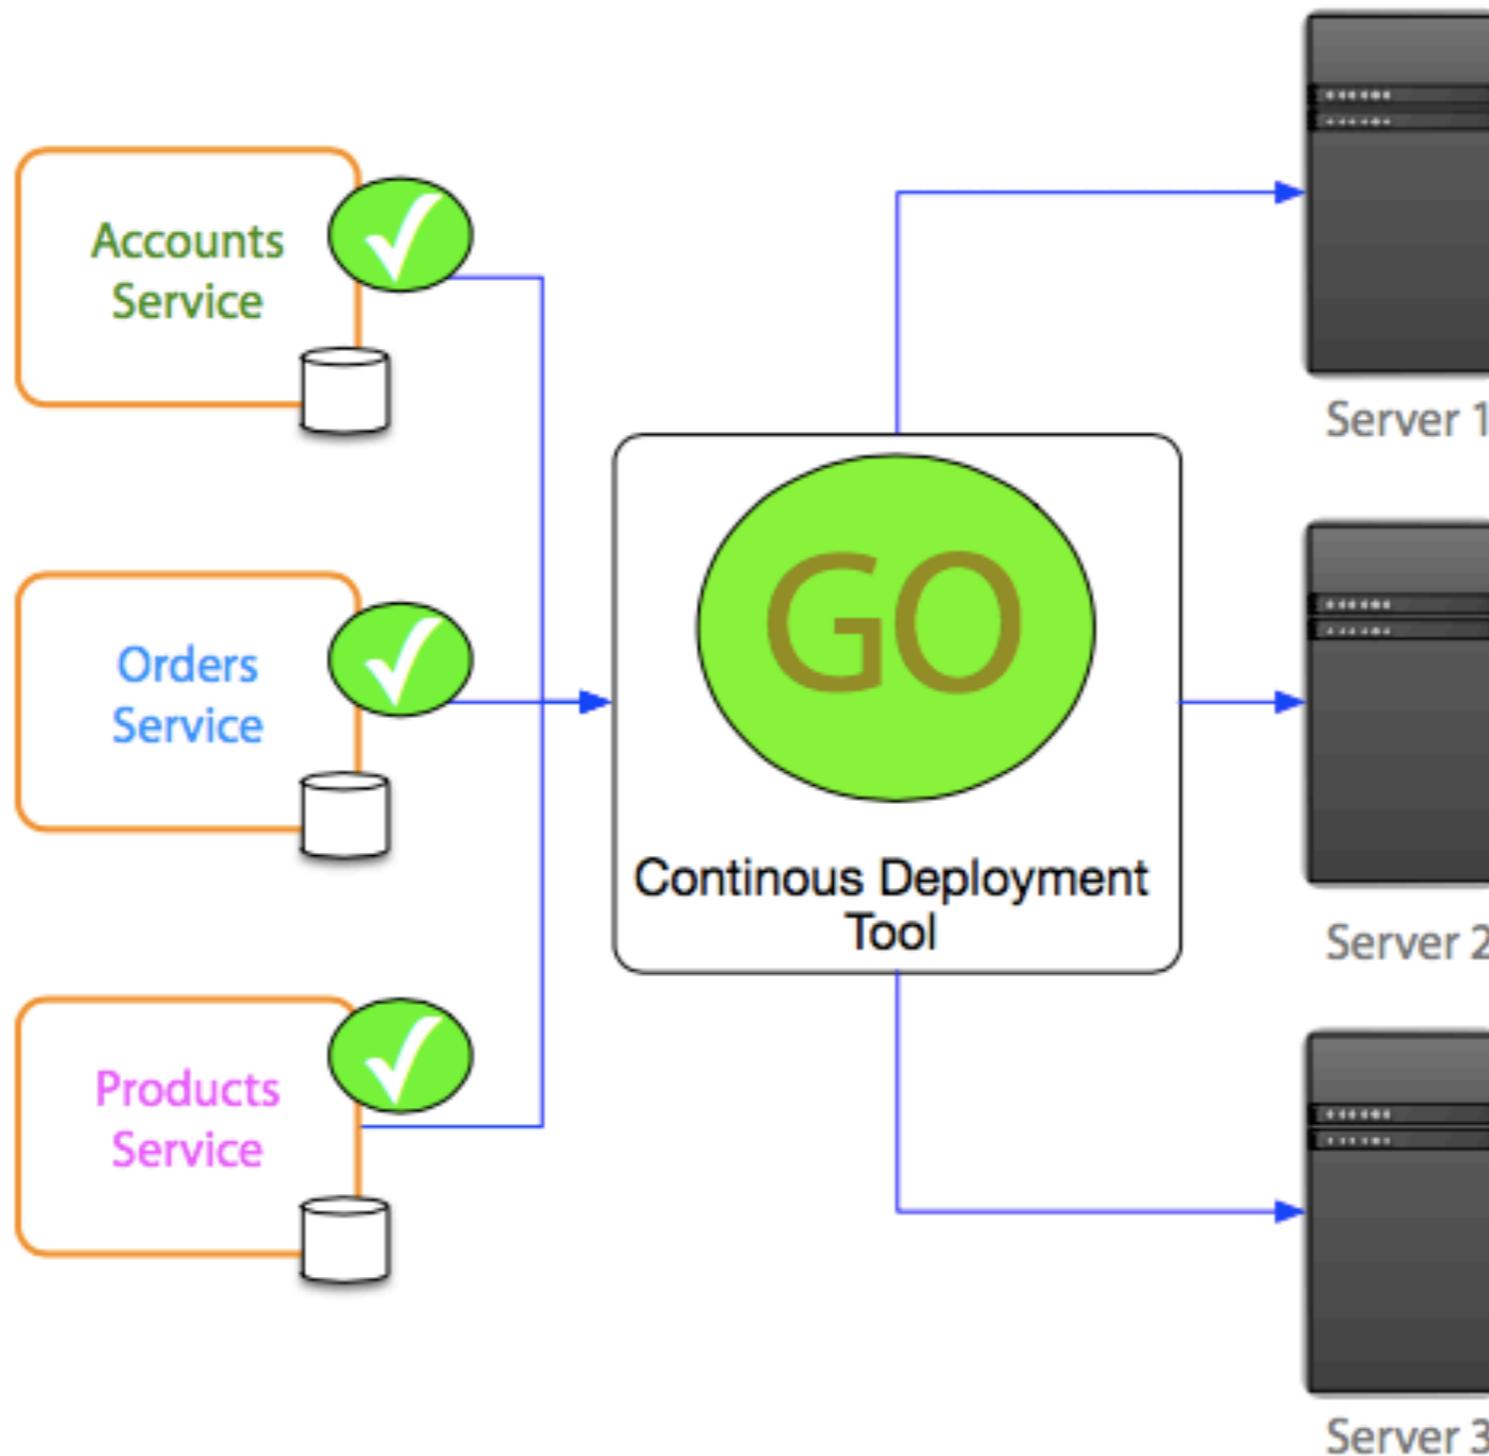
Continuous Integration

- Urgency to fix quickly
- Creation of build
- Build ready for test team
- Build ready for deployment

Continuous Deployment

- Automate software deployment
 - Configure once
 - Works with CI tools
 - Deployable after check in
 - Reliably released at anytime
- Benefits
 - Quick to market
 - Reliable deployment
 - Better customer experience

Continuous Deployment



Design Principle

High Cohesion
Single thing done well
Single focus

Autonomous
Independently changeable
Independently deployable

Business Domain Centric
Represent business function or represent a business domain

Resillience
Embrace failure
Default or Degrade functionality

Observable
See System Health
Centralize logging and monitoring

Automation
Tools for testing and feedback
Tools for deployment

Approach

High Cohesion

Keeps splitting service until it only has one reason to change

Autonomous
Loosely Coupled system
Versioning strategy
Microservice Ownership by team

Business Domain Centric

Course grain business domain
Subgroups into functions and areas

Resillience

Design for known failures
Design for fail fast (recover fast)

Observable
Real-time centralized monitoring
Centralized logging

Automation

Continuous Integration
Cntinuous Deployment

Technology for MicroServices

Communication

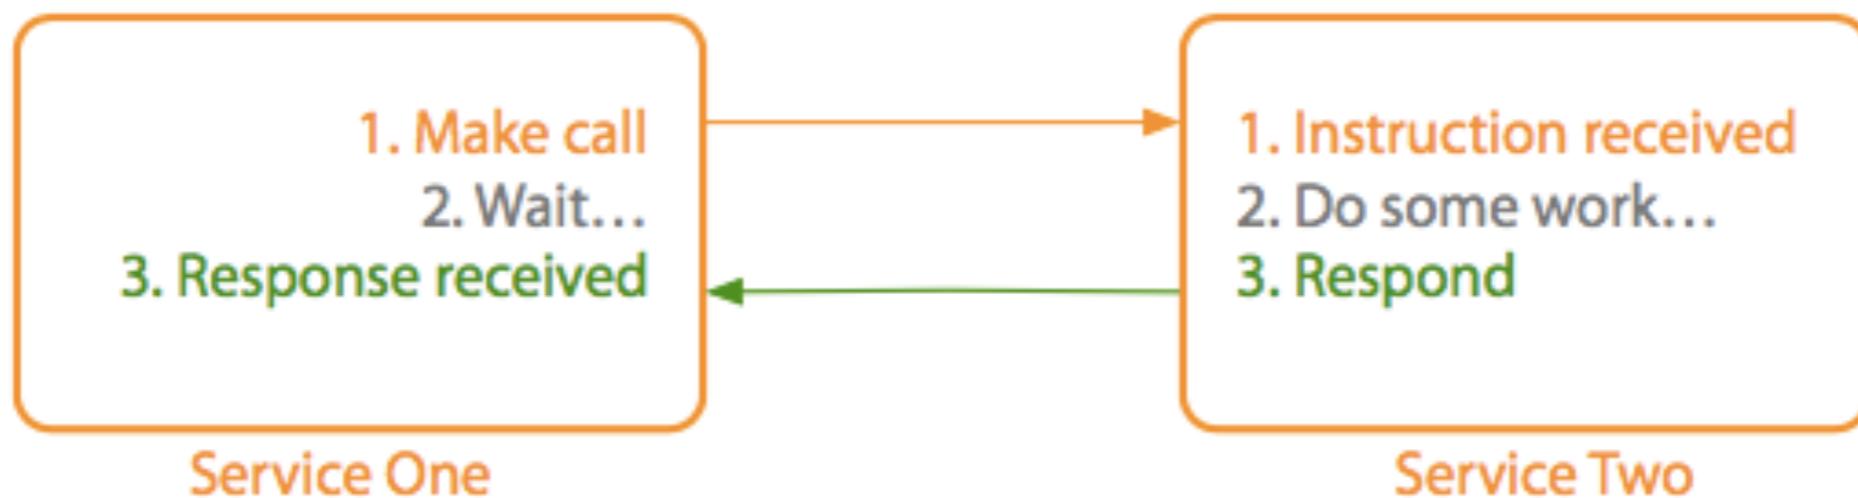
Inter-Process Communications

Synchronous
One-To-One

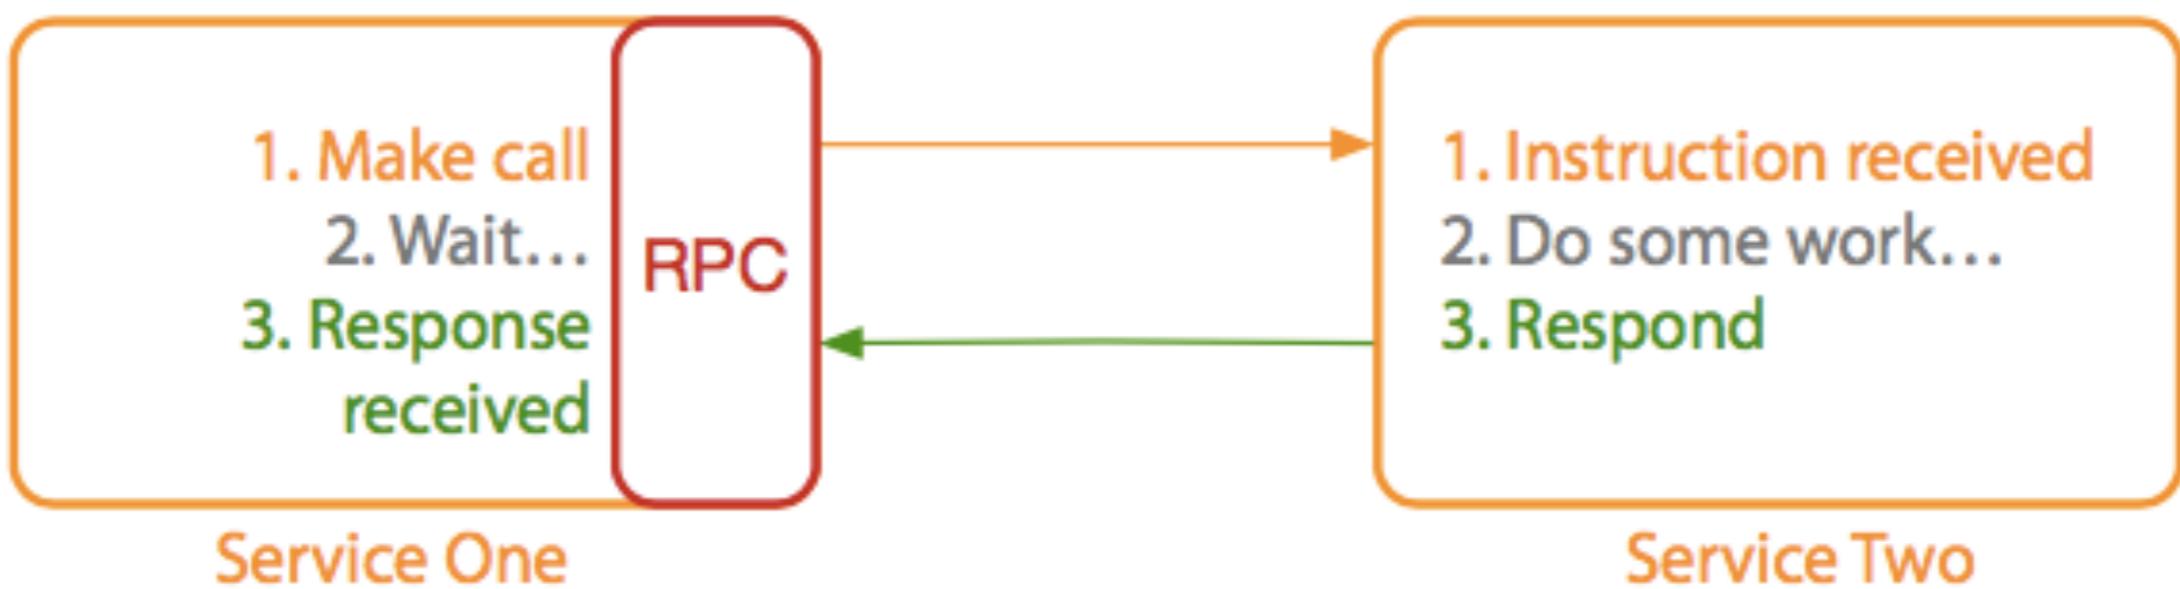
Asynchronous
One-To-Many

Synchronous

- Request response communication
 - Client to service
 - Service to service
 - Service to external



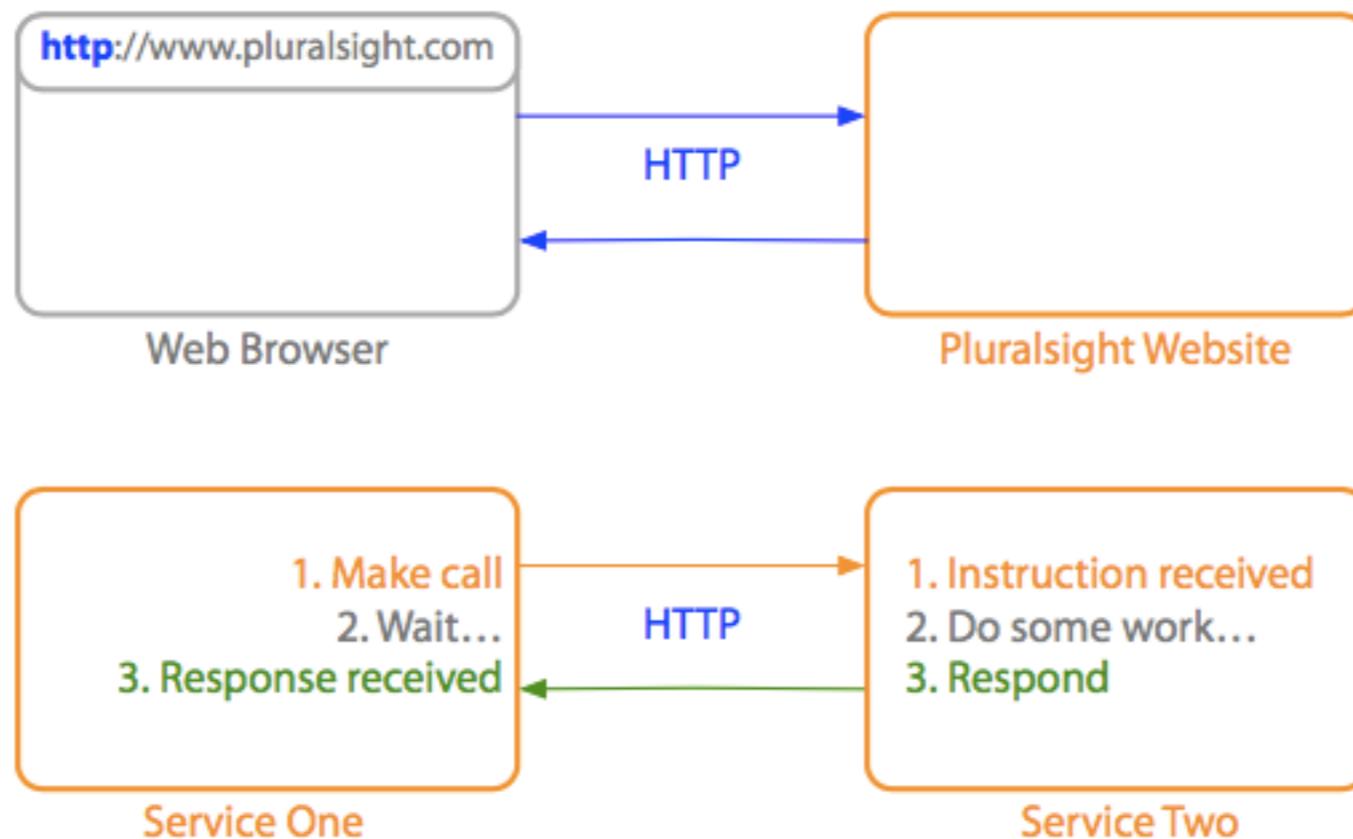
Remote Procedure Call



Sensitive to Change

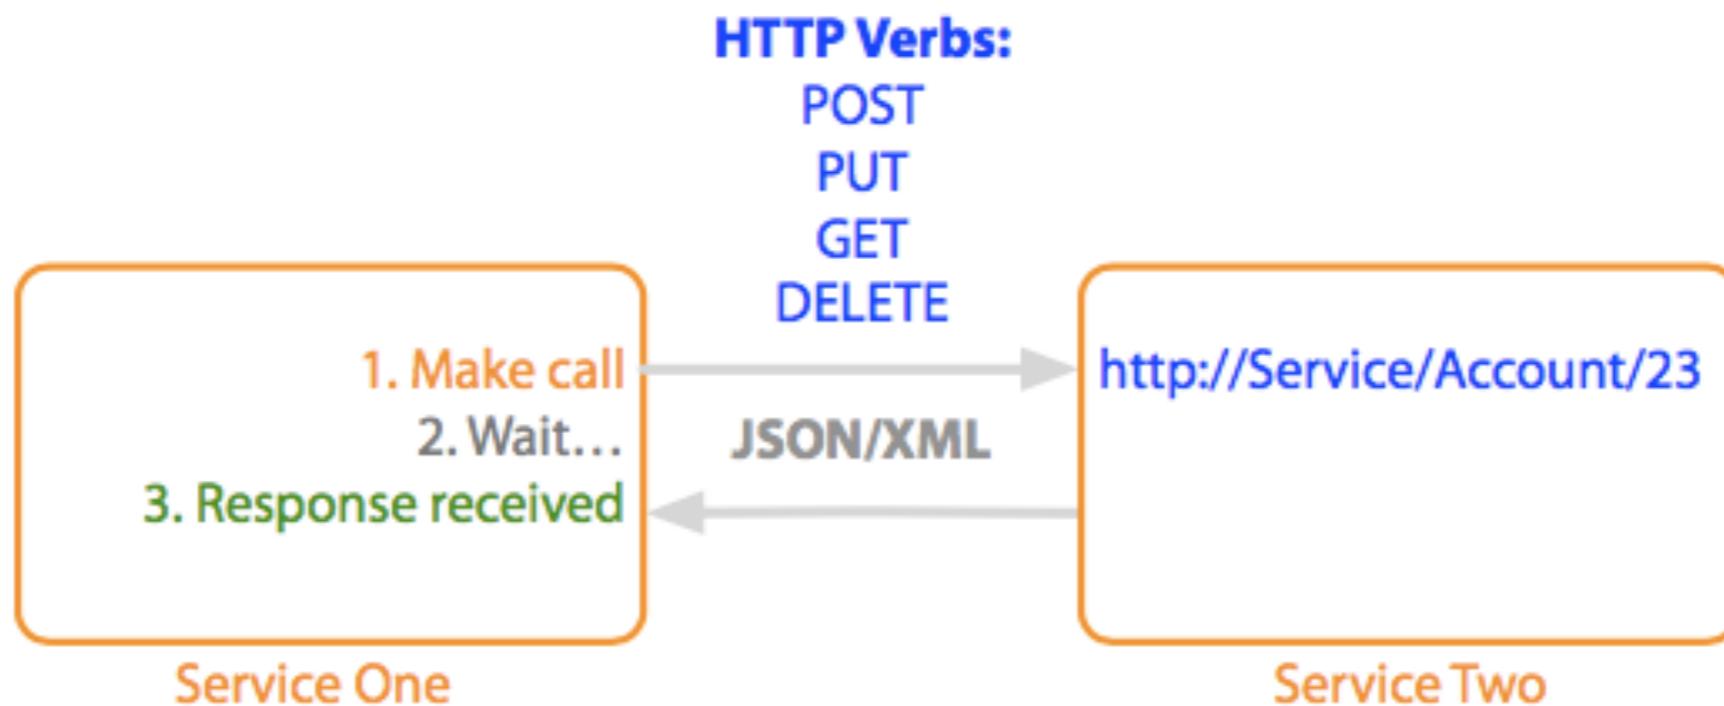
HTTP

- Work across the internet
- Firewall friendly

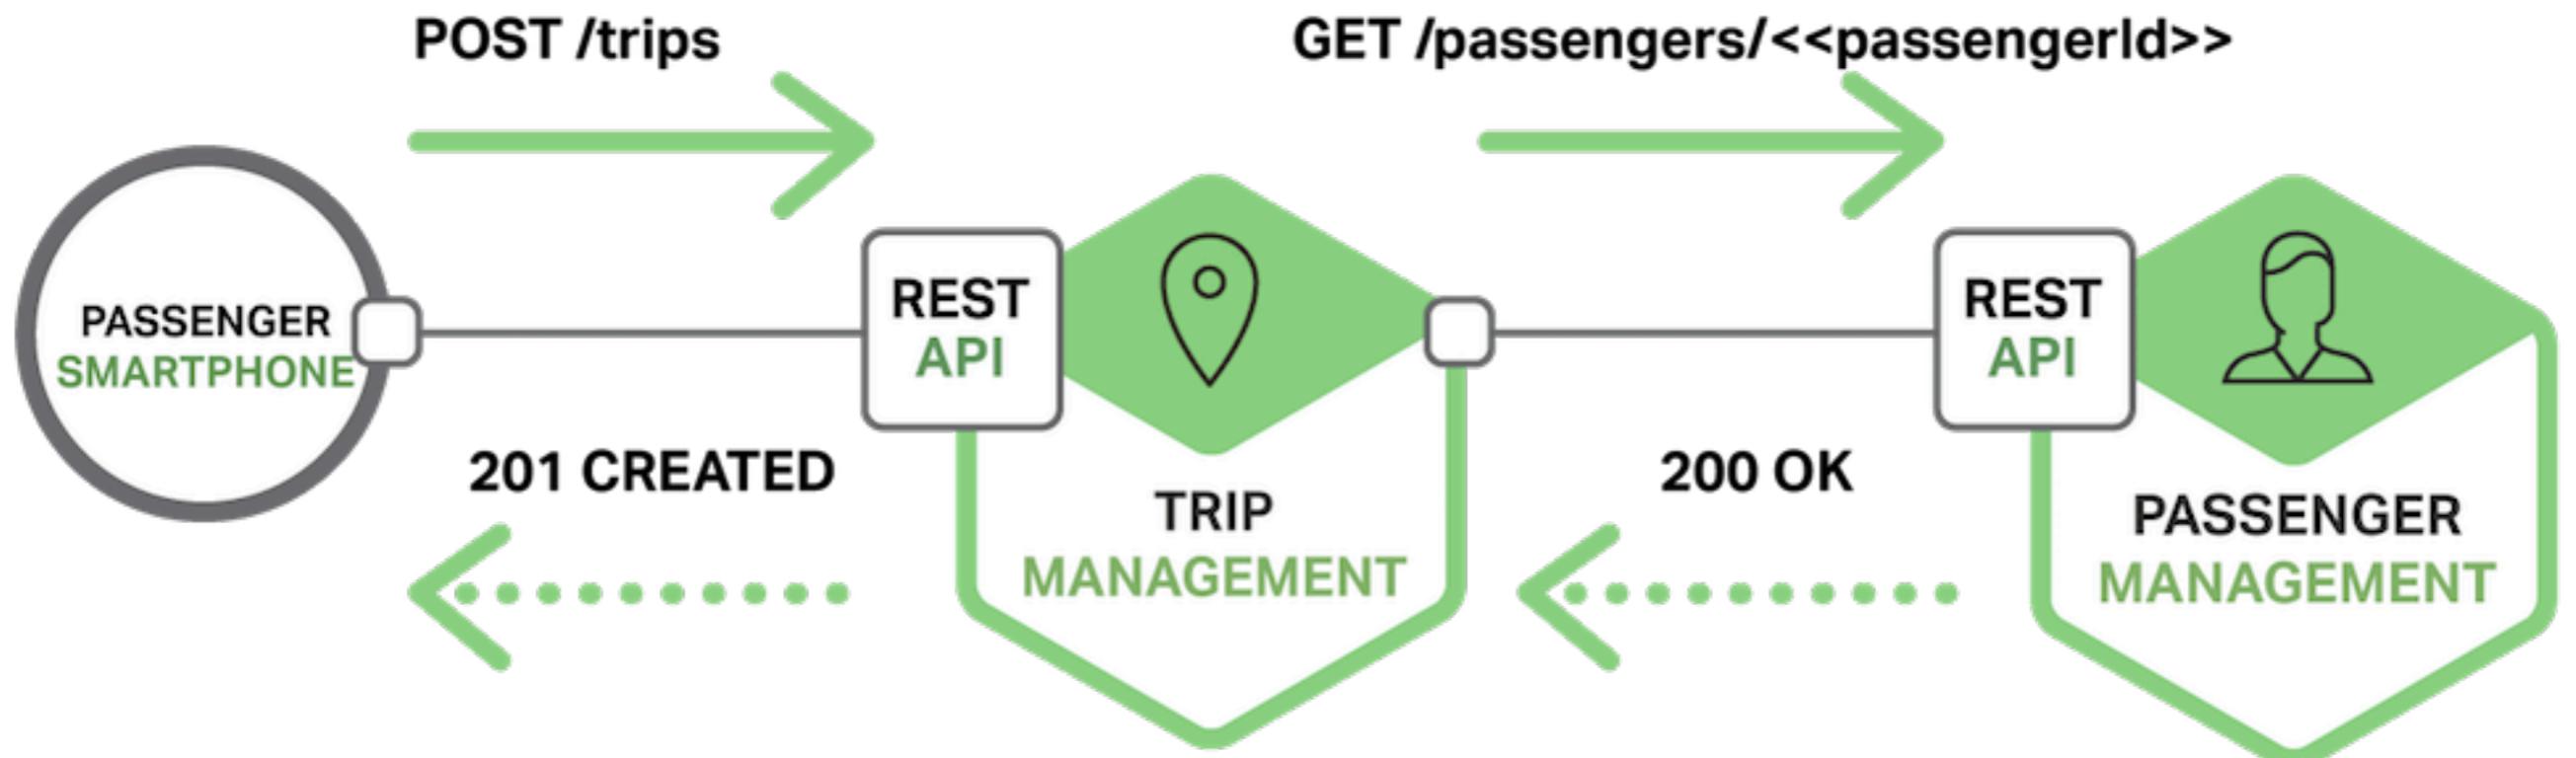


RESTful Service

- CRUD using HTTP verbs
- Natural decoupling
- Open communication protocol



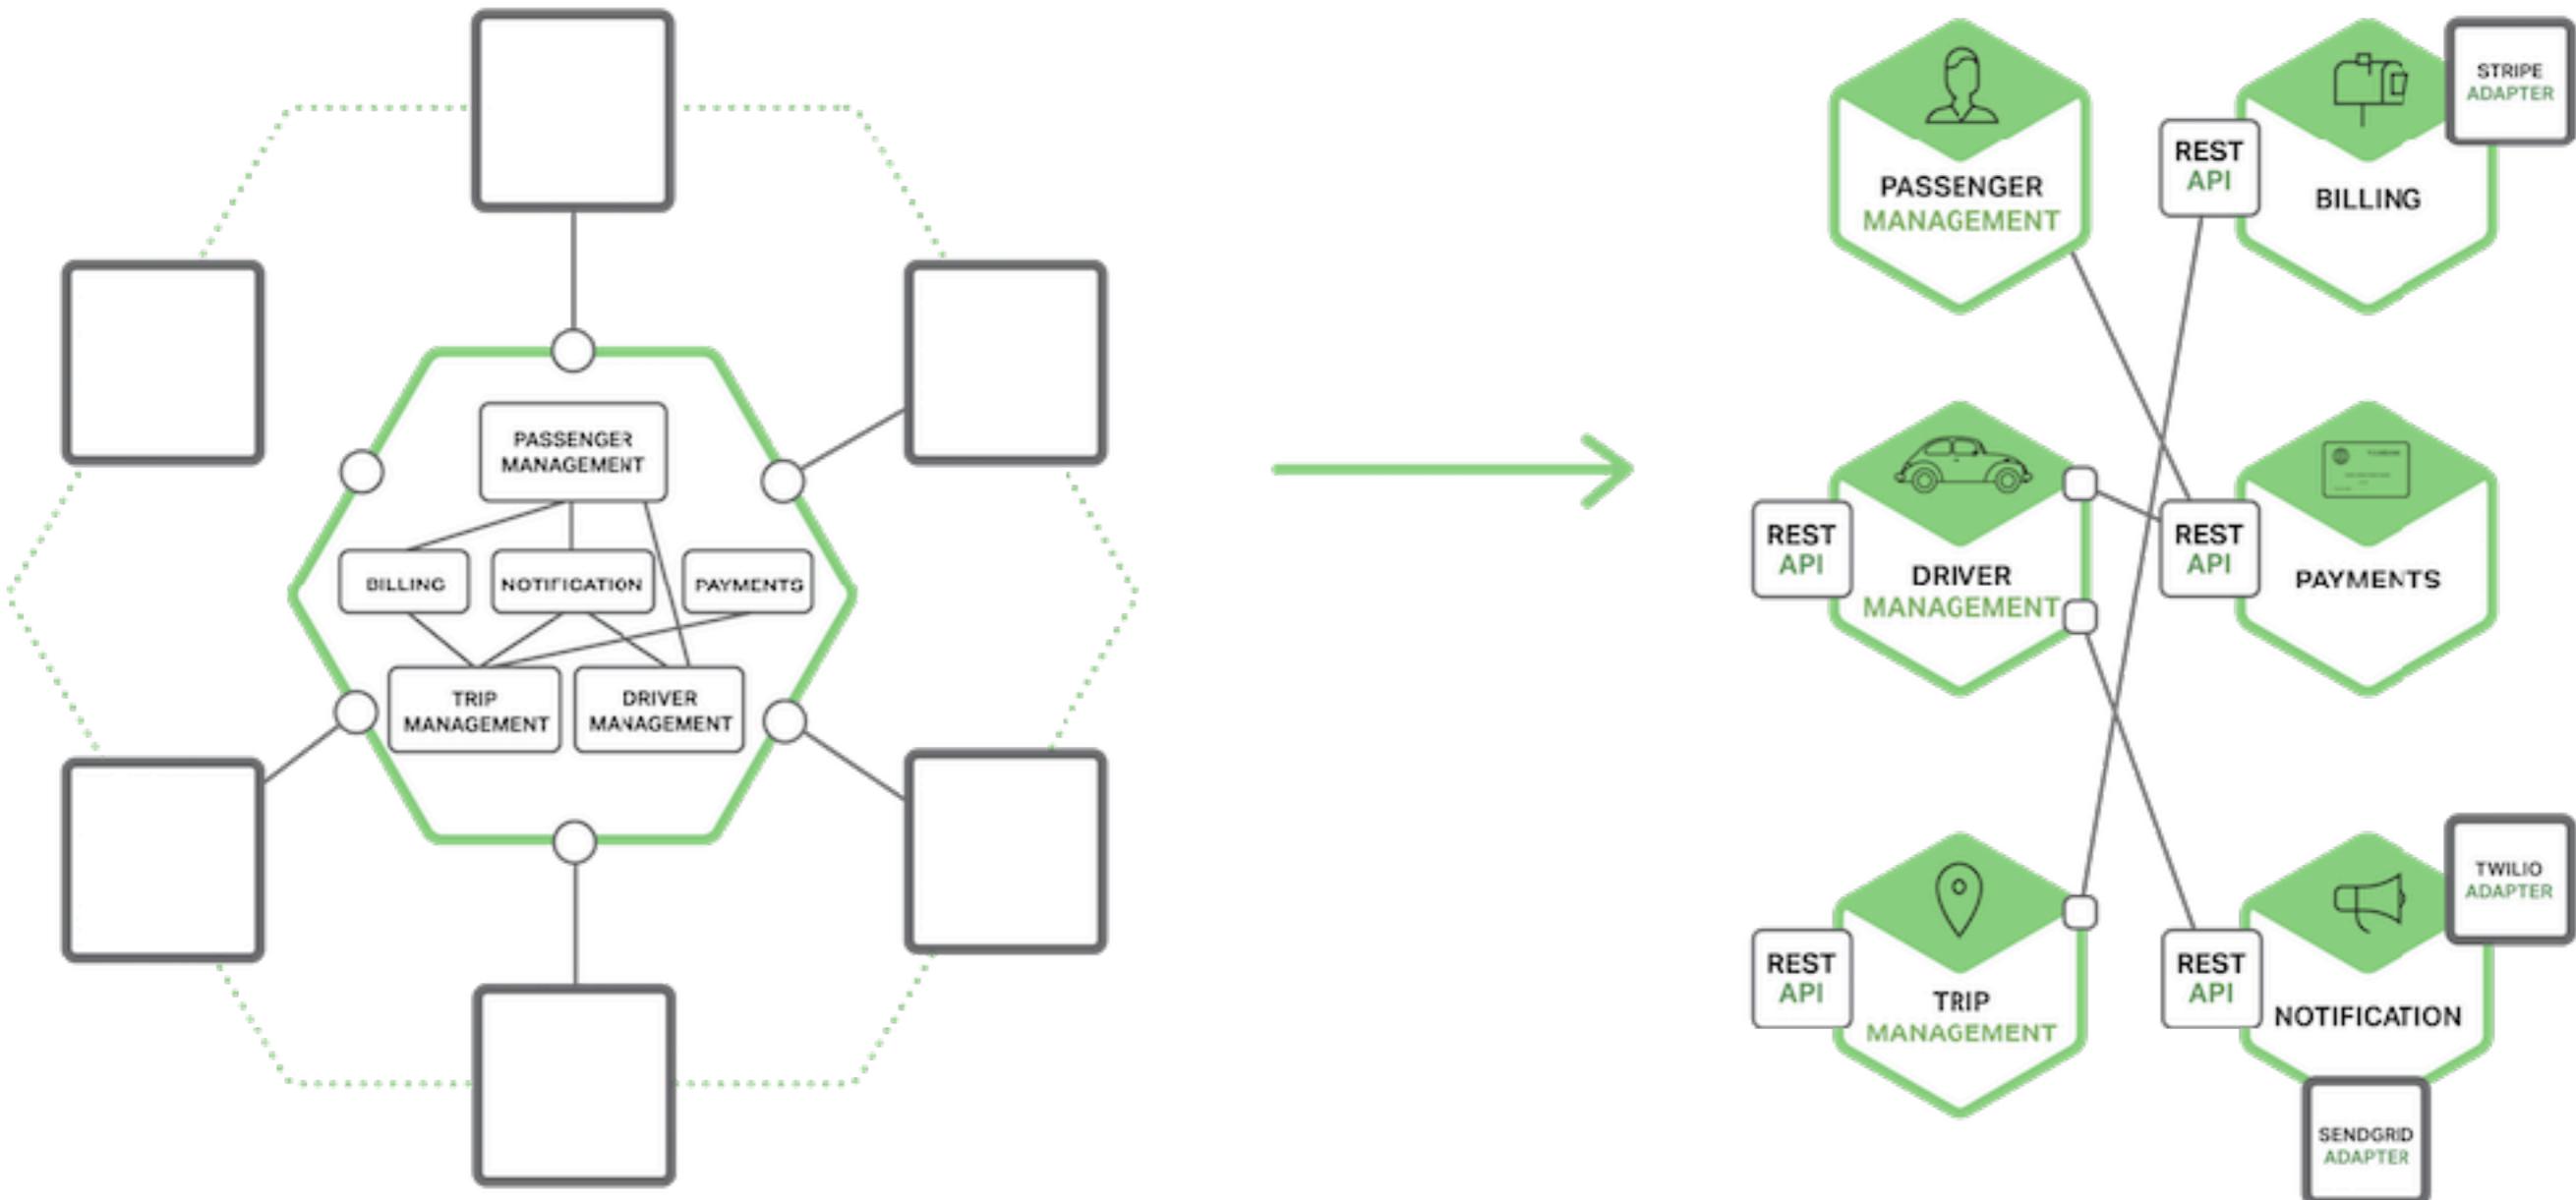
Taxi-Haling REST



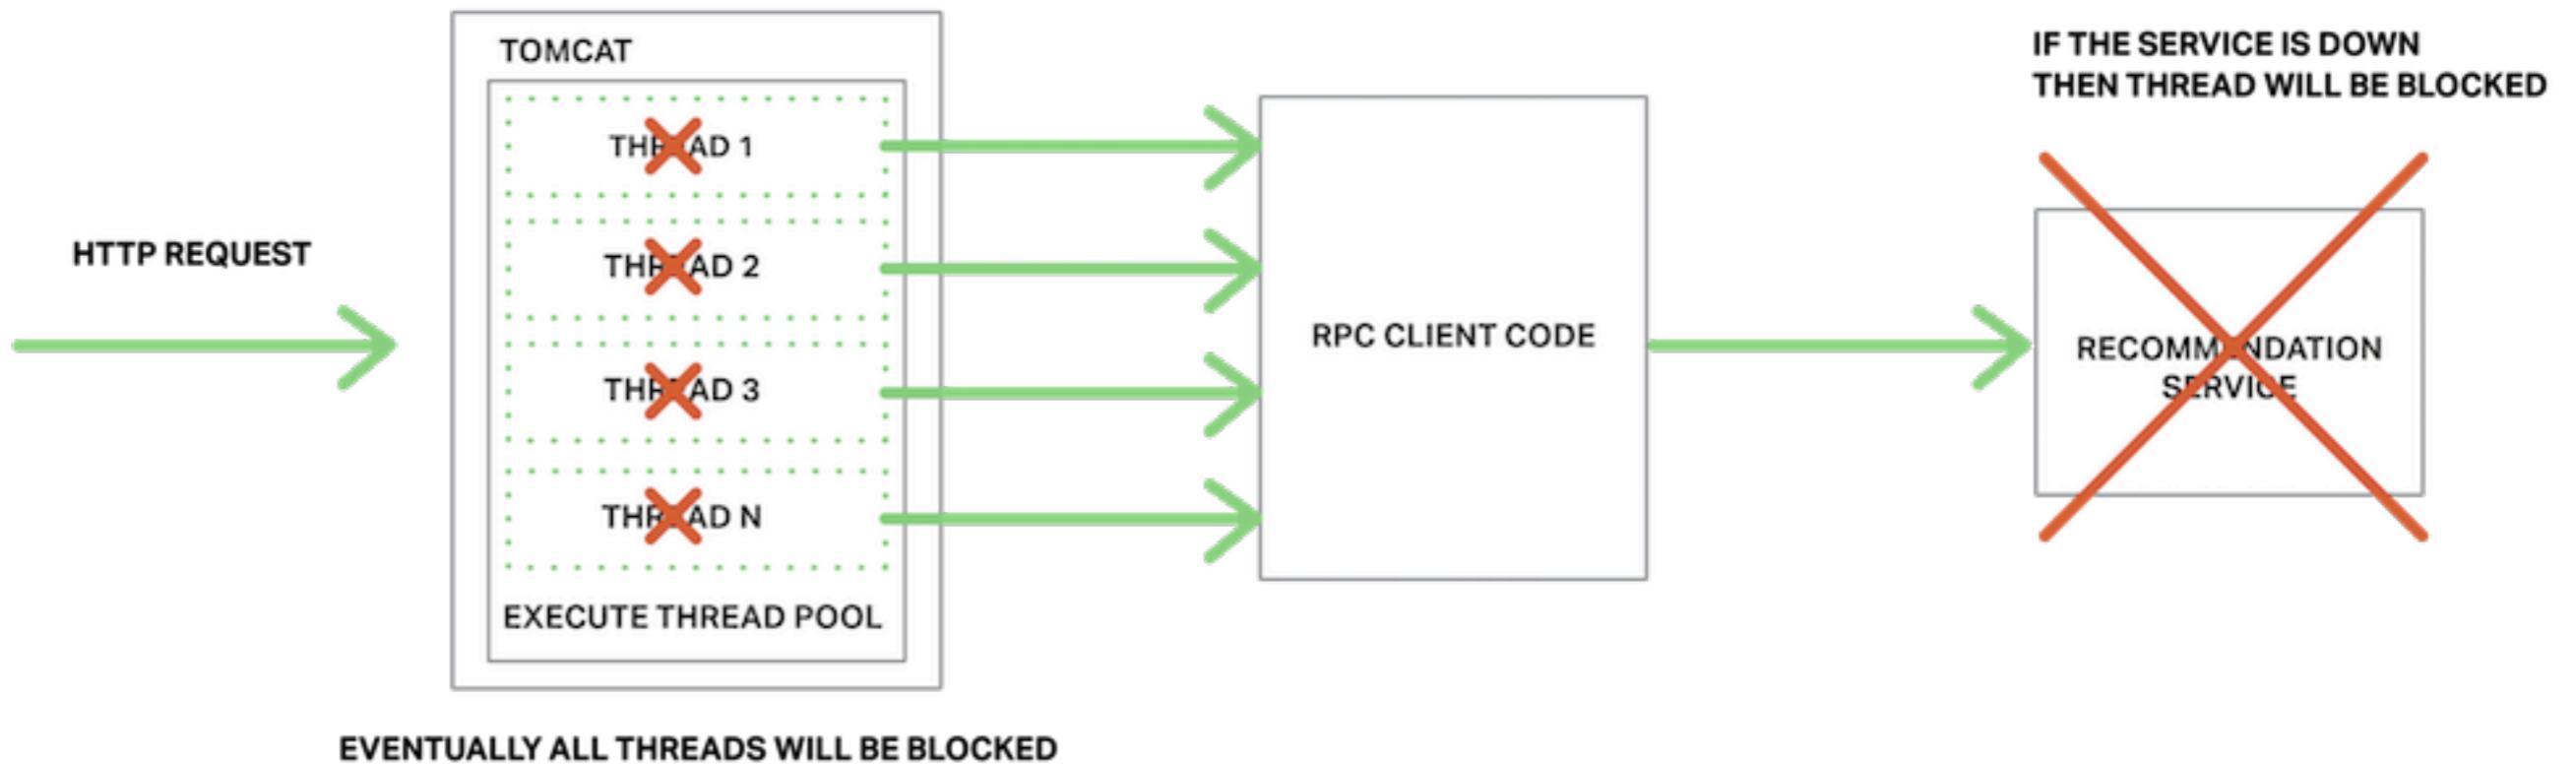
Synchronous Issues

- Both parties have to be available
- Performance subject to network quality
- Clients must know location of service (host\port)

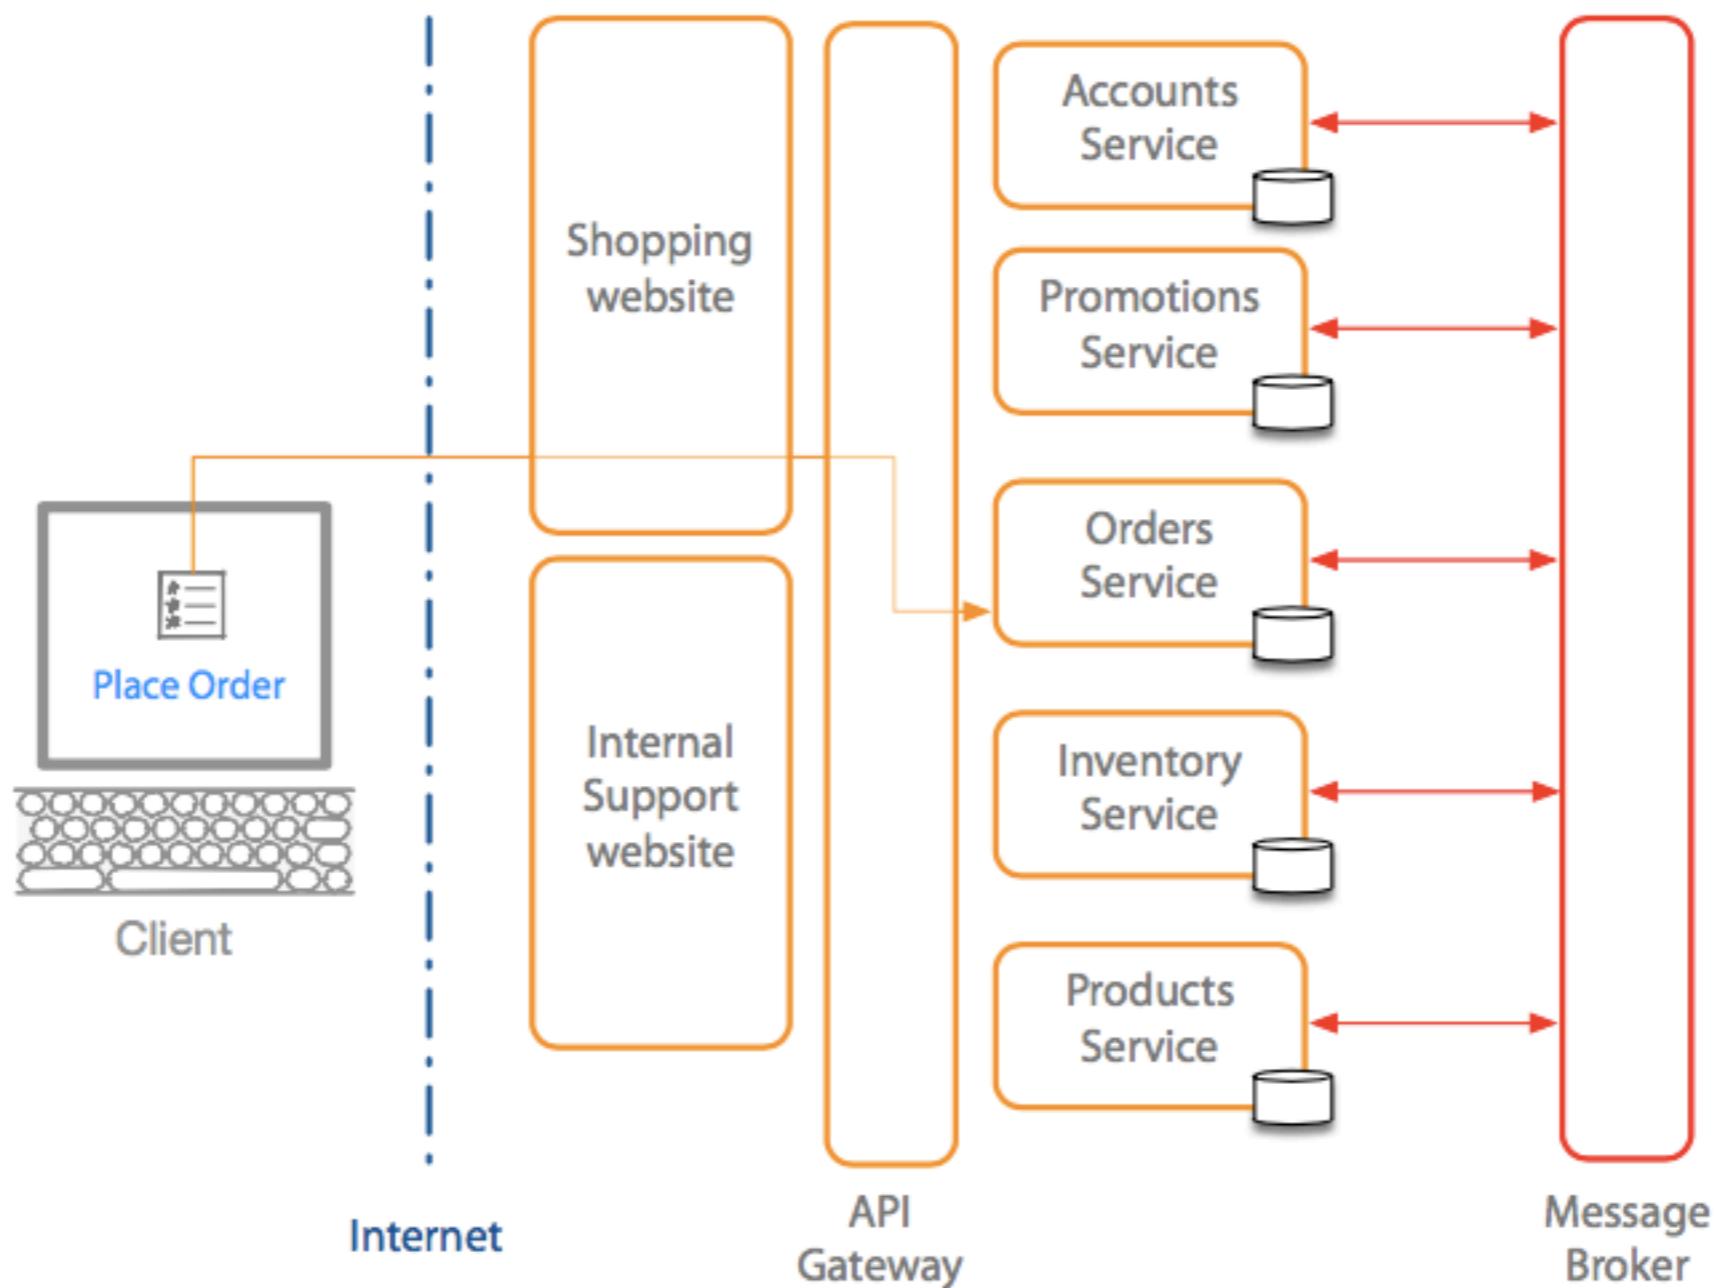
One-To-One



Problem



Asynchronous

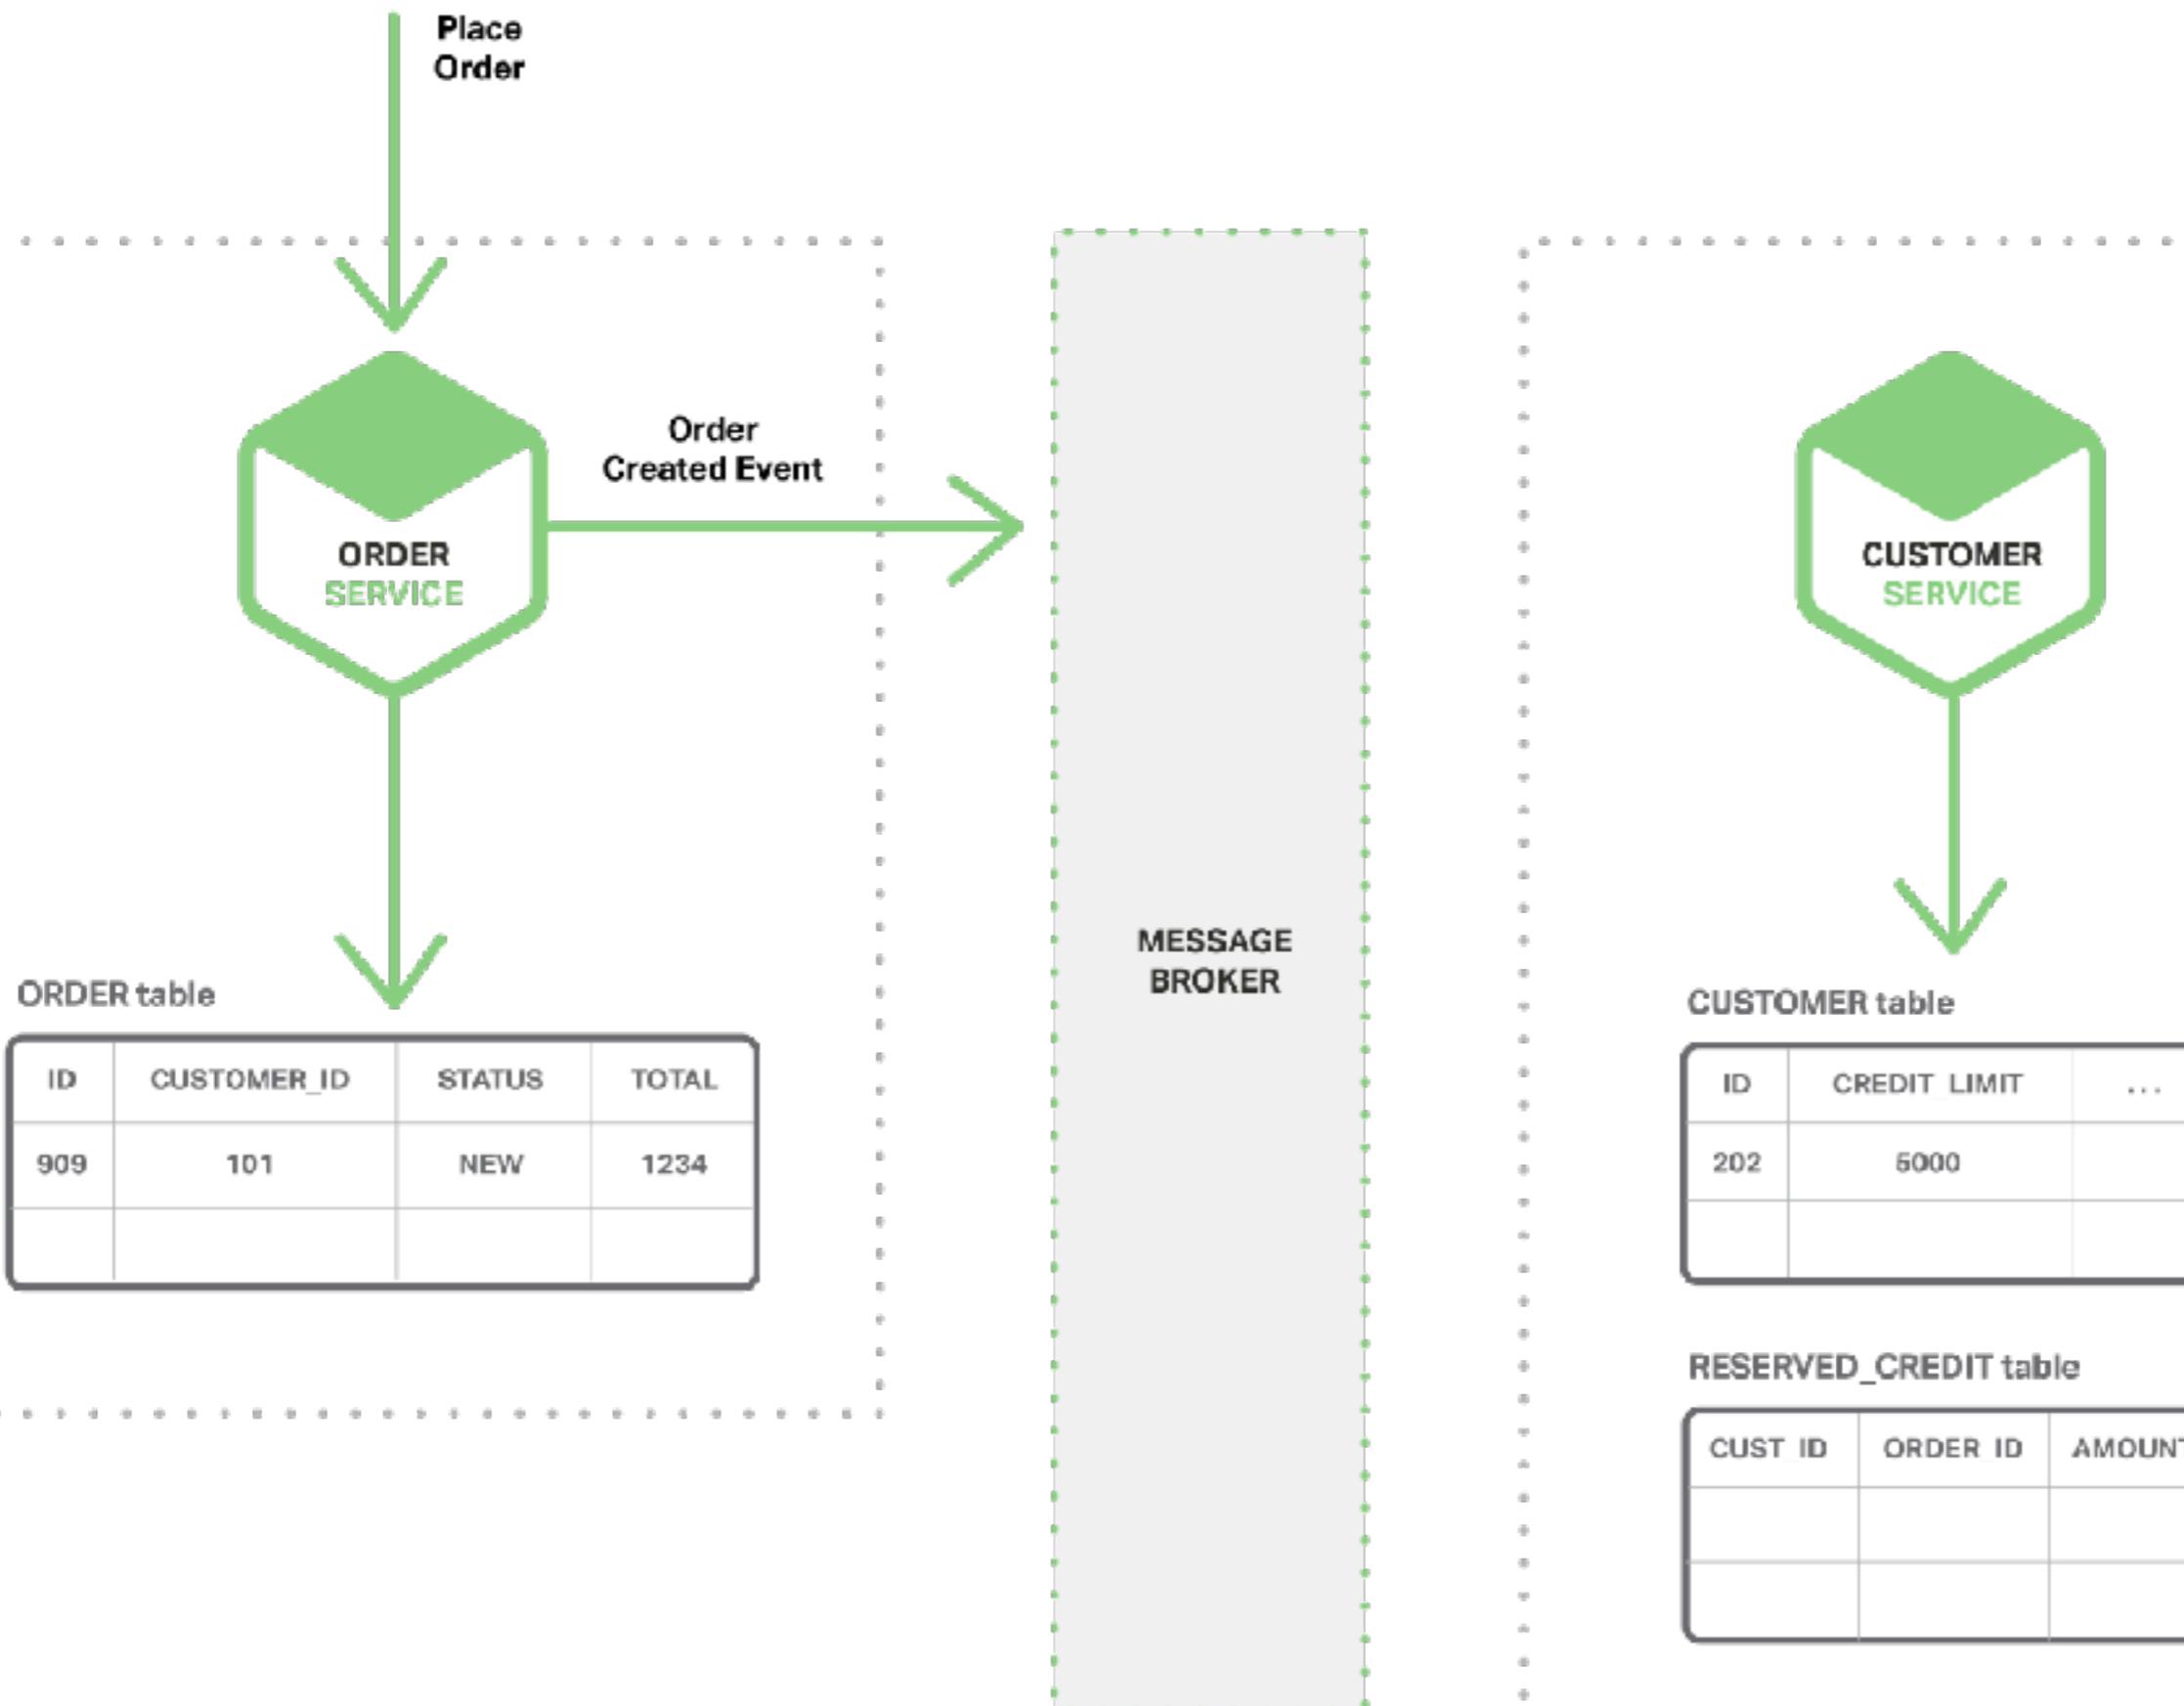


Asynchronous

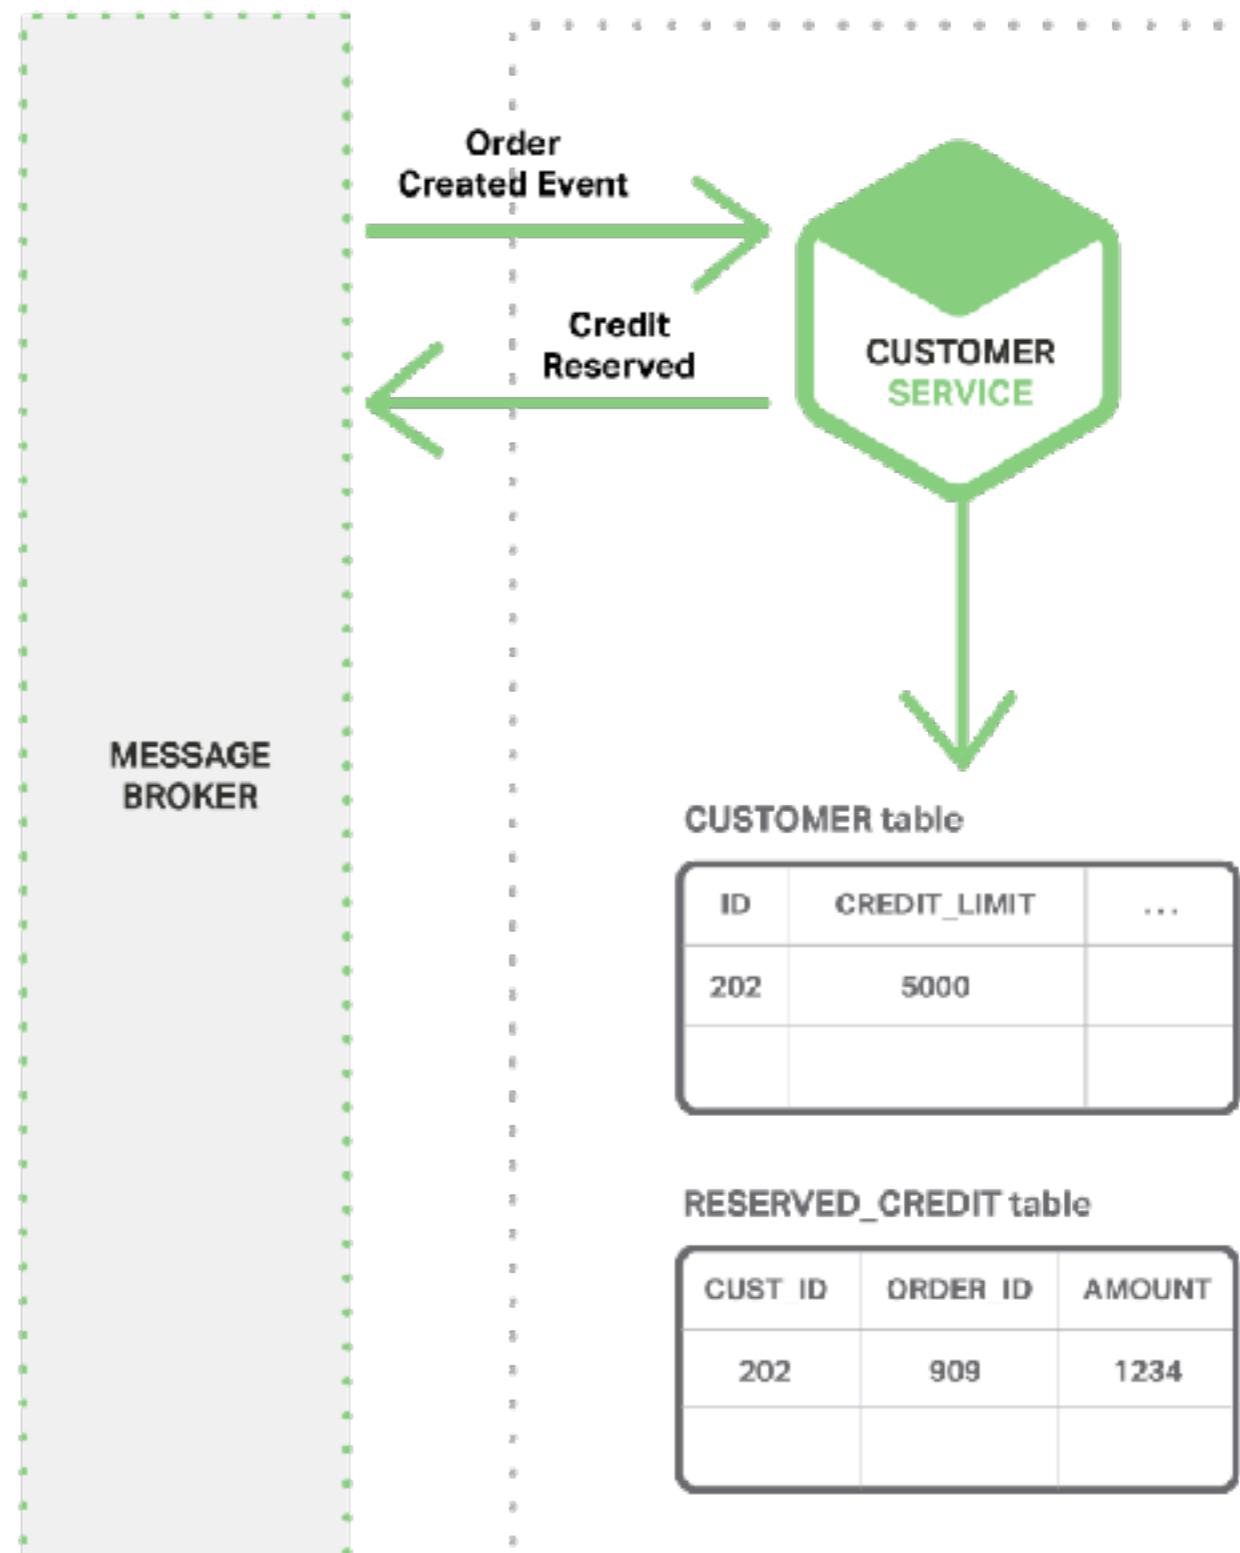
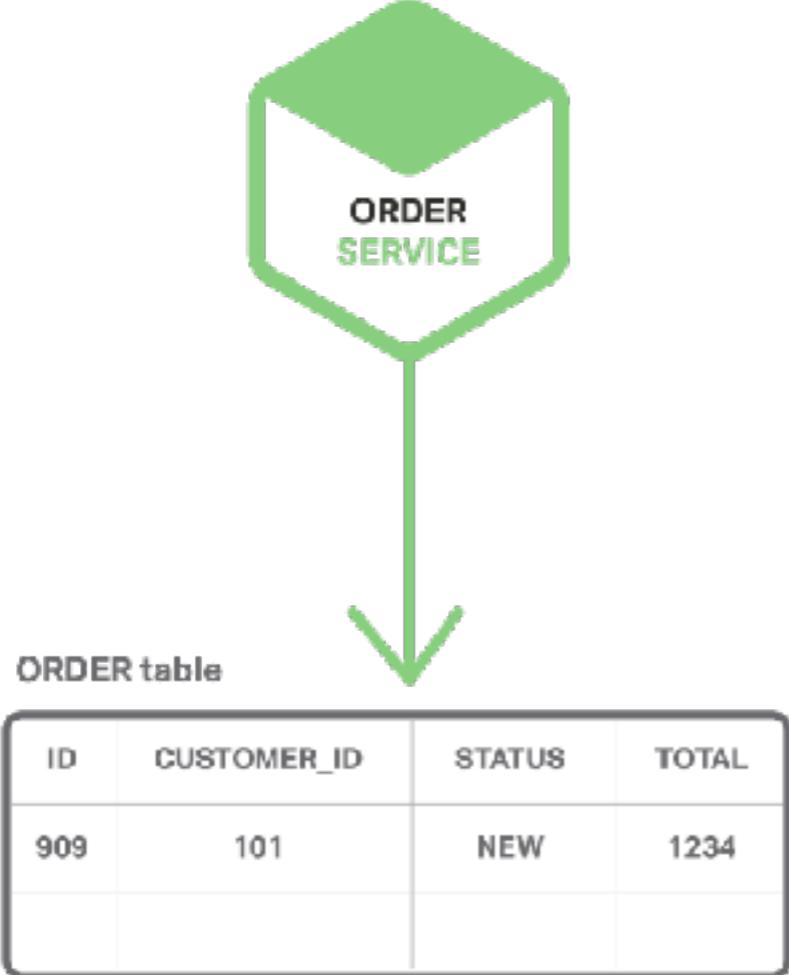
- Event based
 - Mitigates the need of client and service availability
 - Decouples client and service
- Message queueing protocol
 - Message Brokers
 - Subscriber and publisher are decoupled
 - Microsoft message queuing (MSMQ)
 - RabbitMQ
 - ATOM (HTTP to propagate events)

Asynchronous

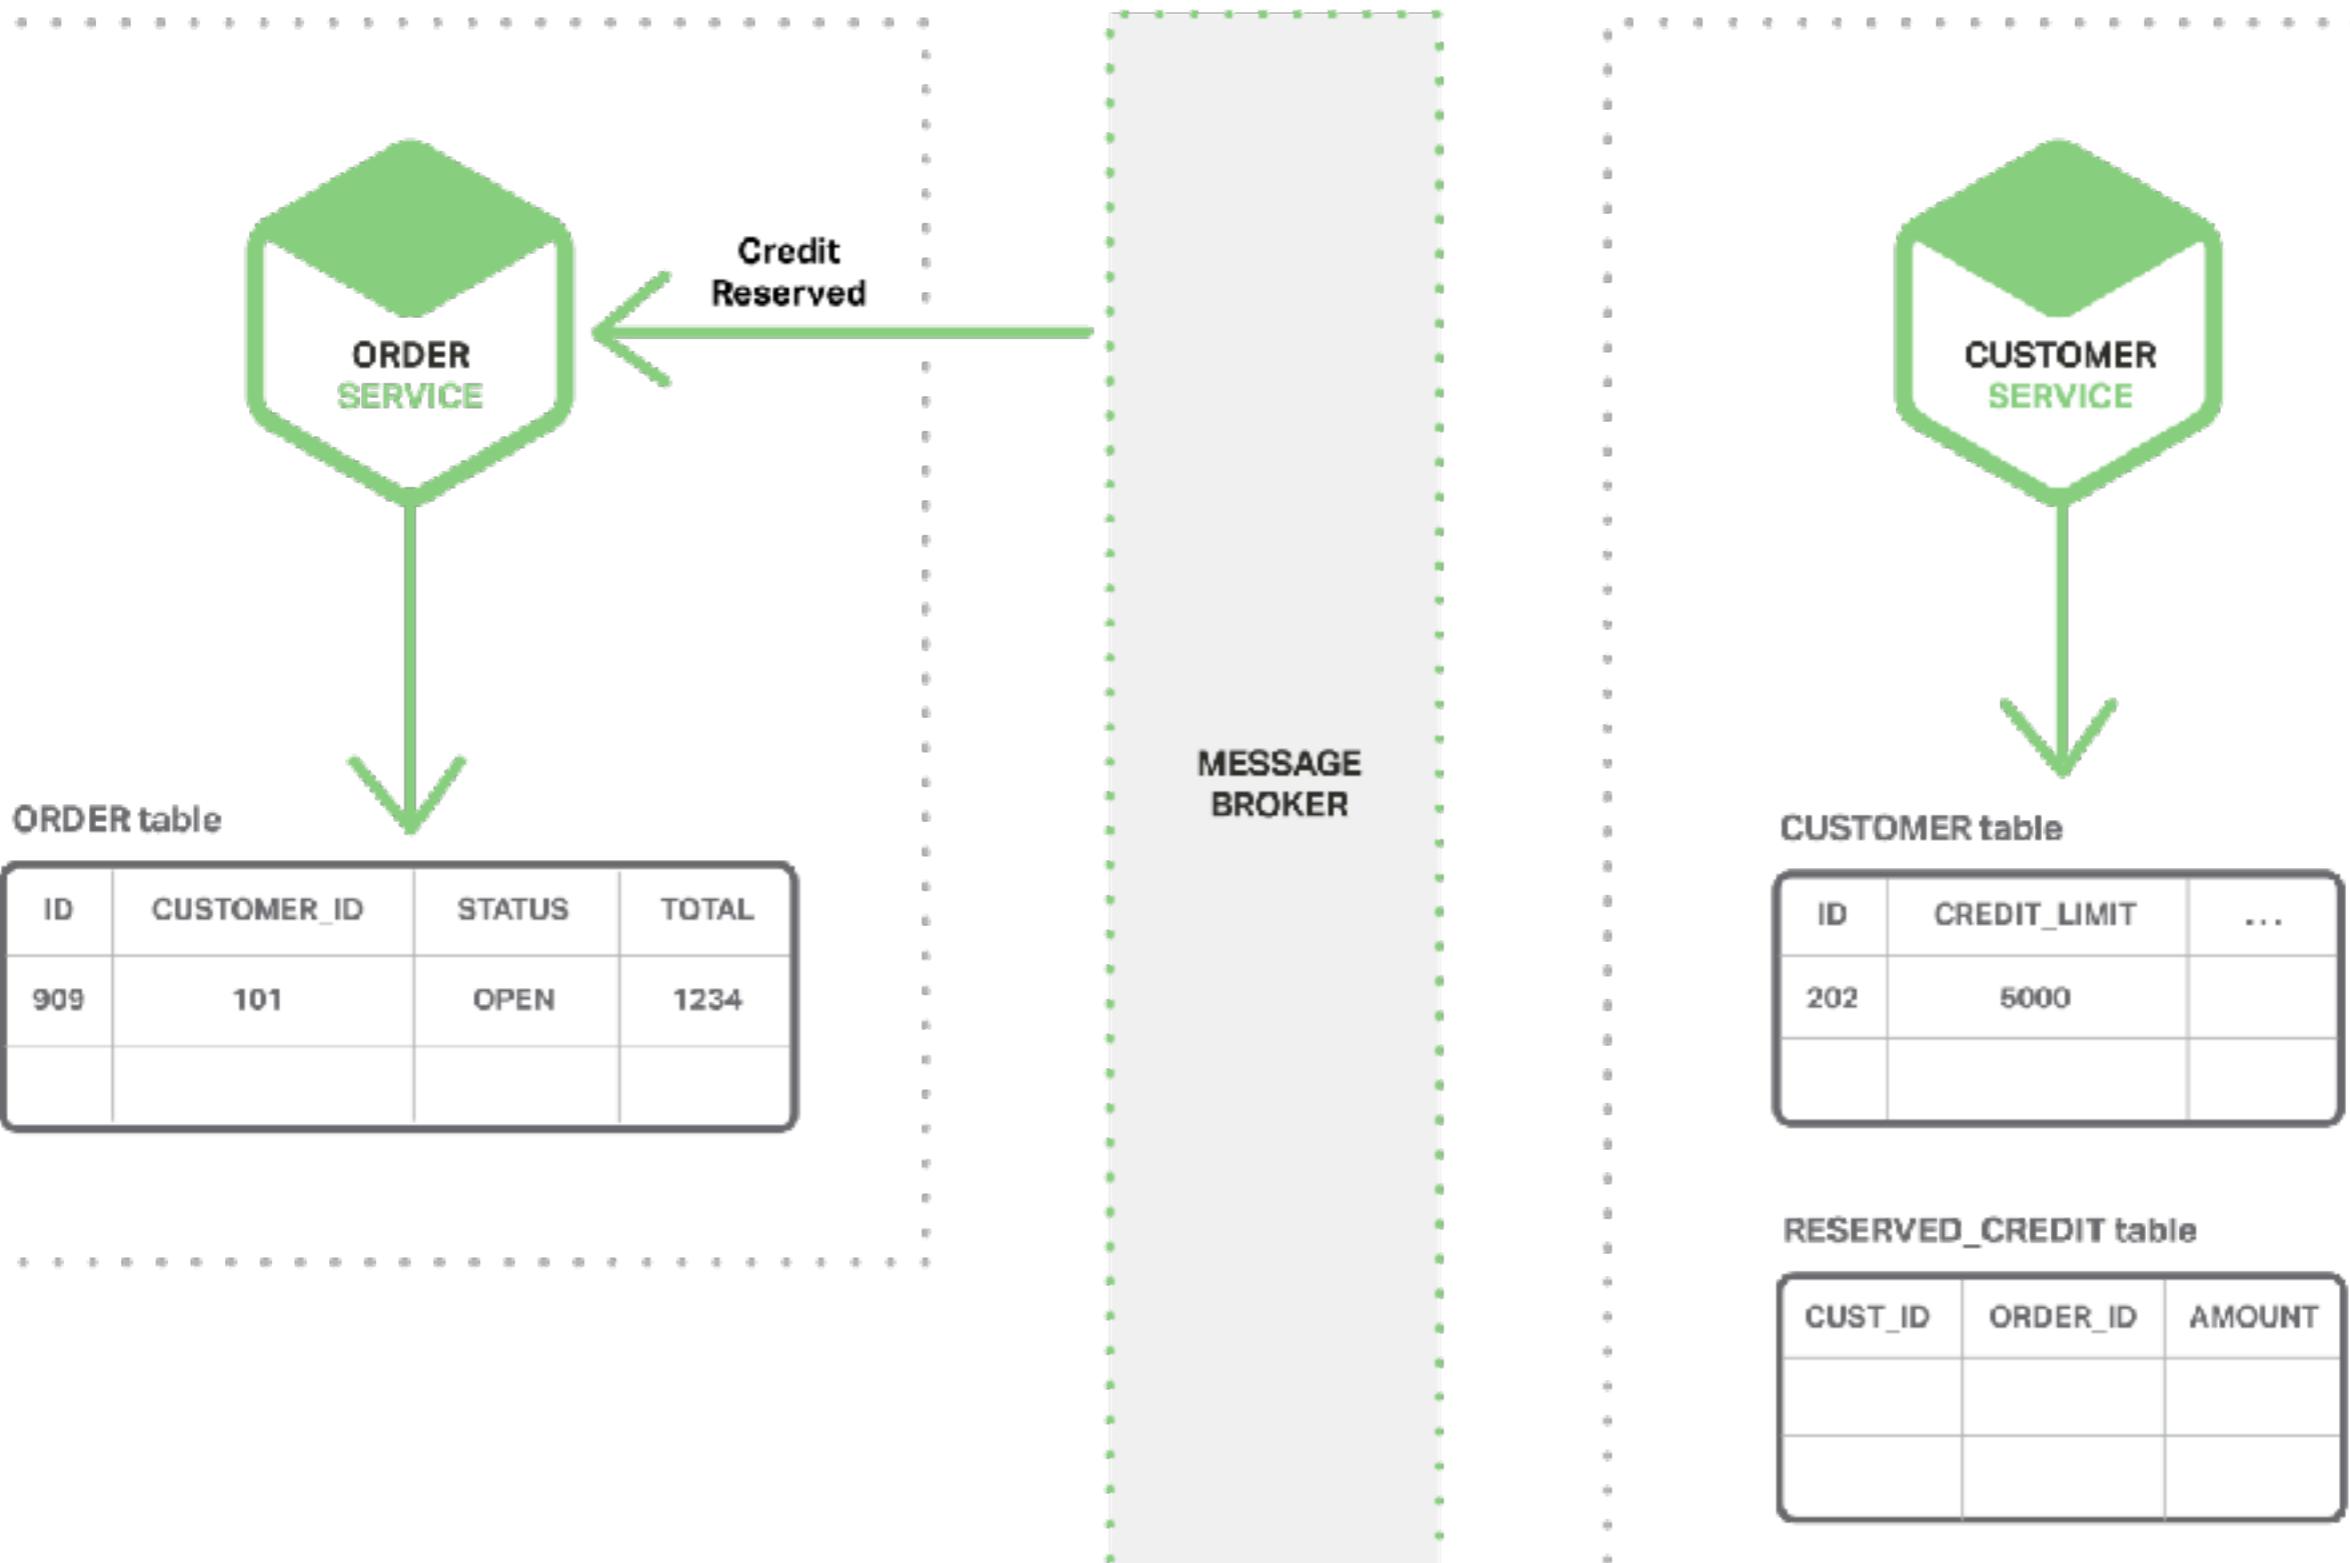
- Asynchronous challenge
 - Complicated
 - Reliance on message broker
 - Visibility of the transaction
 - Managing the messaging queue
- Real world systems
 - Would use both synchronous and asynchronous



Event-Driven Architecture

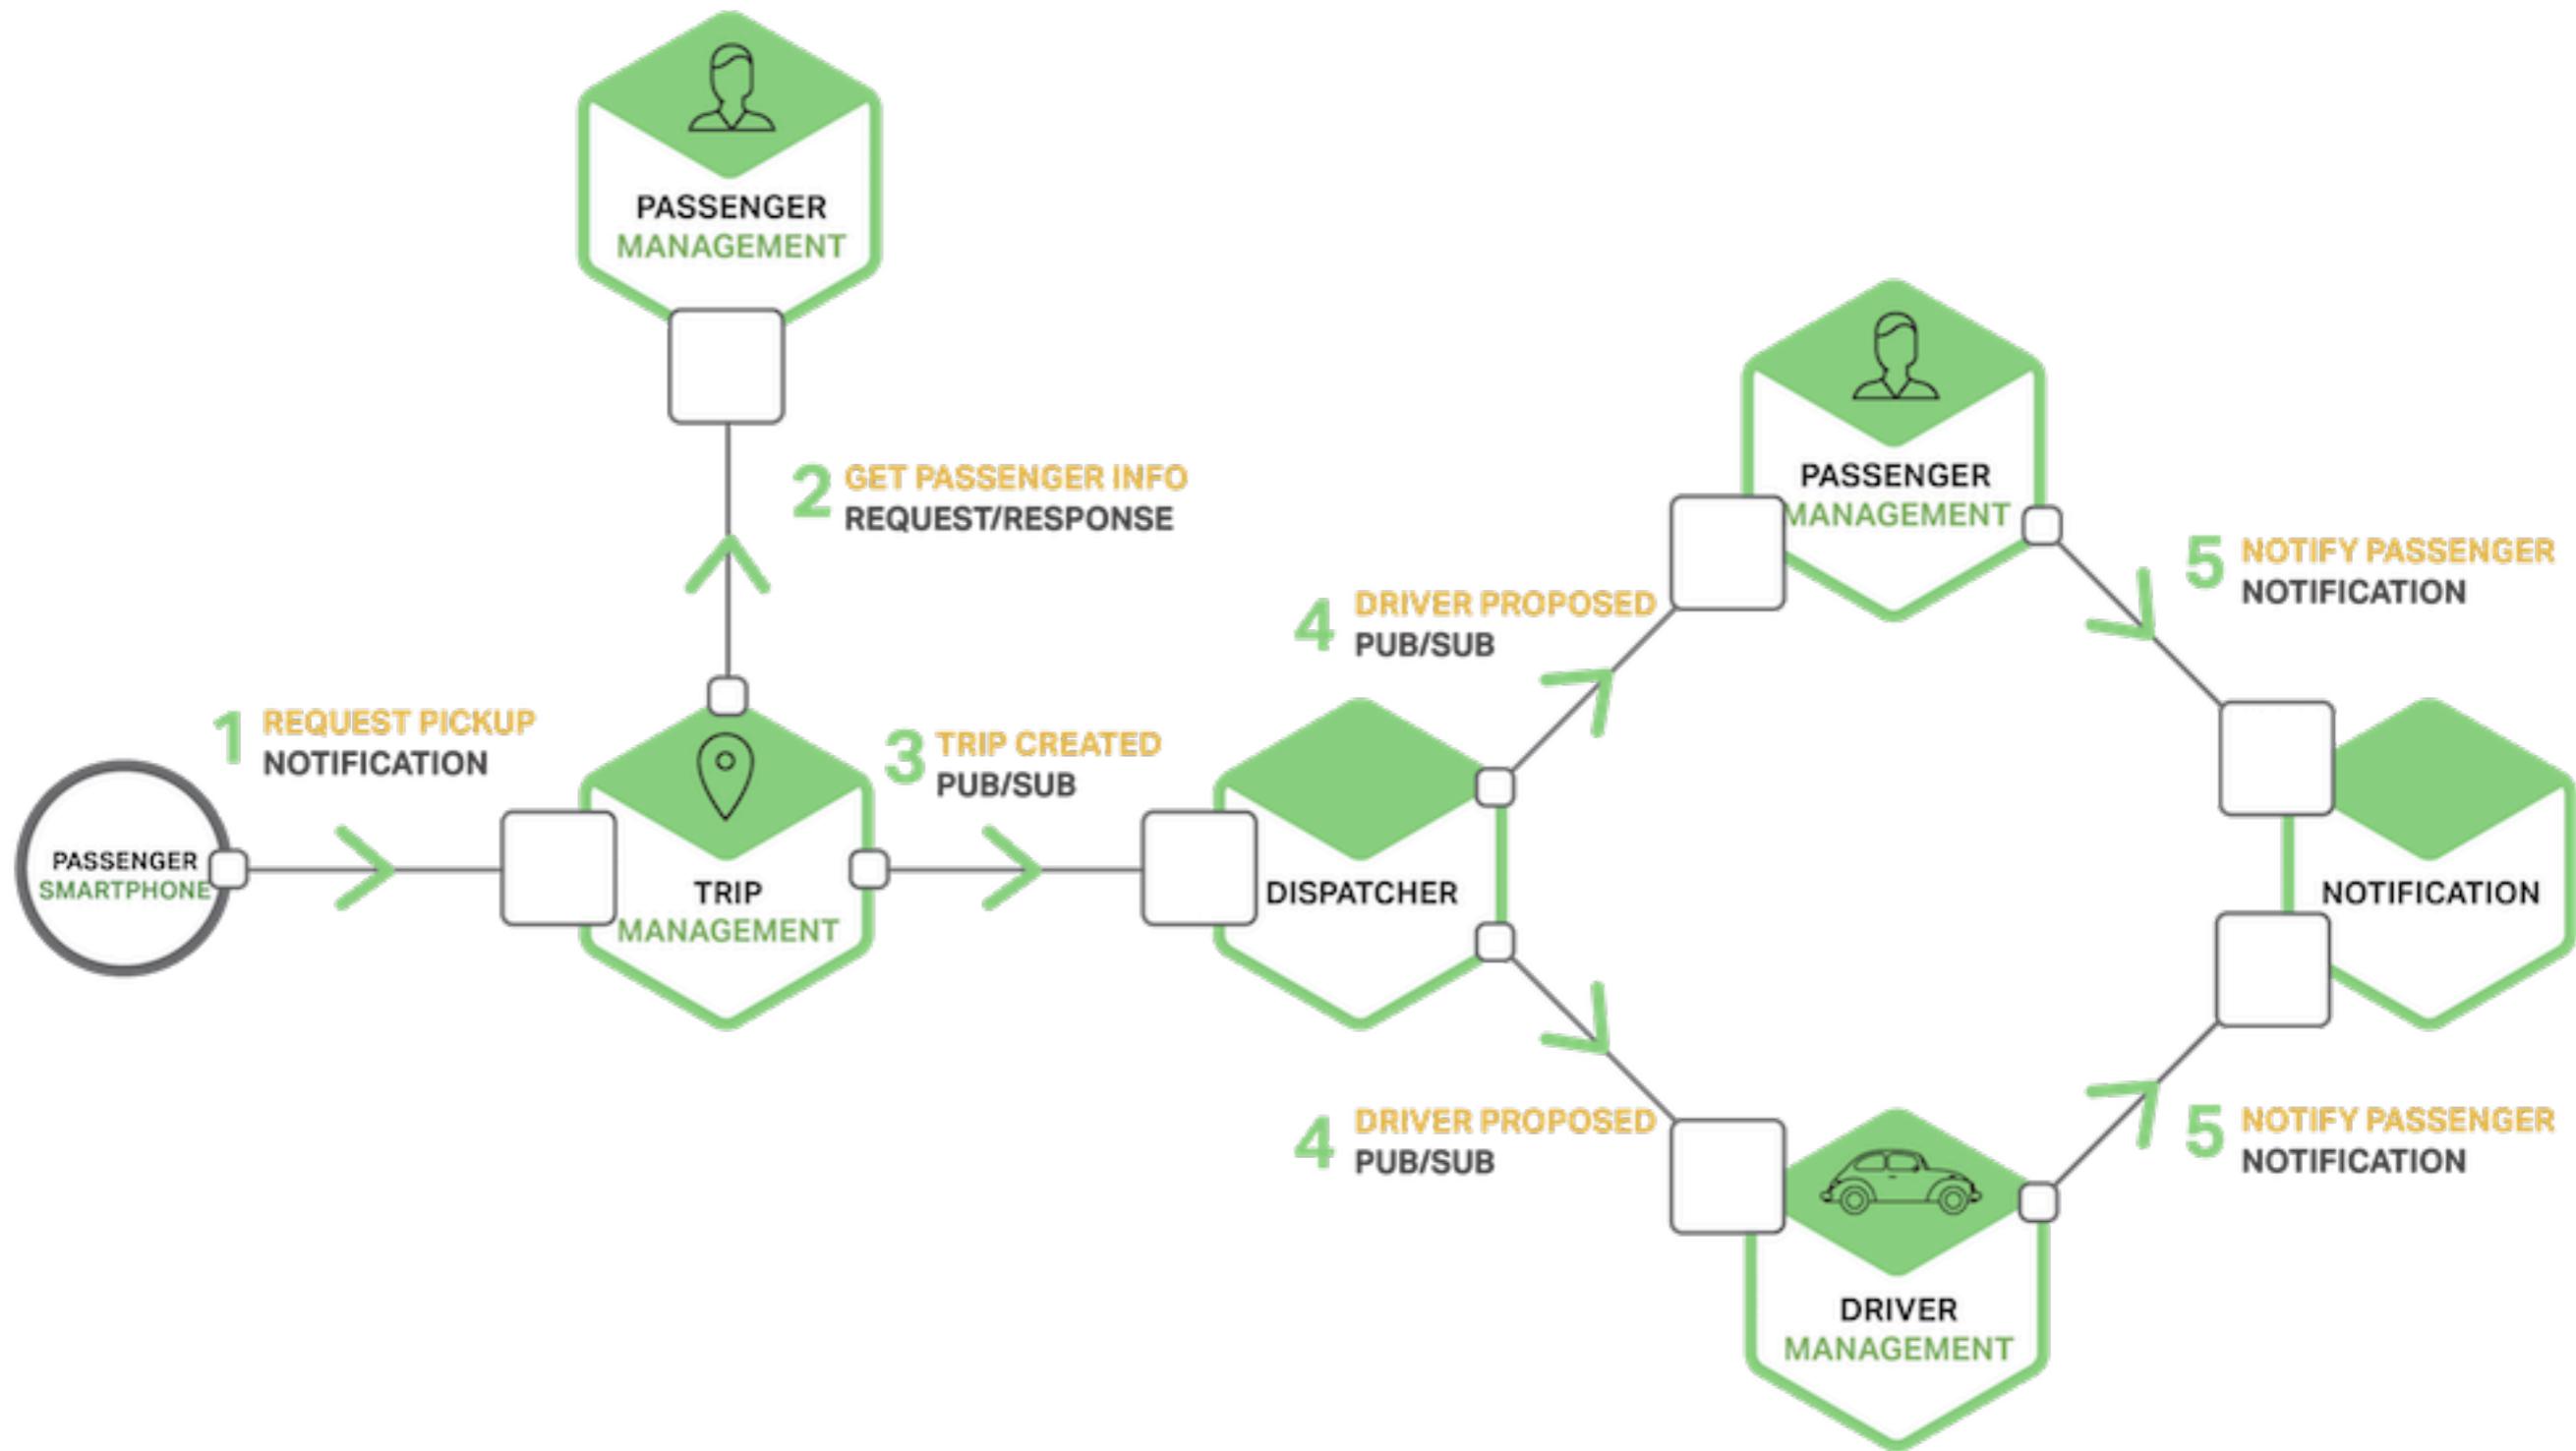


Event-Driven Architecture



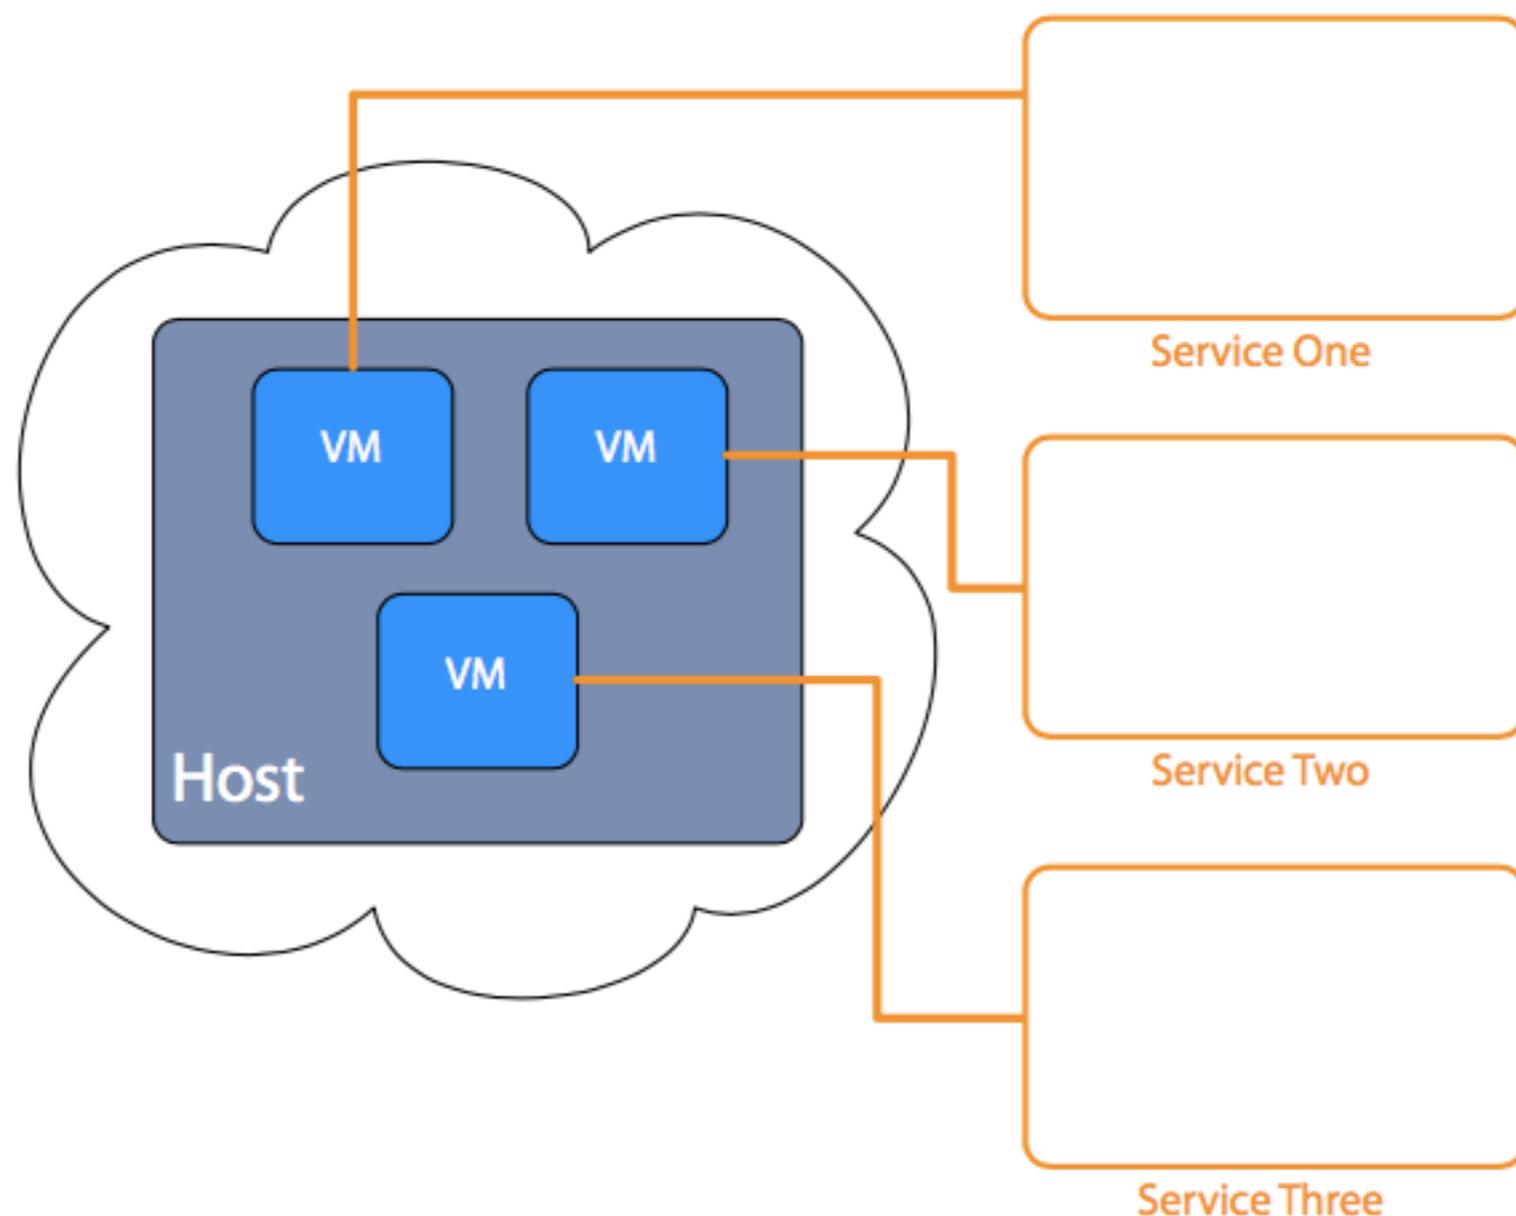
Event-Driven Architecture

One-To Many



Hosting Platforms

Virtualization



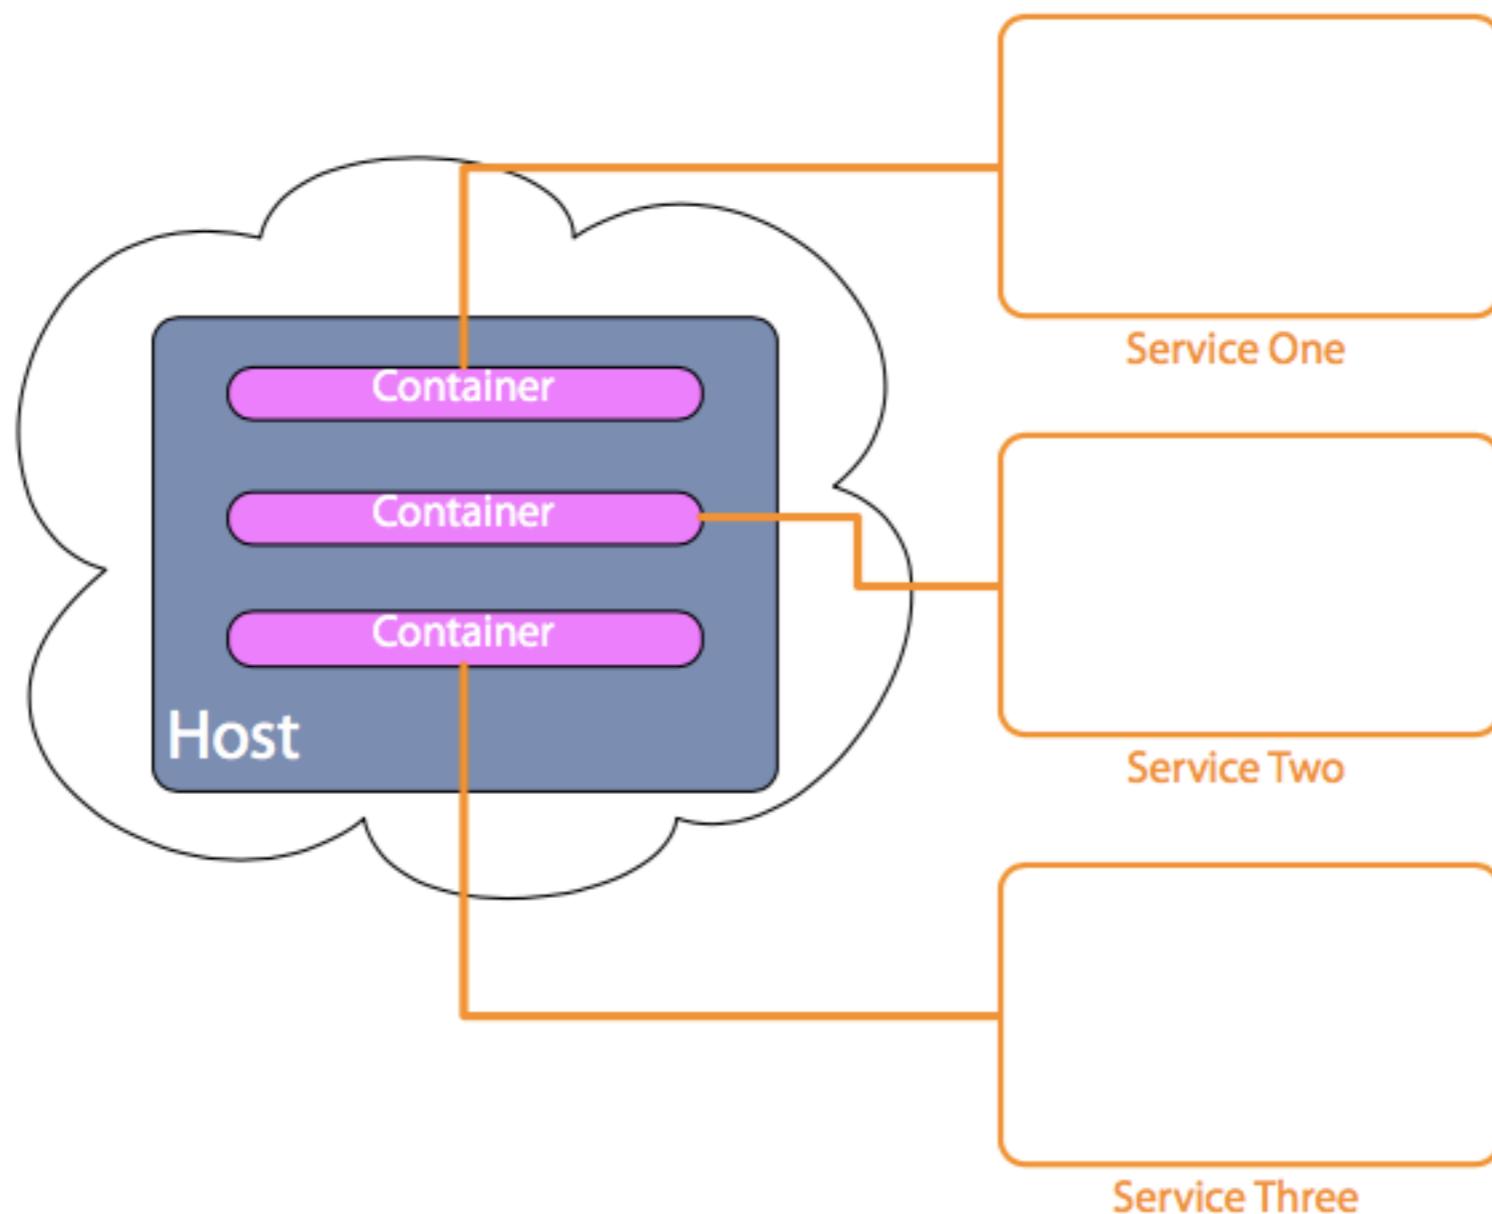
Virtualization

- A virtual machine as a host
- Foundation of cloud platforms
 - Platform as a service (PAAS)
 - Microsoft Azure
 - Amazon web services
 - Your own cloud (for example vSphere)

Virtualization

- Could be more efficient
 - Takes time to setup
 - Takes time to load
 - Take quite a bit of resource
- Unique features
 - Take snapshot
 - Clone instances

Container



Container

- Type of virtualization
- Isolate services from each other
- Single service per container
- Different to a virtual machine
 - Use less resource than VM
 - Faster than VM
 - Quicker to create new instances

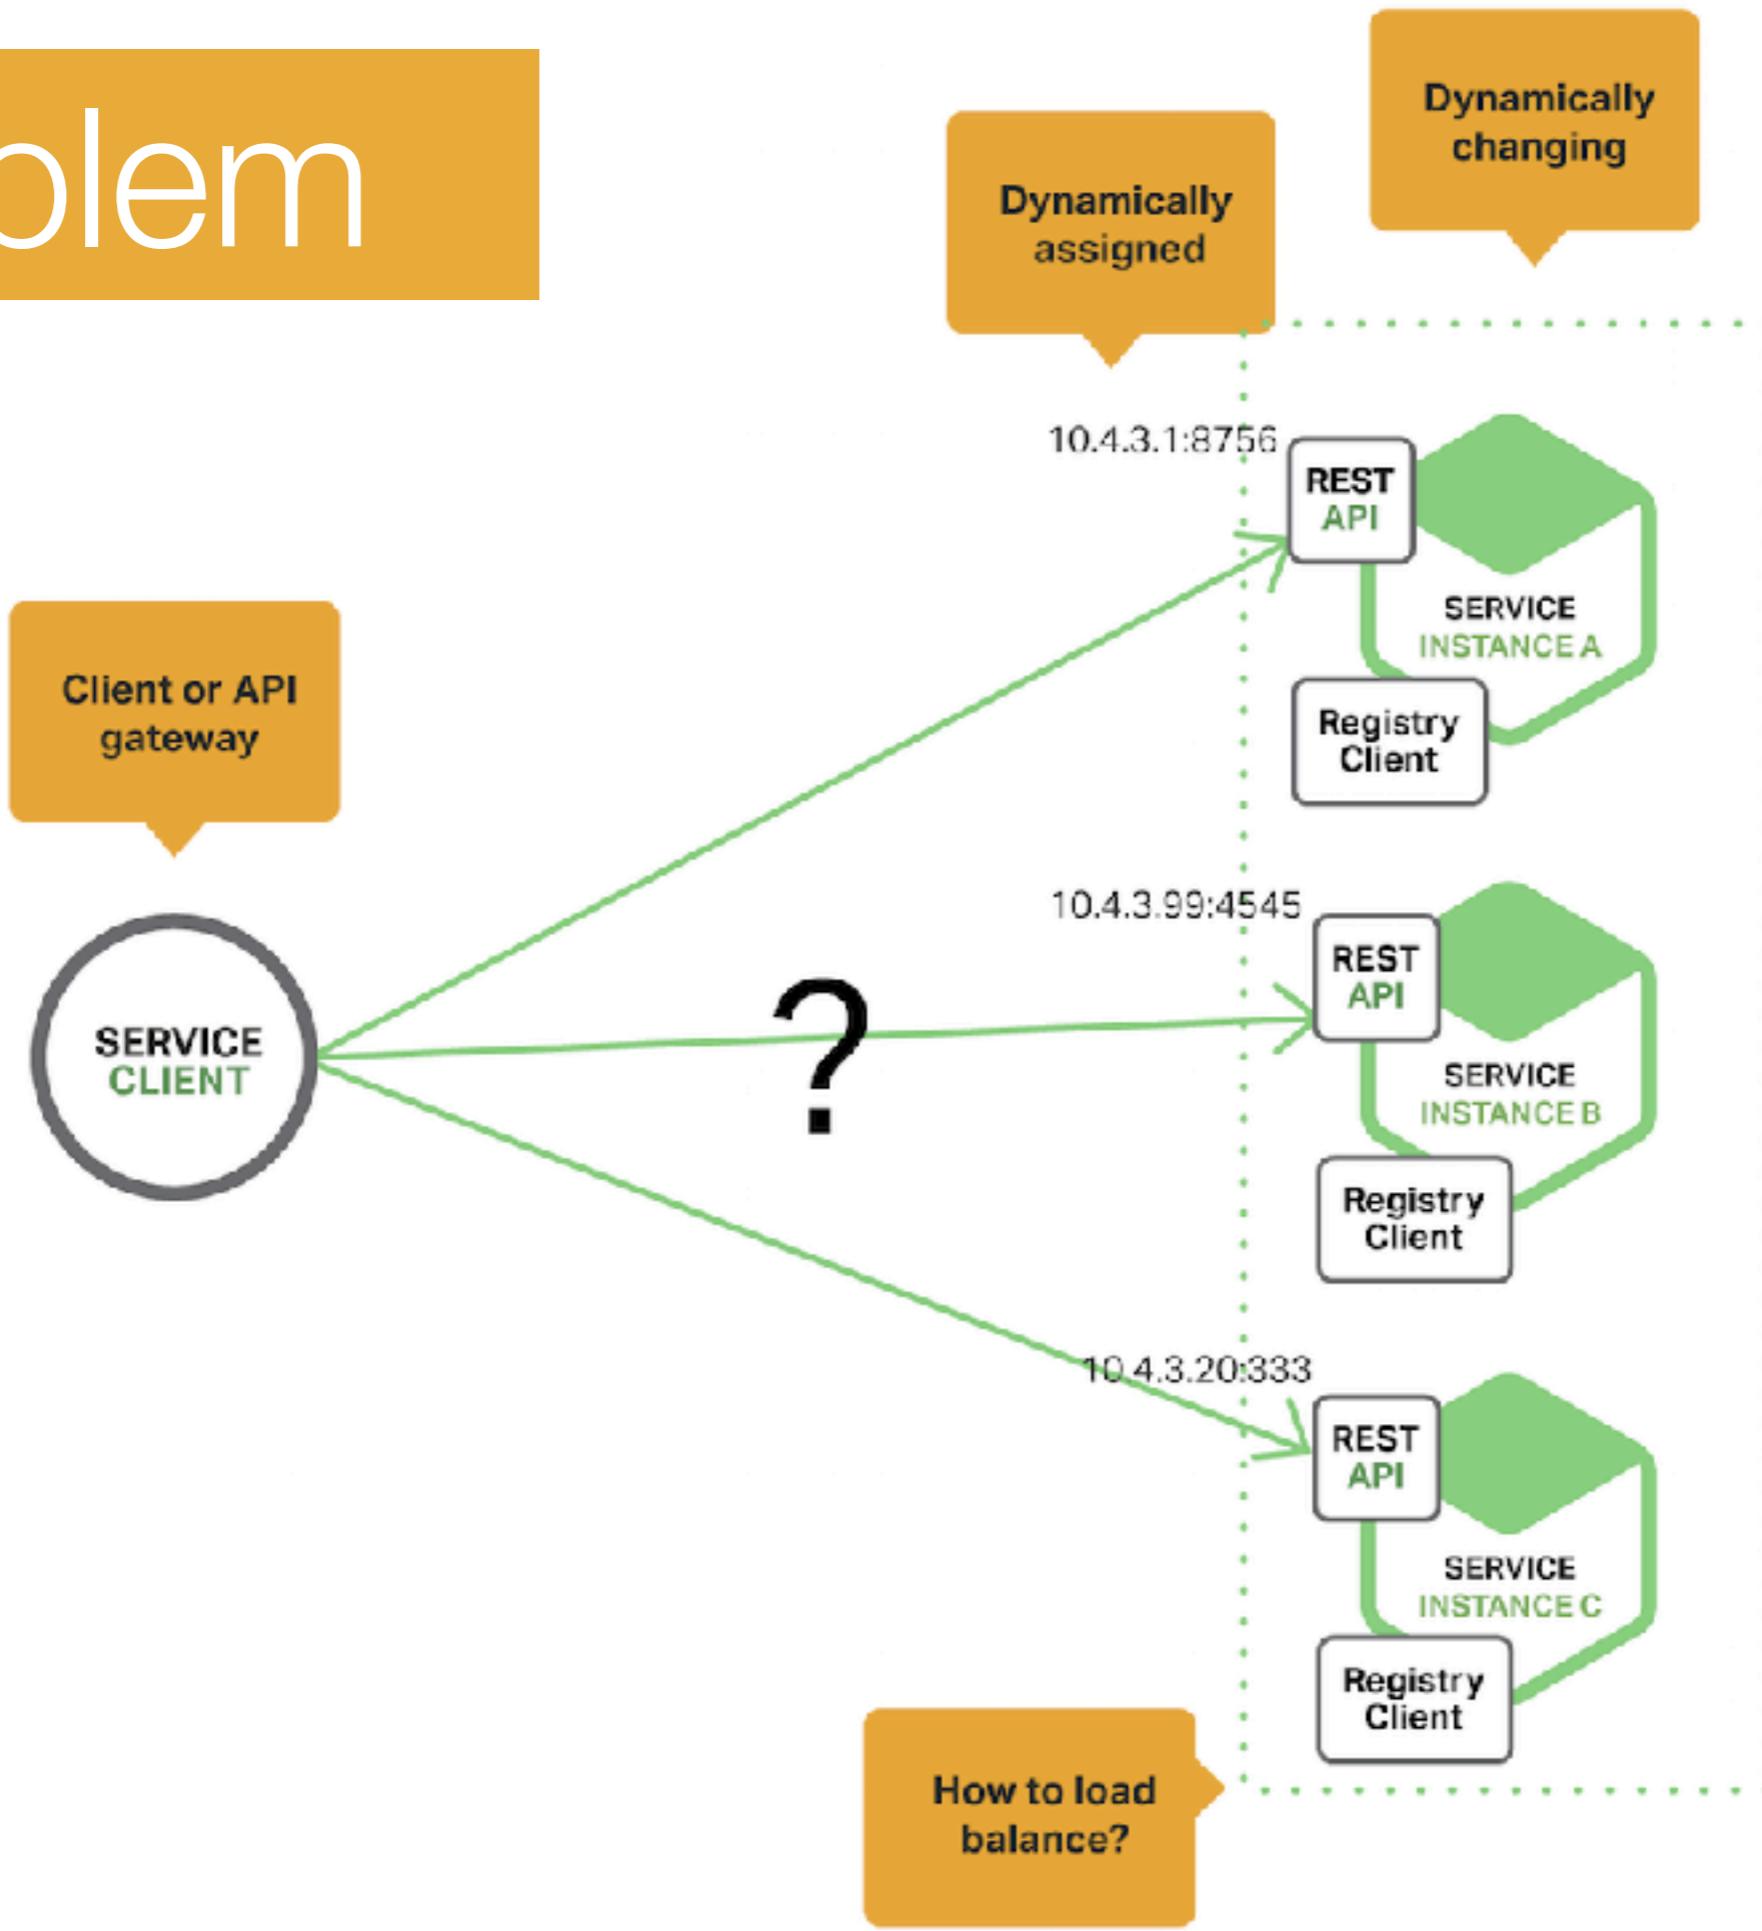
Container

- Future of hosted apps
- Cloud platform support growing
- Mainly Linux based
- Not as established as virtual machines
 - Not standardised
 - Limited features and tooling
 - Infrastructure support in its infancy
 - Complex to setup

Self-Hosting

- Implement your own cloud
 - Virtualization platform
 - Implement containers
- Use of physical machines
 - Single service on a server
 - Multiple services on a server
- Challenges
 - Long-term maintenance
 - Need for technicians
 - Training
 - Need for space
 - Scaling is not as immediate

Problem

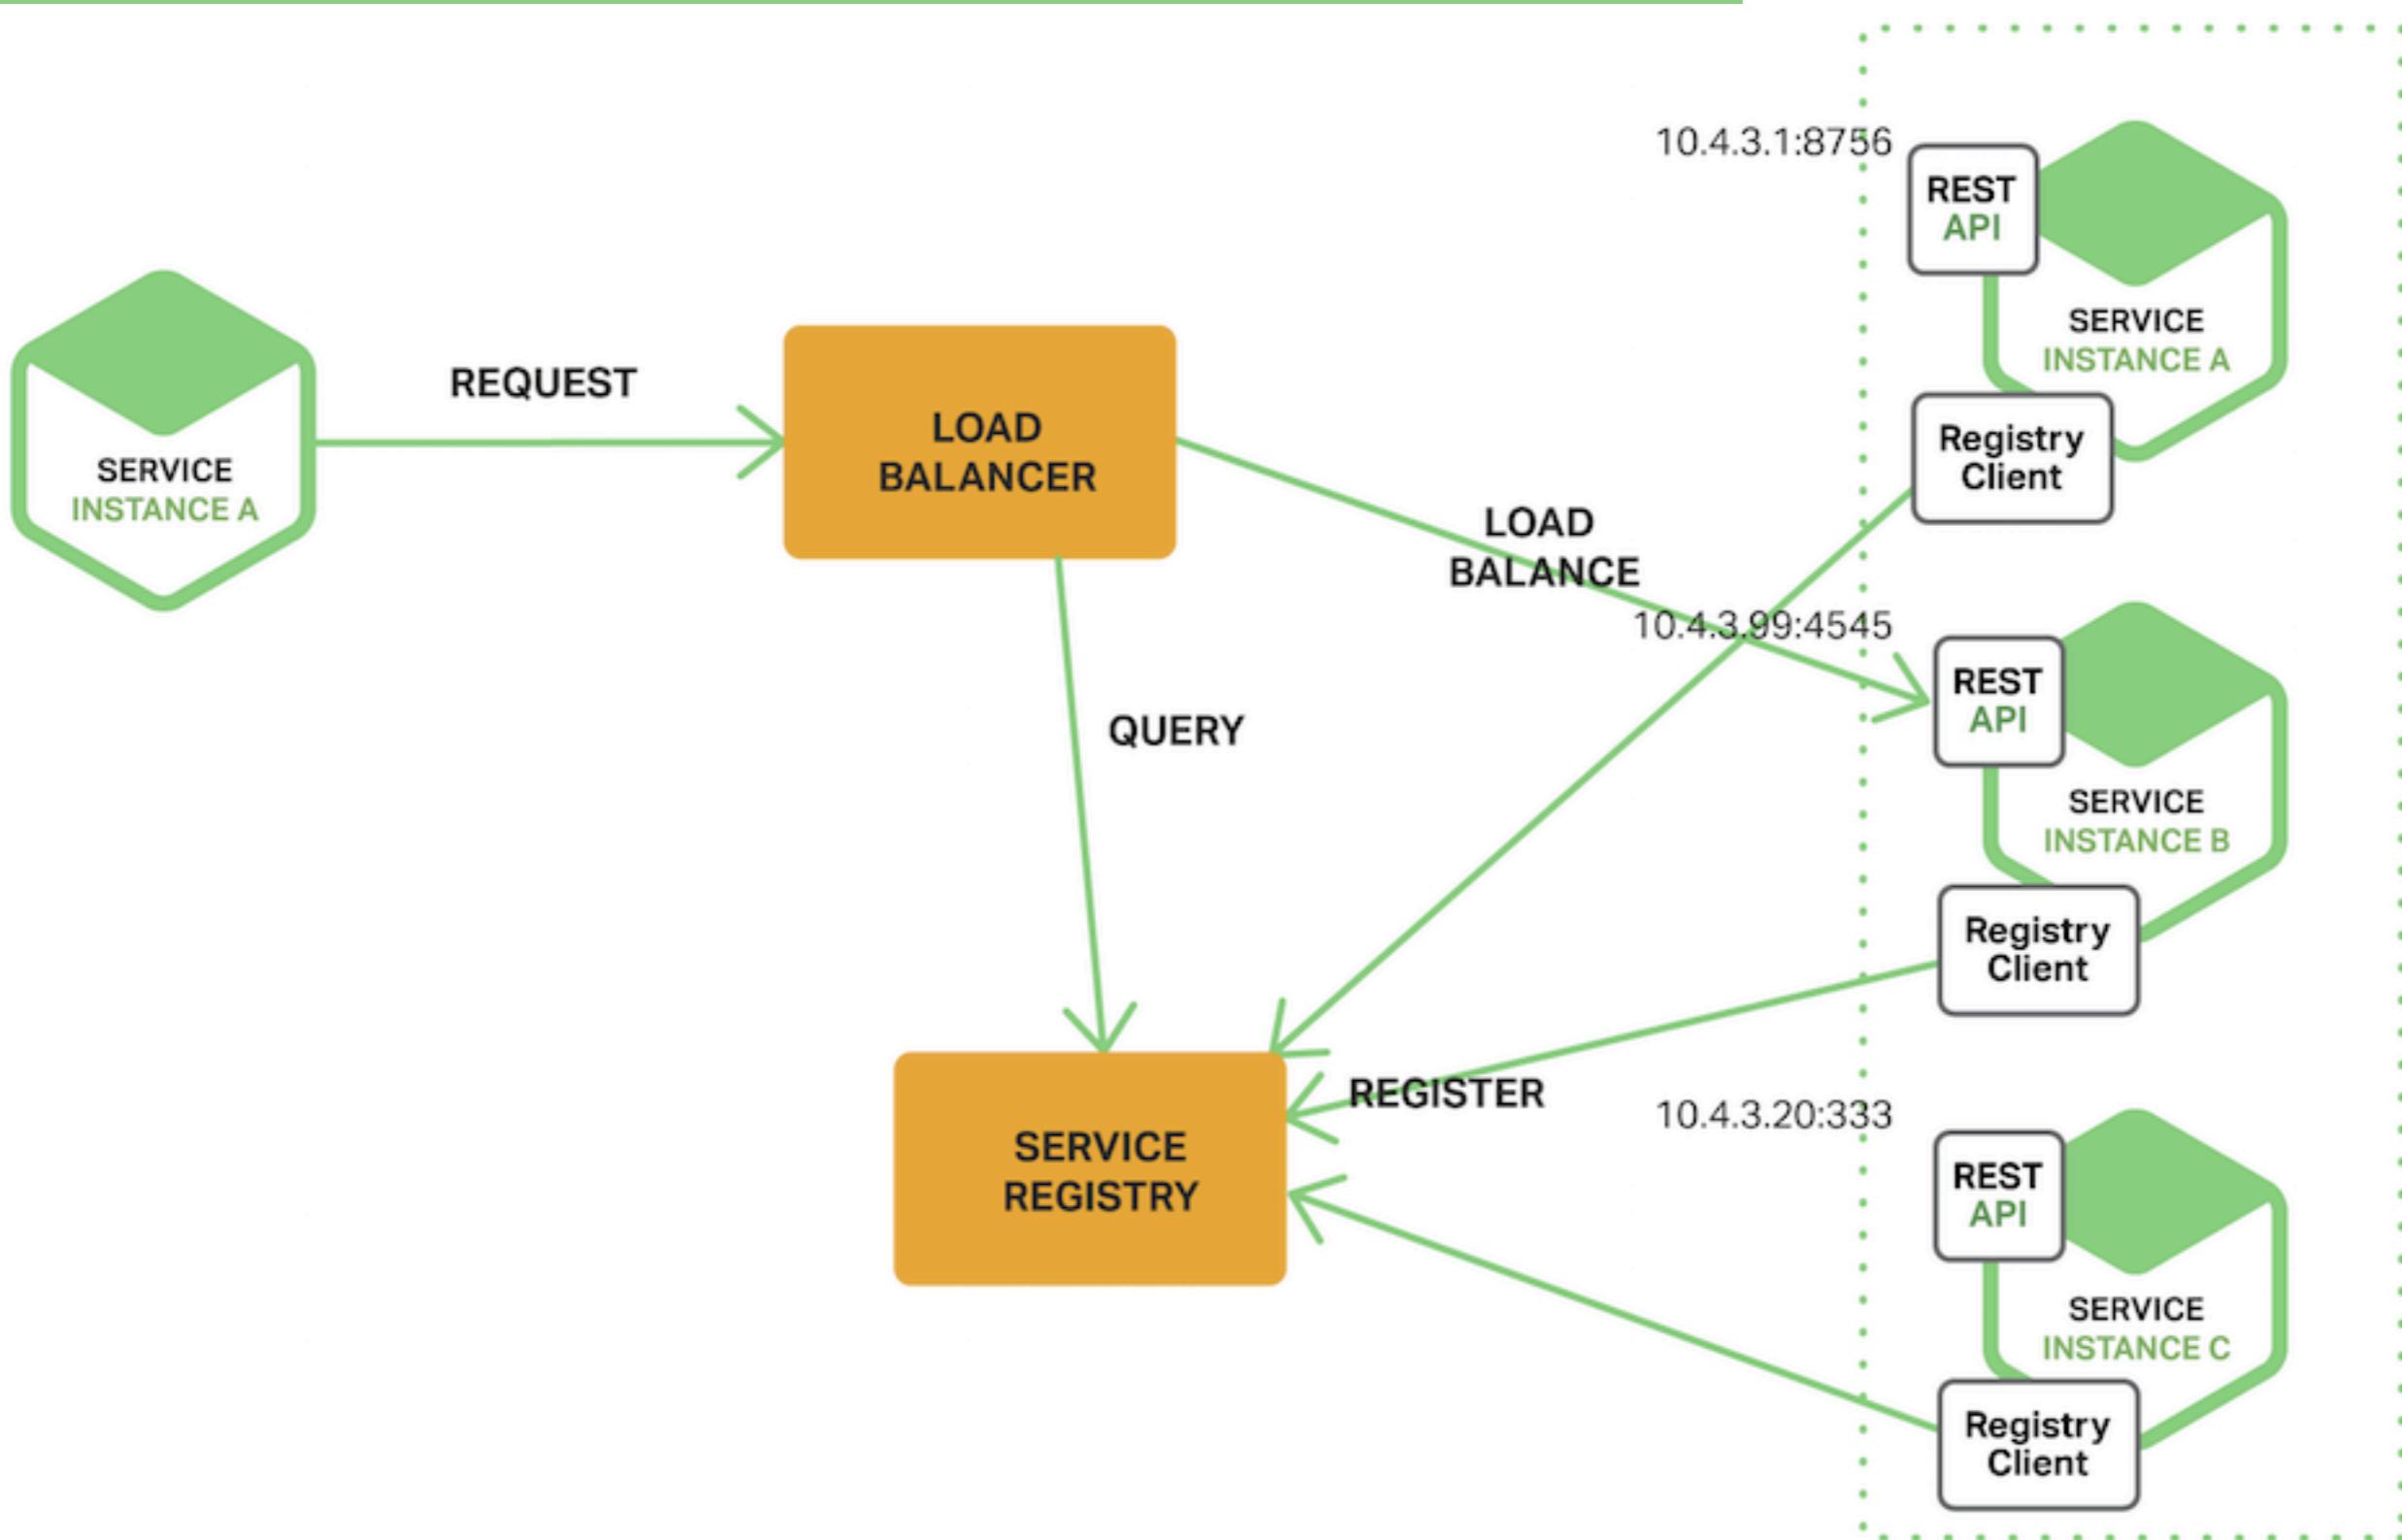


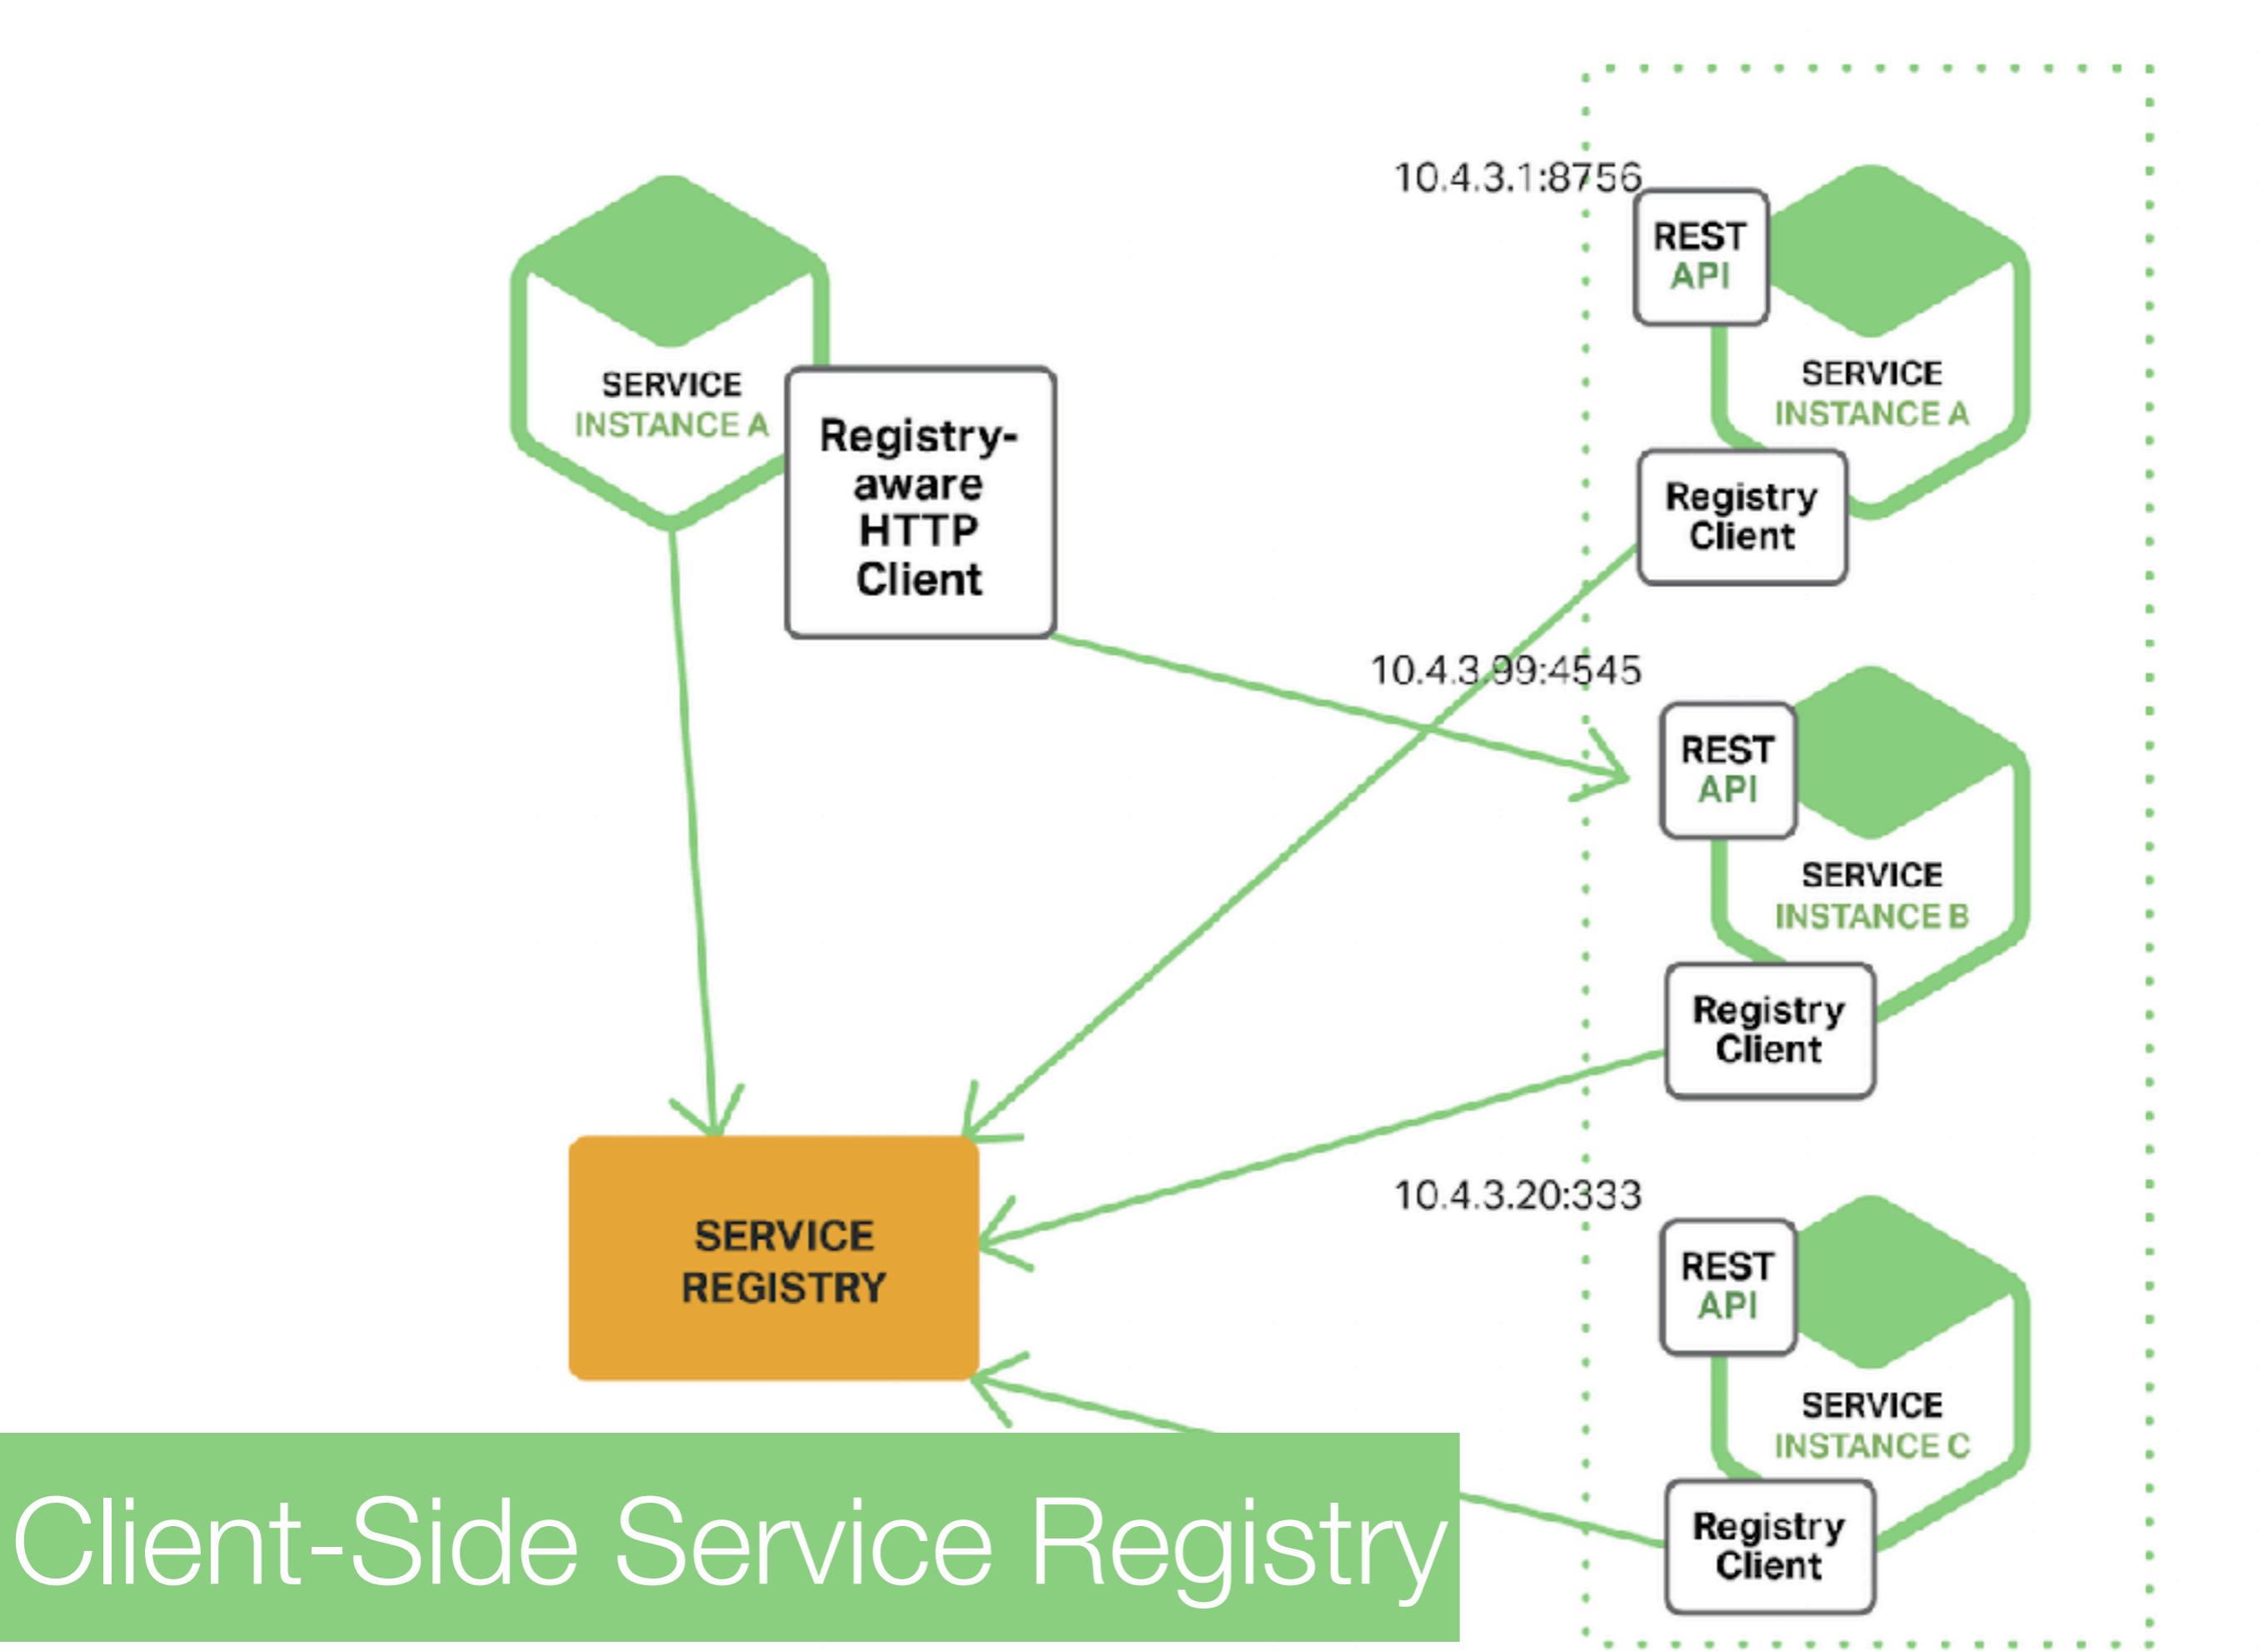
Registration and Discovery

- Service registry database
- Register on startup
- Deregister service on failure
- Cloud platforms make it easy
- Local platform registration options
 - Self registration
 - Third-party registration
- Local platform discovery options
 - Client-side discovery
 - Server-side discovery

Host, Port, Version

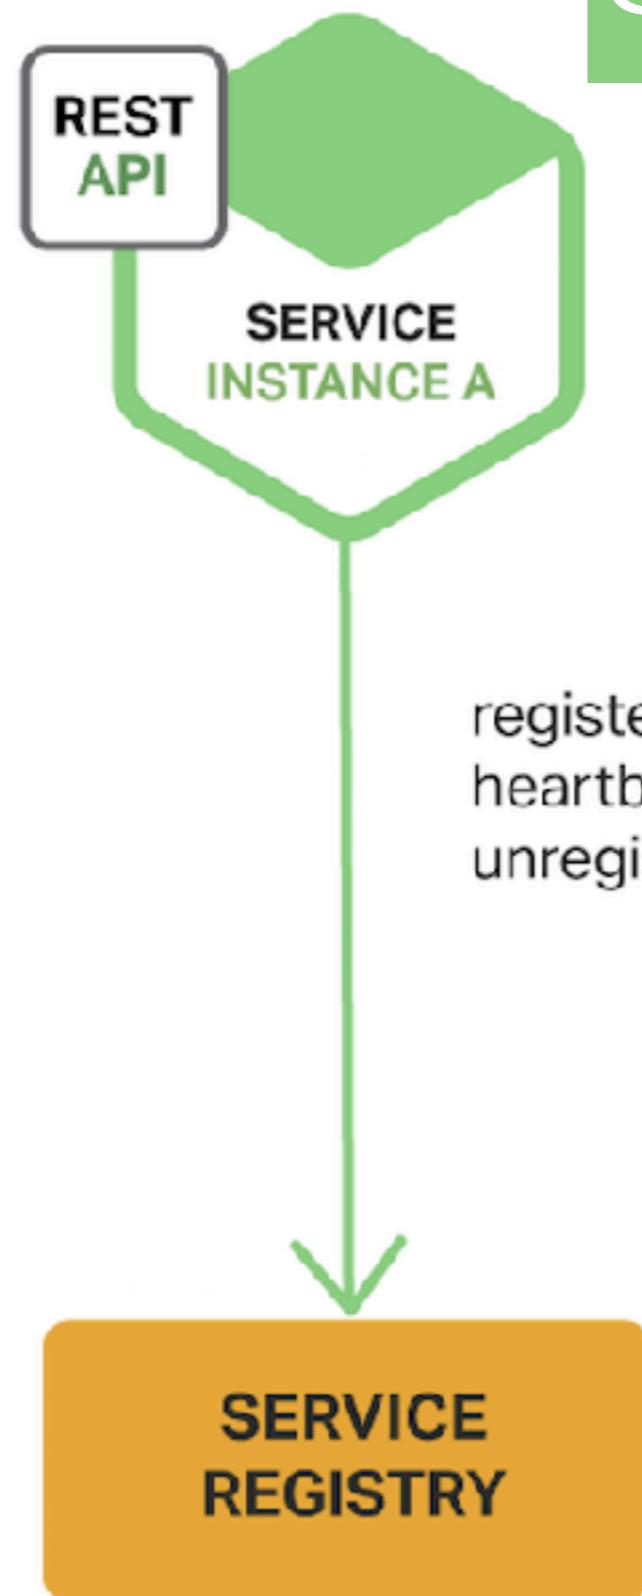
Server-Side Service Registry



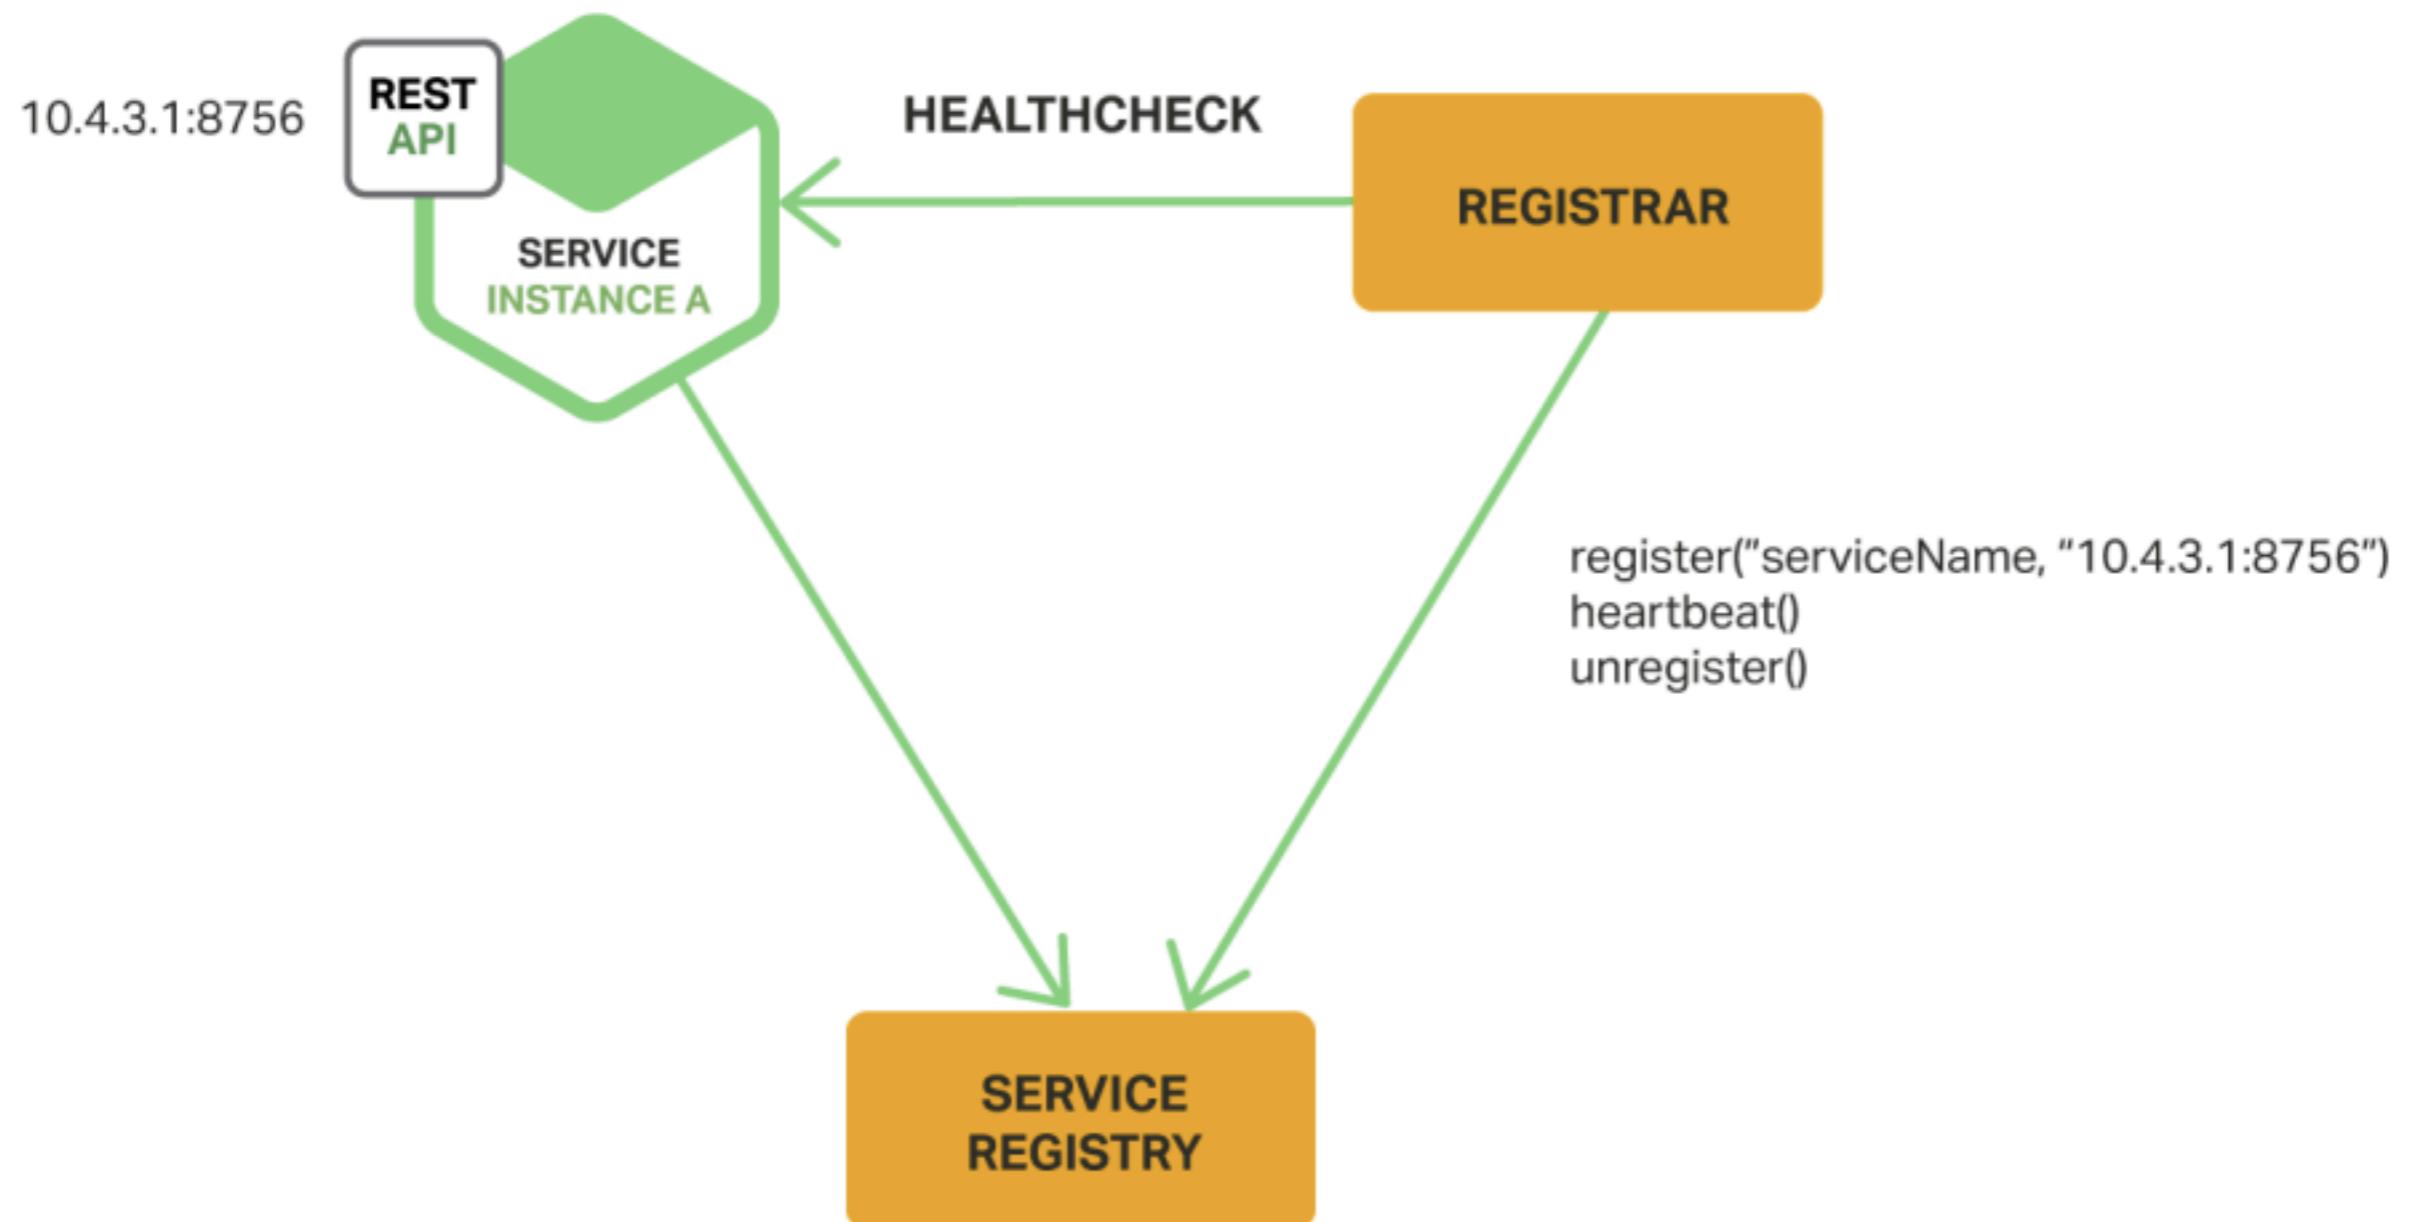


Self-Registration Pattern

10.4.3.1:8756



Third-Party Registration Pattern



Observable

Observable Microservices

Logging Tech

Monitoring Tech

Monitoring Tools

- **cAdvisor**
This container-monitoring tool, developed by Google, natively supports Docker containers. The cAdvisor daemon, which collects, aggregates, processes, and exports information about containers, is available in a graphical version
- **Heapster**
Use Heapster to monitor Kubernetes container clusters. It collects and interprets various signals such as compute resource usage and lifecycle events, supporting multiple sources of data
- **Prometheus**
Prometheus is an open-source service-monitoring system and time-series database
- **Grafana**
A metrics dashboard, Grafana integrates with many data sources, including Graphite, CloudWatch, Prometheus, Elasticsearch, and InfluxDB

Features

- Metrics across servers
- Automatic or minimal configuration
- Client libraries to send metrics
- Test transactions support
- Alerting

Microservice Monitoring

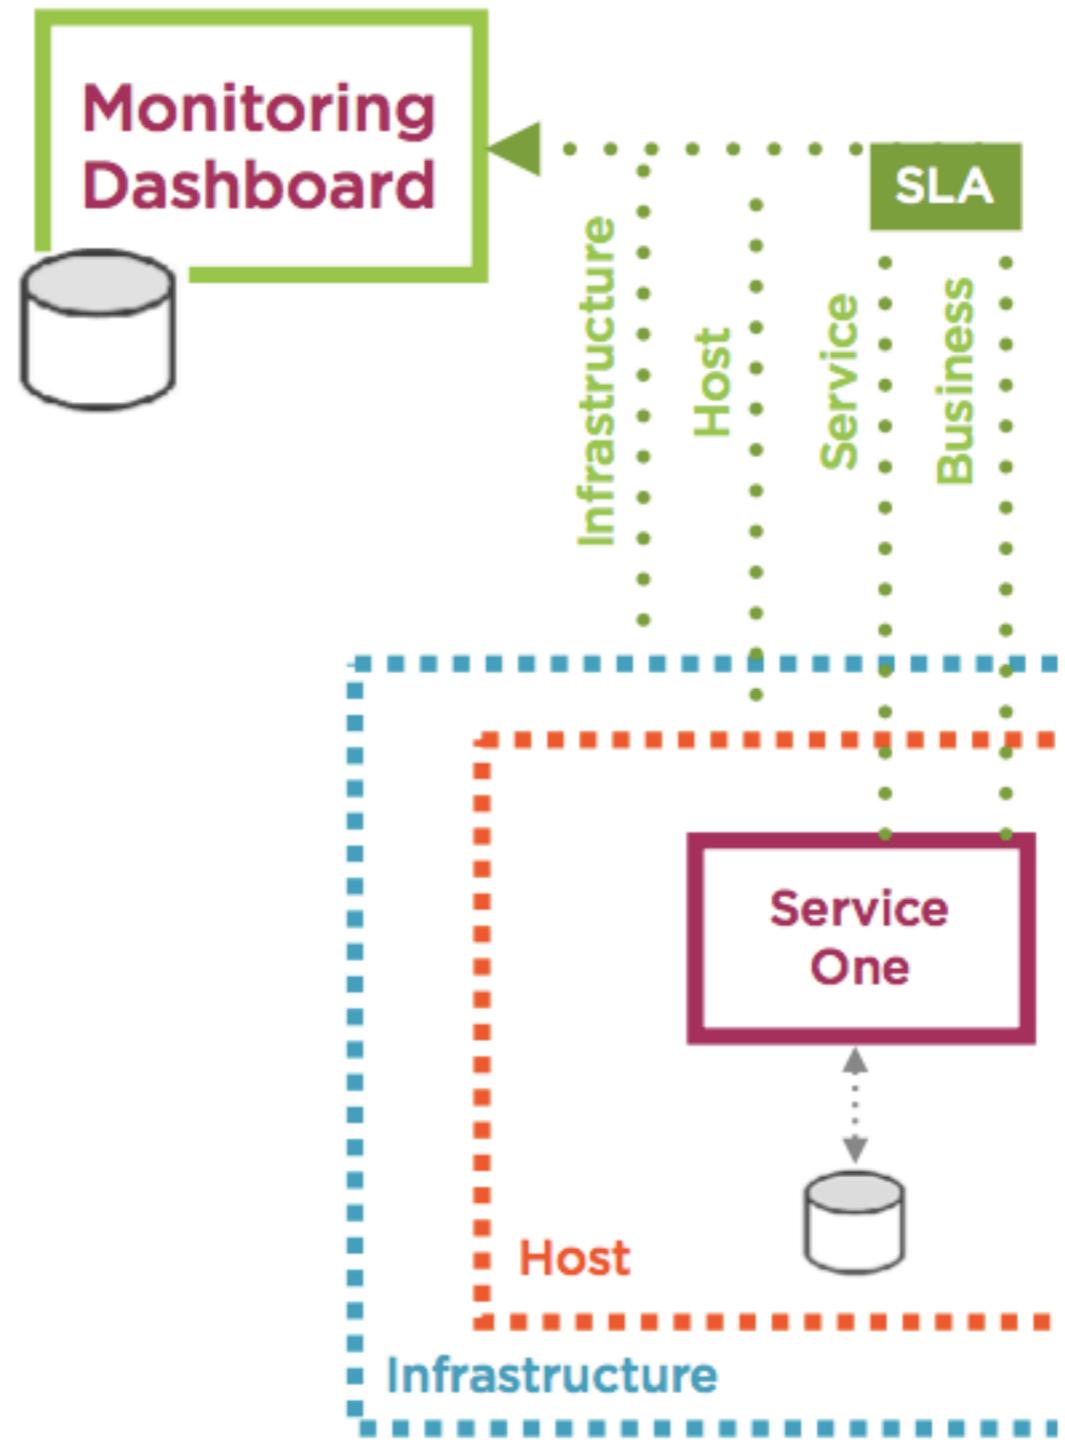
- Why the need to monitor
 - Microservices complexity
 - High availability is a requirement
- Causes of microservices outages
 - Deployment issues
 - Delayed and long repair
 - Downstream dependency issues
- Monitoring is the **proactive solution**
 - Know the state of all components
 - Know the state at every stage
 - Collect all statistics using metrics

Key Metrics

- Host and infrastructure metrics
 - State of hosts and infrastructure
 - State affects microservices
- Business metrics
 - Hidden issues and failures
 - Sudden changes in behavior
- Monitor microservice metrics
 - Availability
 - Response rate
 - Success or failures of end Business metrics
 - Hidden issues and failures
 - Sudden changes in behavior points
 - Errors, exceptions and timeouts
 - Health of dependencies
 - Metrics at different stages

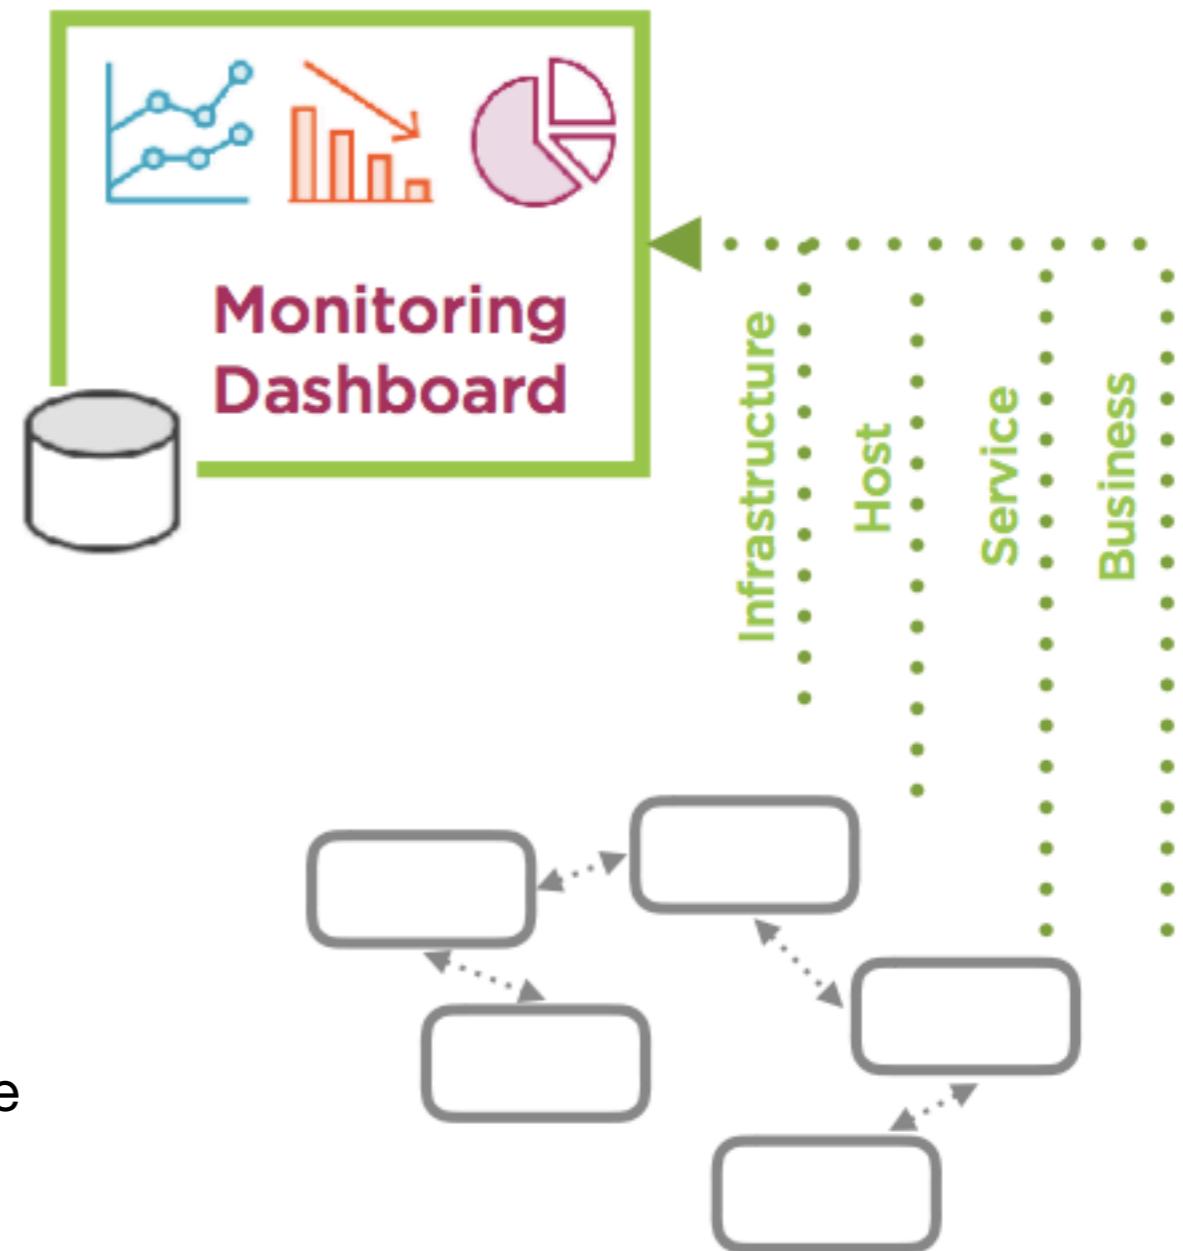
Monitor SLA Metrics

- Monitor SLAs at all levels
 - Service level
 - Service endpoint level
- Examples of SLA metrics
 - Throughput and uptime
 - Concurrent clients
 - Average response times
- Impact of dependencies on SLAs
 - Monitor dependencies
- Use SLA monitoring to alert
 - Actionable alerts to fix SLA issues
- Communicate SLA issues early



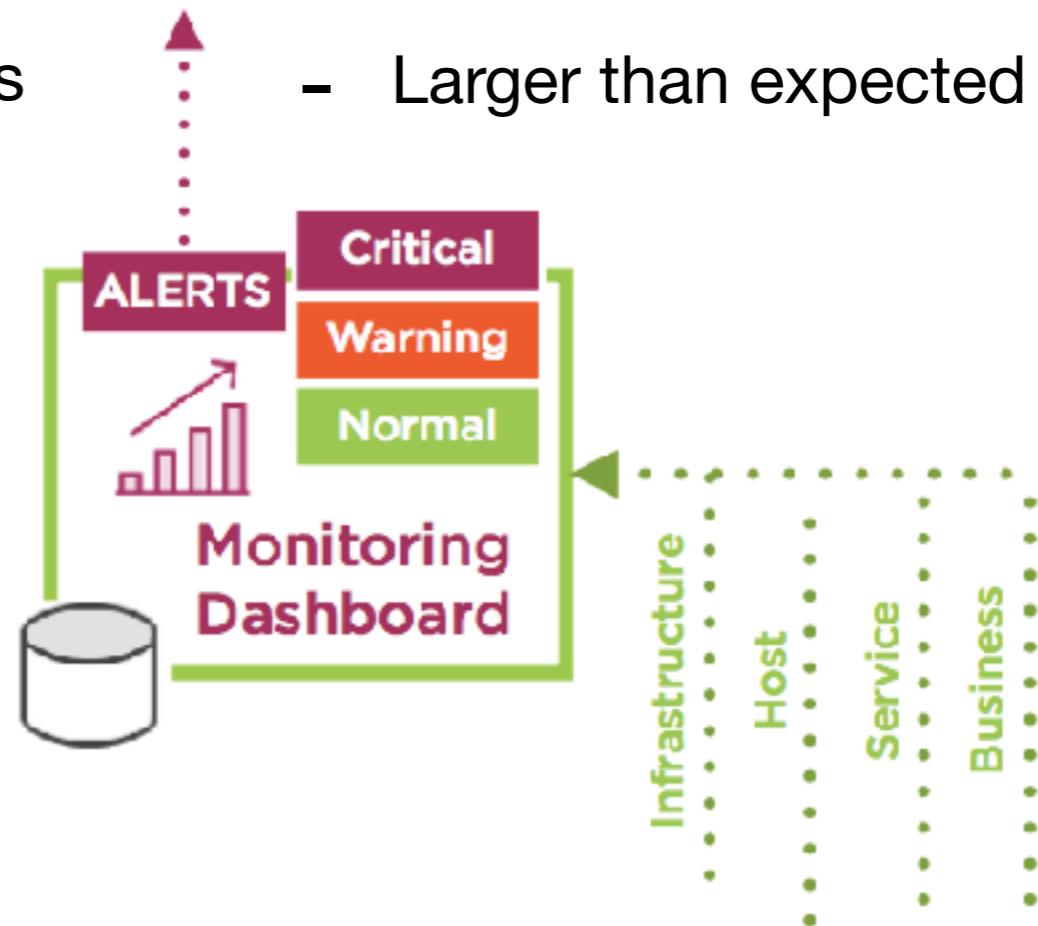
Monitoring Dashboard

- Display and collect metrics
- Real-time
- Centralized, accessible and standardized
- Simple, visual and minimal information
Graphs showing metrics over time
- Highlights effectiveness of monitoring
- What to display
 - Metrics and their alerts
 - Metrics for each environment and phase
 - Correlate metrics to events
- Use to help define thresholds for alerts

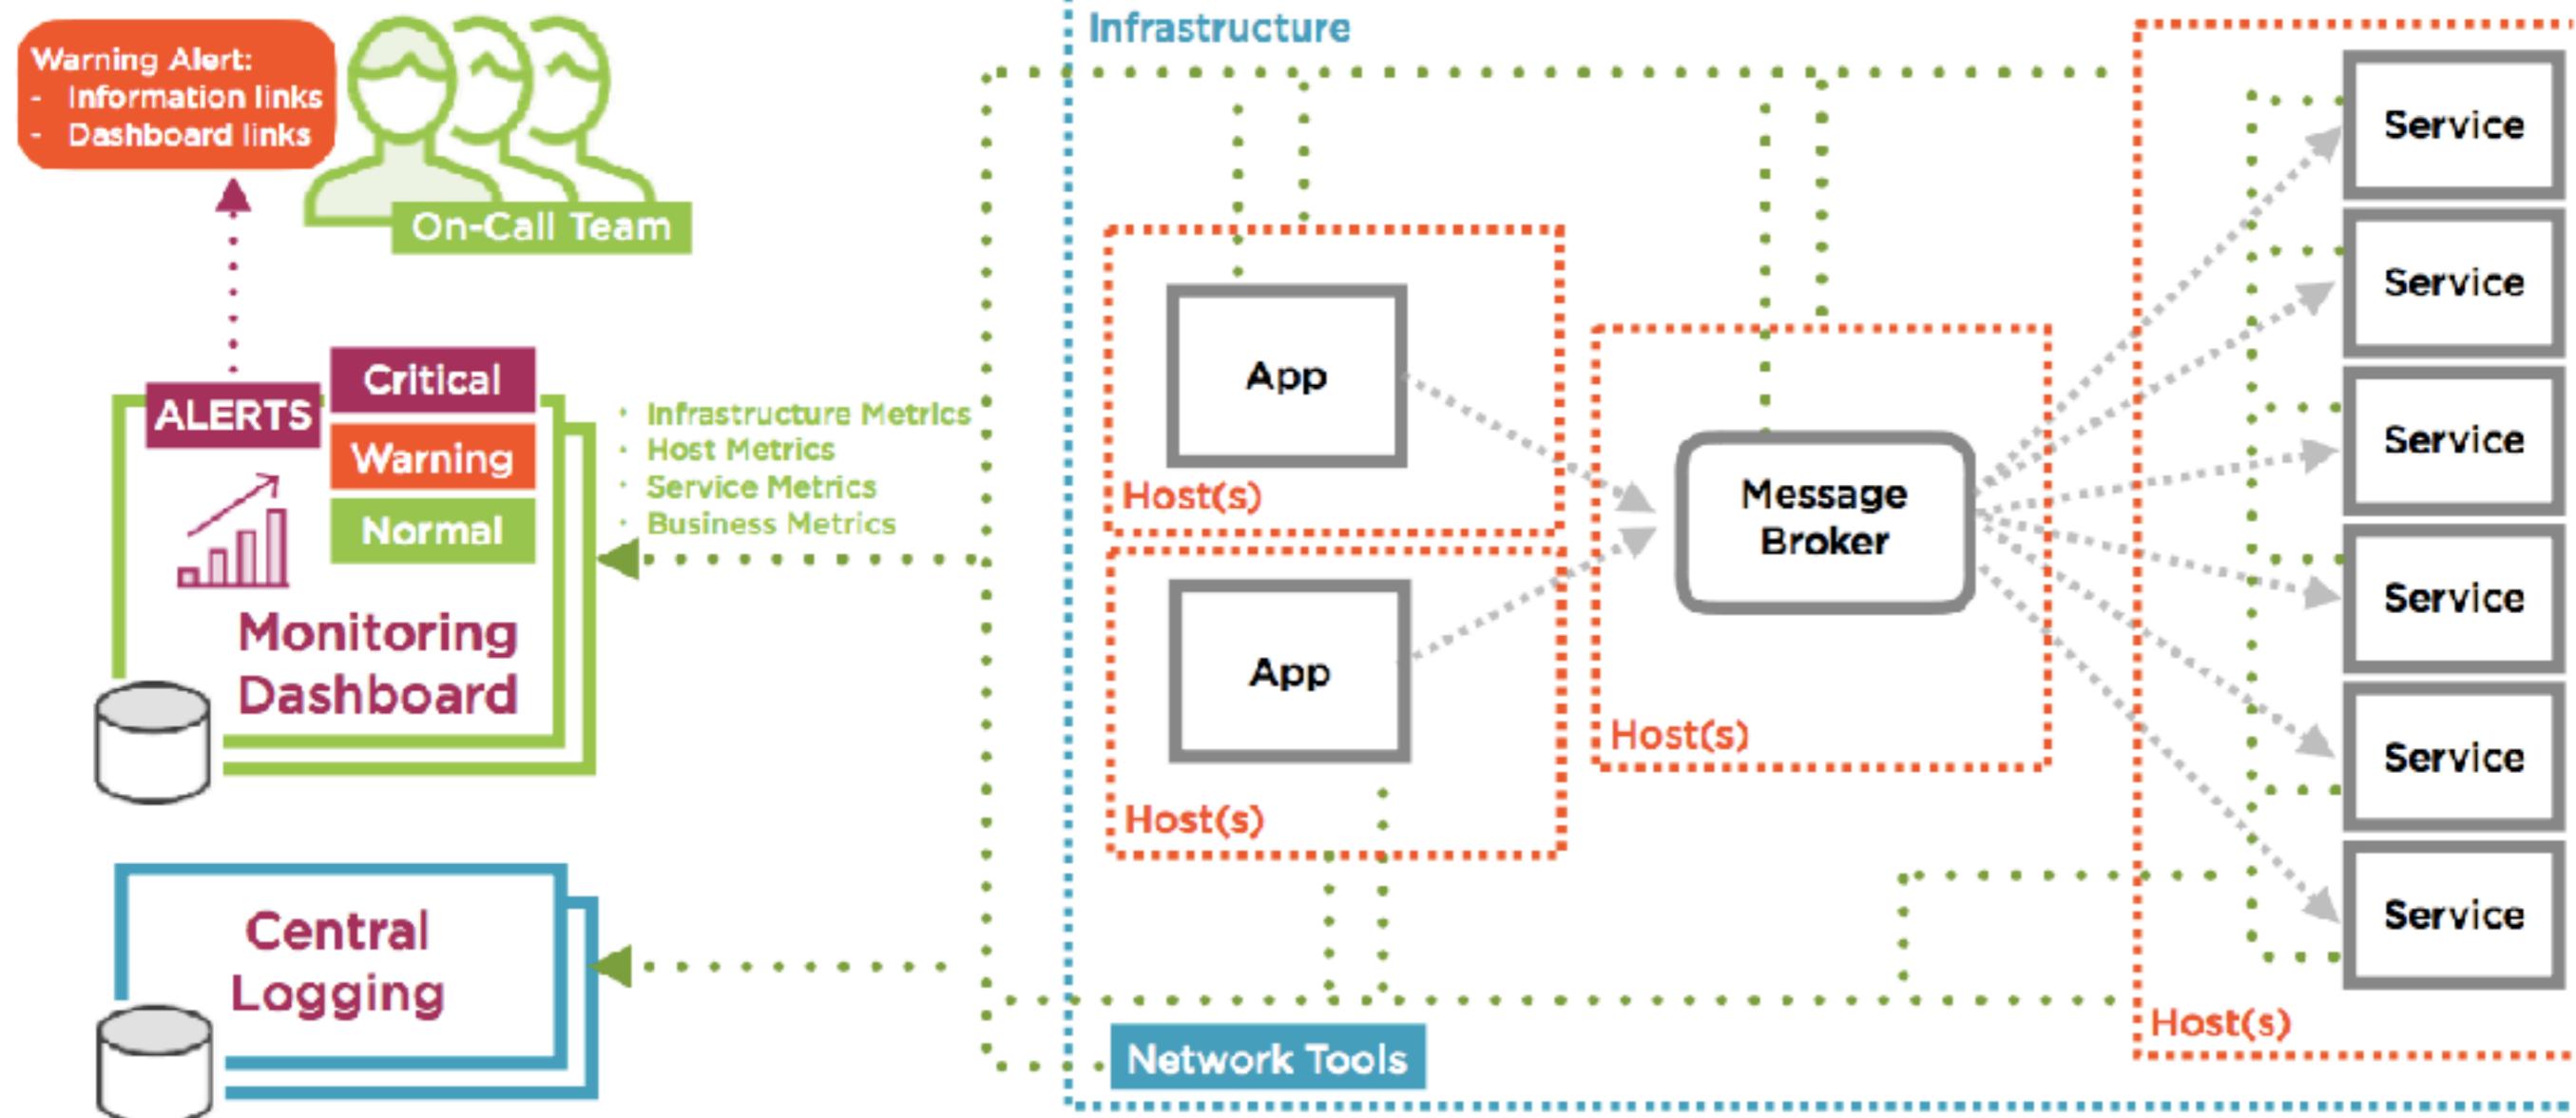


Defining Thresholds for Alerts

- Ranges required for effective alerting
- Types of thresholds
 - Normal state
 - Warning state
 - Critical state
- Warning alerts prevent outages
- Effective levels for thresholds
 - Use historical data
 - Performance testing and load testing data
 - Expected traffic for normal thresholds
 - Larger than expected traffic



MicroServices Monitoring Patterns



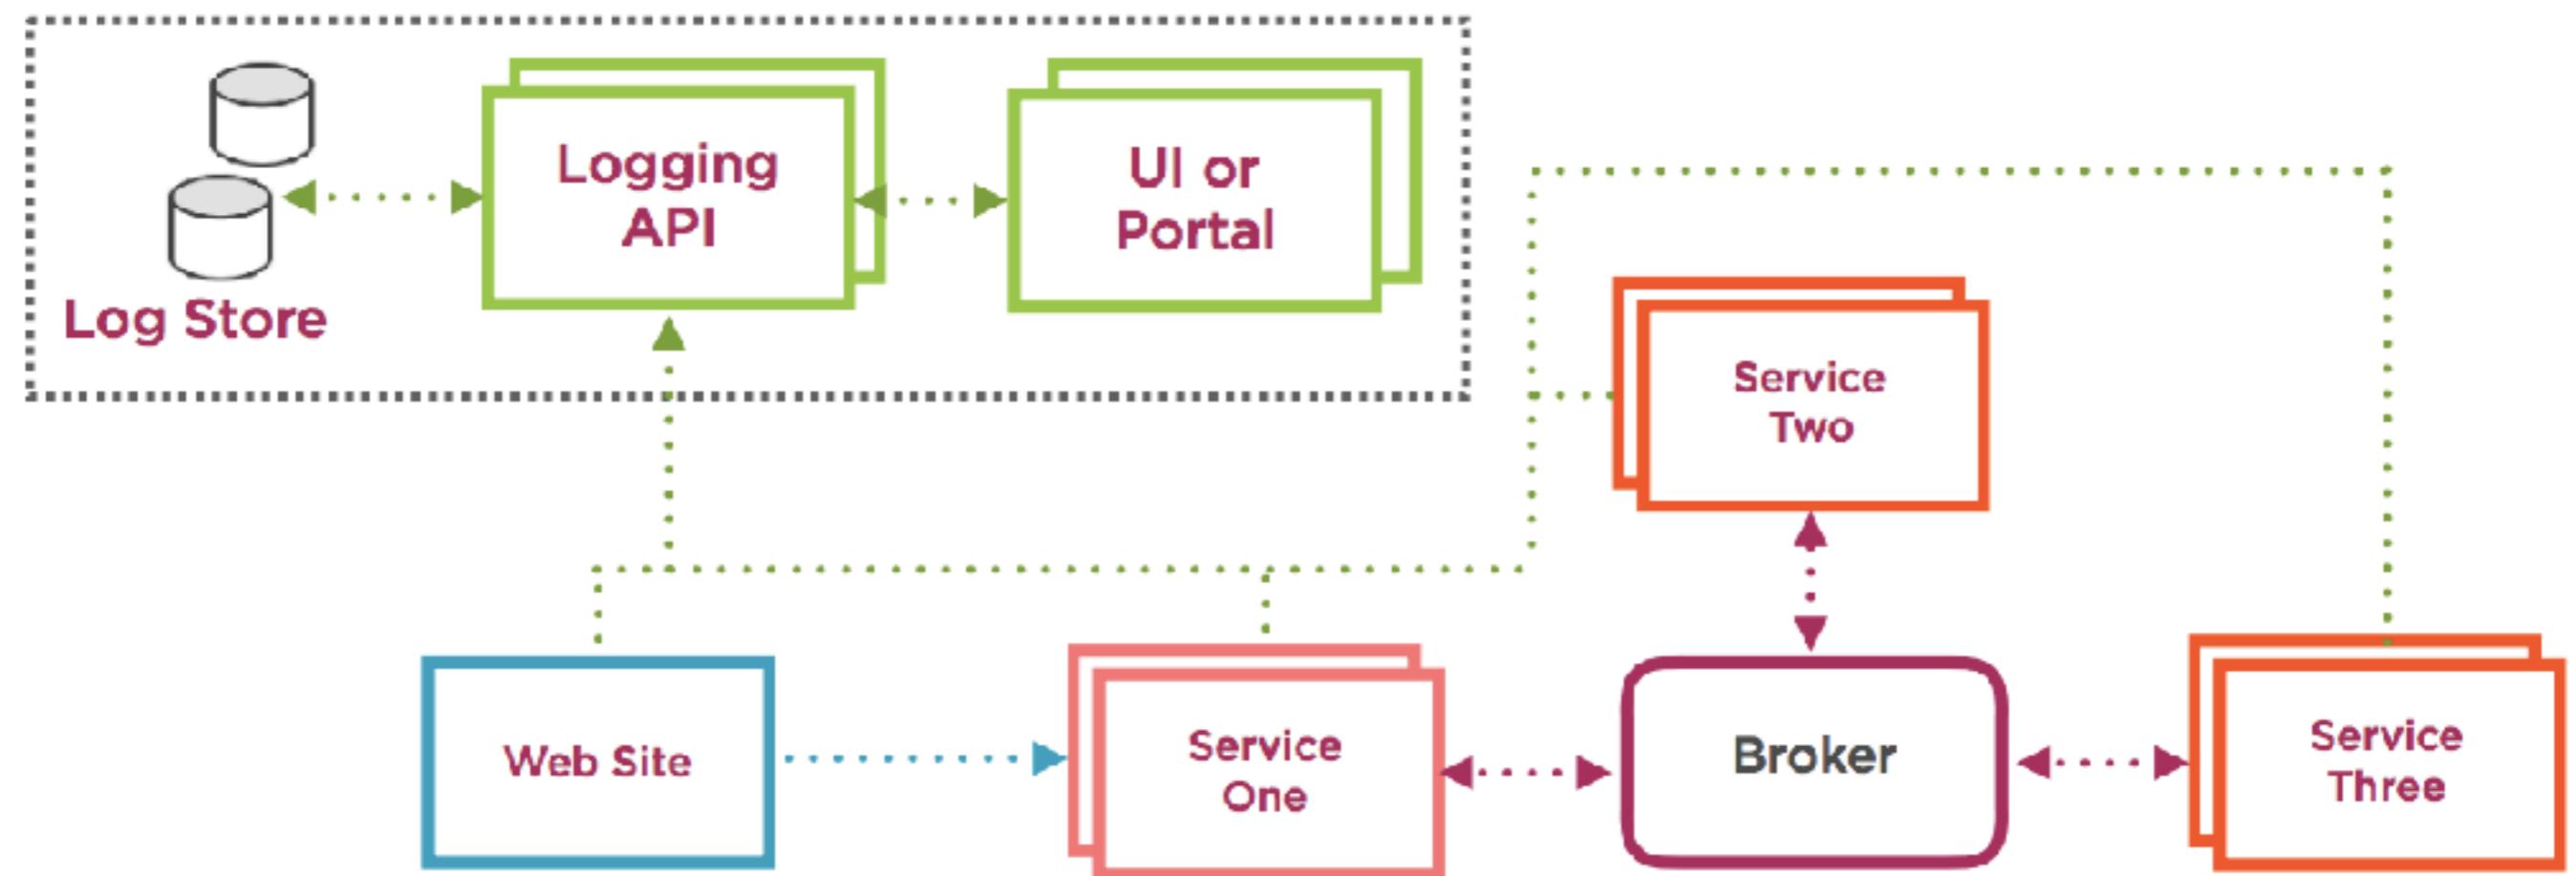
Logging Tools

- Logstash
- Kibana
- Splunk
- Graphite

Features

- Structured logging
- Logging across servers
- Automatic or minimal configuration
- Correlation\Context ID for transactions

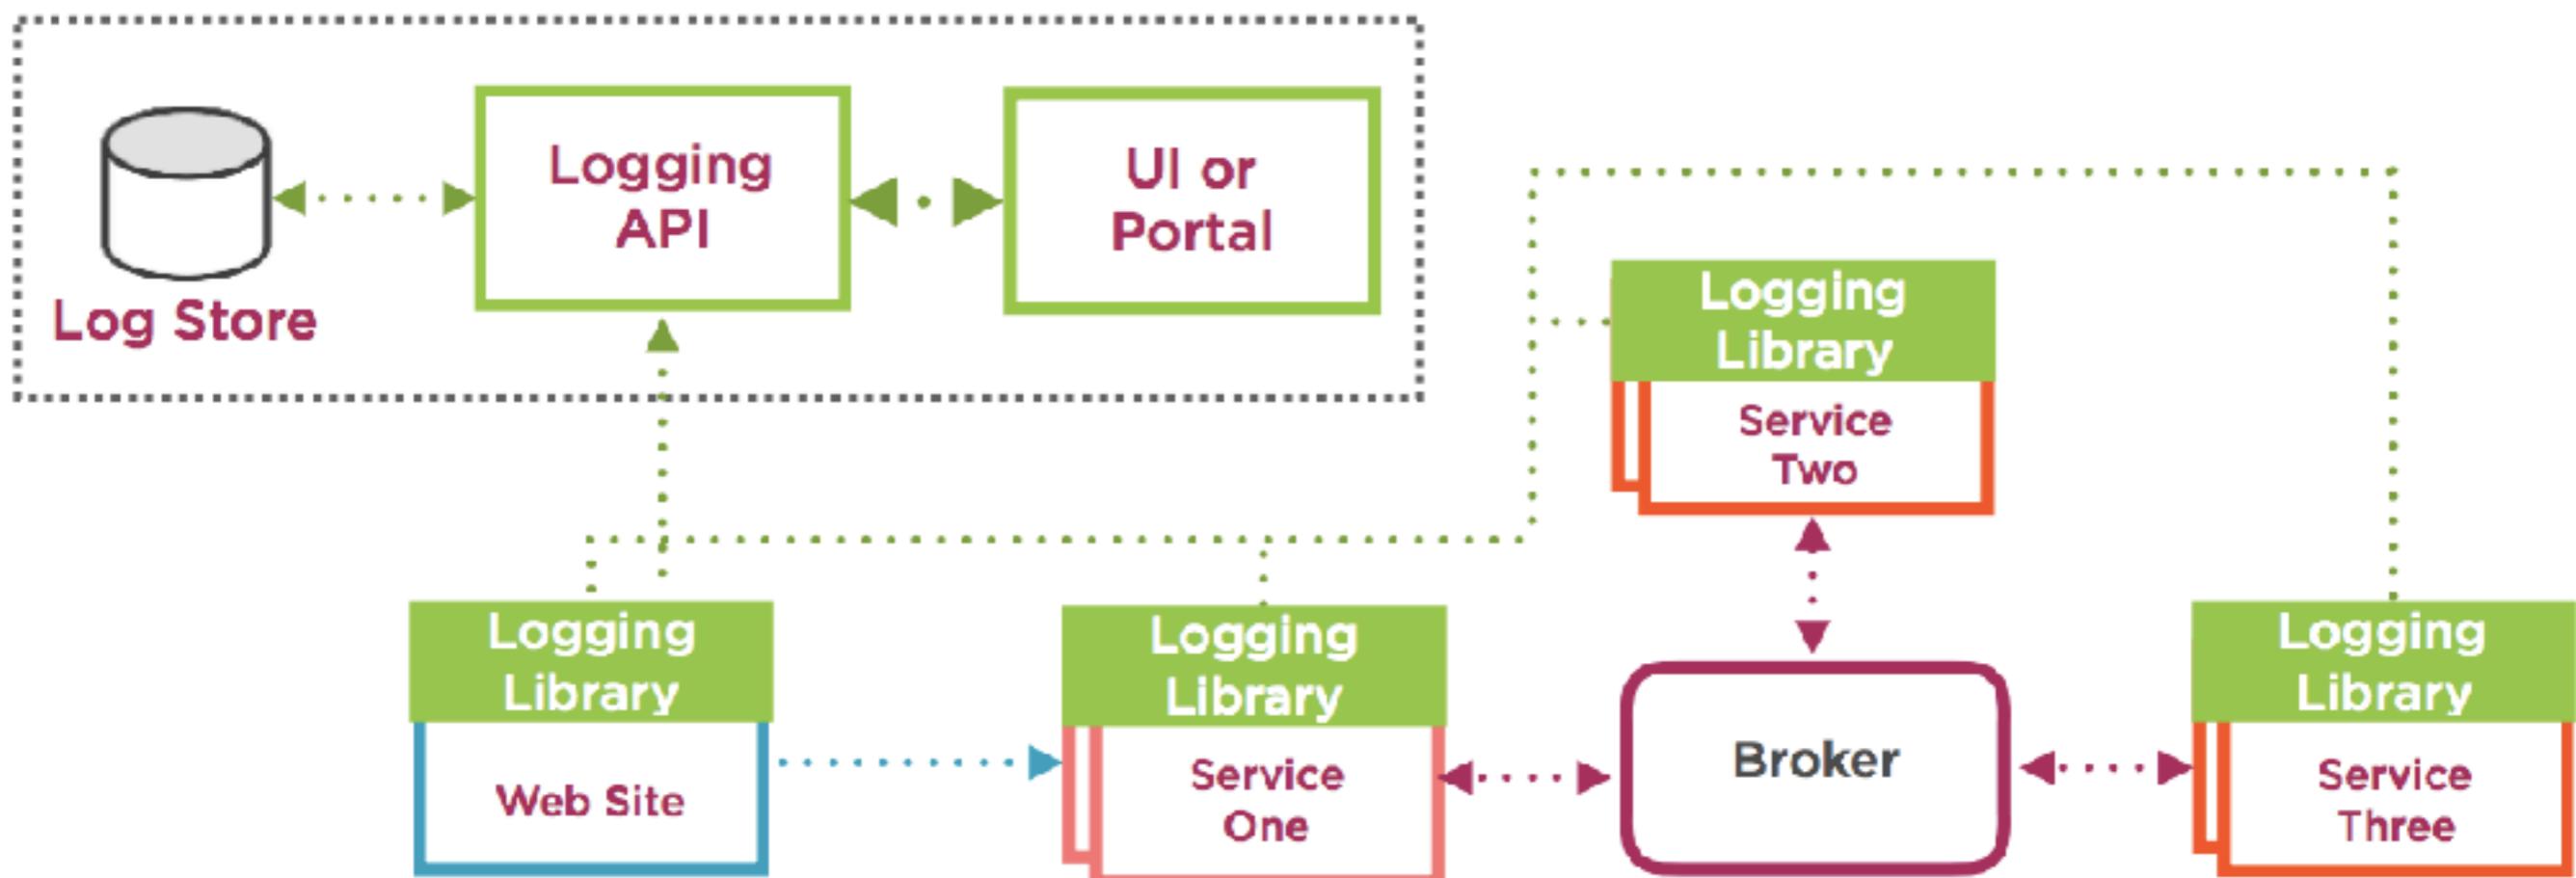
Centralize Logging



Consistency Logging Format

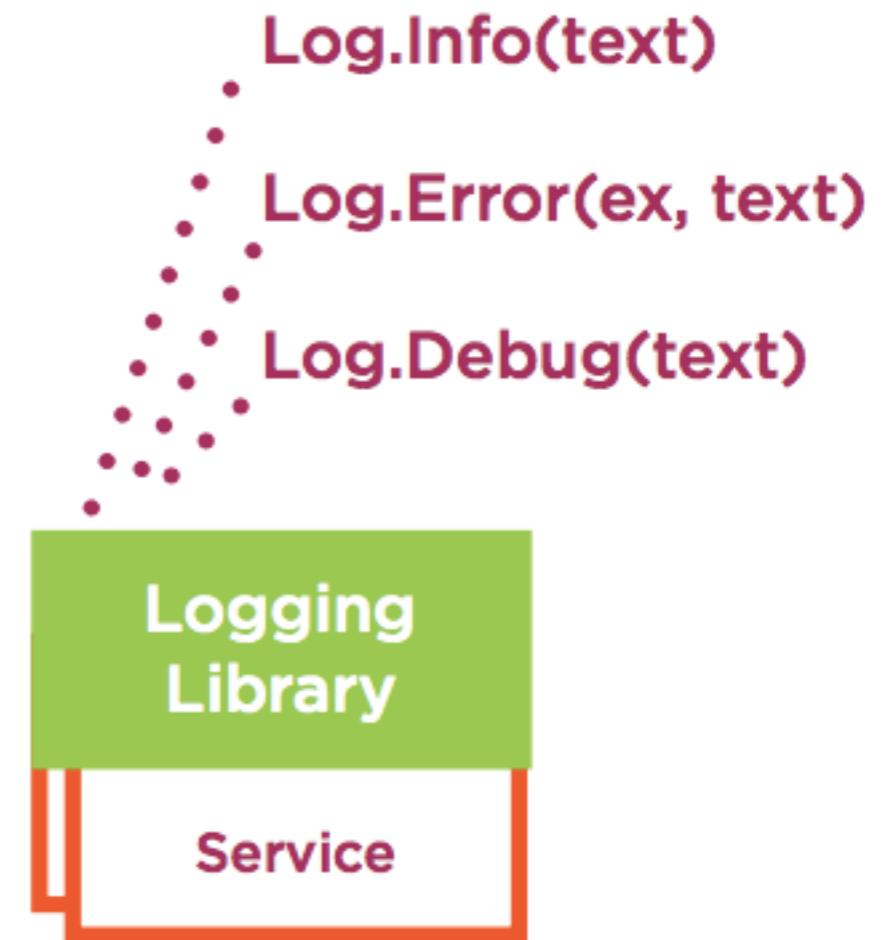
- Across the architecture
- Structured format
 - Levels
 - Date and time
 - Correlation ID
 - Host and app information - Message
- Shared library
- Decoupled from implementation

Consistency Logging Format



Logging Level

- Information level
 - Non-technical descriptive message
 - Simple enough for entire business
 - Key transaction milestones
- Debug level
 - Technical milestones
 - Calls and parameters
 - Response times and timeout stats
- Error level
 - Information on exceptions raised
 - Error information and number
 - Call stack



Performance

Scaling

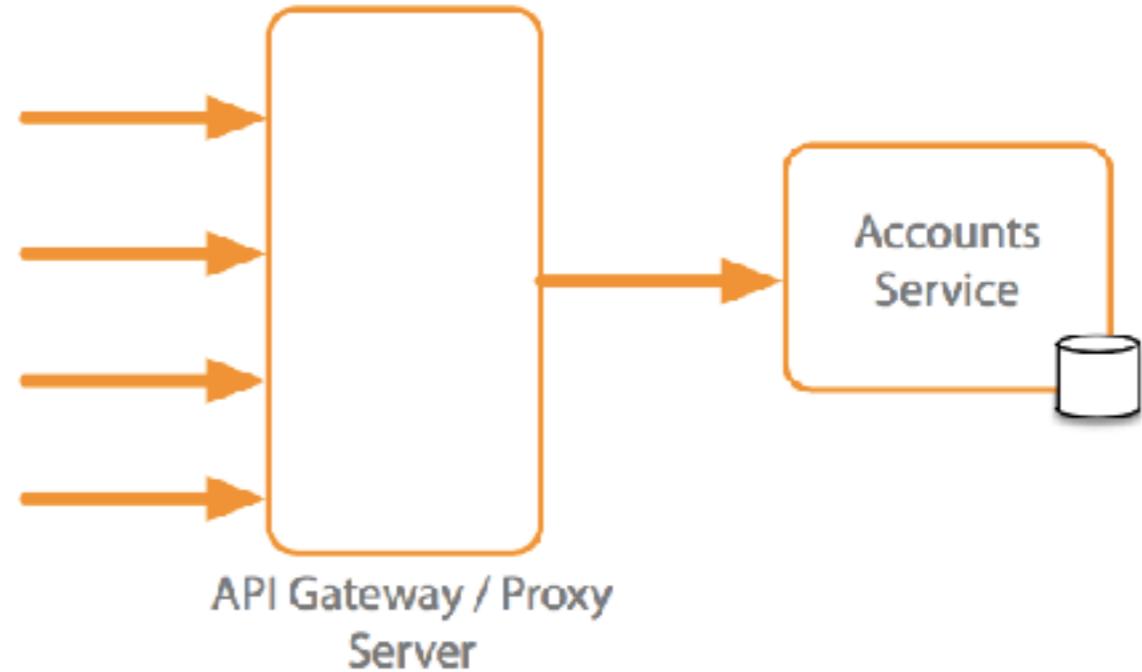
- How
 - Creating multiple instances of service
 - Adding resource to existing service
- Automated or on-demand
- PAAS auto scaling options
- Virtualization and containers

Scaling

- Physical host servers
- Load balancers (API Gateway)
- When to scale up
 - Performance issues
 - Monitoring data
 - Capacity planning

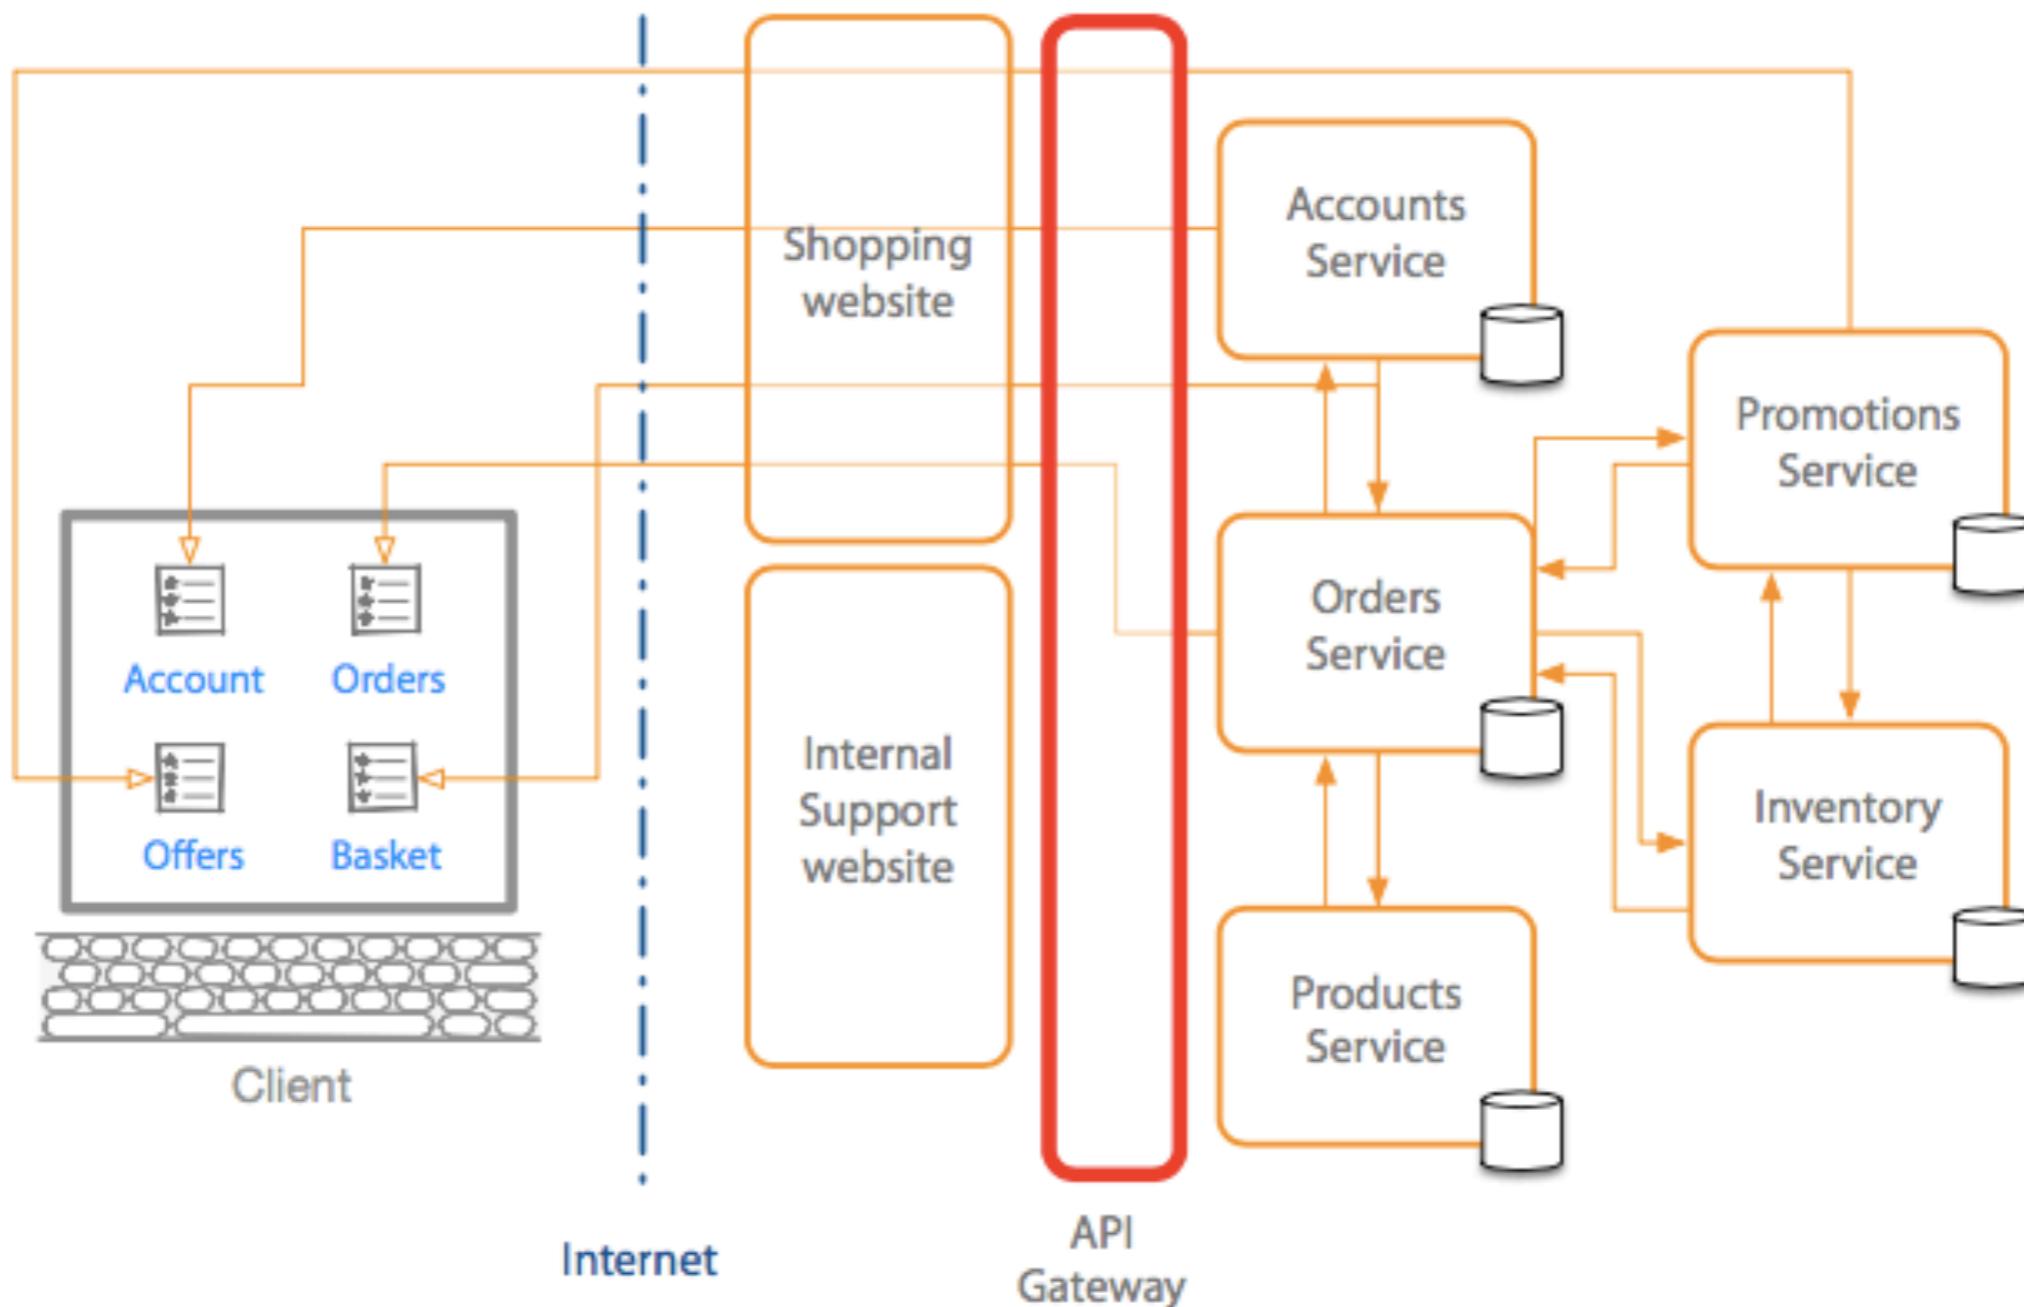
Caching

- Caching to reduce
 - Client calls to services
 - Service calls to databases
 - Service to service calls
- API Gateway\Proxy level
- Client side
- Service level
- Considerations
 - Simple to setup and manage
 - Data leaks

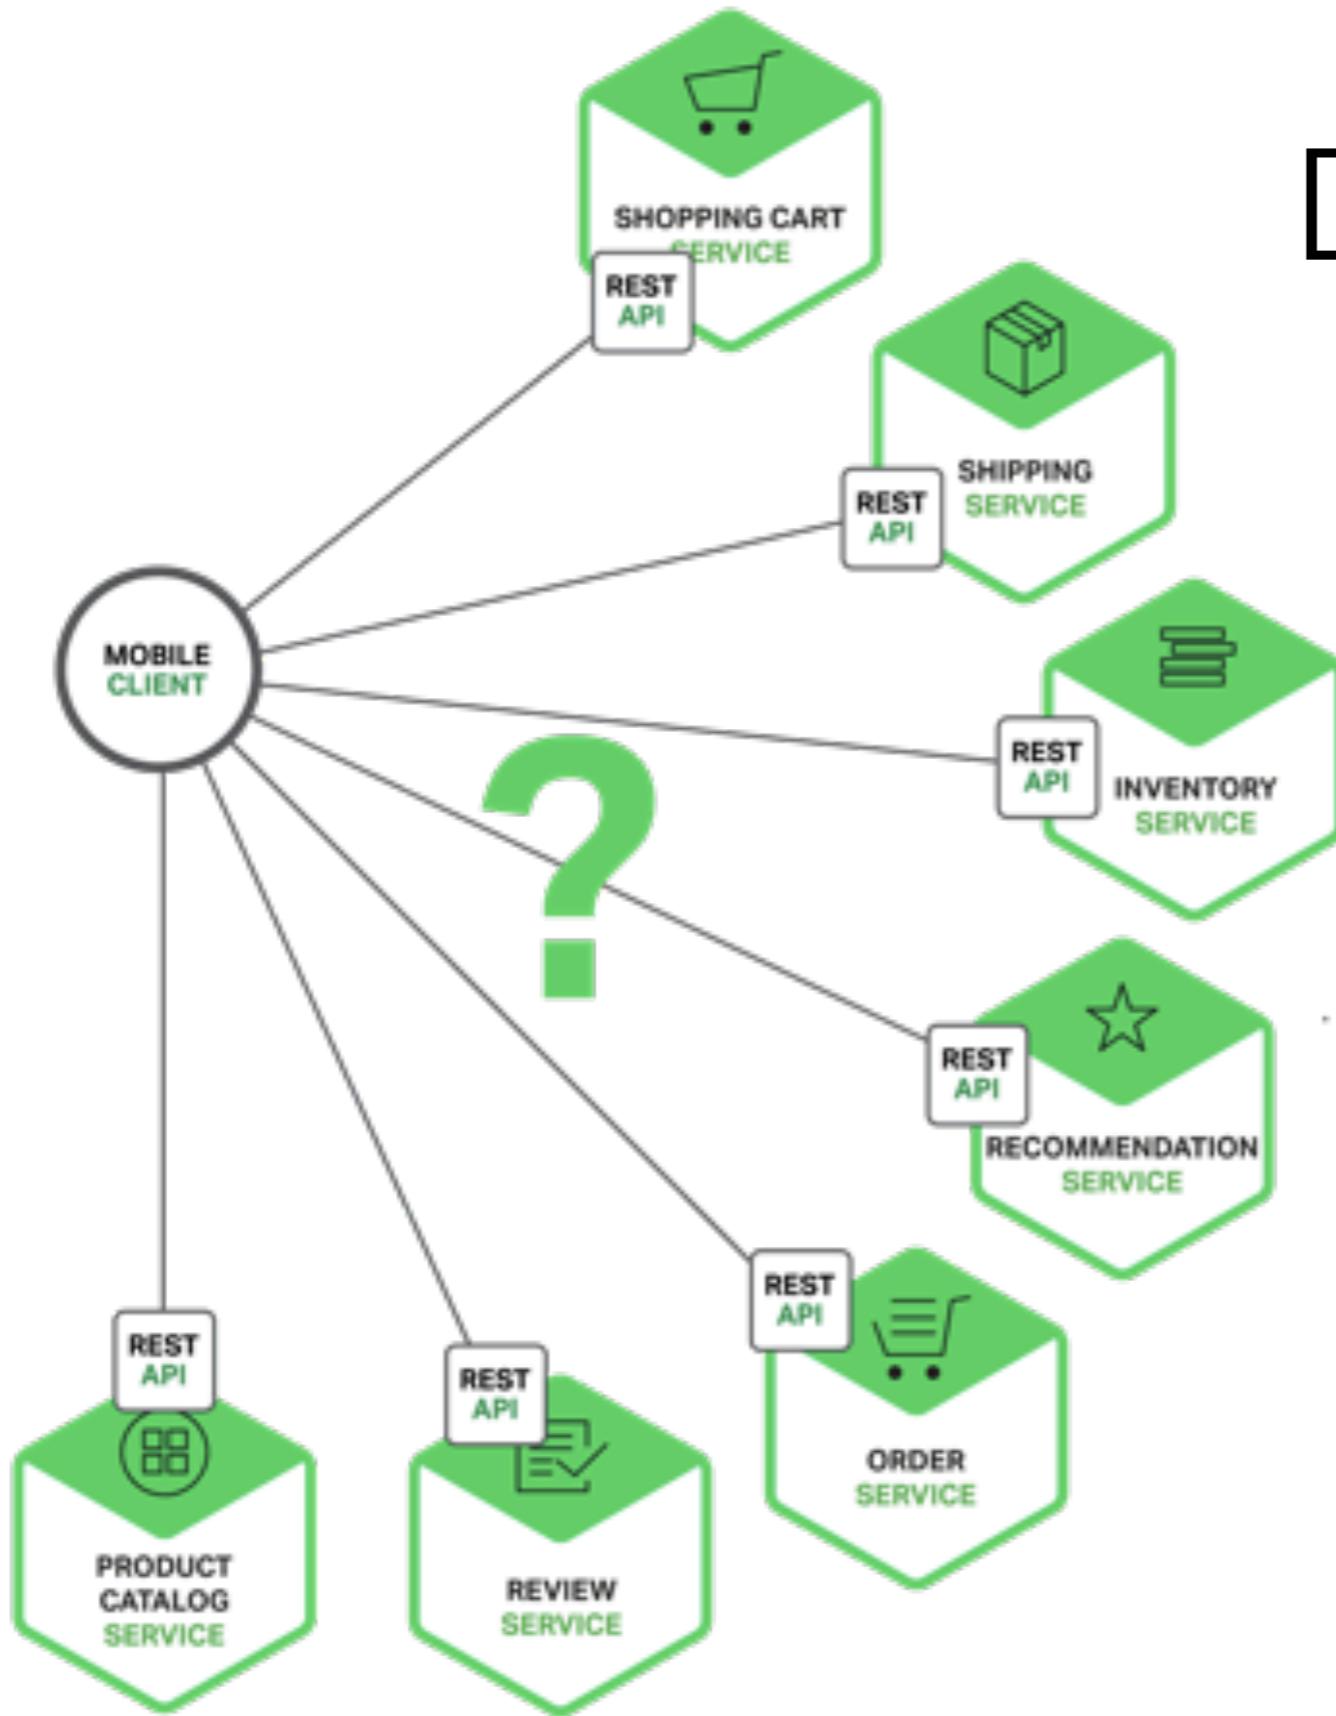


API Gateway

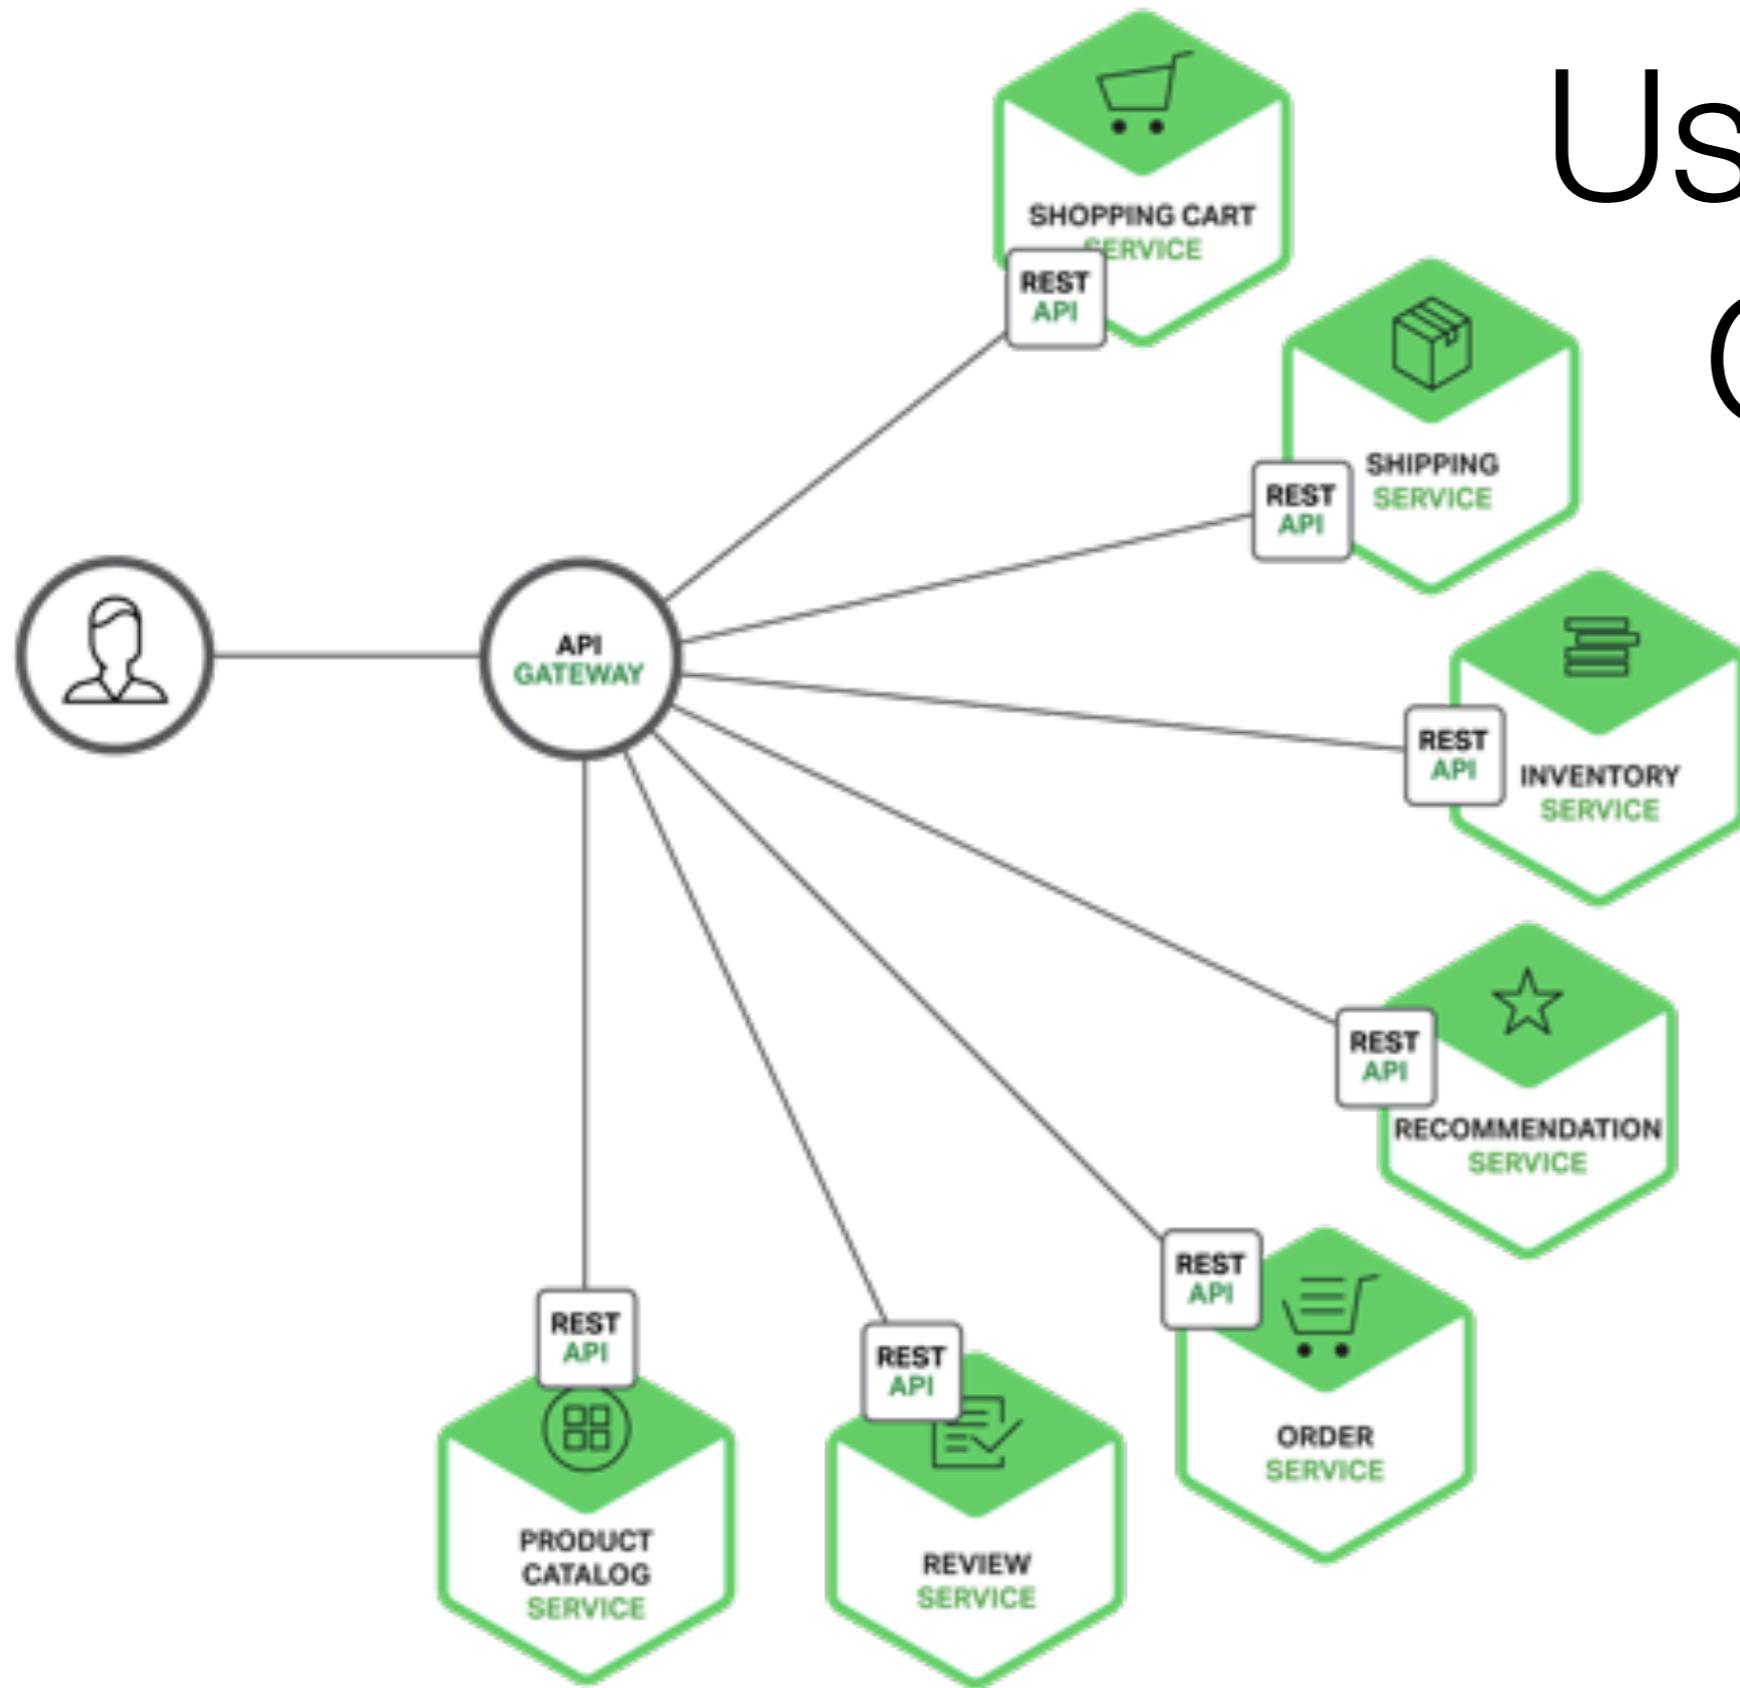
API Gateway



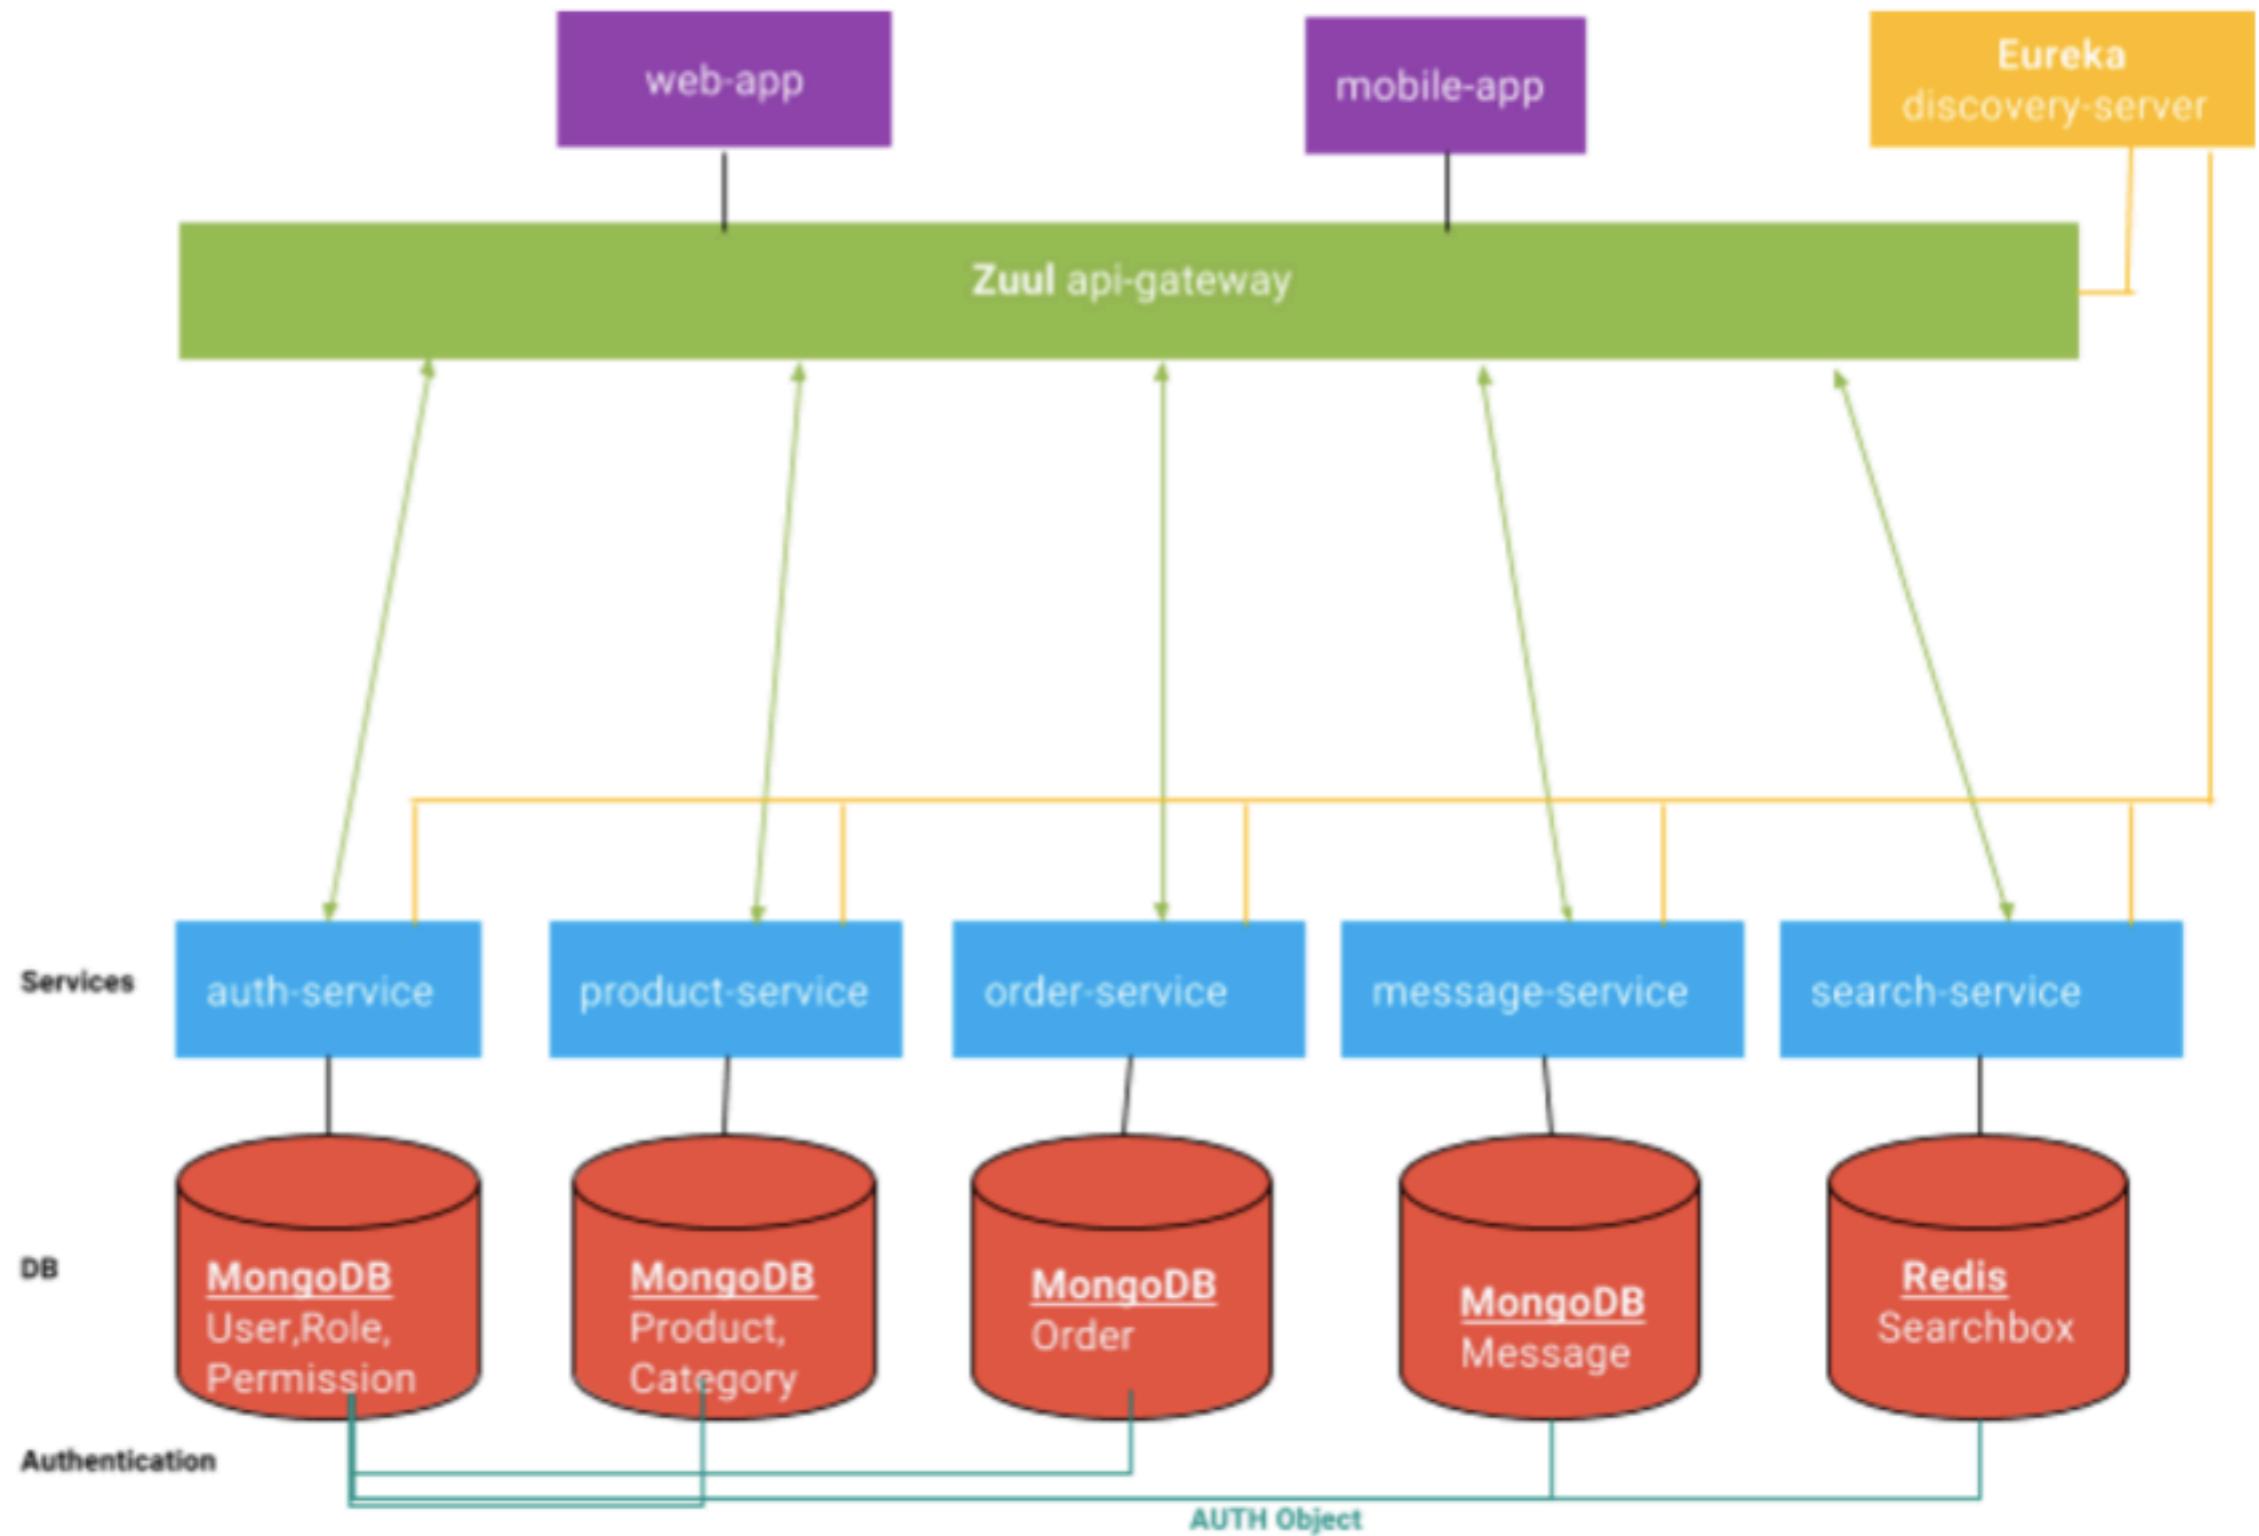
Direct Client to Microservices



Using an API Gateway



Netflix



Help with

- Creating central entry point
- Exposing services to clients
- One interface to many services
- Dynamic location of services
- Routing to specific instance of service
- Service registry database

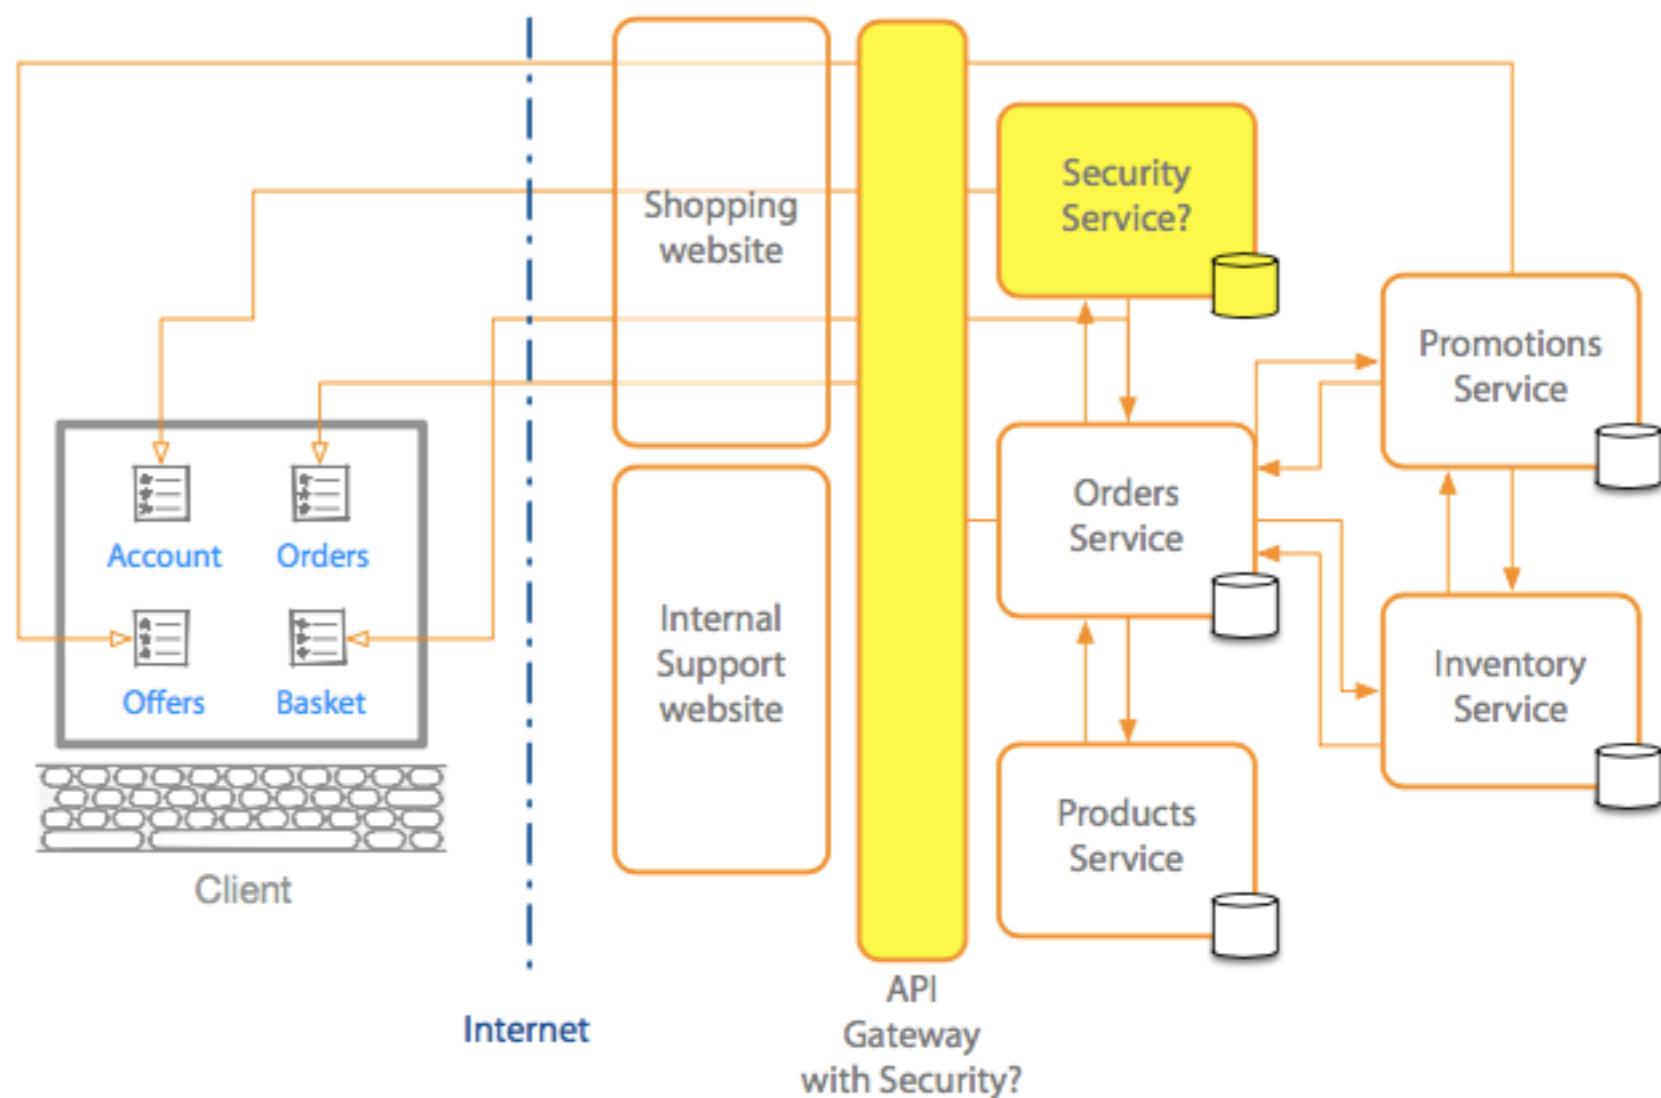
Help with Performance

Load Balancing

Caching

Help with Security

- Dedicated security service
- Central security vs service level



Automation Tools

Automation Tools

Continuous
Integration

Continuous
Deployment

Continuous Integration Tools

- Team Foundation Server
- TeamCity
- Jenkins

Features

- Cross platform
Windows builders, Java builders and others
- Source control integration
- Notifications
- IDE Integration (optional)

Map to CI Build

- Code change triggers build of specific service
- Feedback just received on that service
- Builds and tests run quicker
- Separate code repository for service
- End product is in one place
- CI builds to test database changes
- Both microservice build and database upgrade are ready

Avoid one CI build for
all services

Continuous Deployment

- Jenkins
- Ansible
- Chef & Puppet

Features

- Central control panel
- Simple to add deployment targets
- Support for scripting
- Support for build statuses
- Integration with CI tool
- Support for multiple environments
- Support for PAAS

Moving Forward

Environment

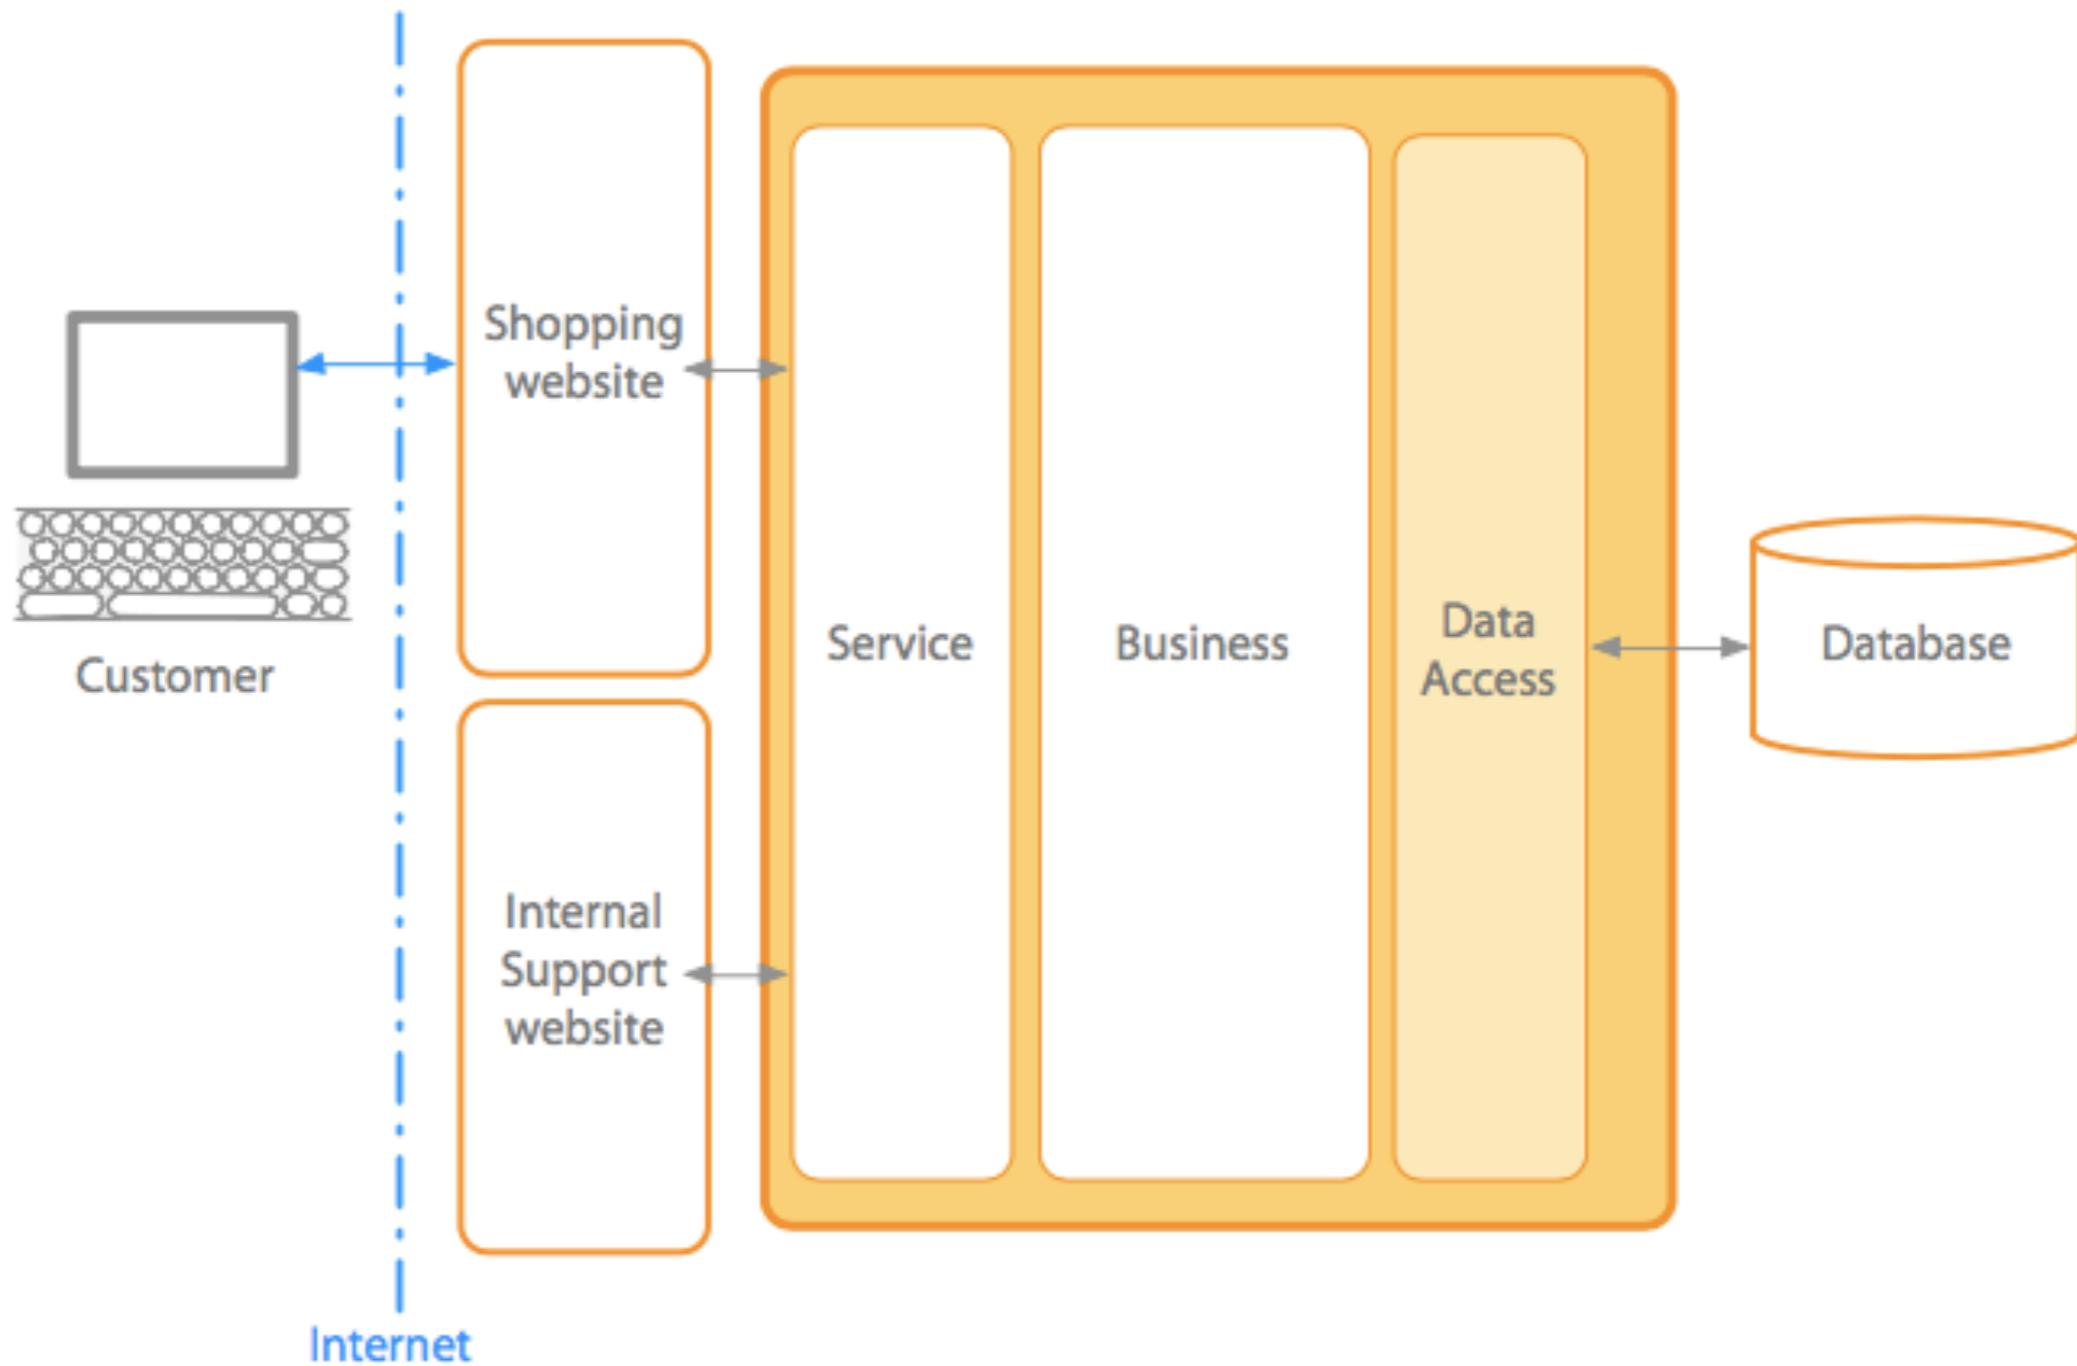
**Brownfield
Microservices**

**Greenfield
Microservices**

Brownfield MicroServices

- Existing system
 - Monolithic system
 - Organically grown
 - Seems to large to split
- Lacks microservices design principles

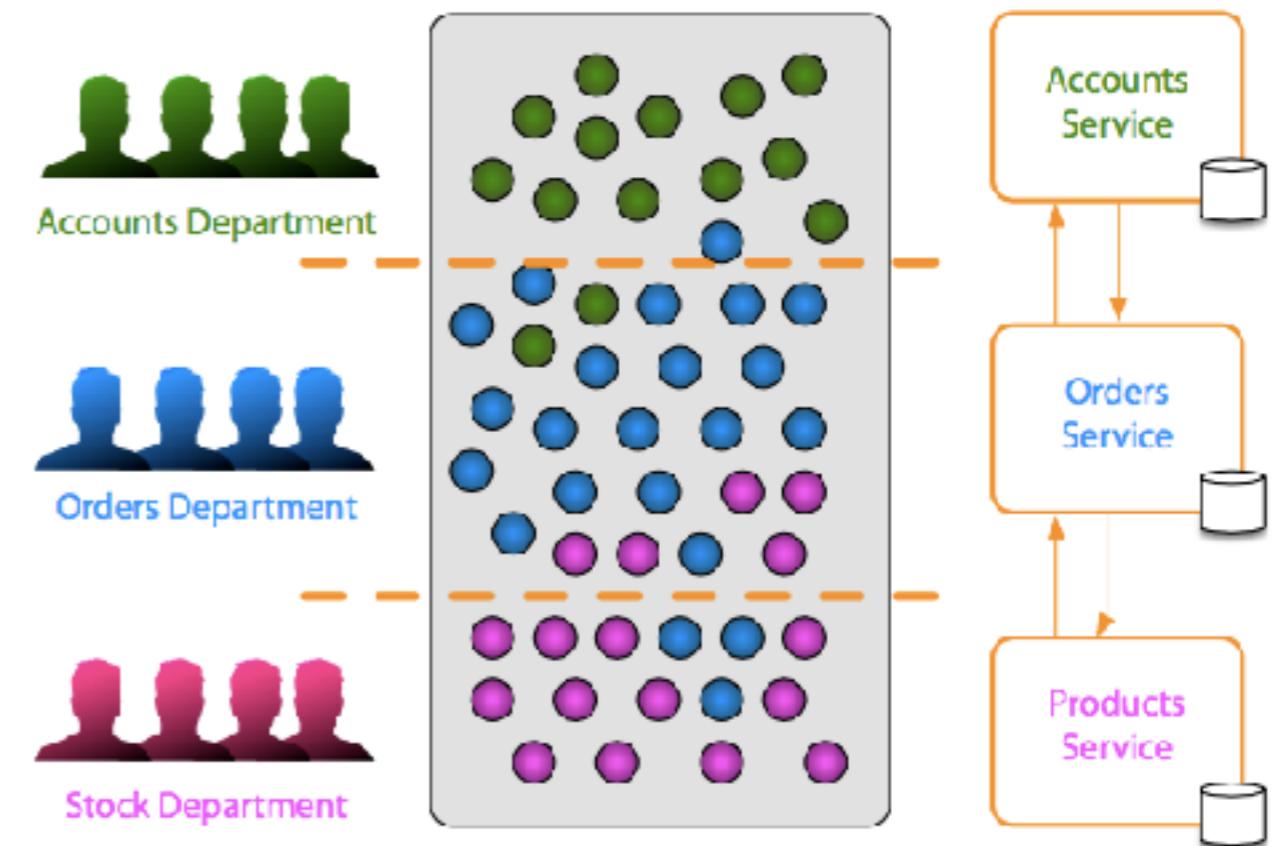
Brownfield MicroServices



Brownfield MicroServices

Identify seams

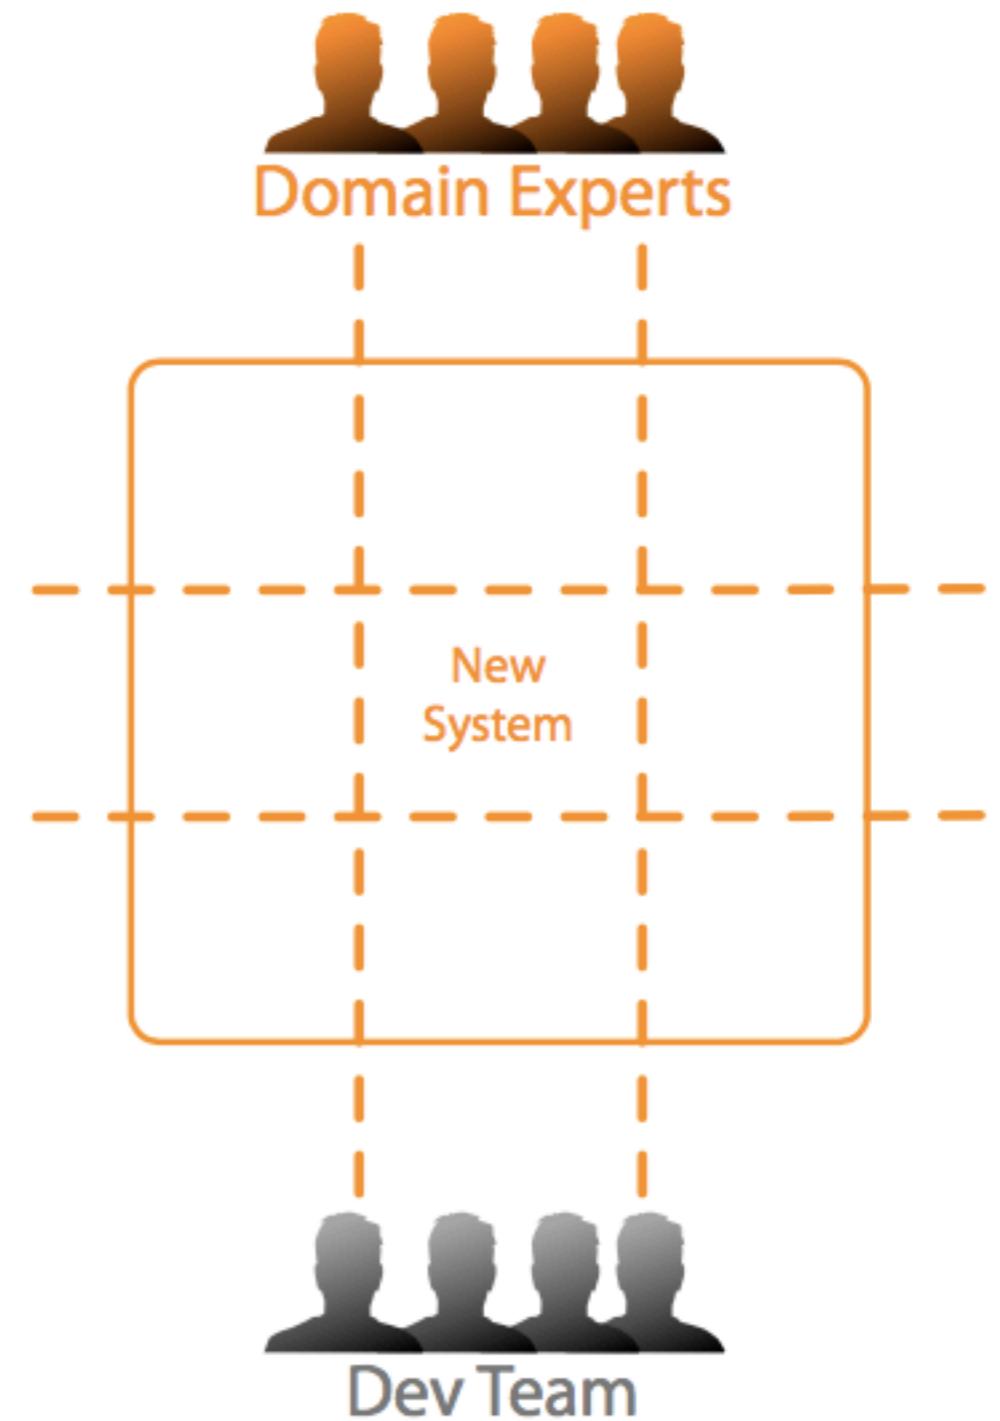
- Separation that reflects domains
- Identify bounded contexts



Seams are future microservice boundaries

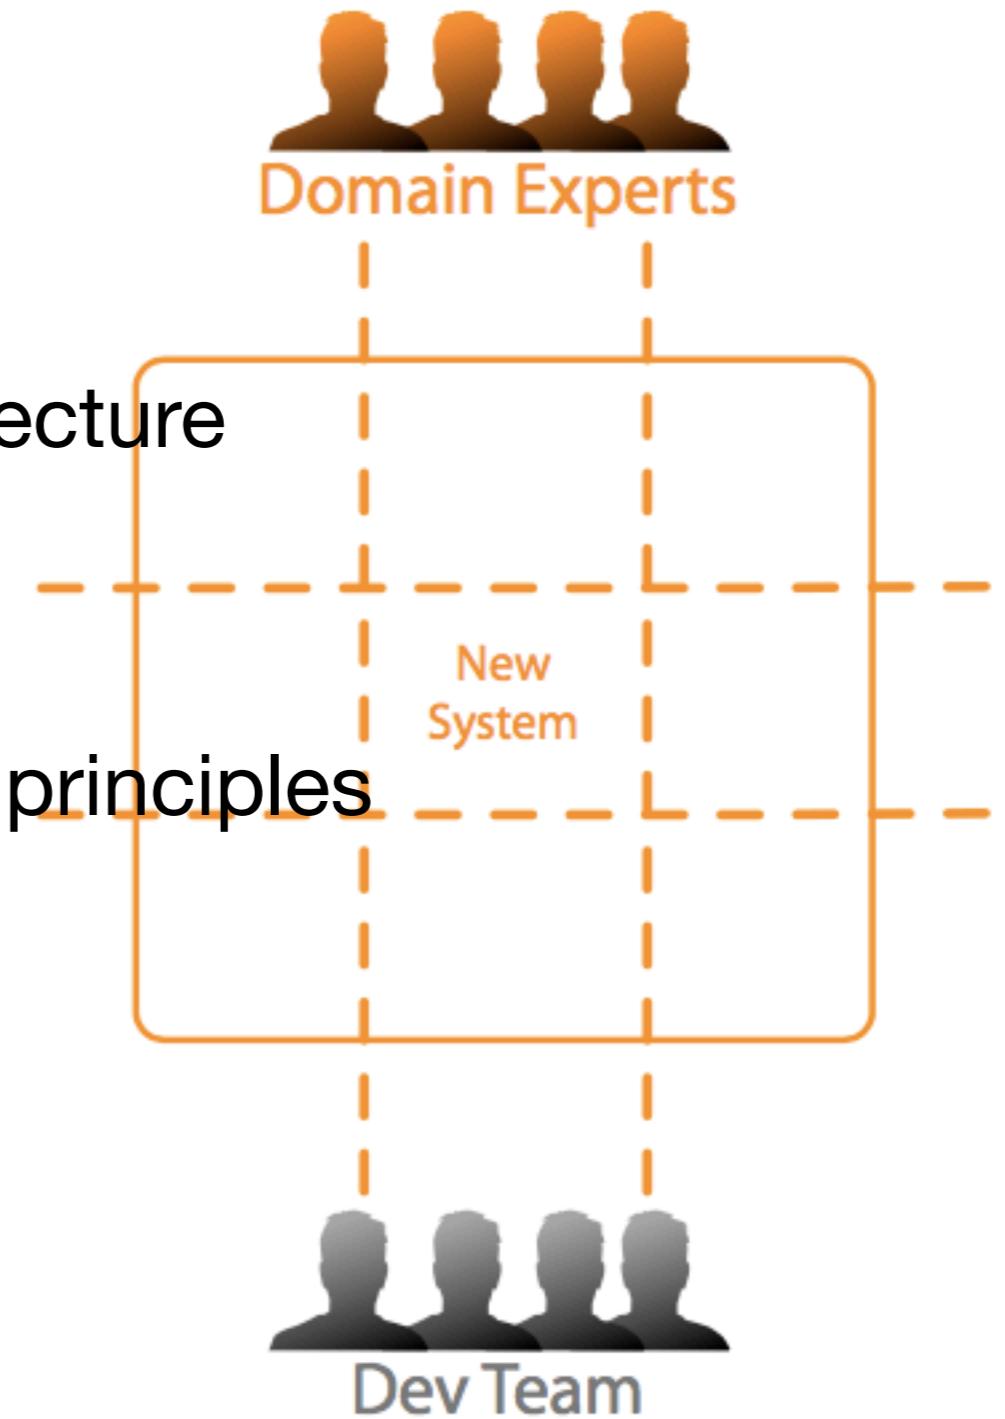
Greenfield MicroServices

- New project
- Evolving requirements
- Business domain
 - Not fully understood
 - Getting domain experts involved
 - System boundaries will evolve
- Teams experience
 - First microservices
 - Experienced with microservices



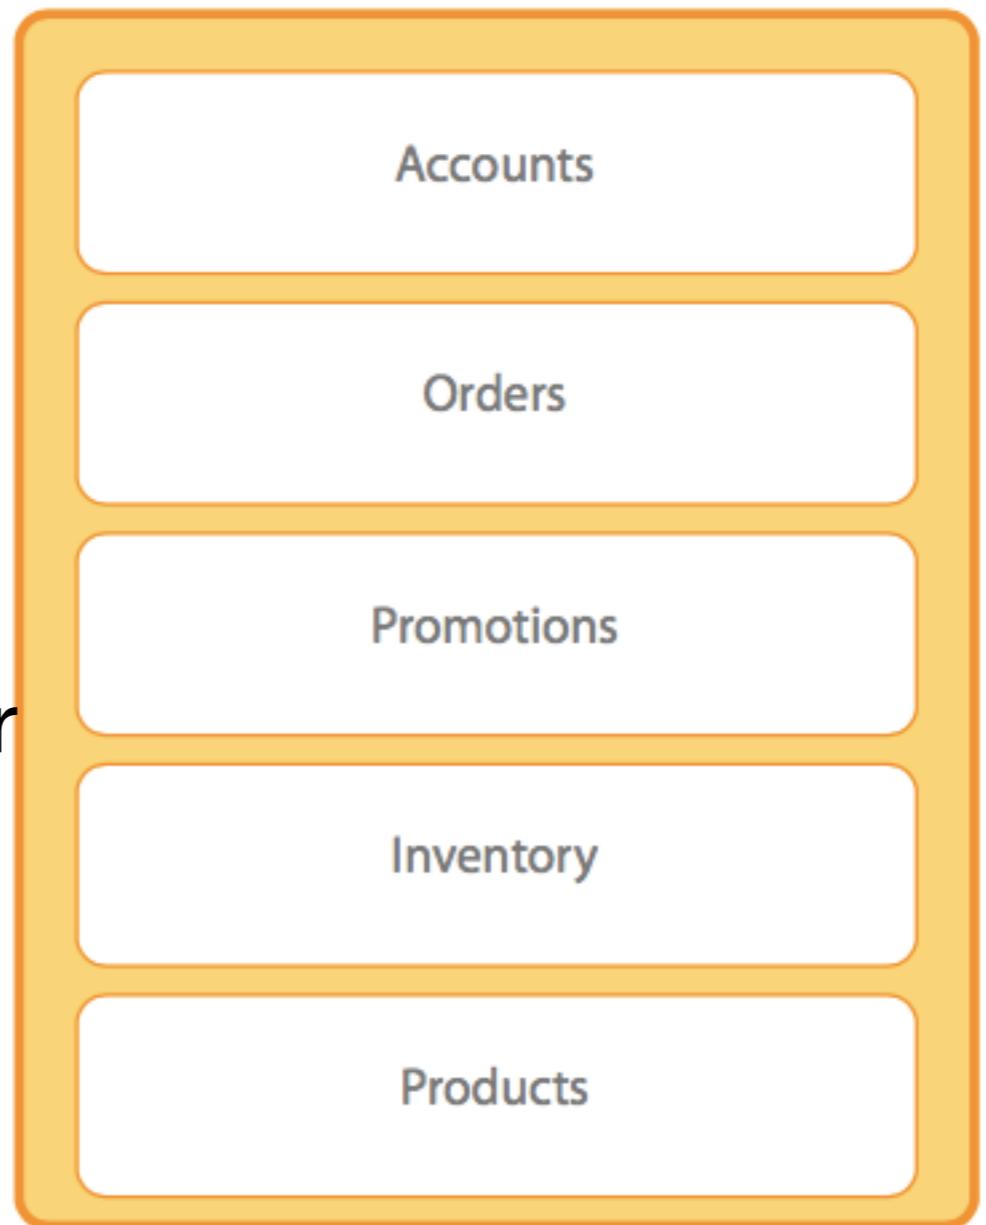
Greenfield MicroServices

- Existing system integration
 - Monolithic system
 - Established microservices architecture
- Push for change
 - Changes to apply microservices principles



Start with Monolithic Design

- High level
- Evolving seams
- Develop areas into modules
- Boundaries start to become clearer
- Refine and refactor design
- Split further when required



Bounded Context

A specific responsibility enforced by an explicit boundary

Core Aspects of Your Domain



Driver
Transport
Shifts
Annual leave

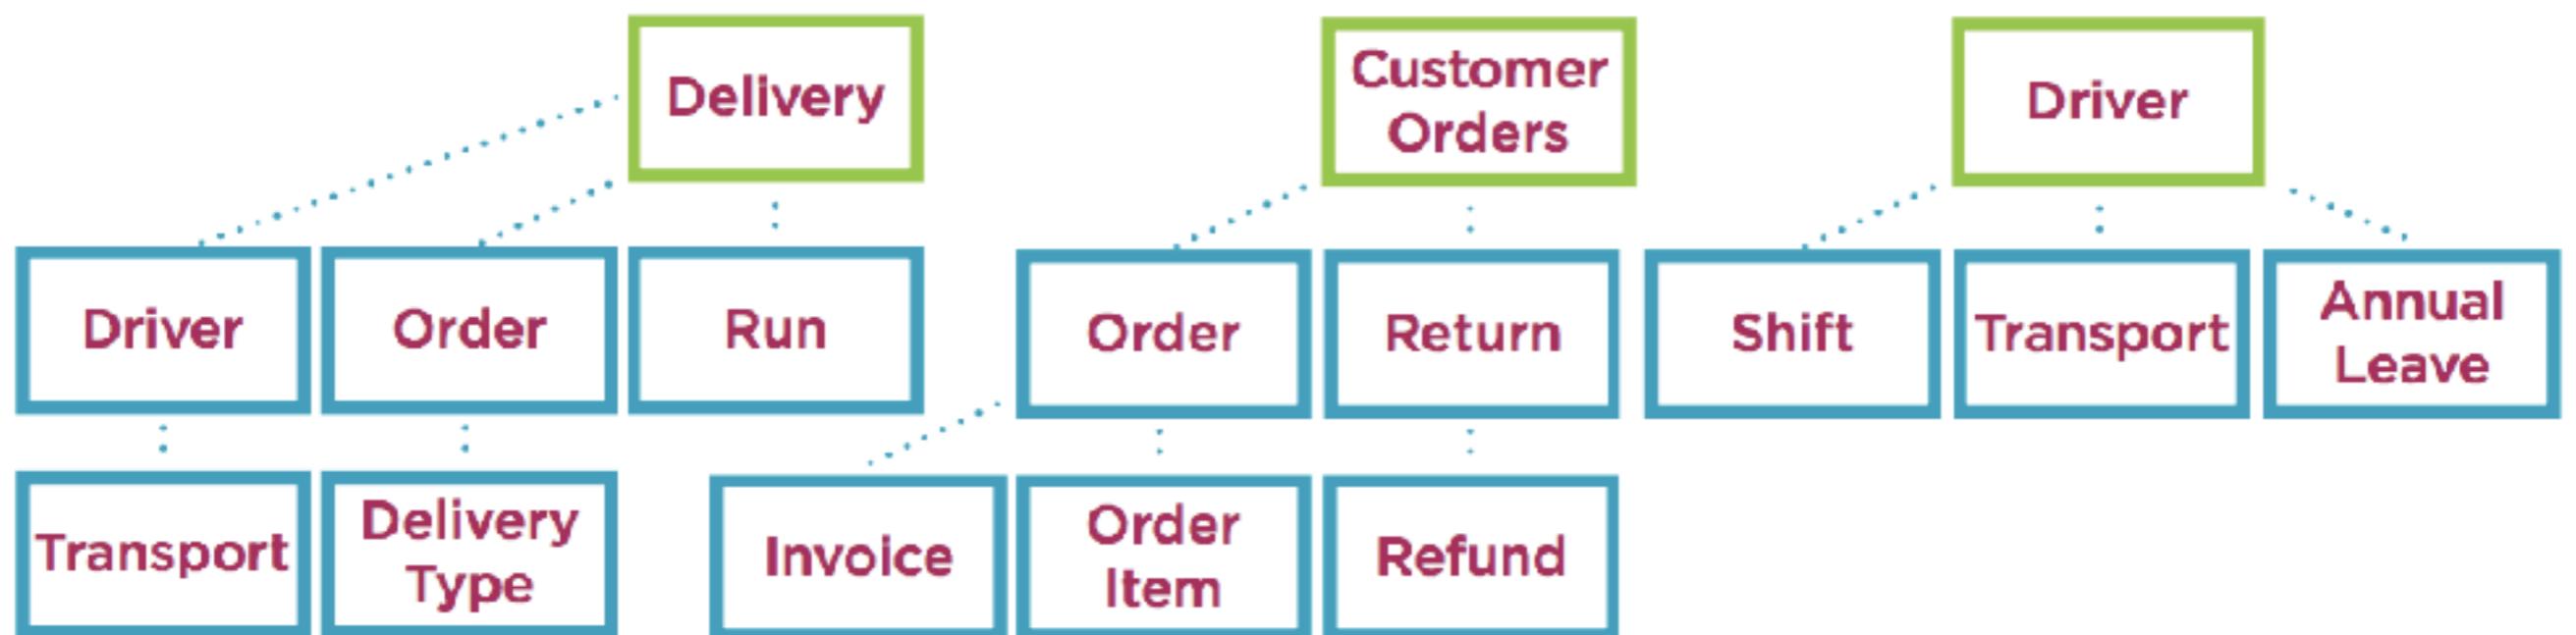


Customer Orders
Orders
Returns

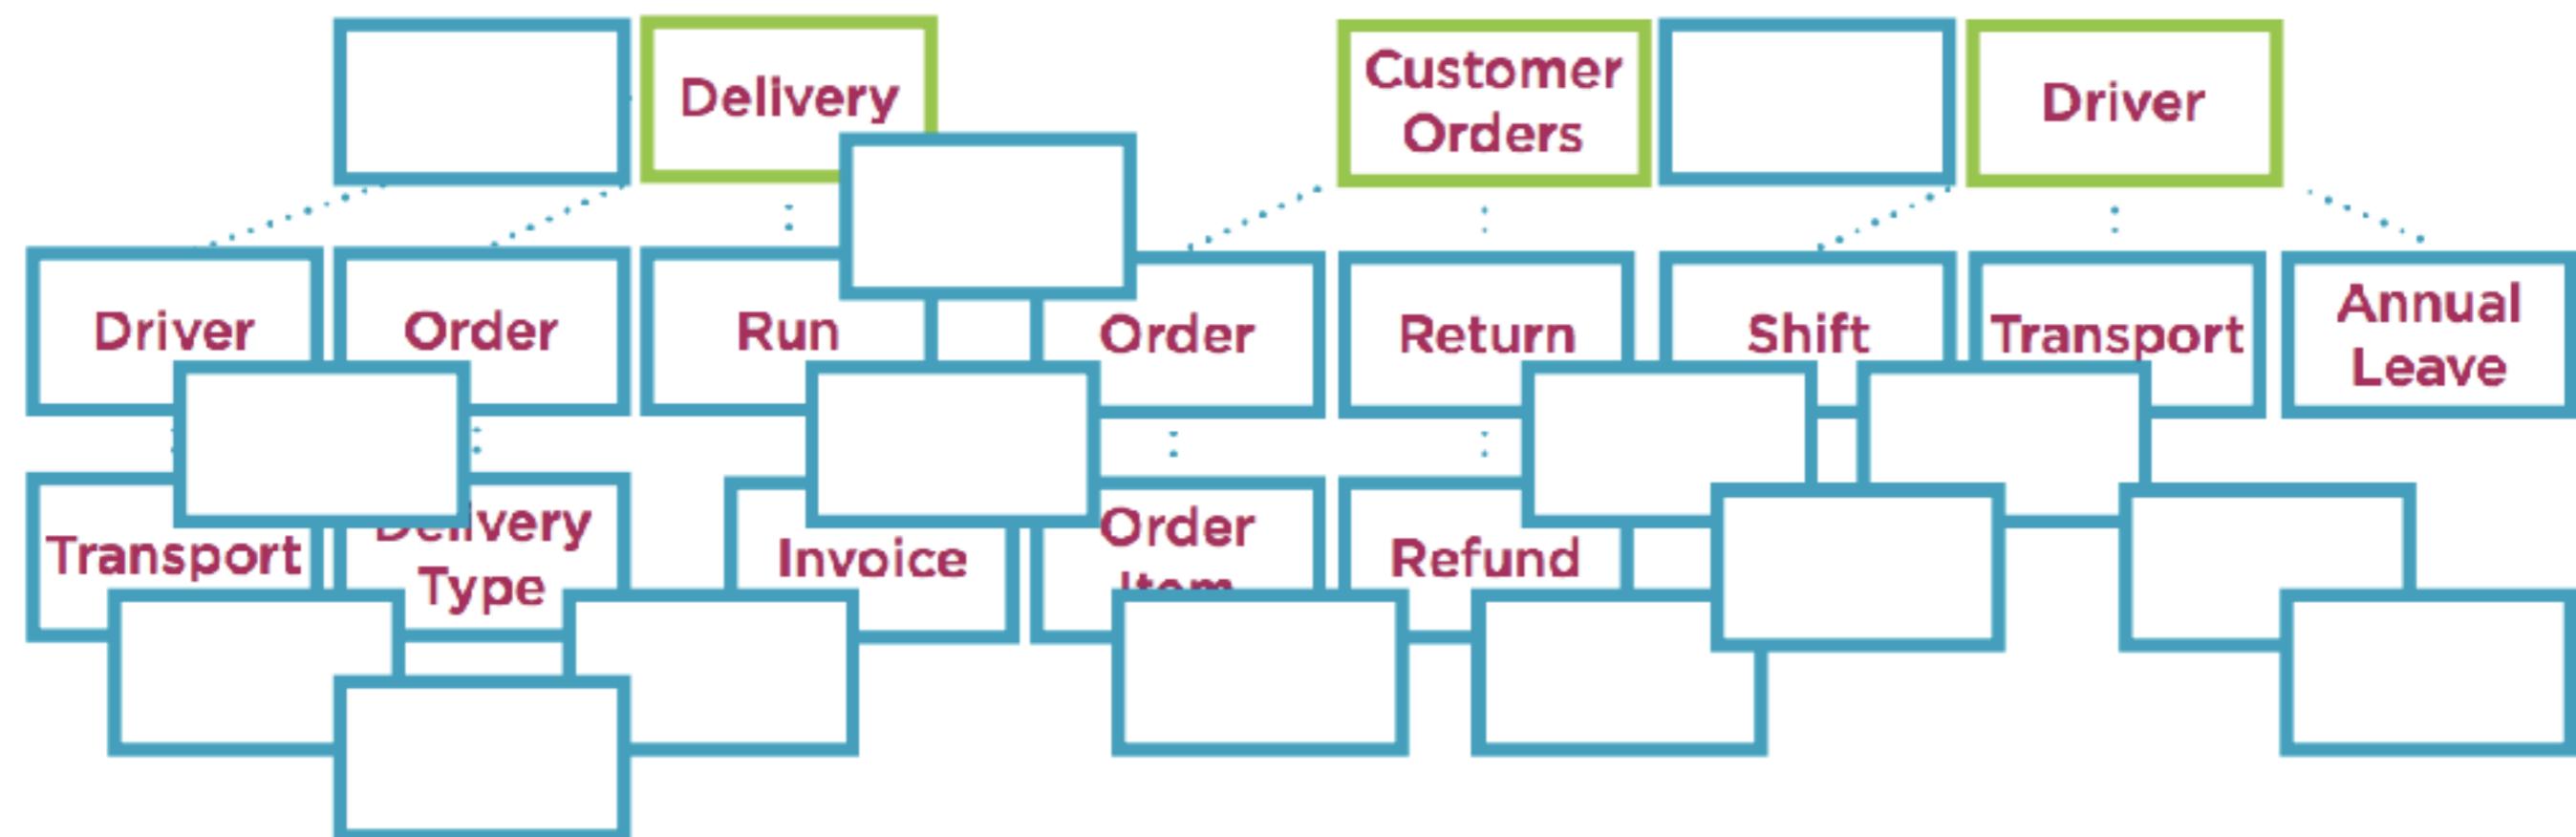


Delivery
Drivers
Orders
Runs

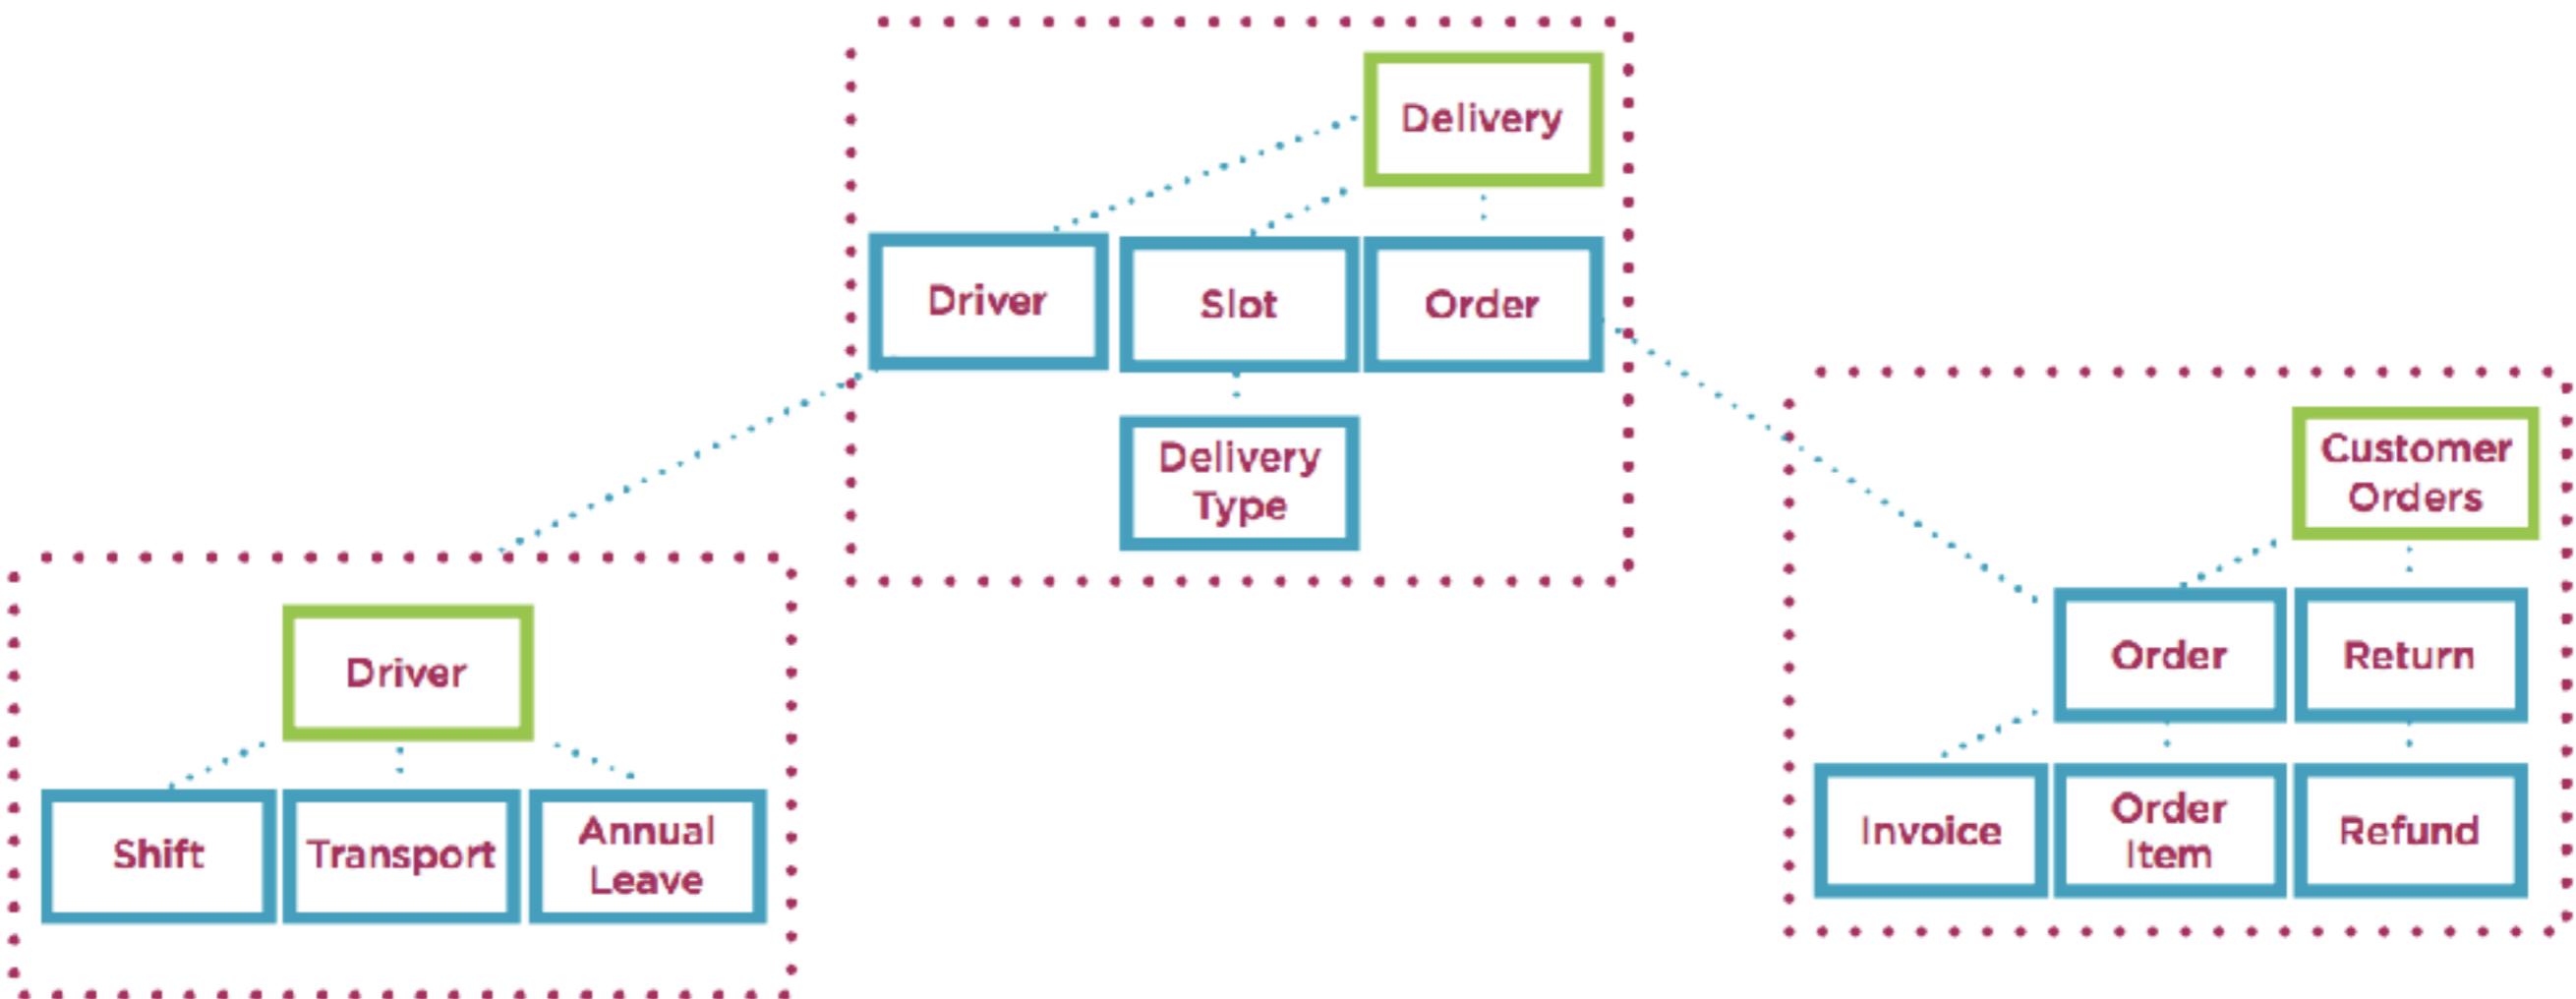
Core Concepts with Supporting Concepts



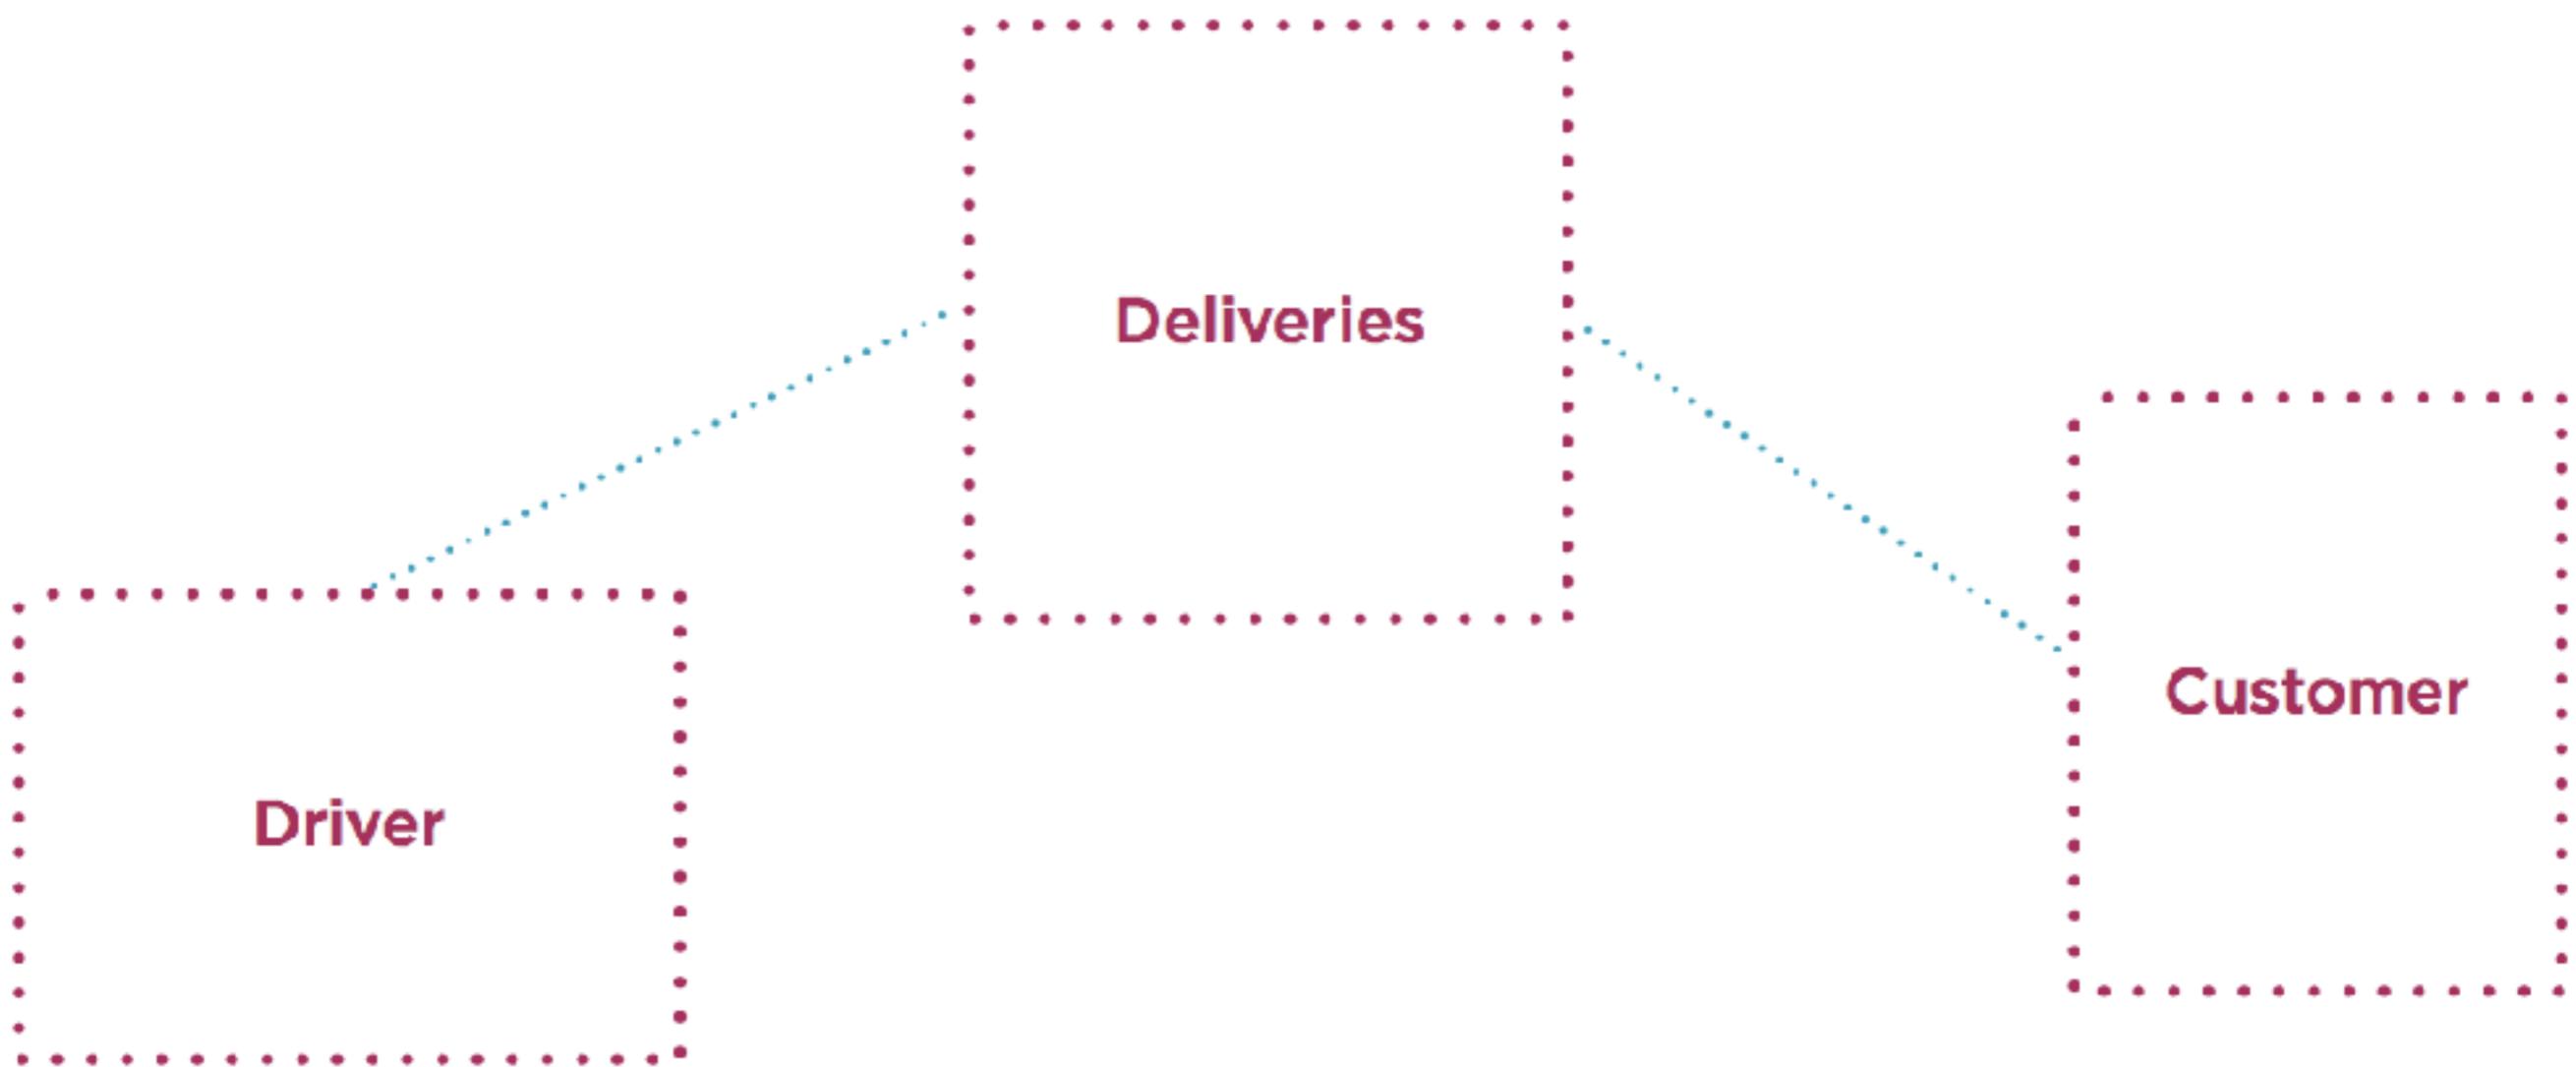
Unbounded Context



Indicate Integration



Become MicroServices



Start Modularizing

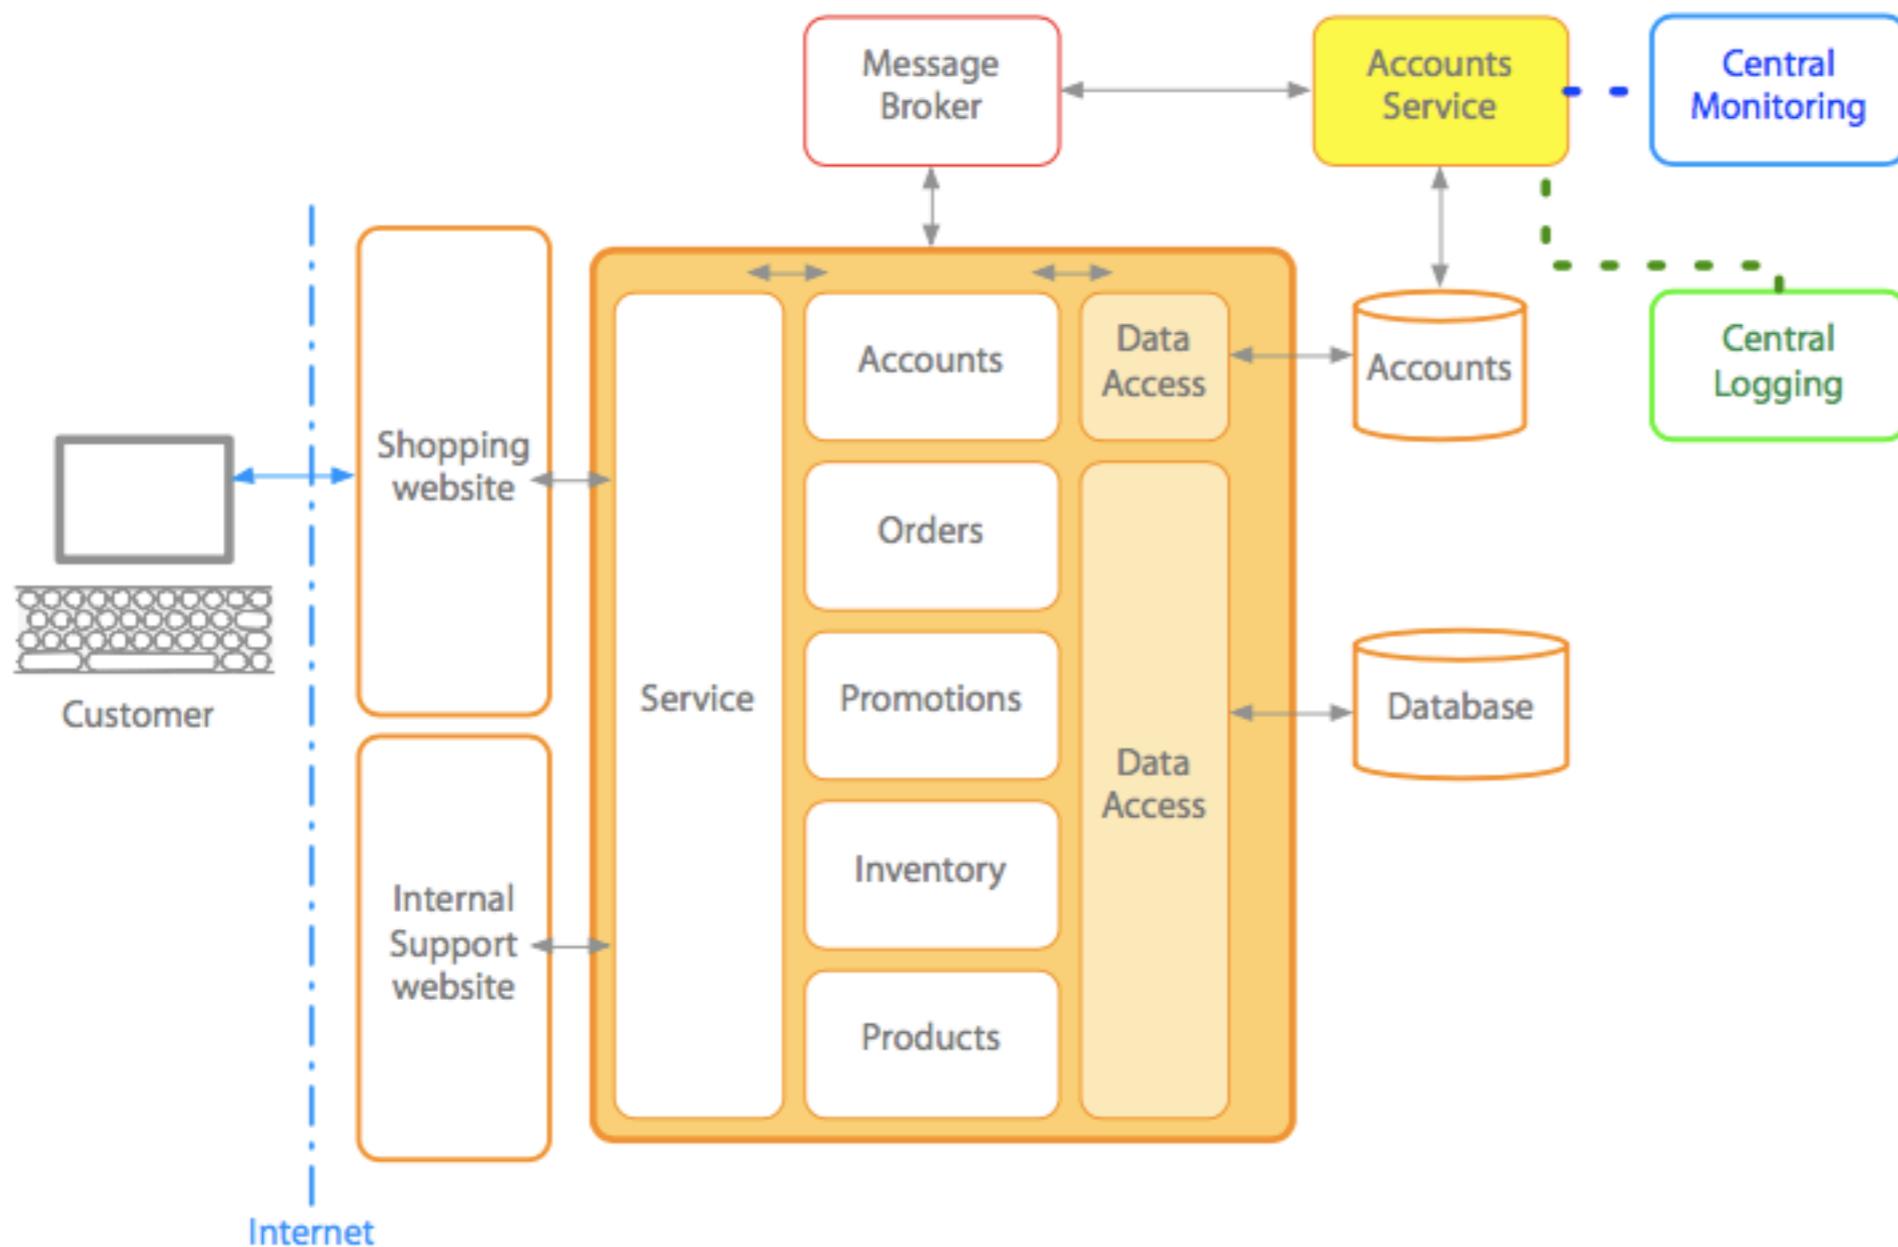
- Move code incrementally
- Tidy up a section per release
 - Take your time
 - Existing functionality needs to remain intact
- Run unit tests and integration tests to validate change
- Keep reviewing

Migration

Migration

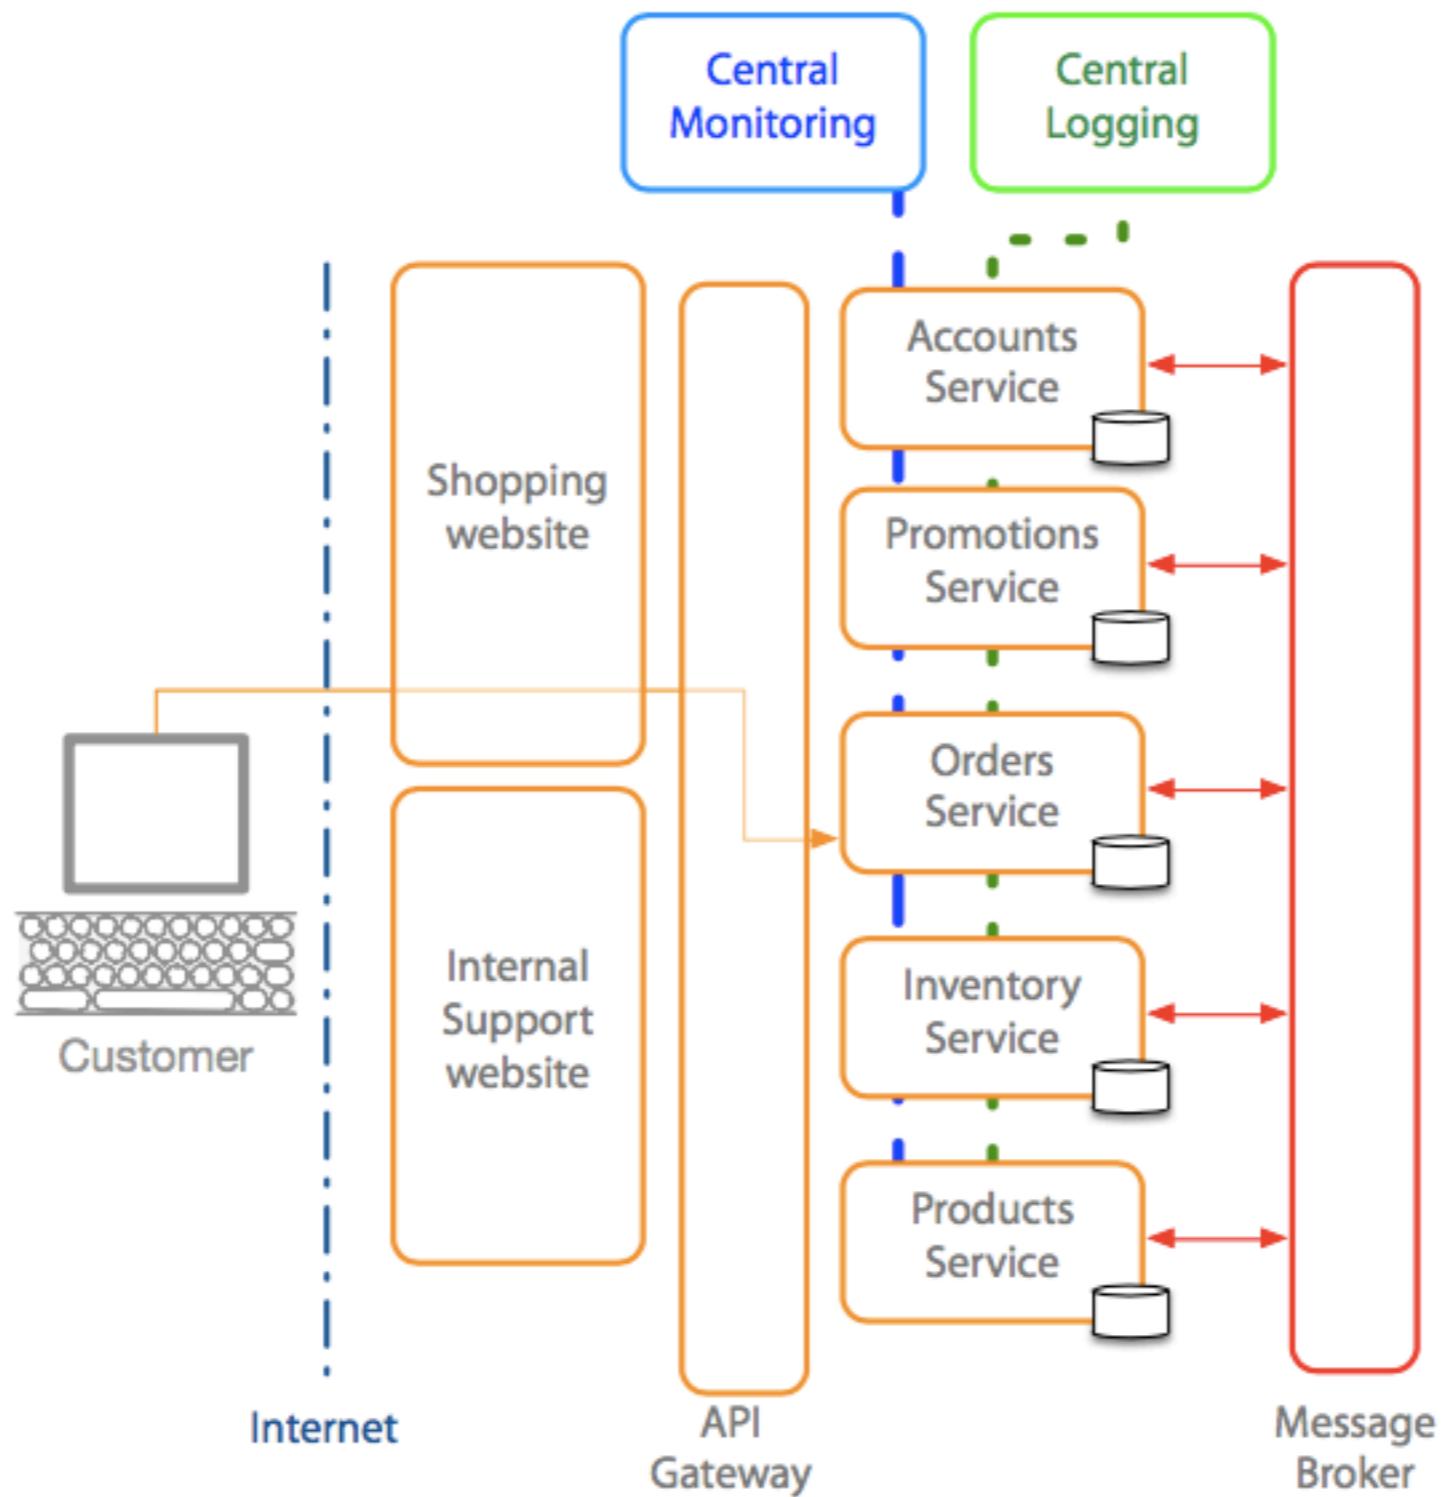
- Code is organized into bounded contexts
 - Code related to a business domain or function is in one place
 - Clear boundaries with clear interfaces between each
- Convert bounded contexts into microservices
 - Start off with one
Use to get comfortable
 - Make it switchable
Maintain two versions of the code

Migration

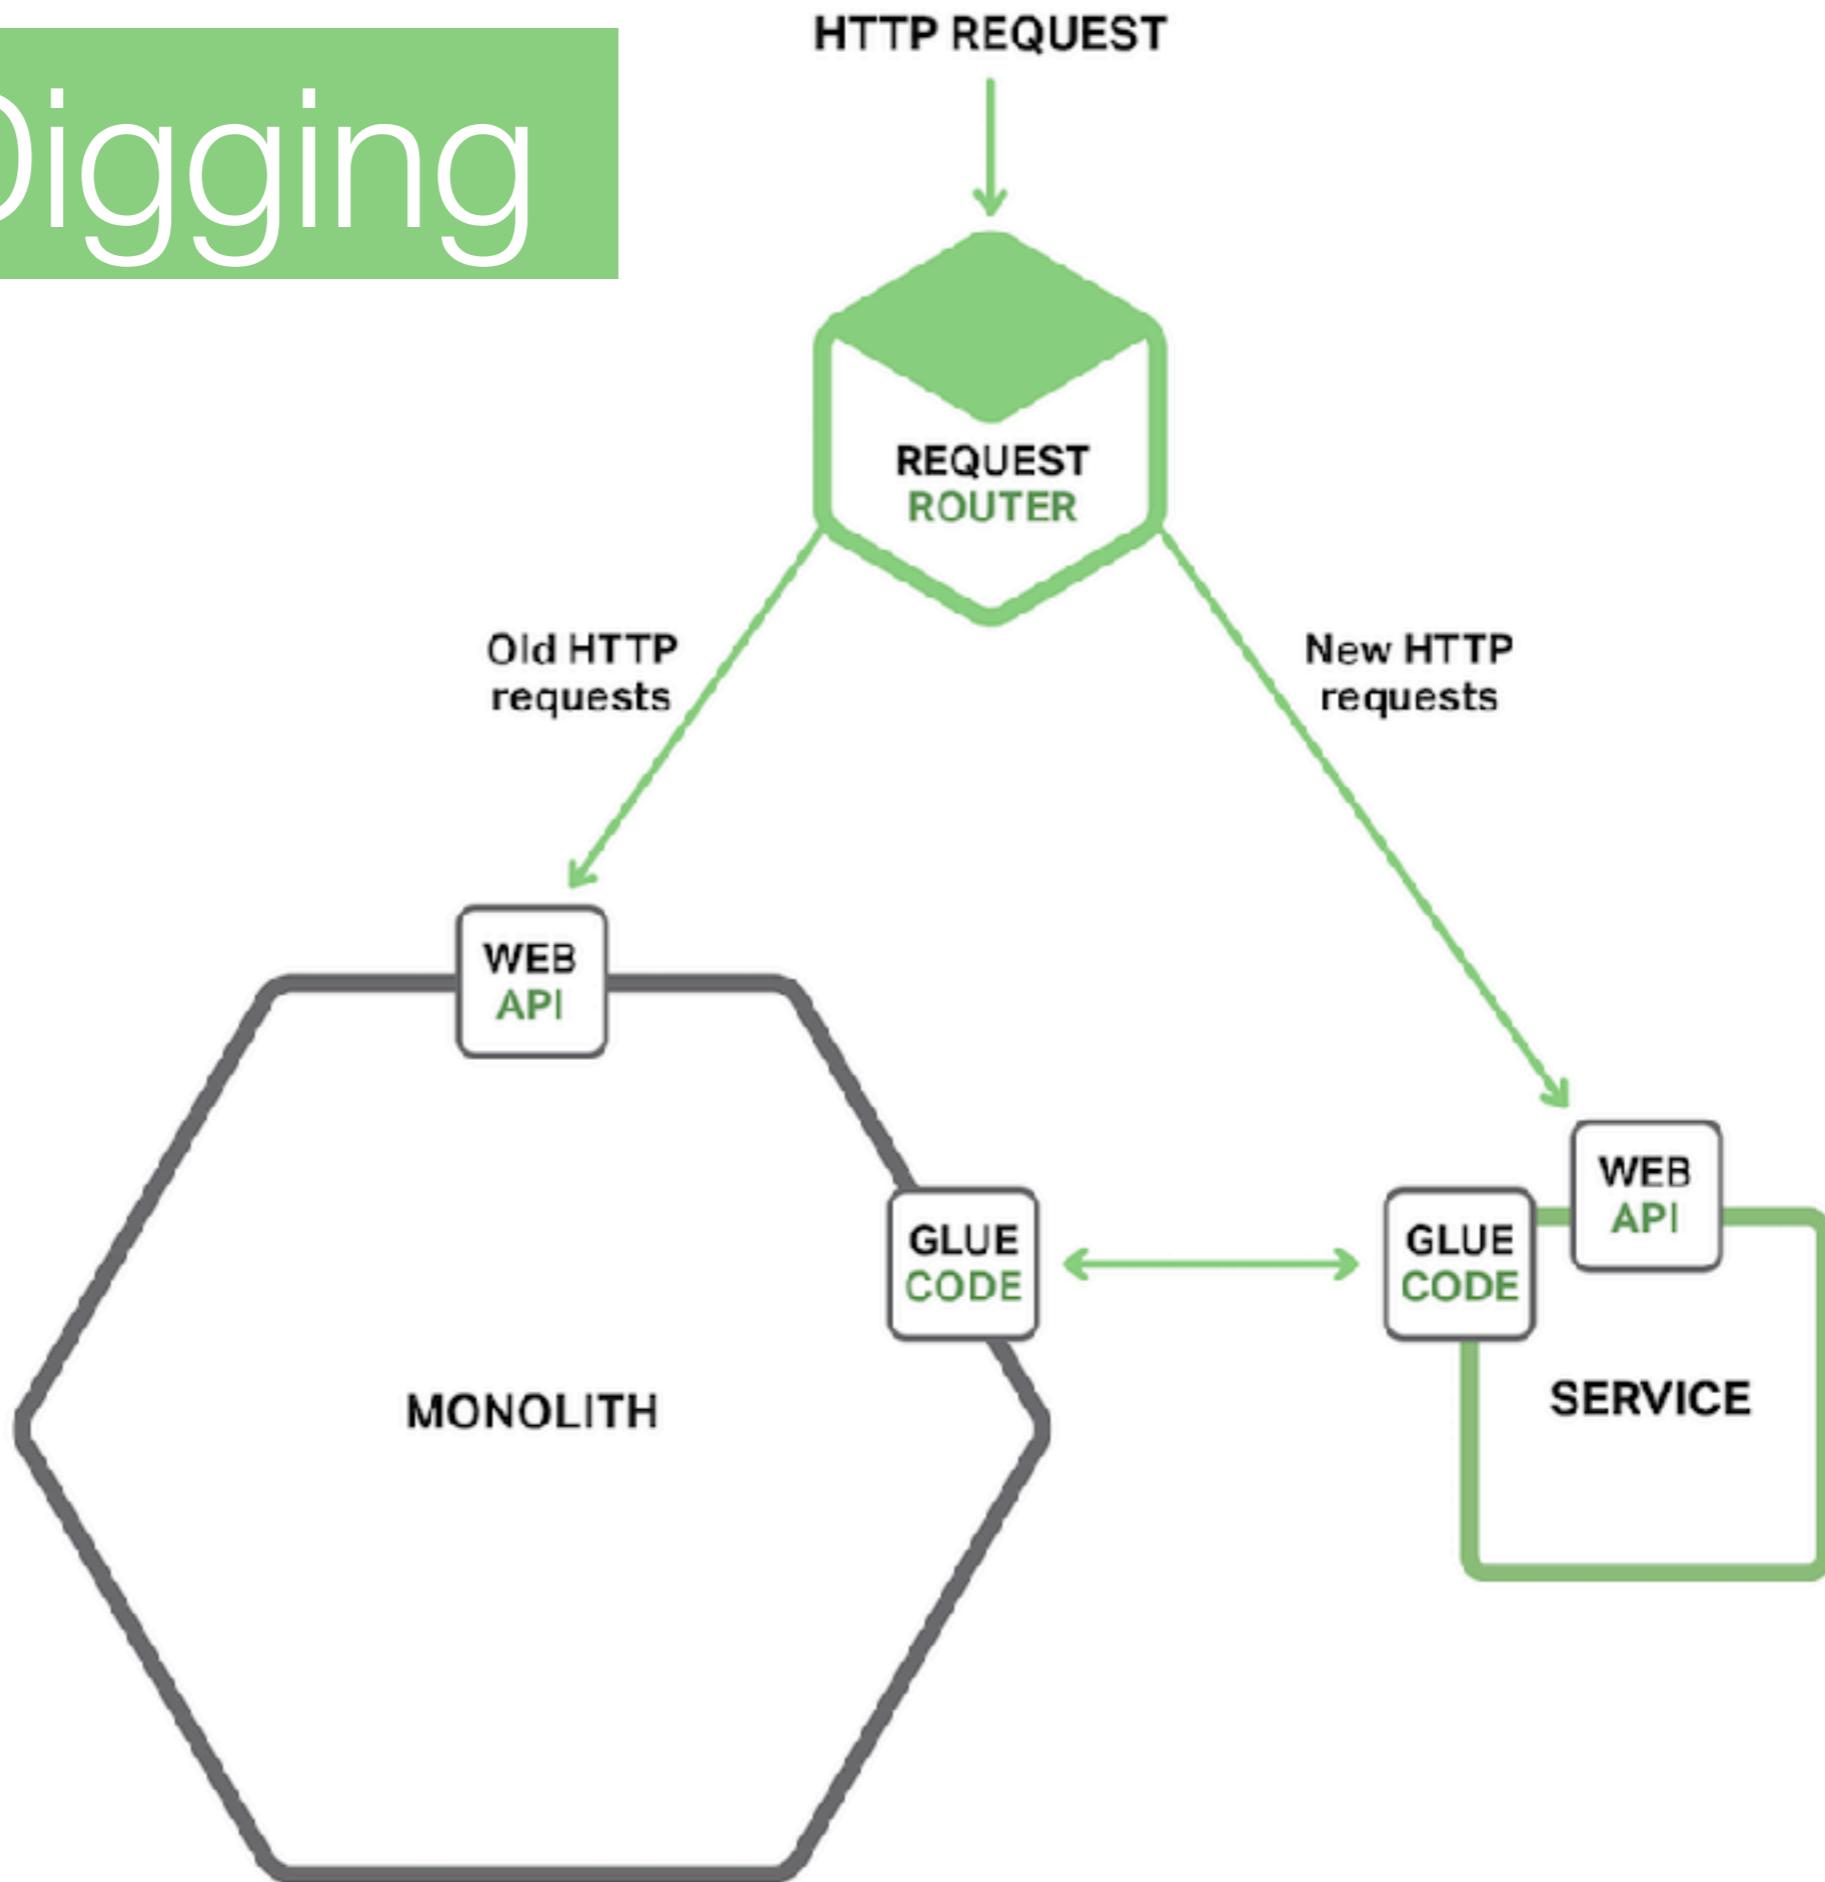


Incremental approach

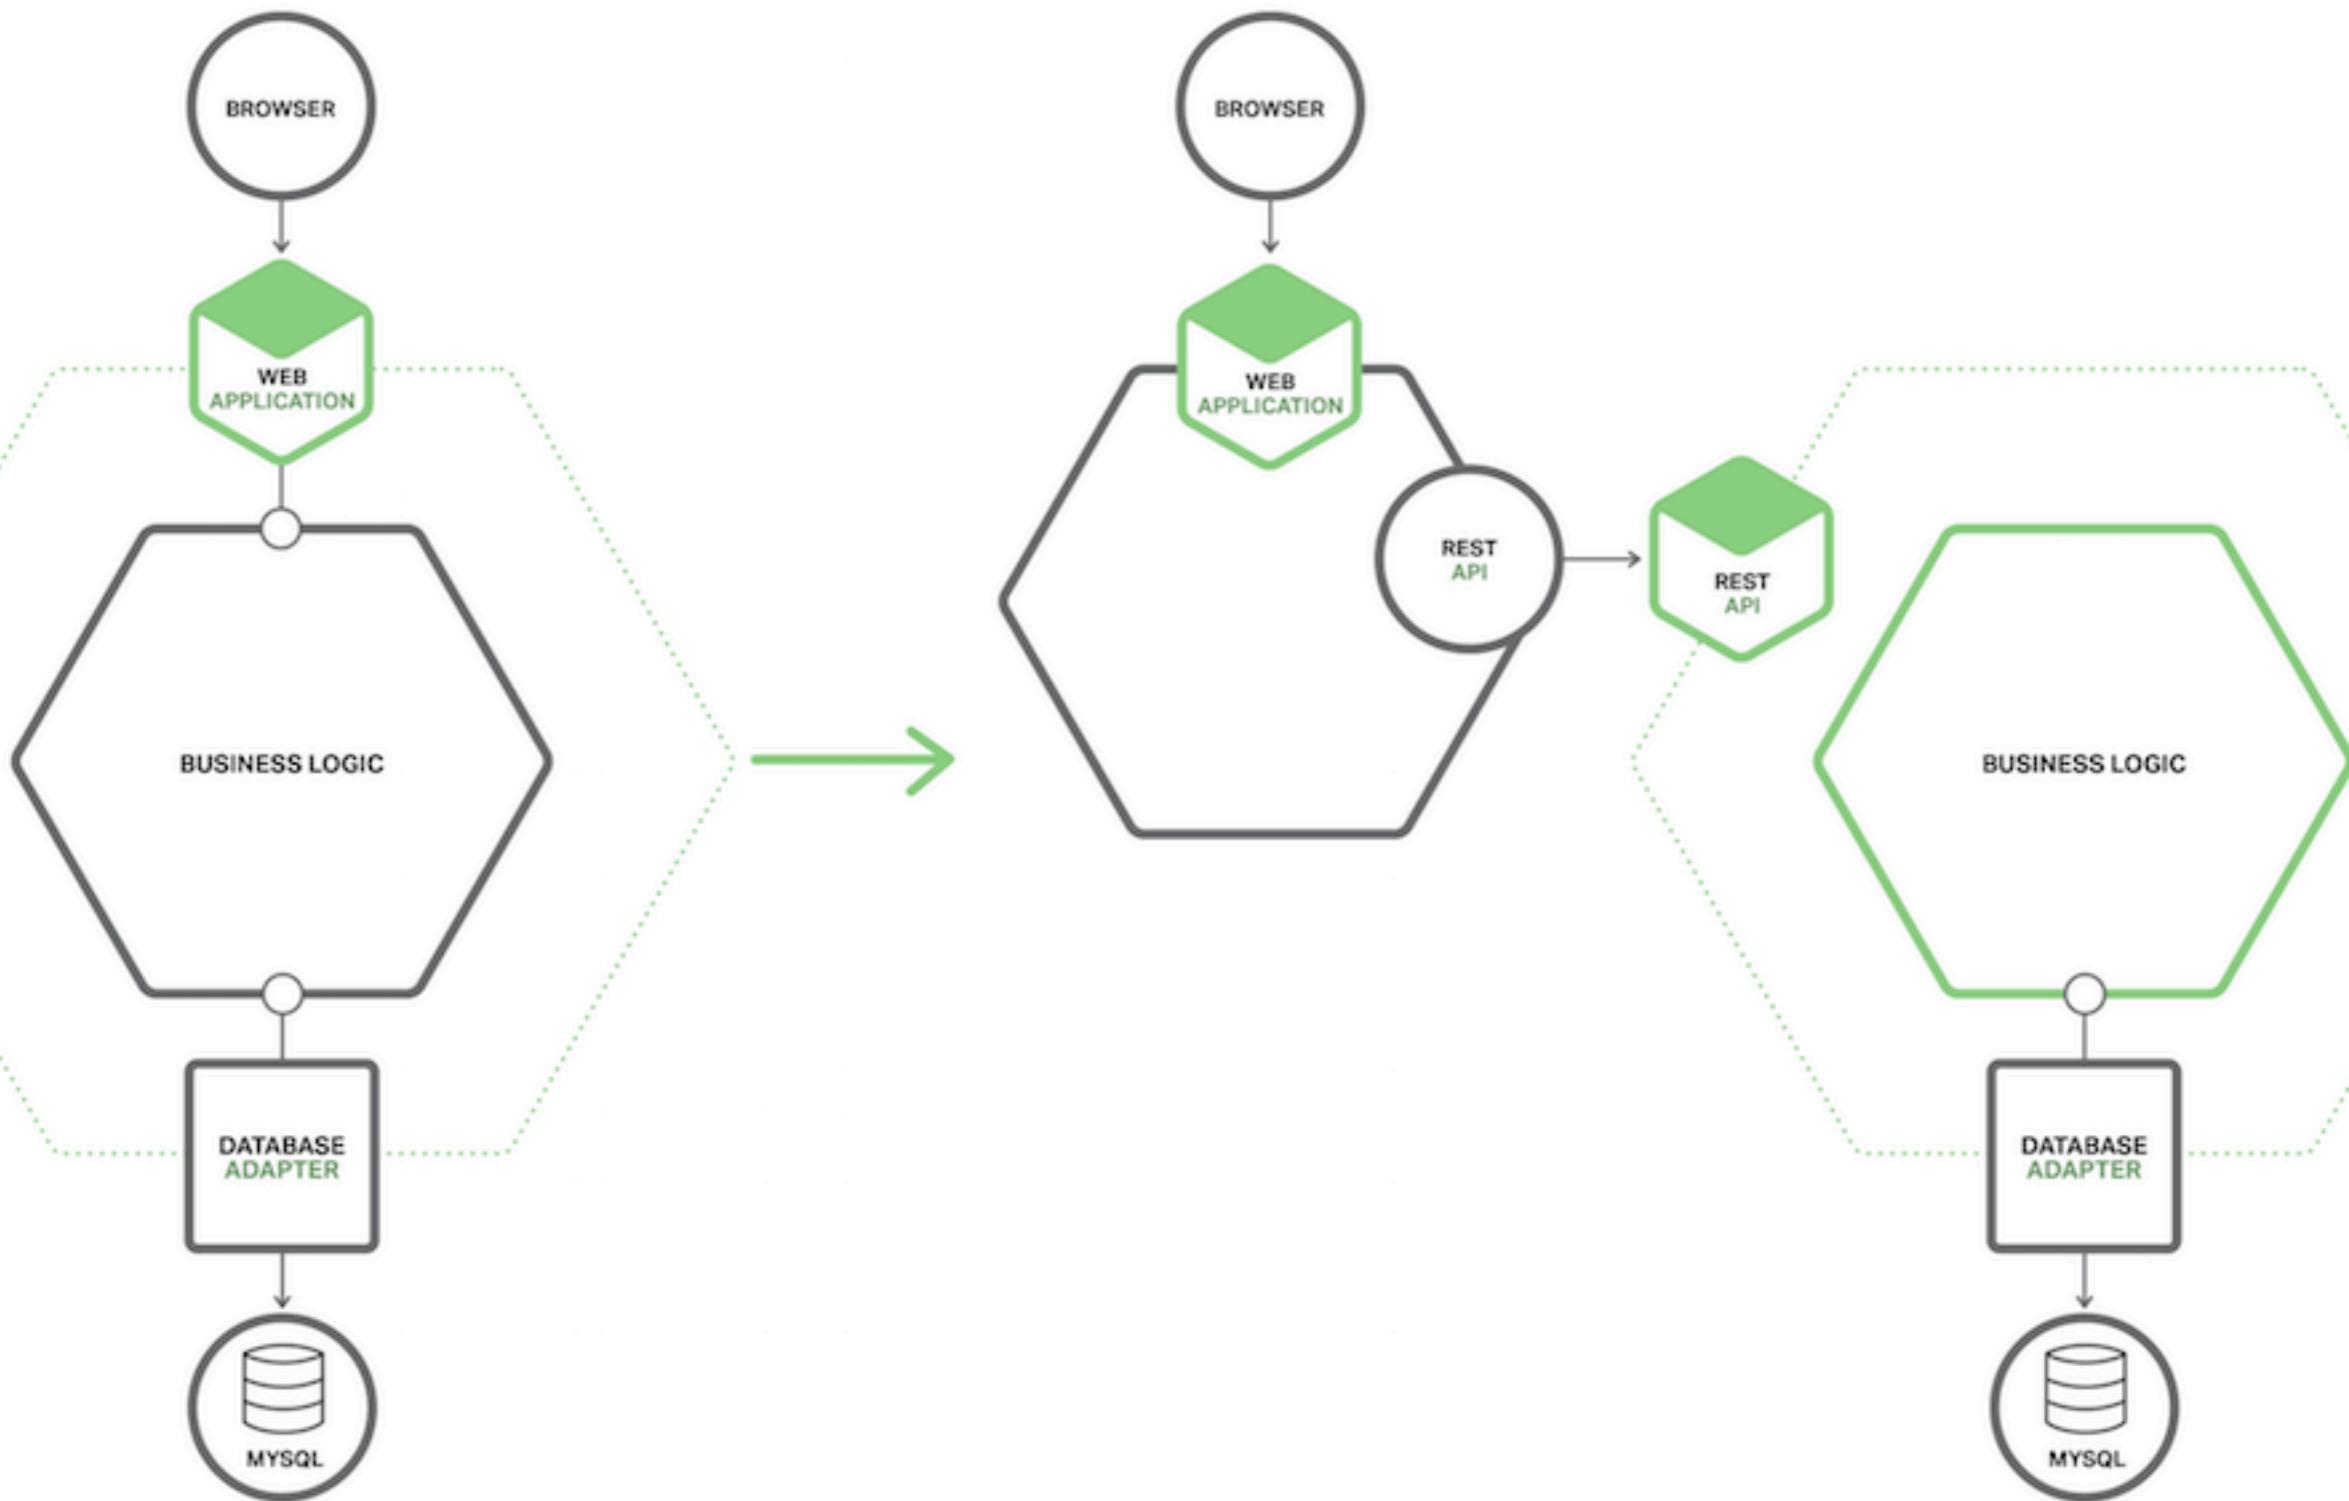
Migration



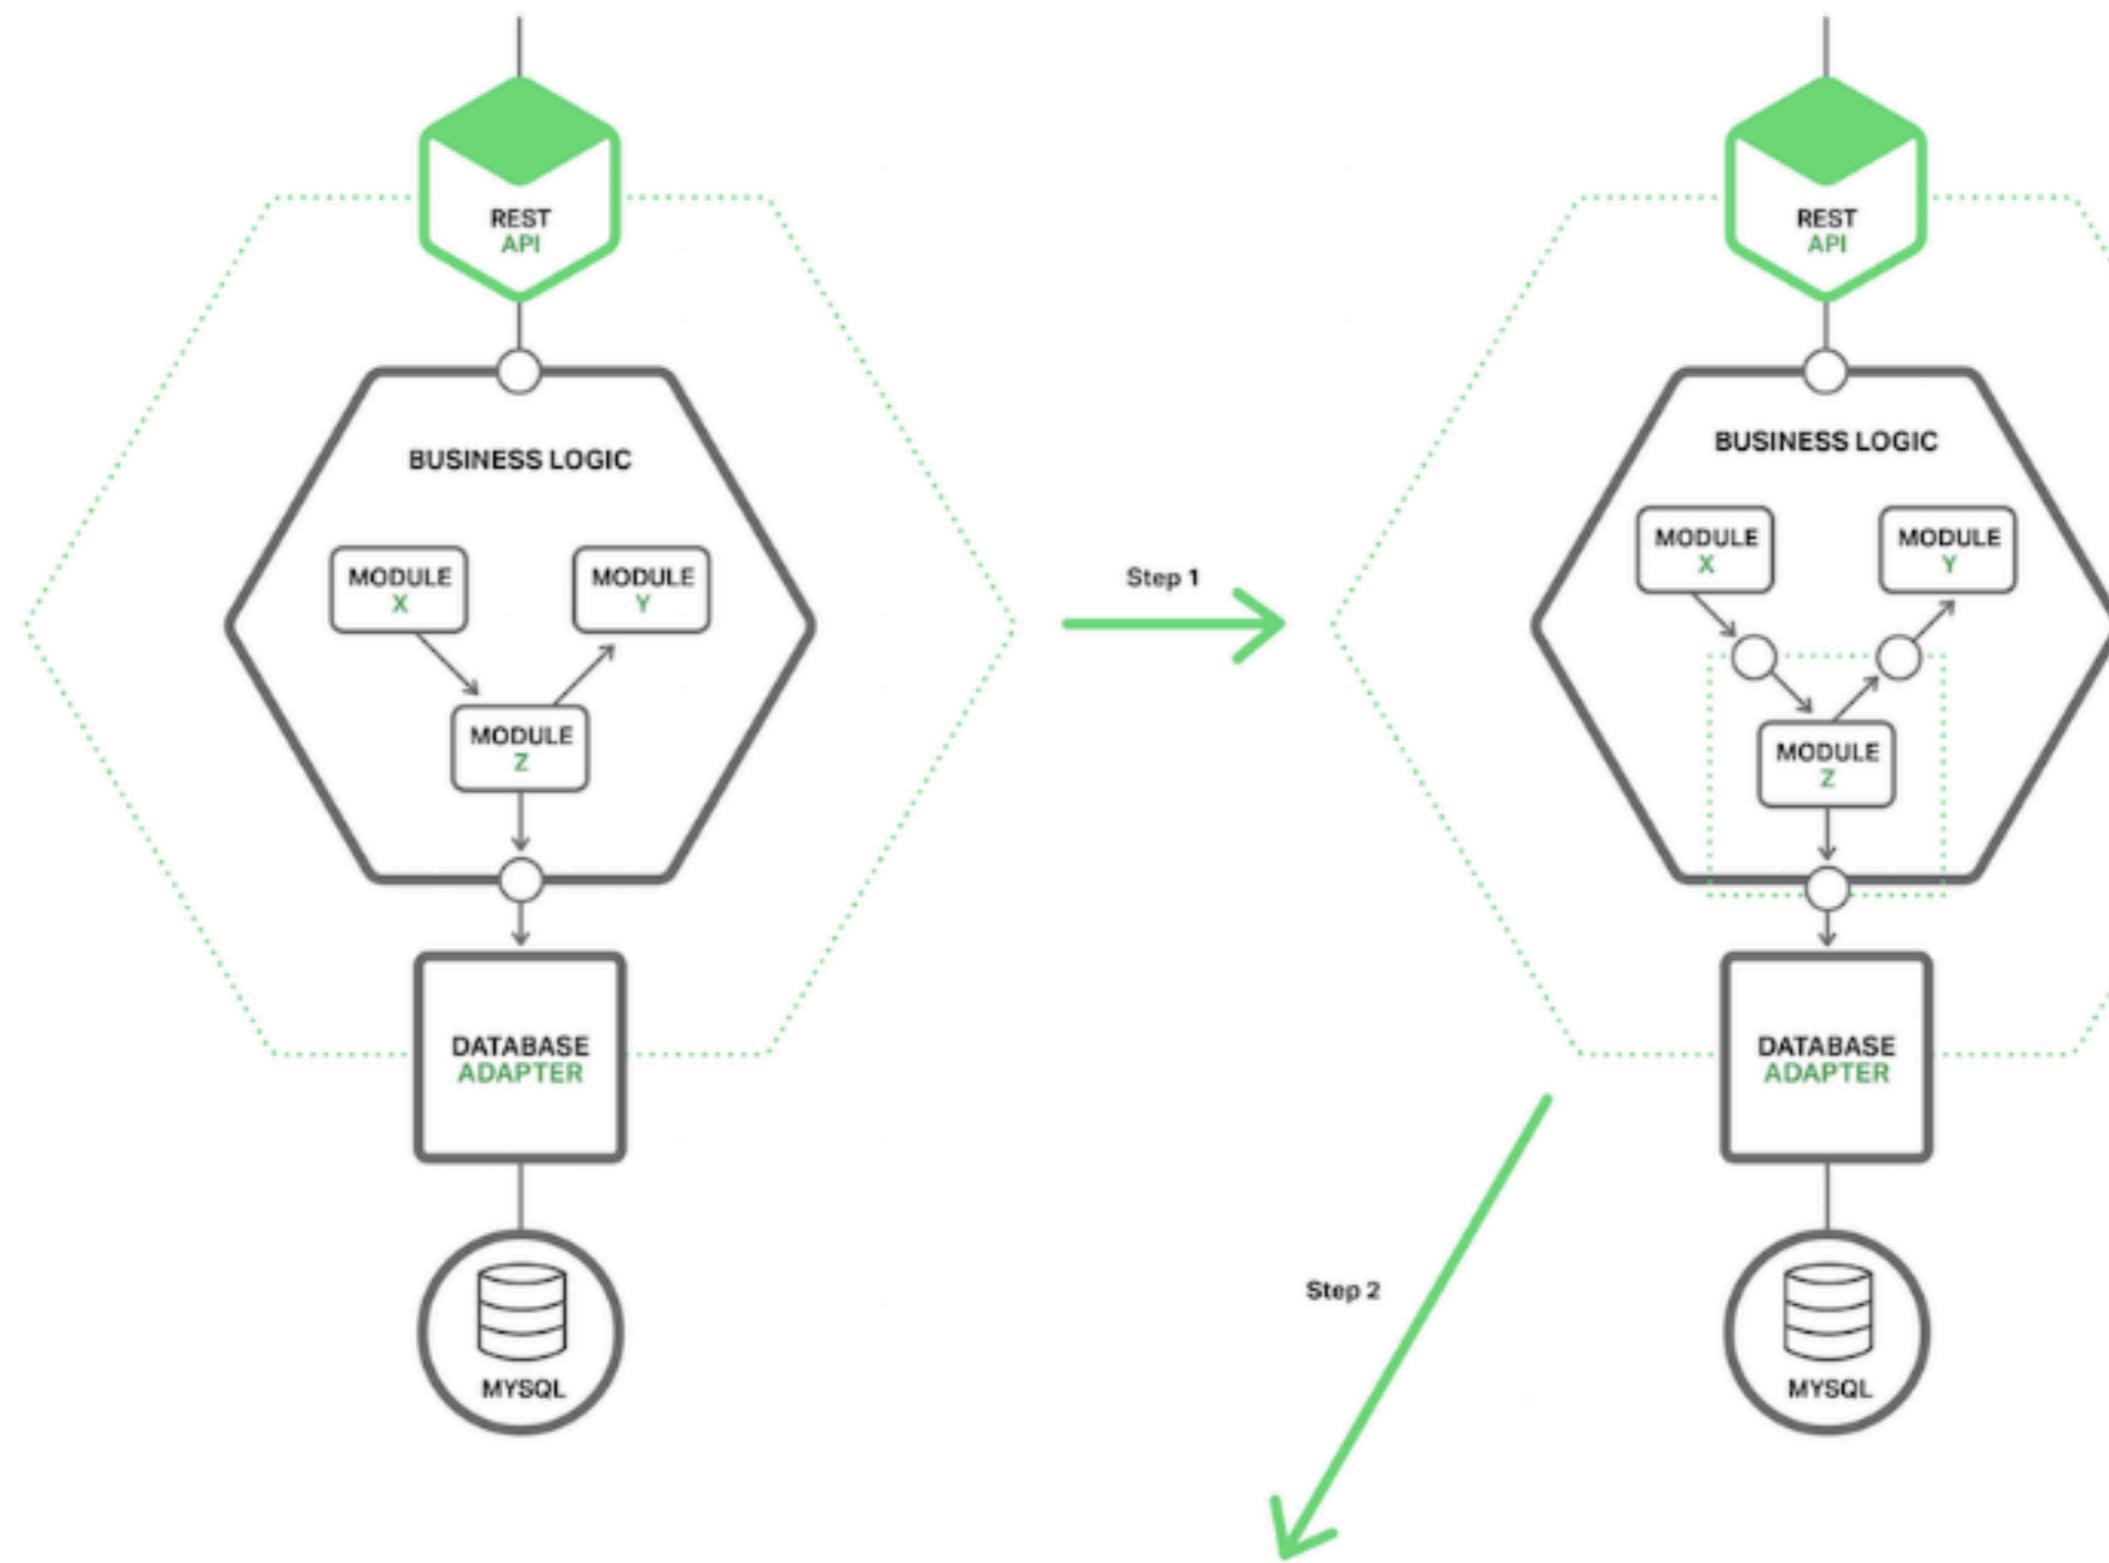
Stop Digging



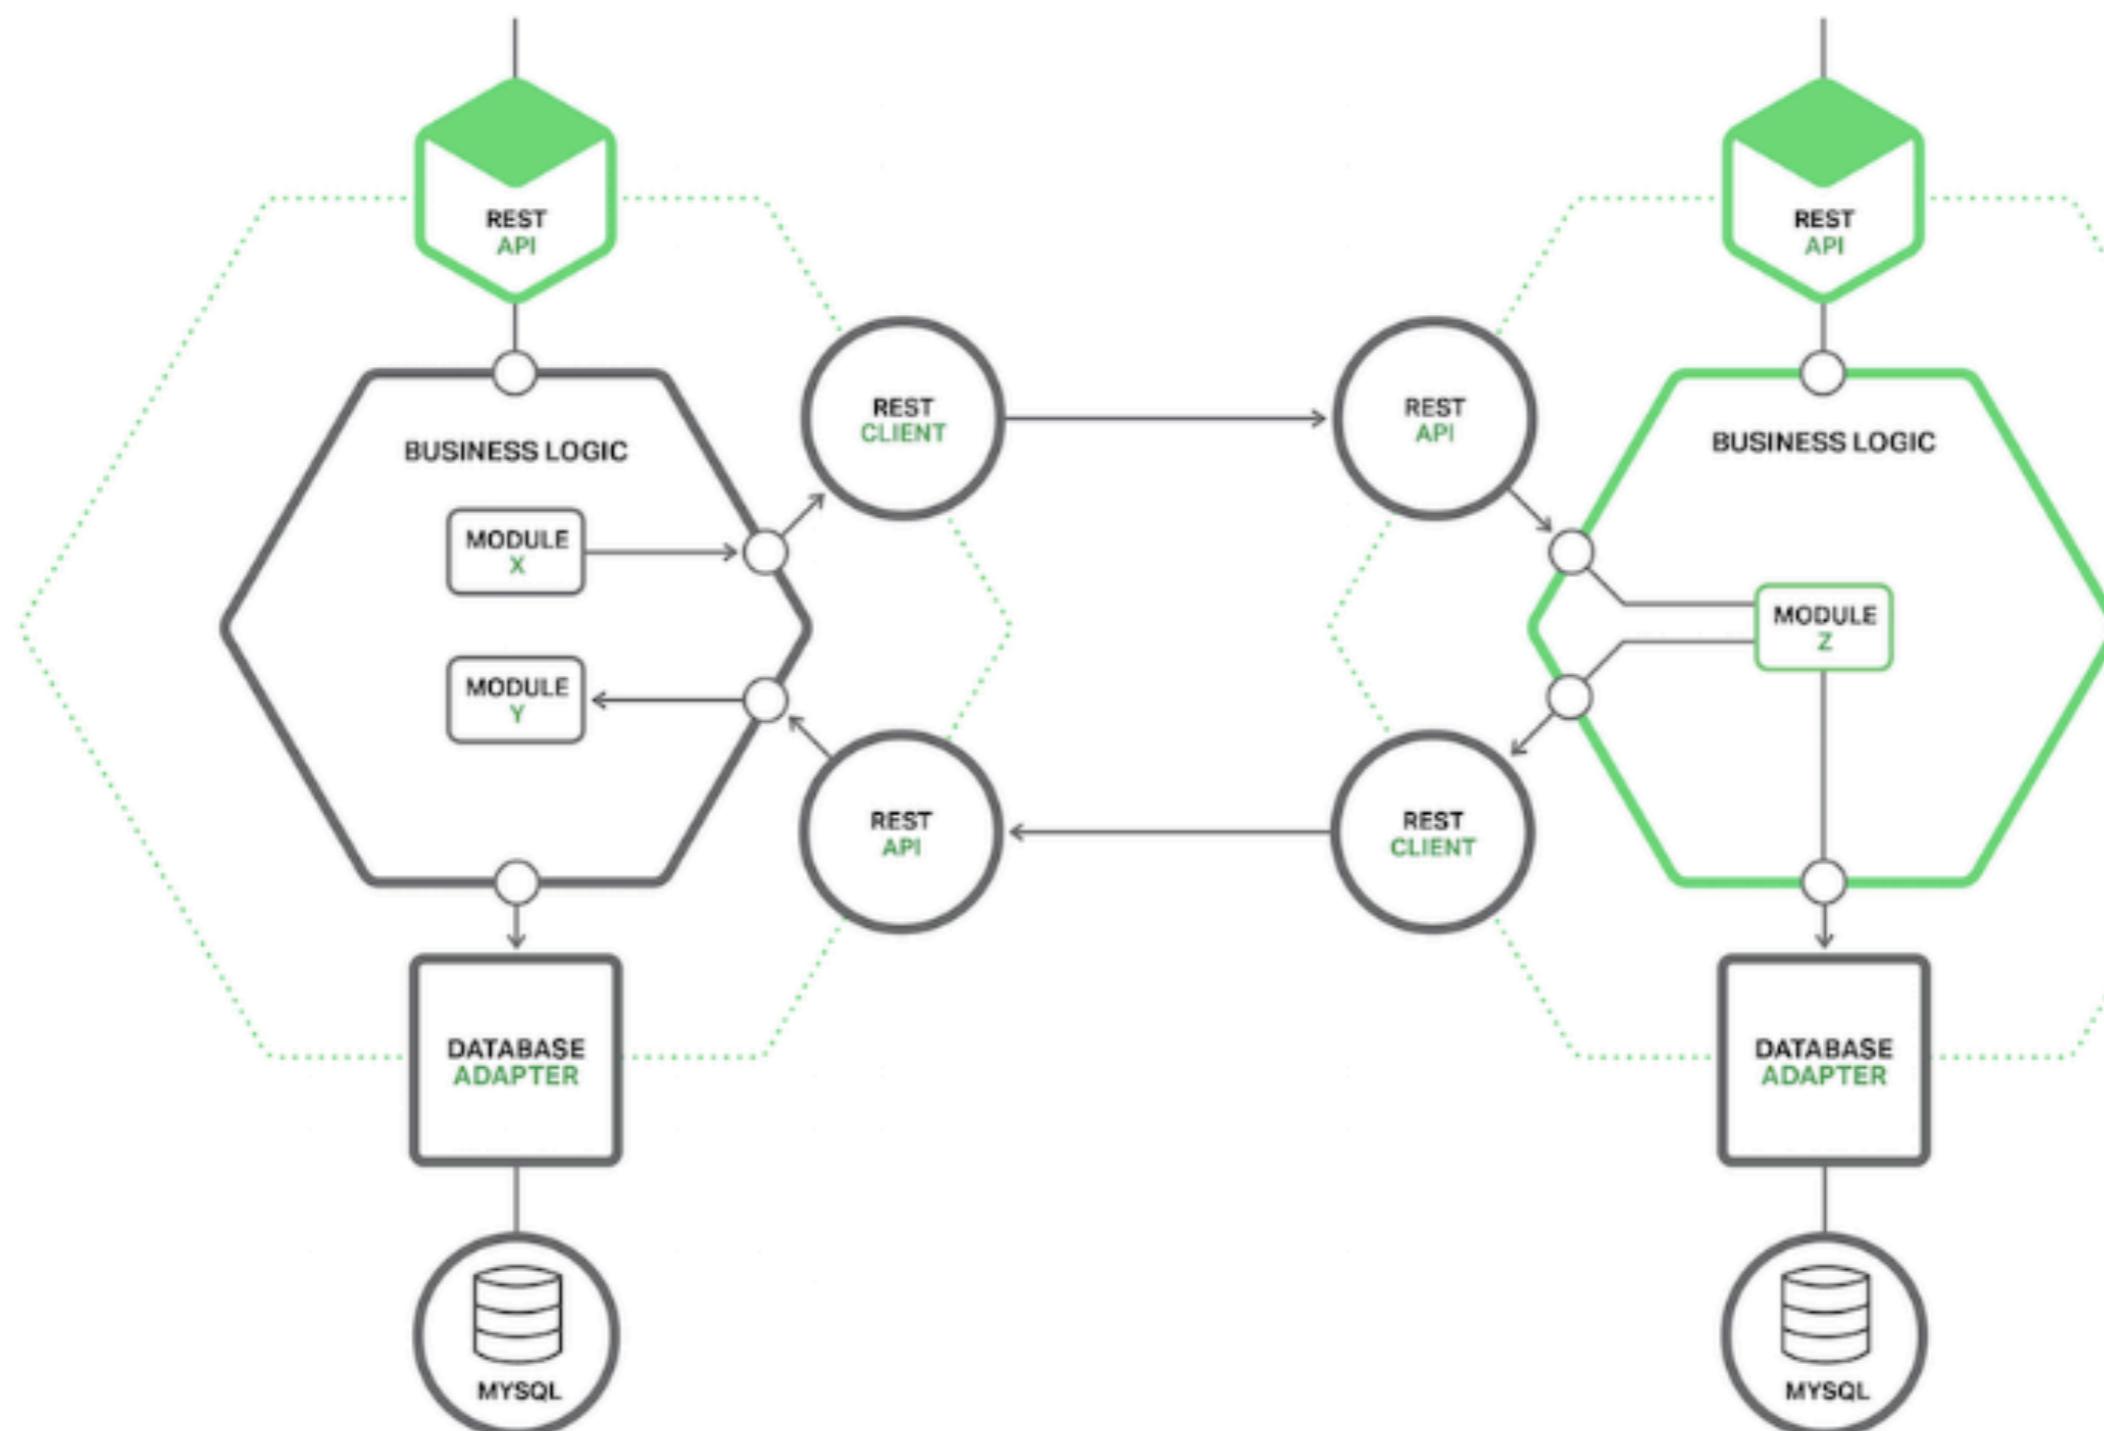
Split Frontend and Backend



Extract Service



Extract Service



How to Prioritize

Risk

Technology

Dependencies

Integrate with Monolith

- Monitor both for impact
- Monitor operations that talk to microservices
- Review and improve infrastructure
- Incrementally the monolithic will be converted

API Architecture Pattern

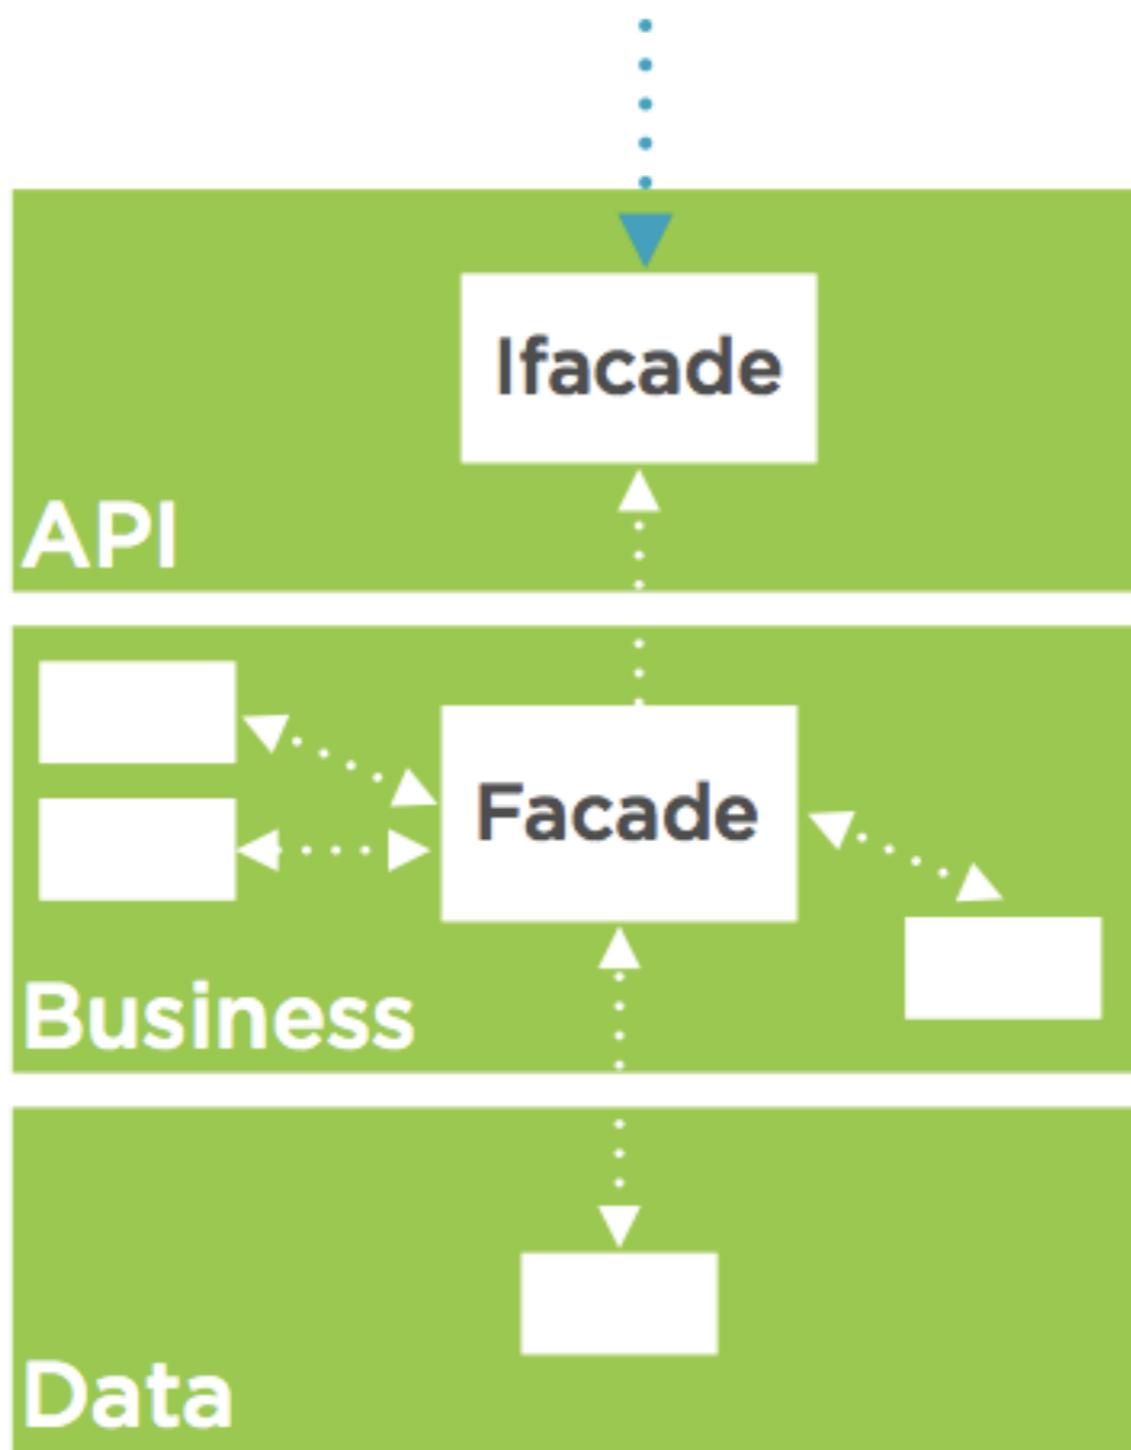
Design Pattern

Facade
design pattern

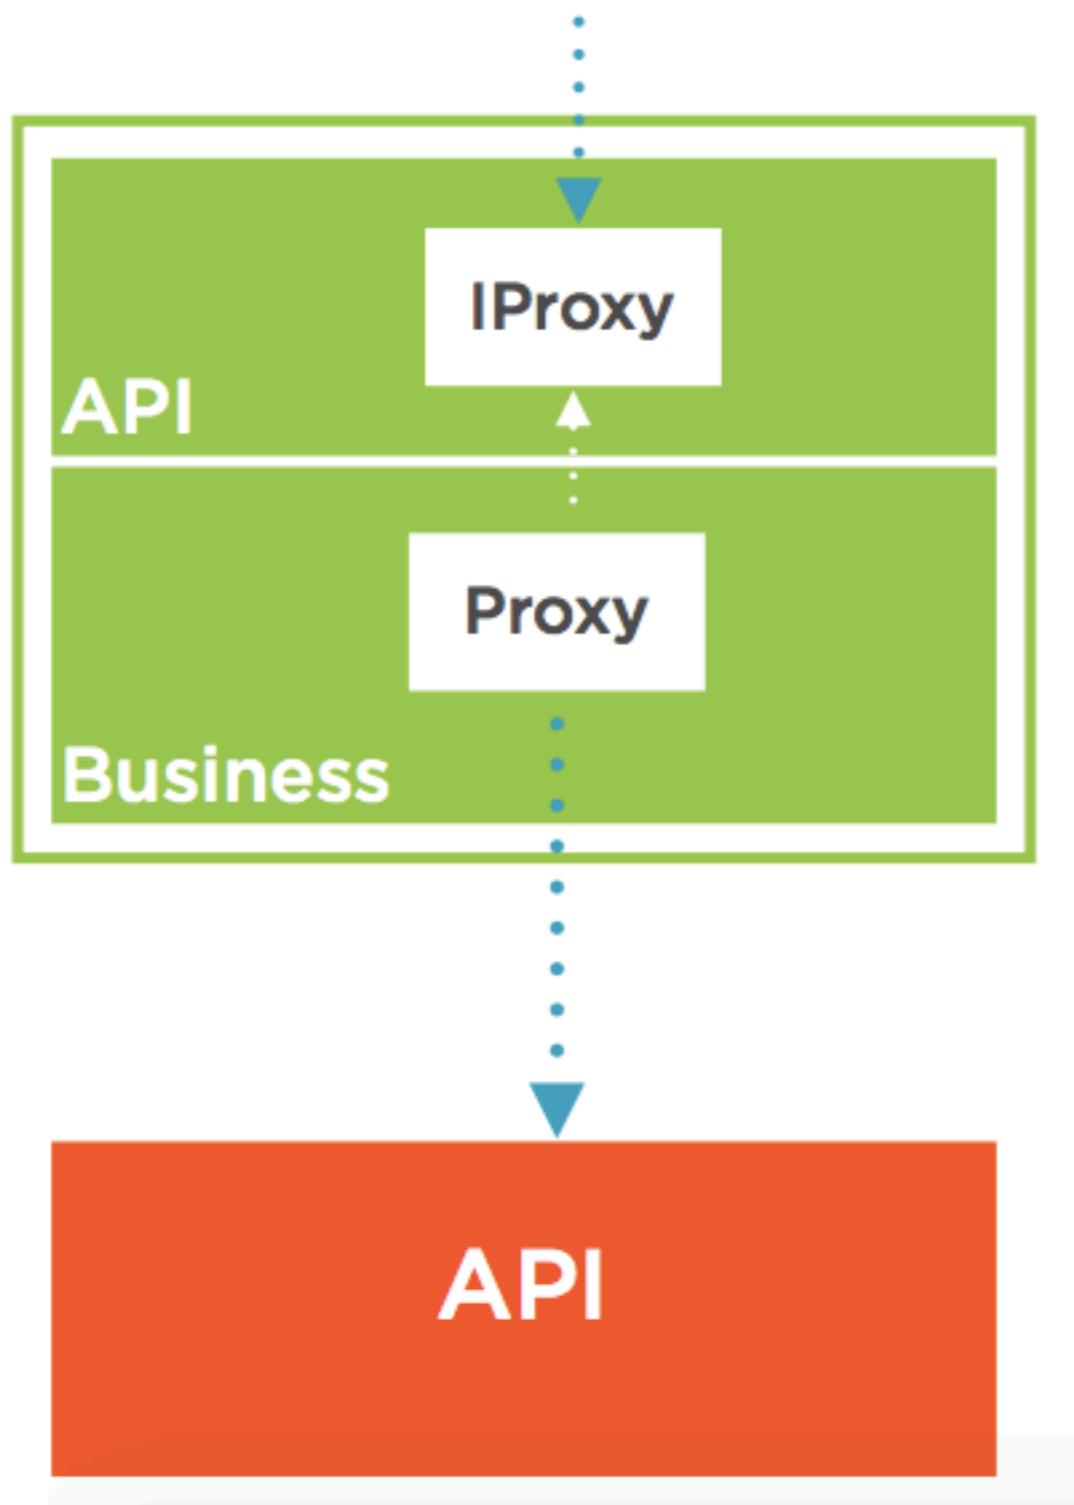
Proxy
design pattern

Stateless
design pattern

Facade Design Pattern



Proxy Design Pattern



- Simplification
- Transformation
- Security
- Validation

Stateless Design Pattern

- What are stateful services
Avoid for microservices
- Stateless Service Pattern
 - No state information maintained
 - Clients maintain state
 - State sent as part of request
- Advantages for microservice architecture
Scalability, performance and availability

Increase Network Traffic

Database Migration

Database Migration

- Avoid shared databases
- Split databases using seams
Relate tables to code seams
- Supporting the existing application
Data layer that connects to multiple database
- Tables that link across seams
API calls that can fetch that data for a relationship

Database Migration

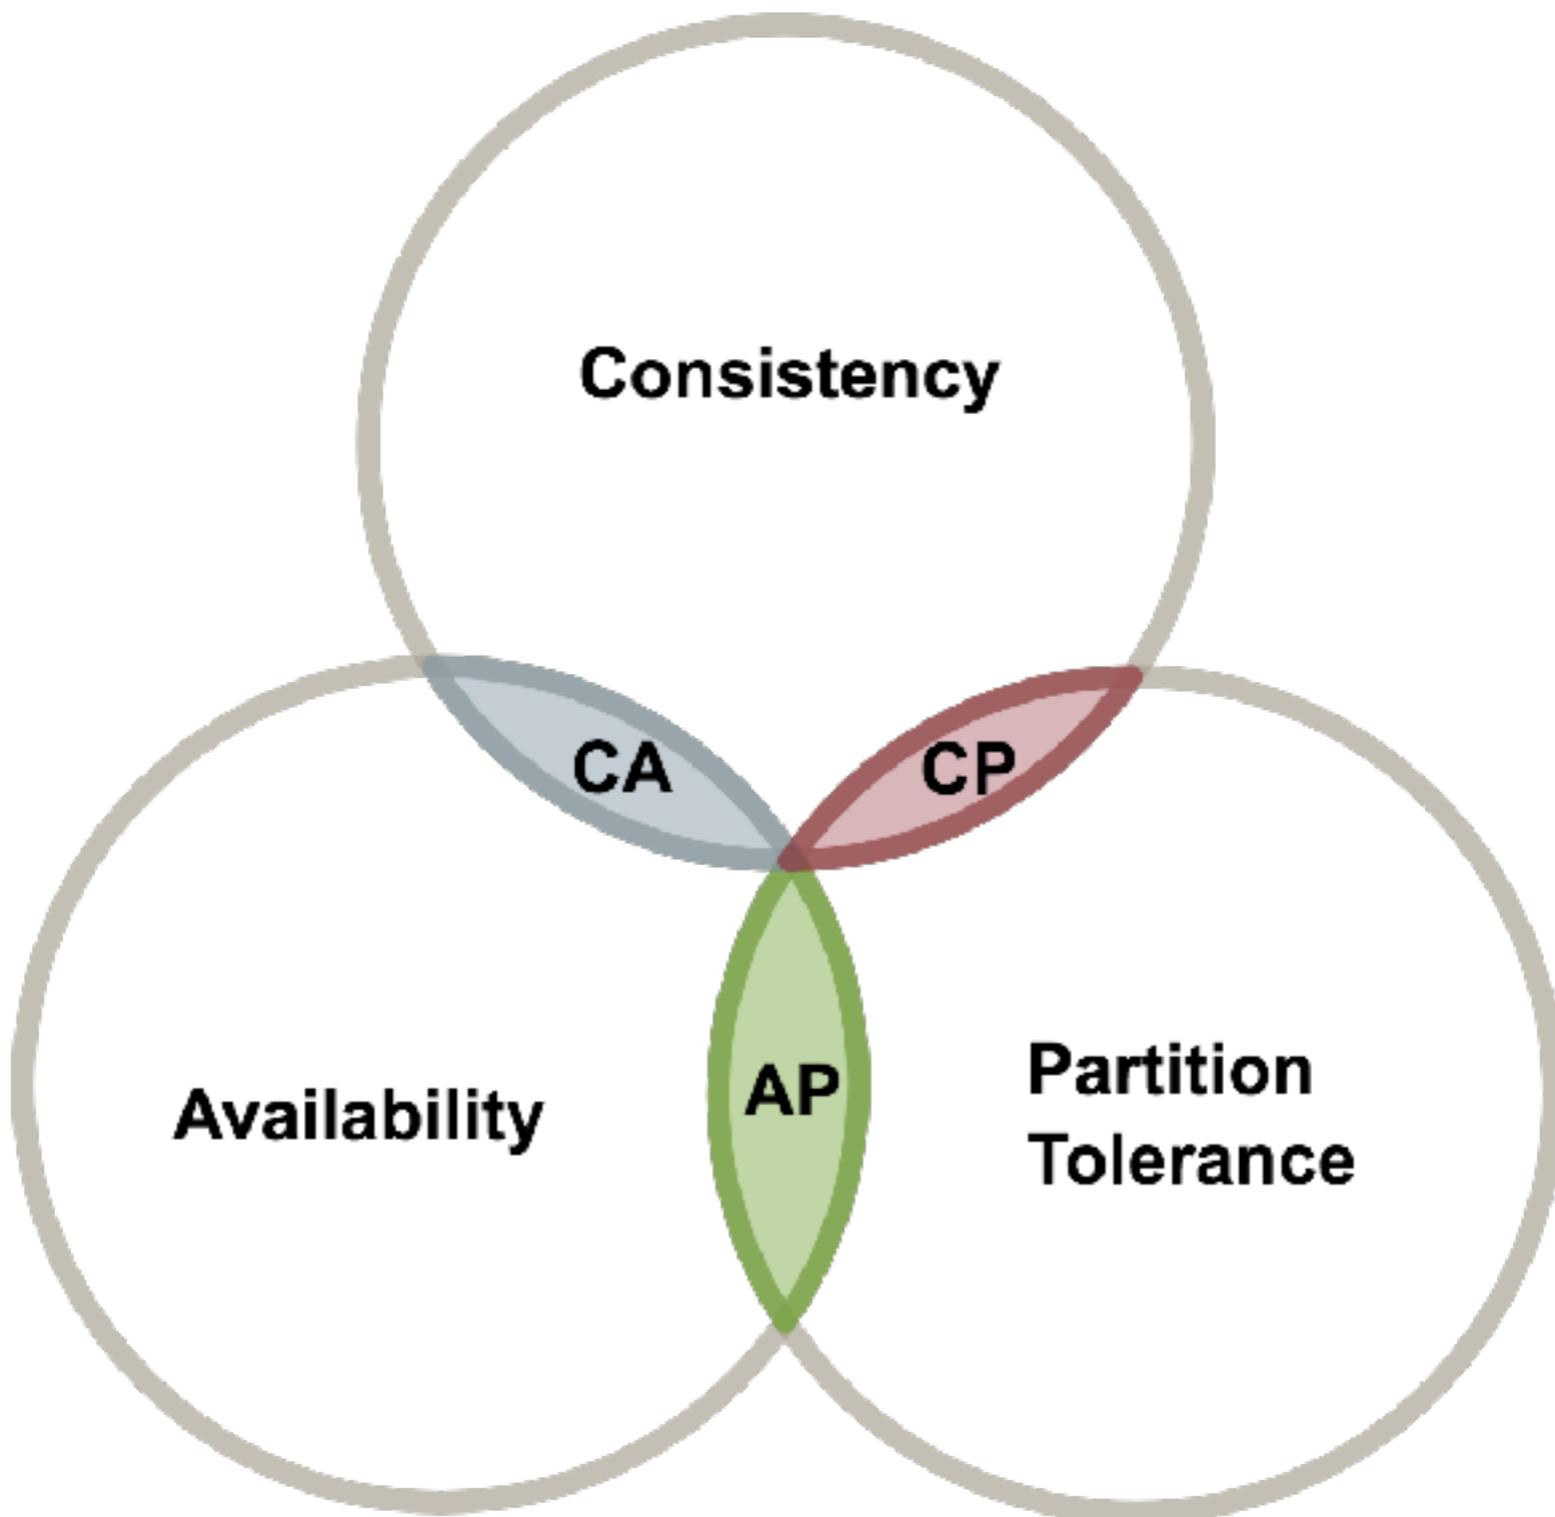
- Refactor database into multiple databases
- Data referential integrity
- Static data tables
- Shared data

Data Consistency

Data Consistency

- Transactions are the traditional approach
- Monolithic system transactions
Single database which is the single truth
- Microservice transactions
 - Distributed architecture
 - Distributed data
 - Distributed transactions
- CAP theorem
 - Network failure will happen
 - Data availability or data consistency?

CAP Theorem



Options

- Traditional ACID transactions
Atomicity, consistency, isolation and durability
- Two phase commit pattern
 - ACID is mandatory
 - CAP Theorem: Choosing consistency
- Saga Pattern
Trading atomicity for availability and consistency
- Eventual consistency pattern
 - Compromise ACID
 - CAP Theorem: Choosing availability

Transactions

ACID

- Atomic
Either succeeds completely or fails completely
- Consistency
Constraints of underlying datastore are enforced
- Isolation
Can't be read by other transactions until it is in a specific state based on isolation rules
- Durability
Once saved, guaranteed to be in the datastore until modified

Transactions

- Transactions ensure data integrity
- Transactions are simple in monolithic applications
- Transactions spanning microservices are complex
 - Complex to observe
 - Complex to problem solve
 - Complex to rollback

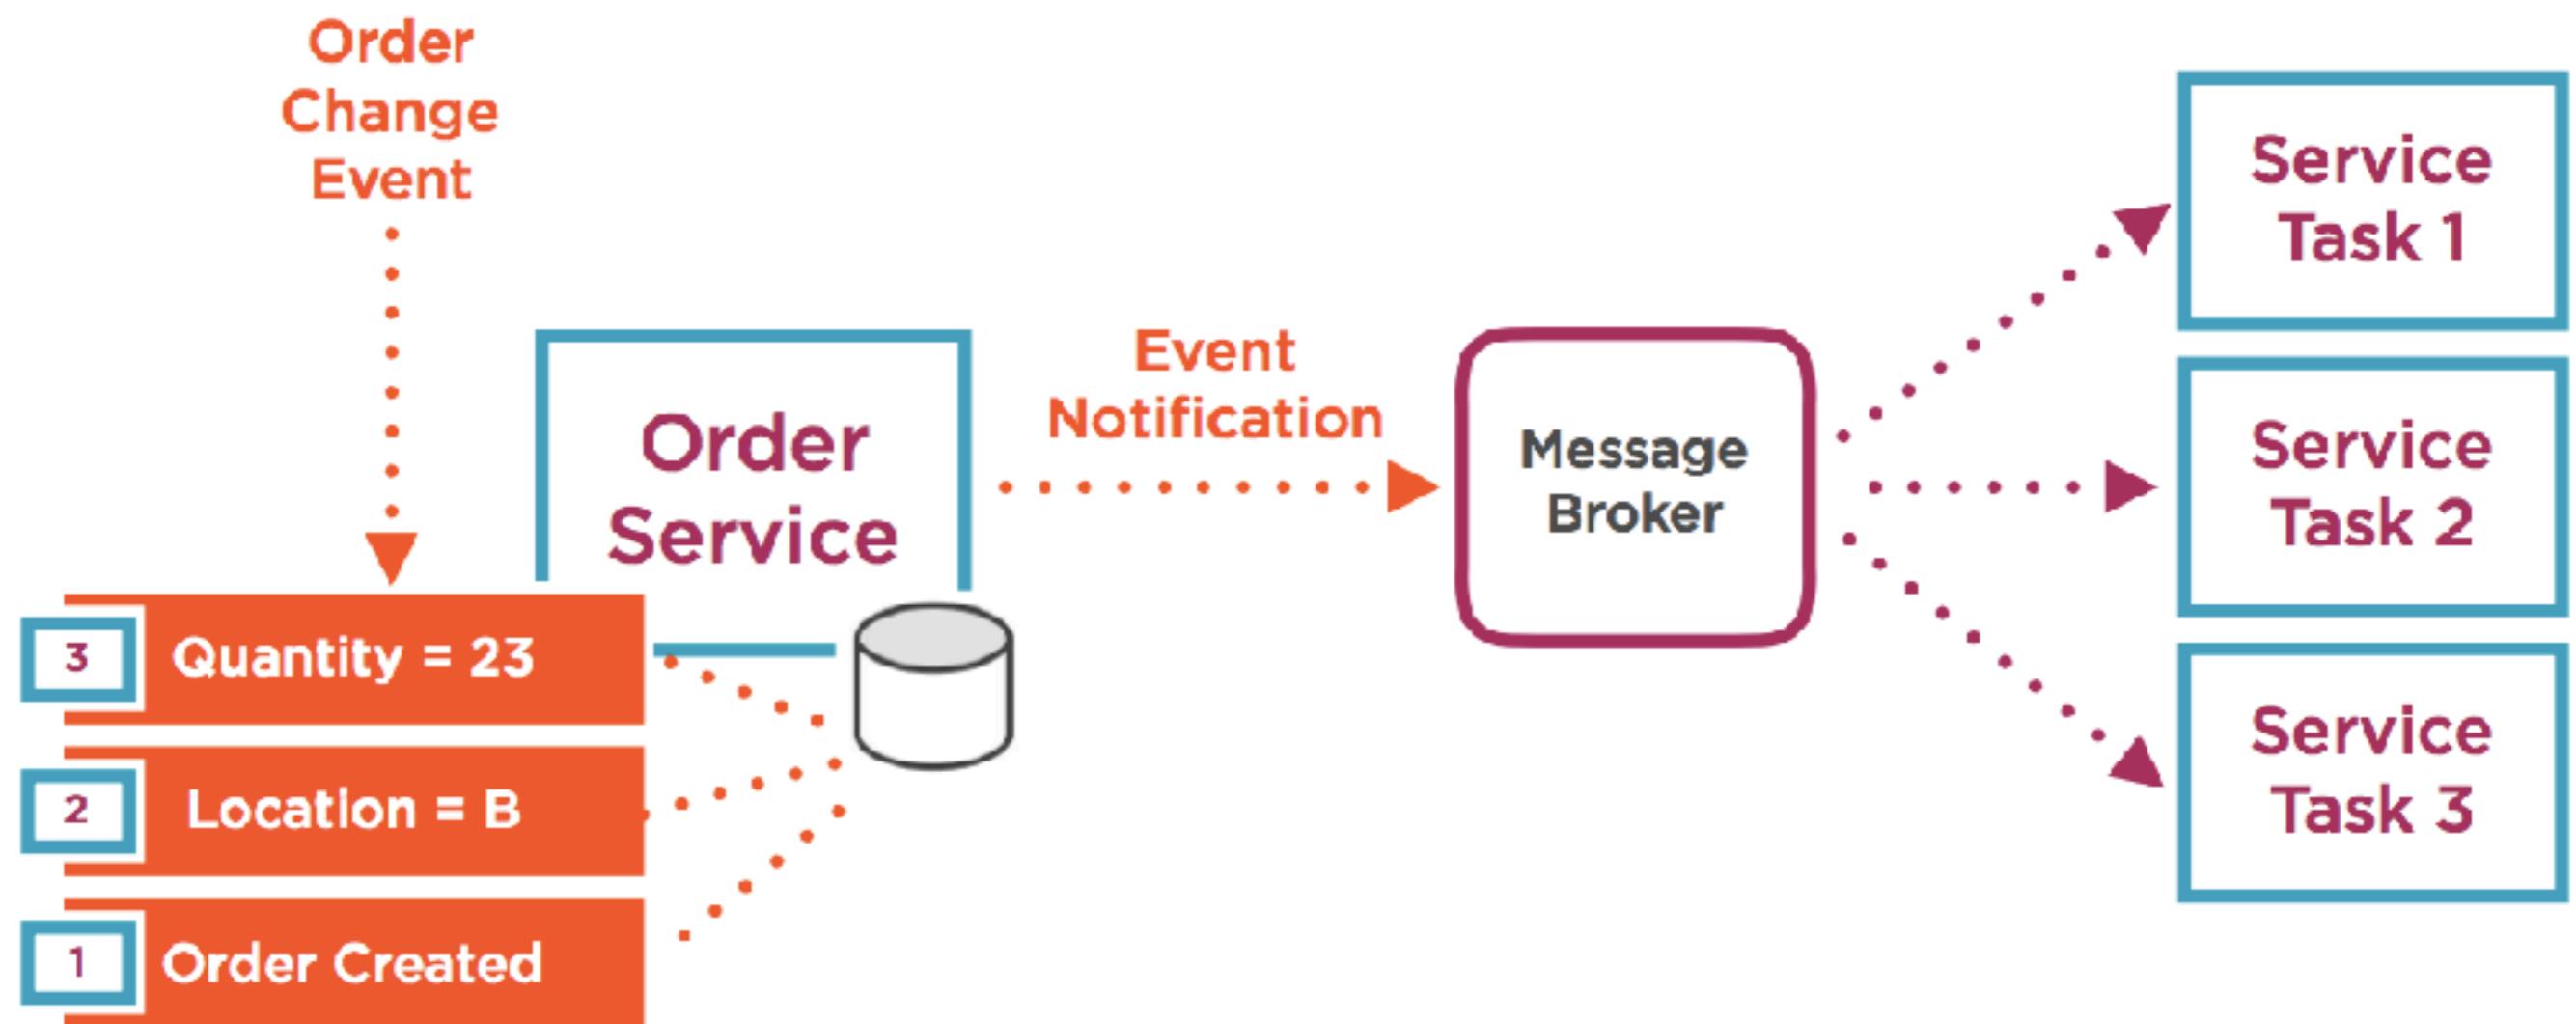
No ACID

Options for Failed

- Try again later
- Abort entire transaction
- Use a transaction manager (Two phase commit)
- Disadvantage of transaction manager
 - Reliance on transaction manager
 - Delay in processing
 - Potential bottleneck
 - Complex to implement

Patterns for Database Design

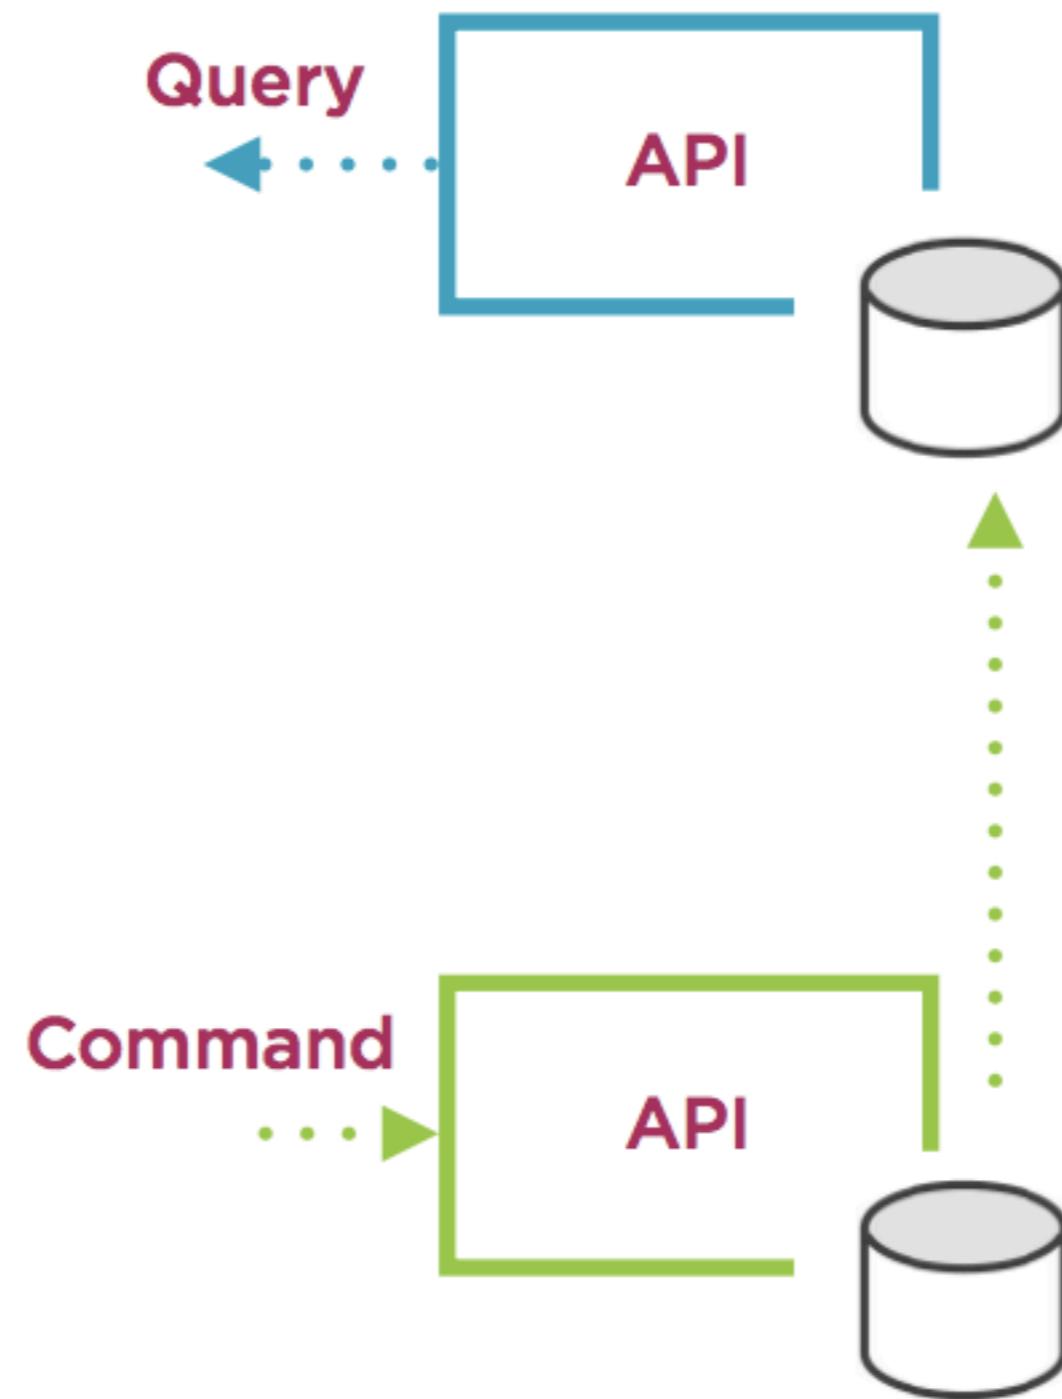
Event Sourcing



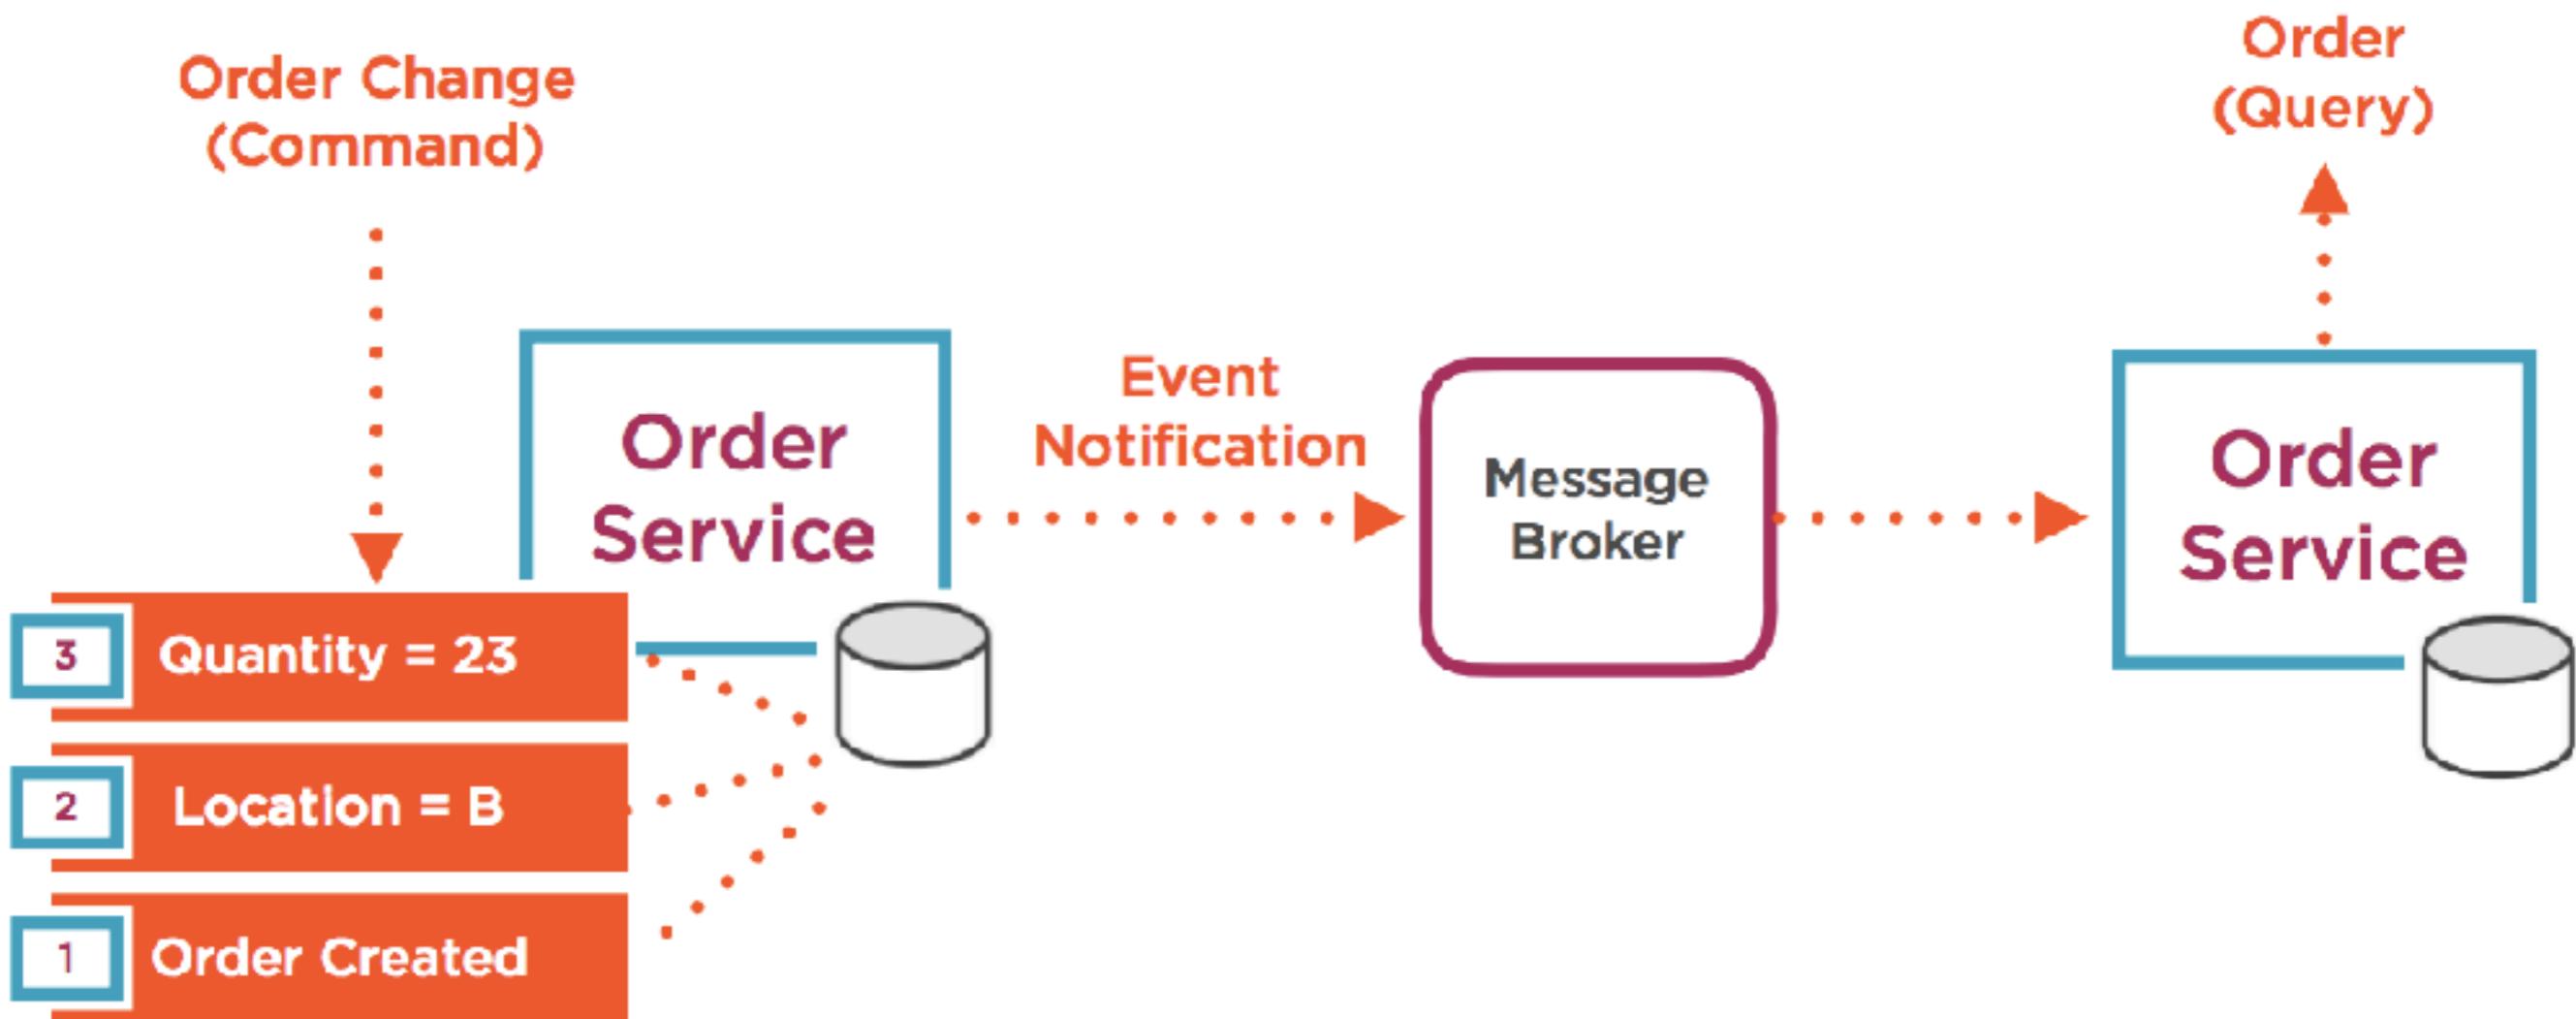
CQRS

- CQRS
 - Command models and or services
 - Query models and or services
- Why
 - Separation of responsibilities
 - Event notifications handled by command - Reporting/ functions handled by query
 - Separation of technologies - Service and storage
- Challenges
 - Command and query database syncing

CQRS

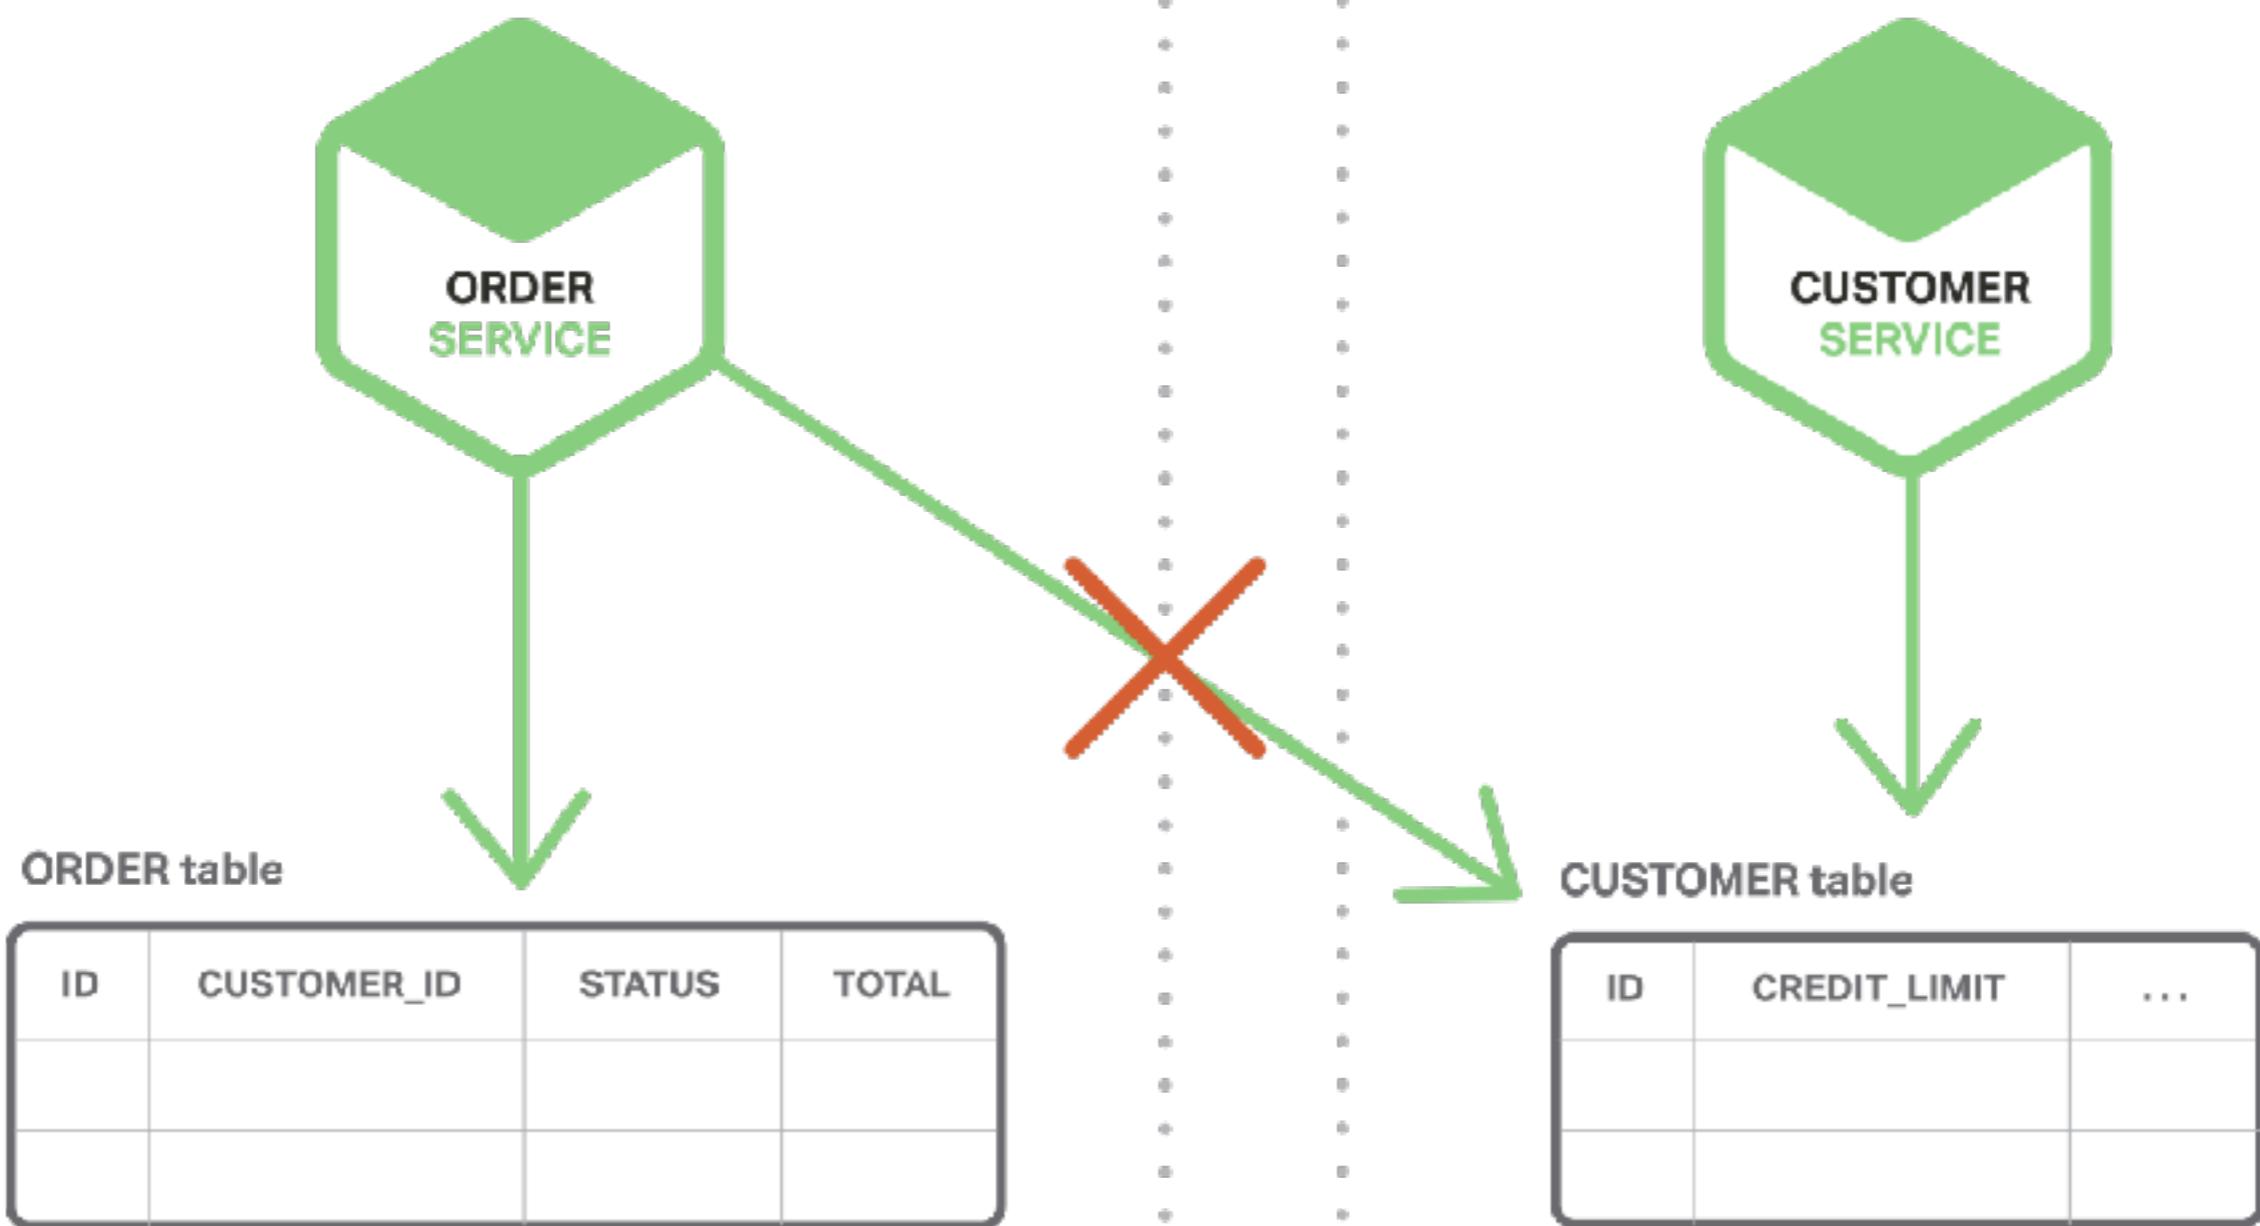


Event Sourcing and CQRS



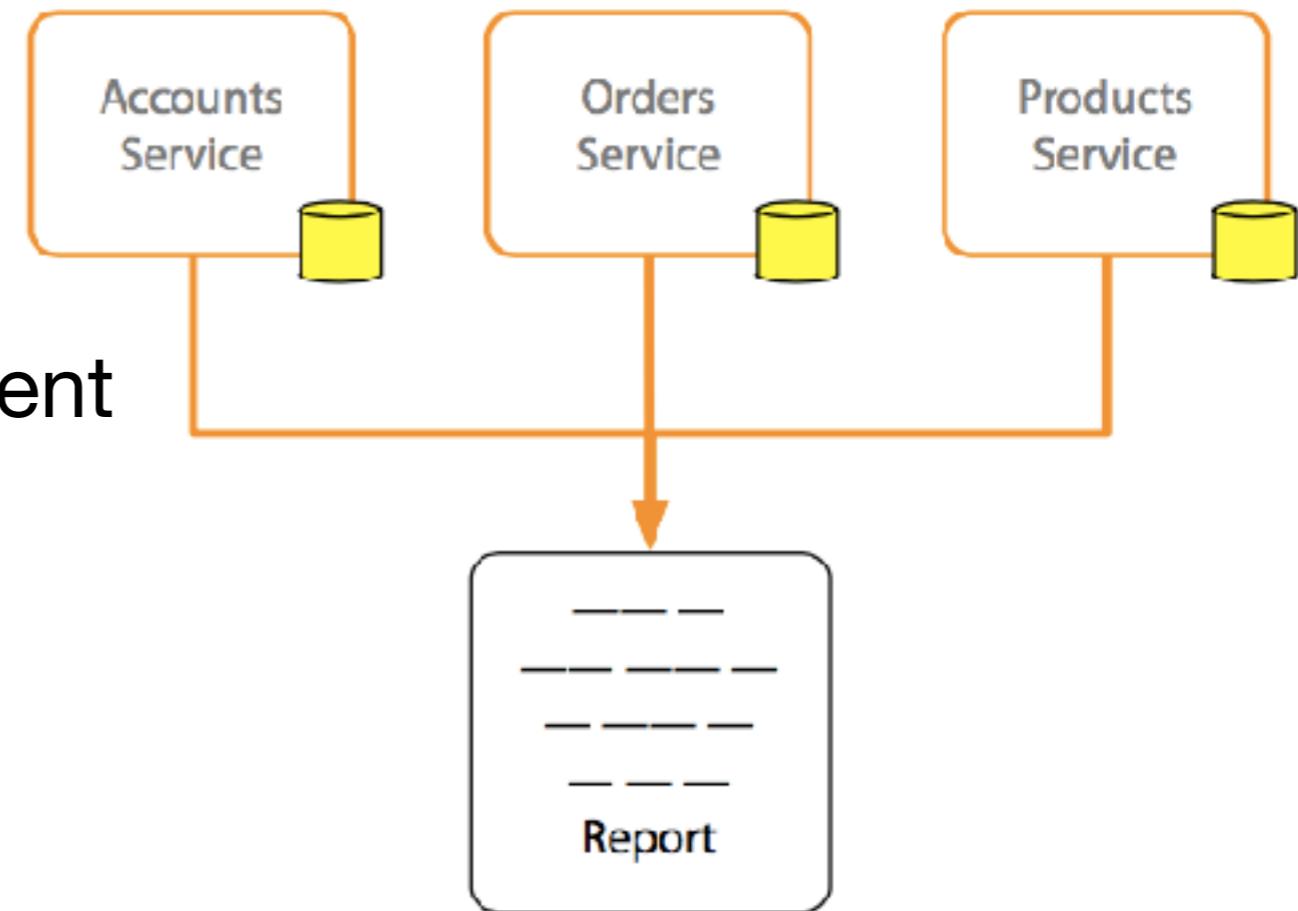
Reporting

No Shared Data



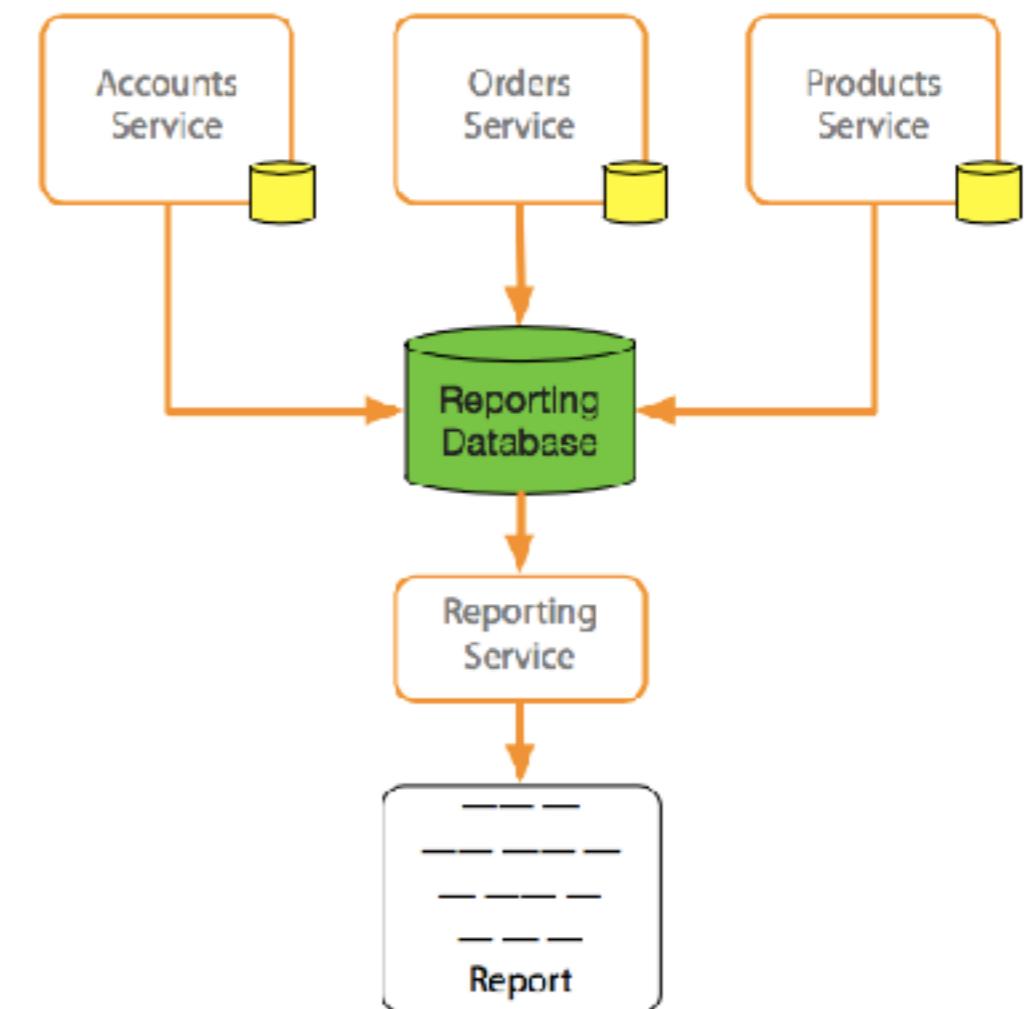
Complicating Report

- Data split across microservices
- No central database
- Joining data across databases
- Slower reporting
- Complicate report development



Possible Solution

- Service calls for report data
- Data dumps (Not good for real-time application)
- Consolidation environment



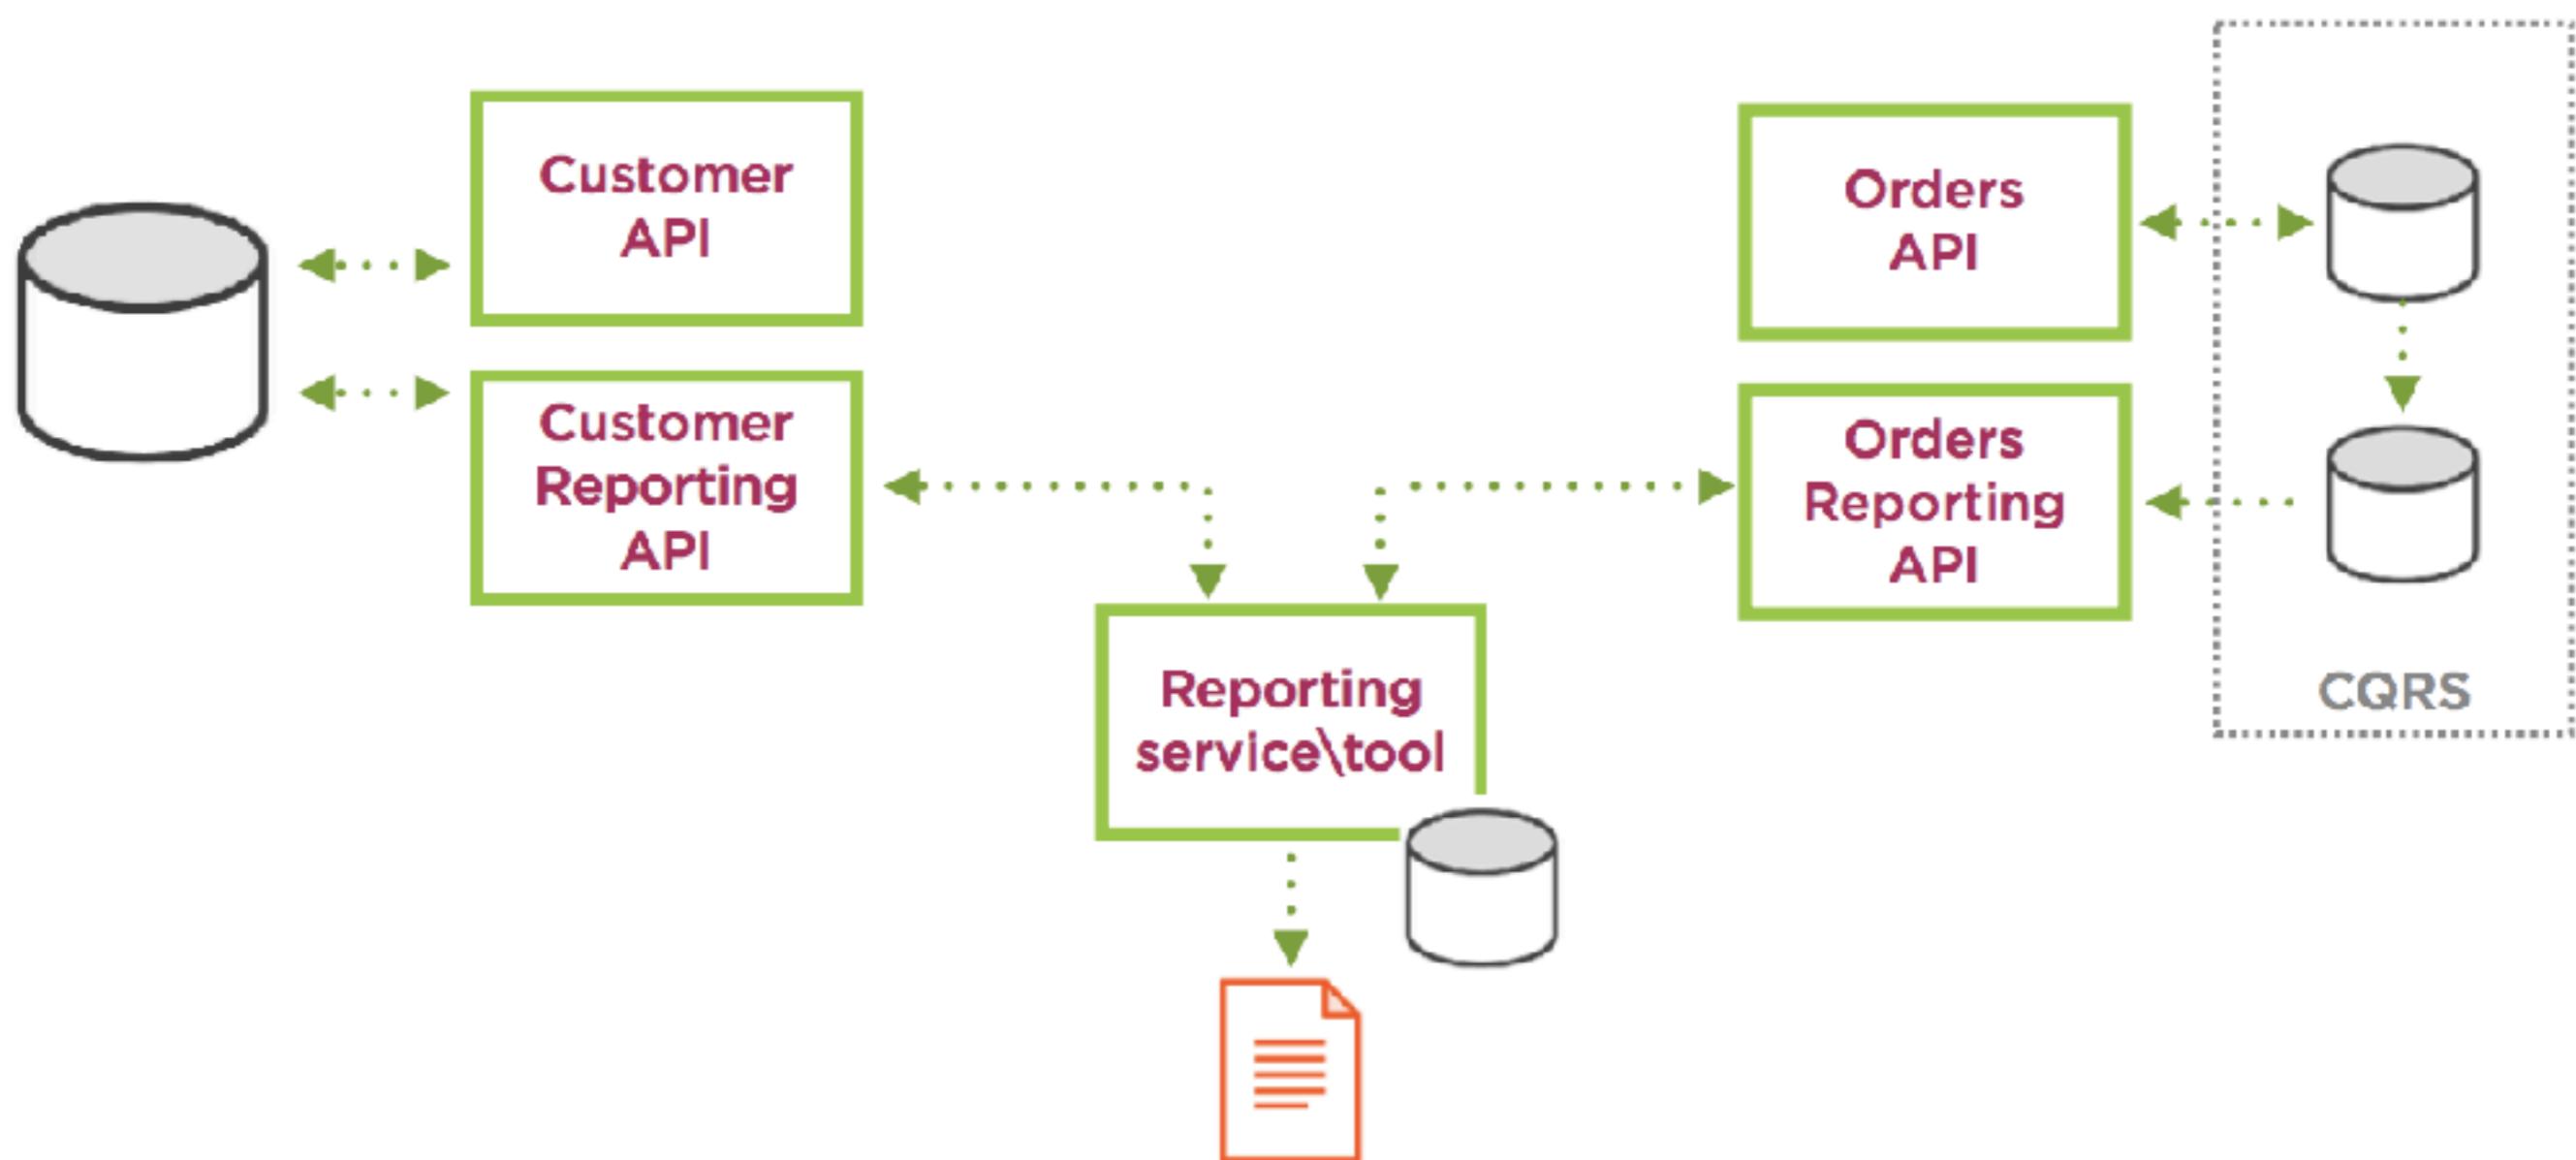
Approach

- Modules become services
- Shareable code libraries promote to service
- Review micro-services principles at each stage
- Prioritize by
 - Minimal viable product (MVP)
 - Customer needs and demand

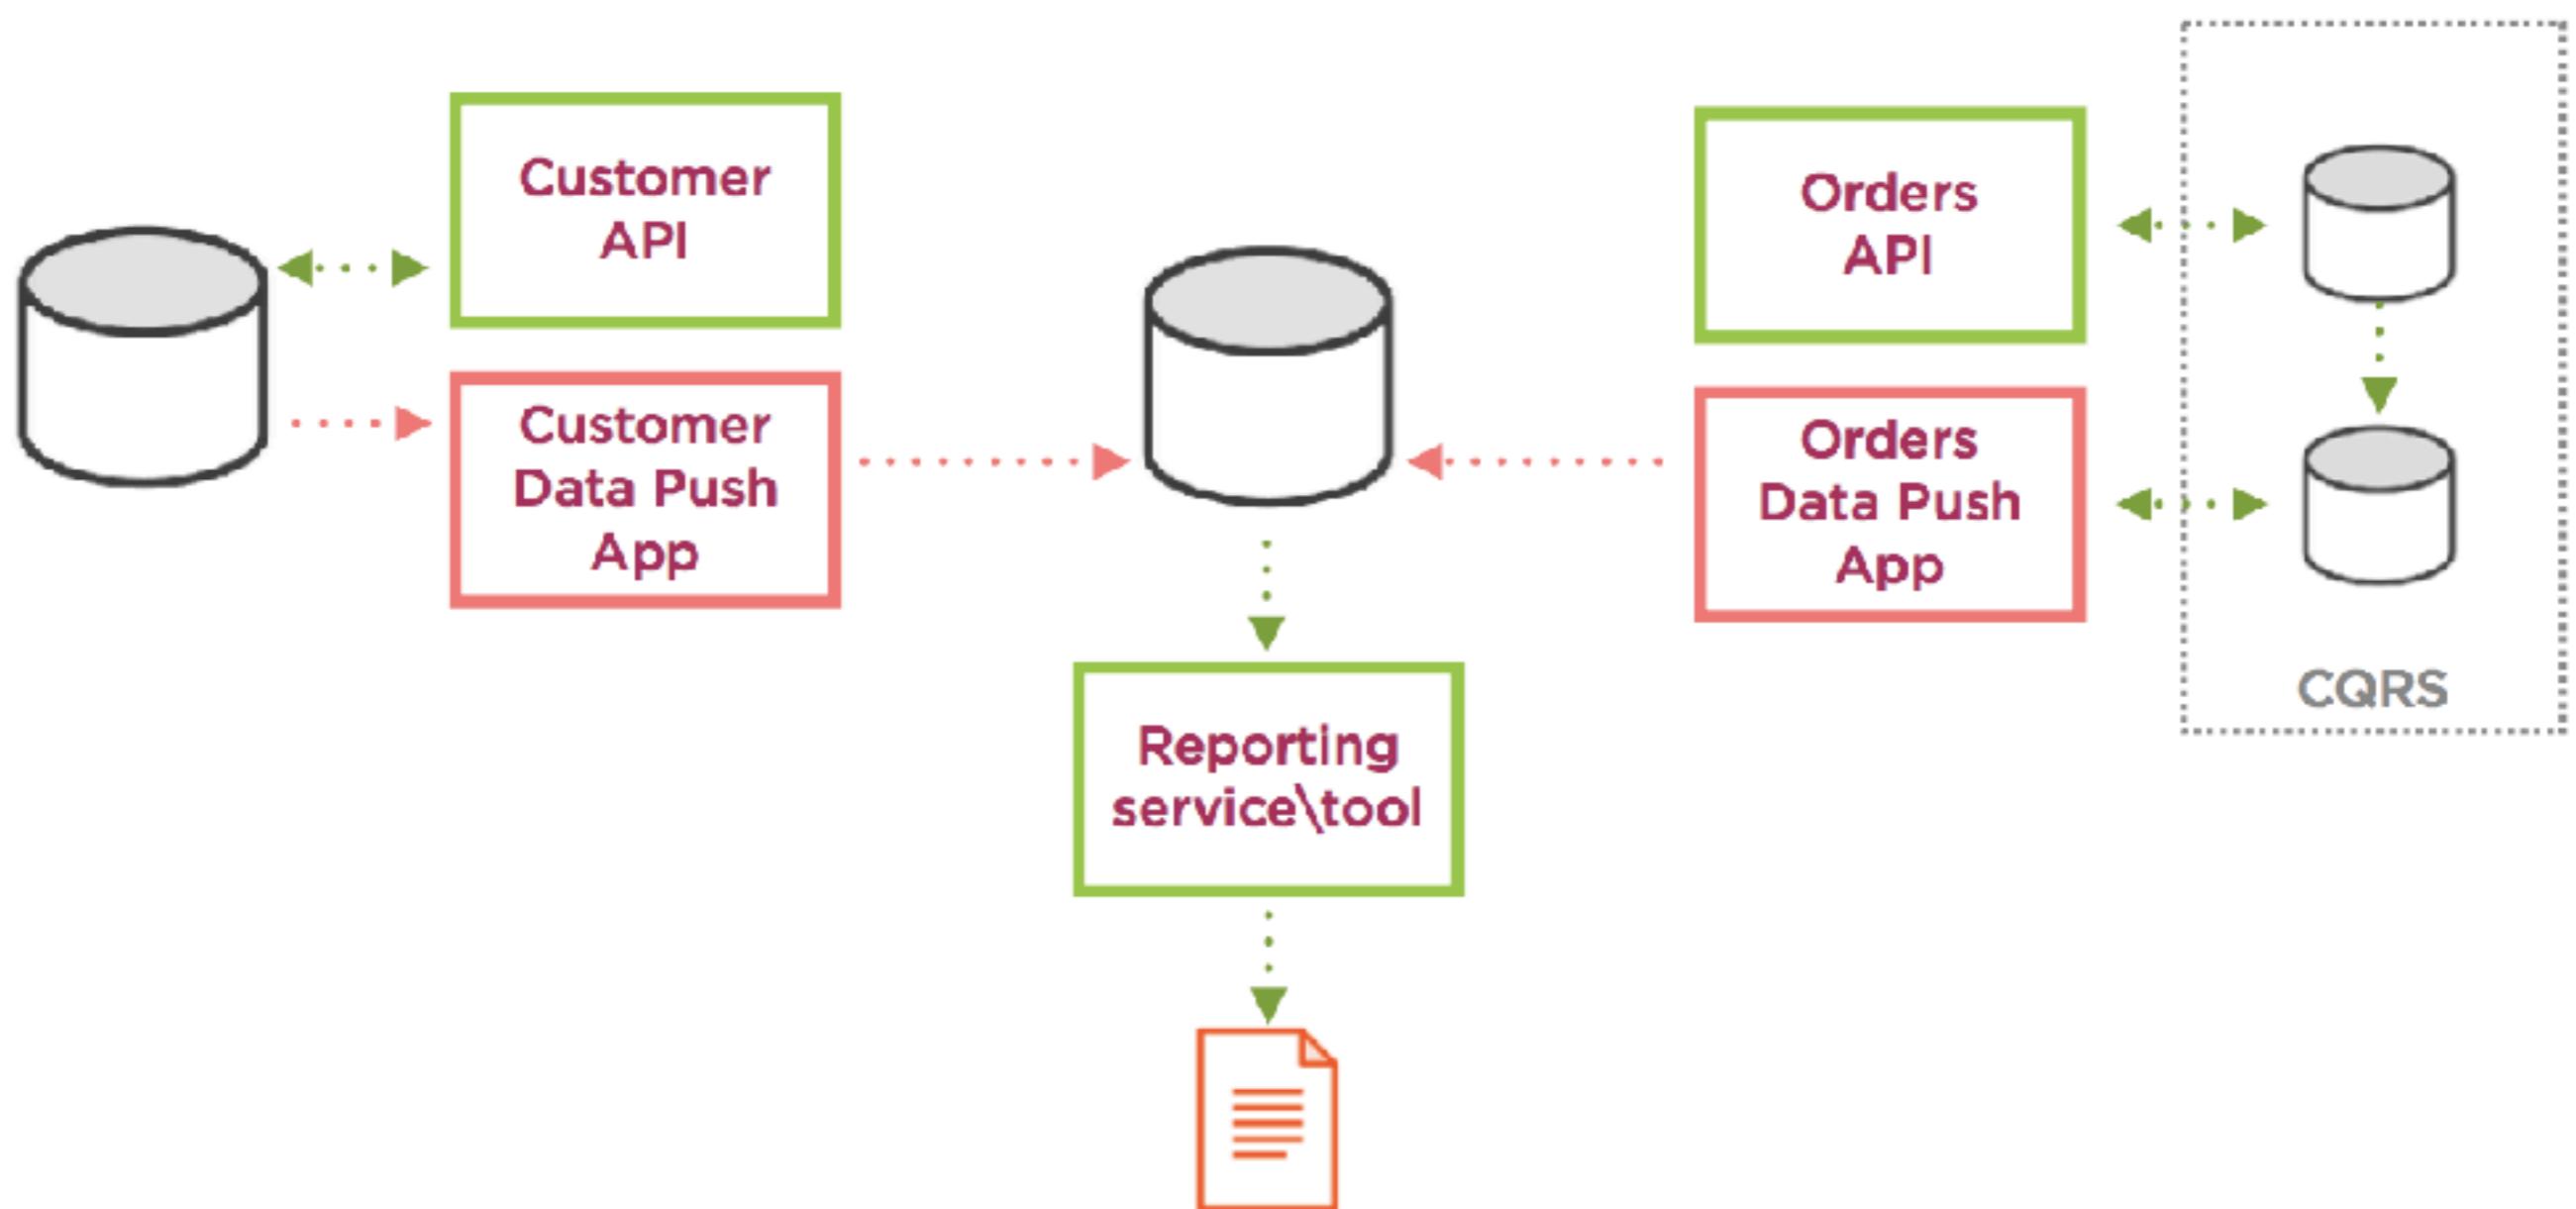
Solutions for Centralize Reporting

- Reporting service calls
- Data push application
- Reporting event subscribers
- Reporting events via gateway
- Using backup imports for reporting
- ETL and data-warehouses

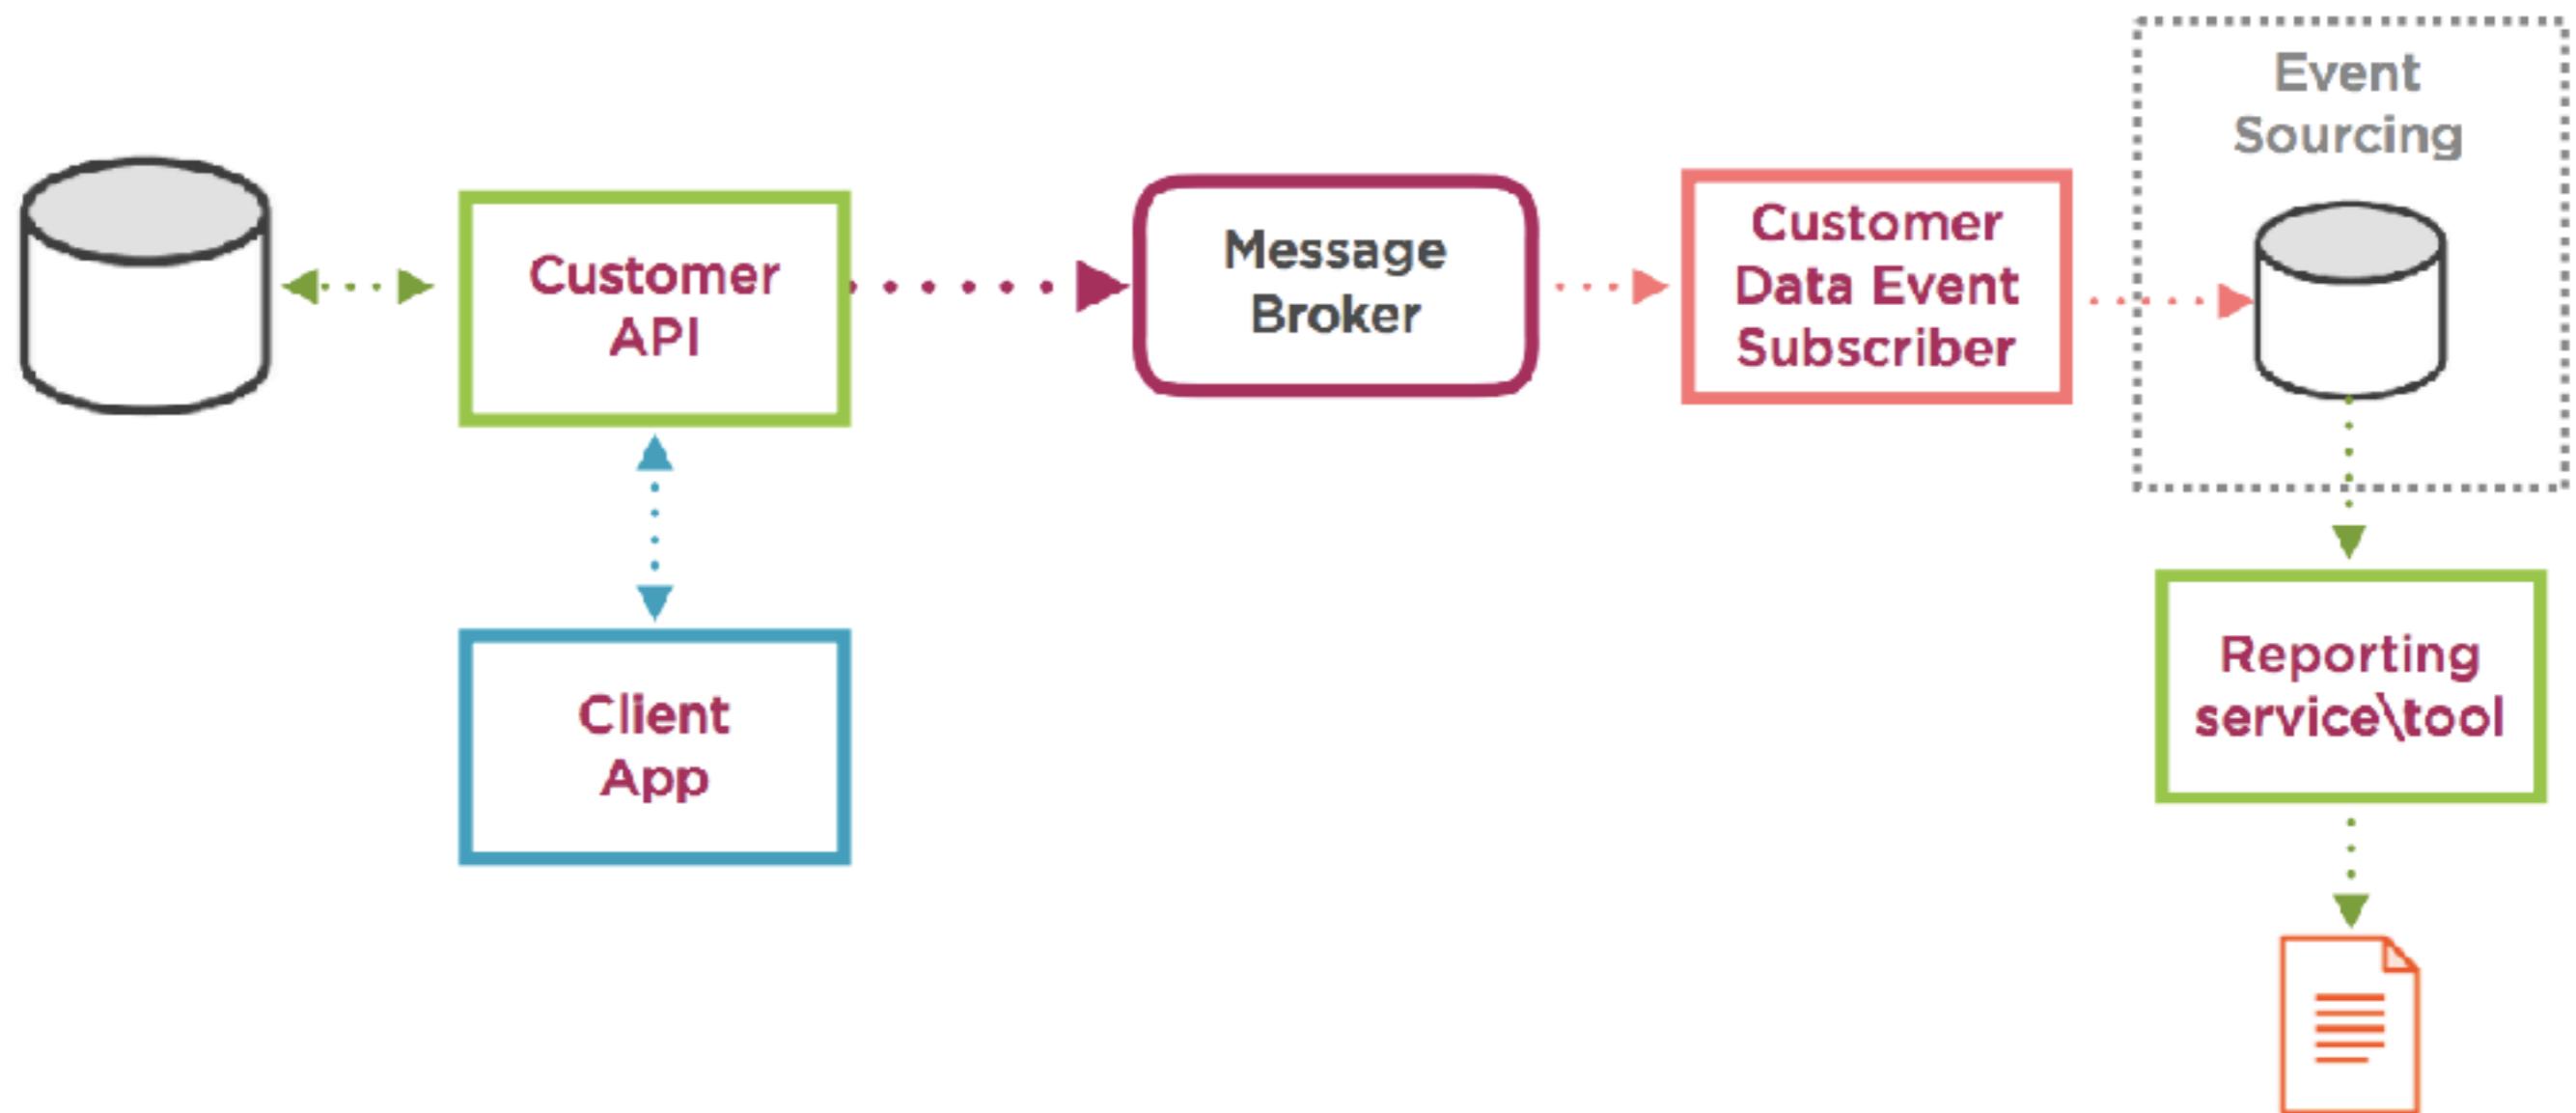
Reporting Service Calls



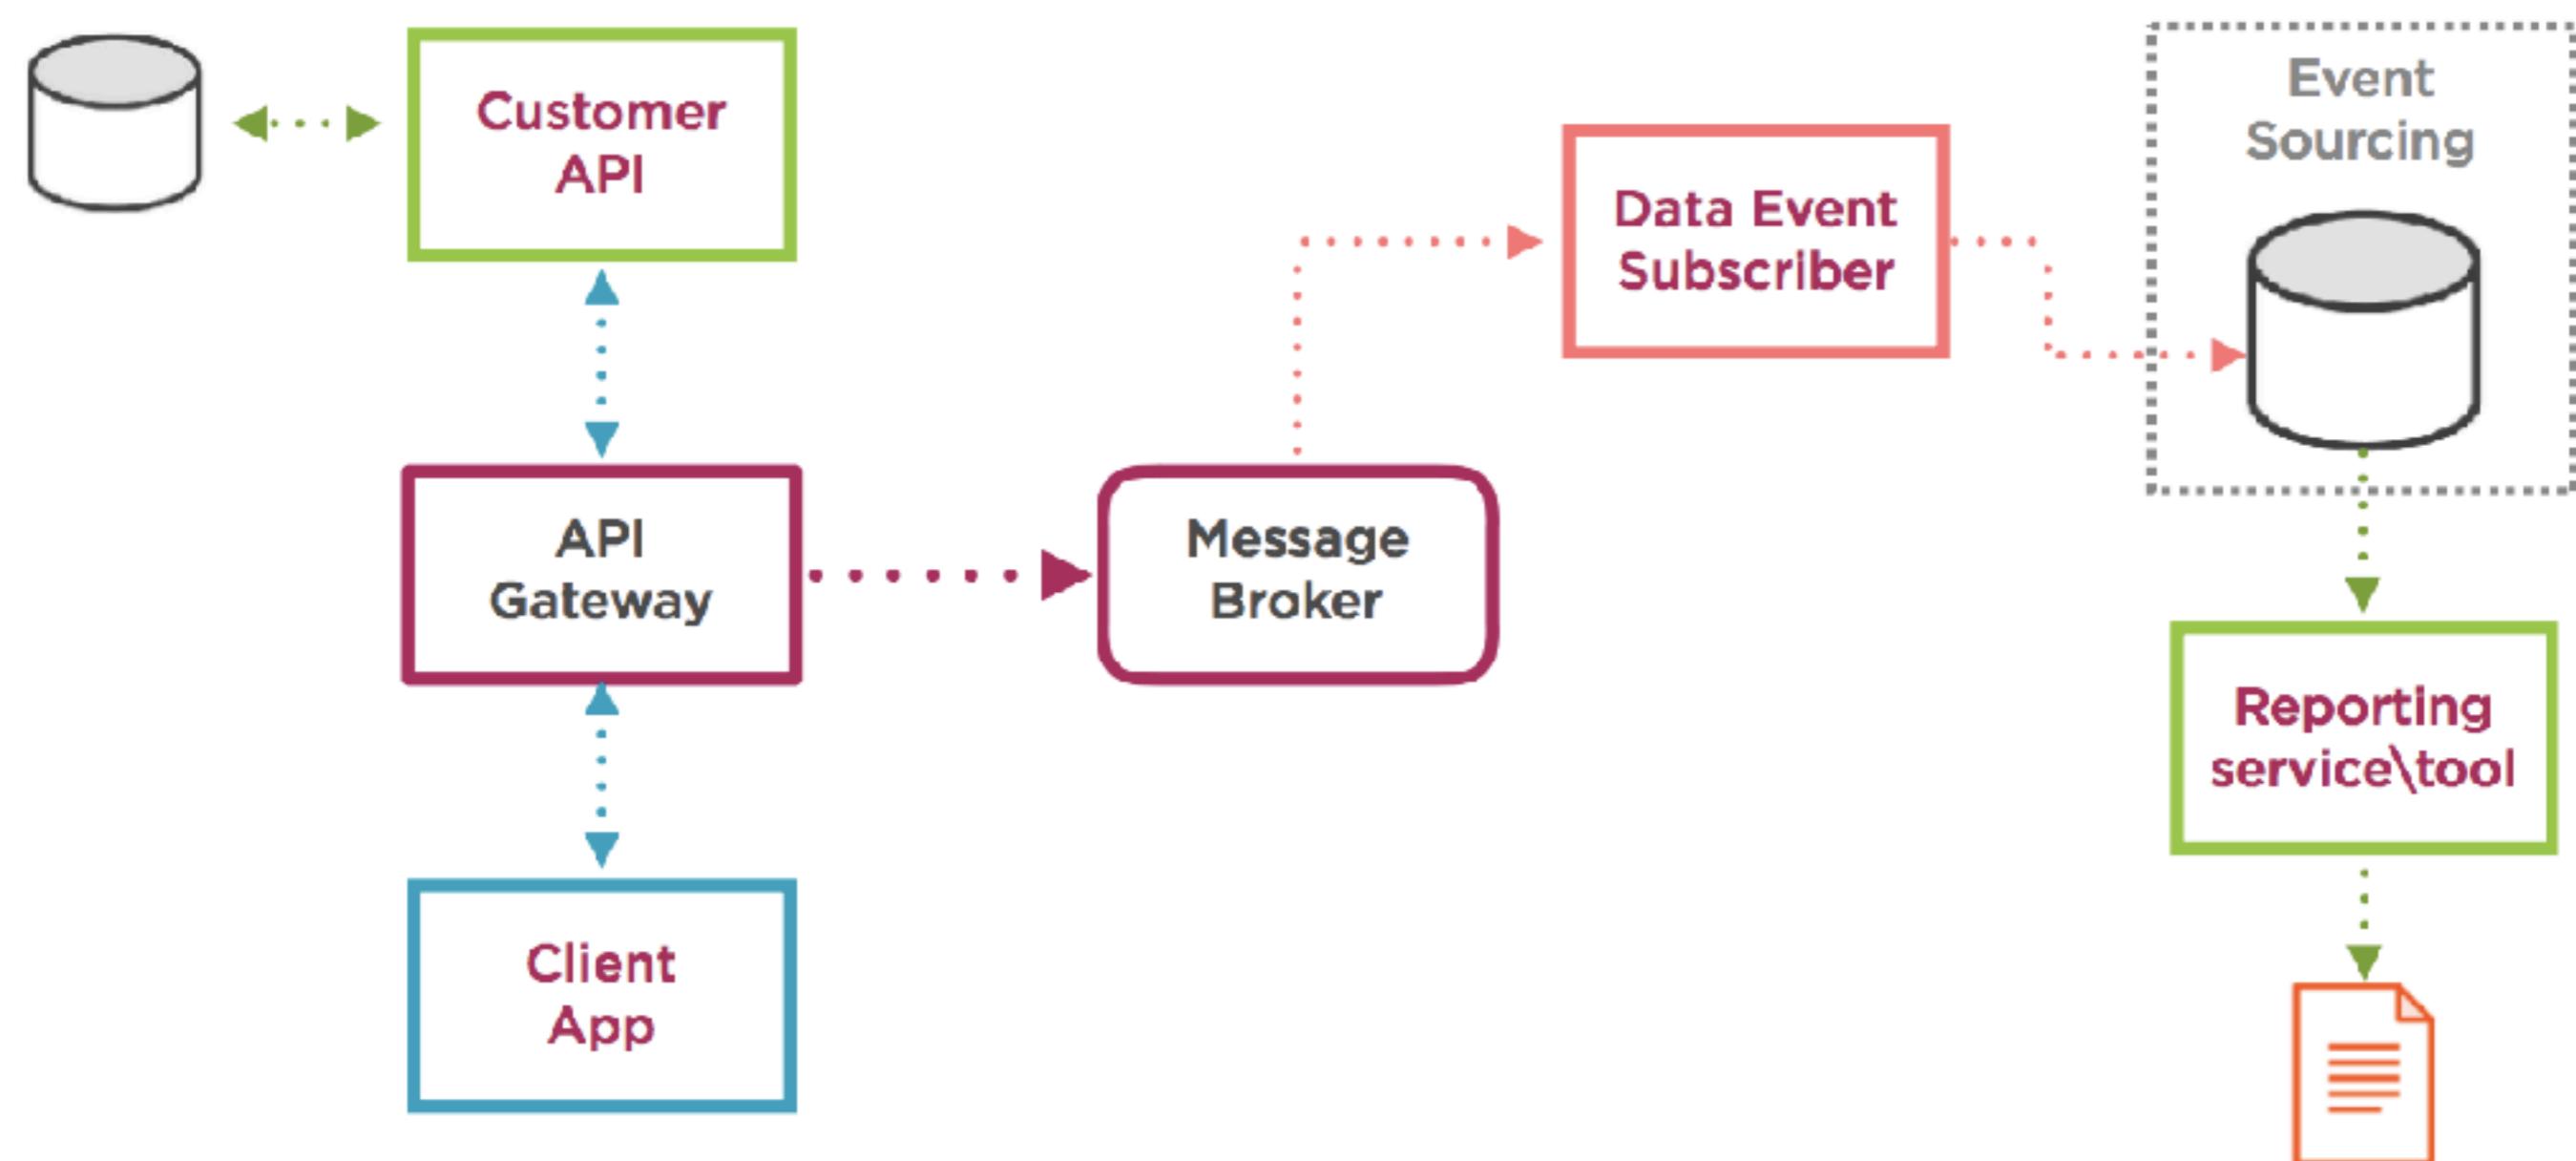
Data Push Application



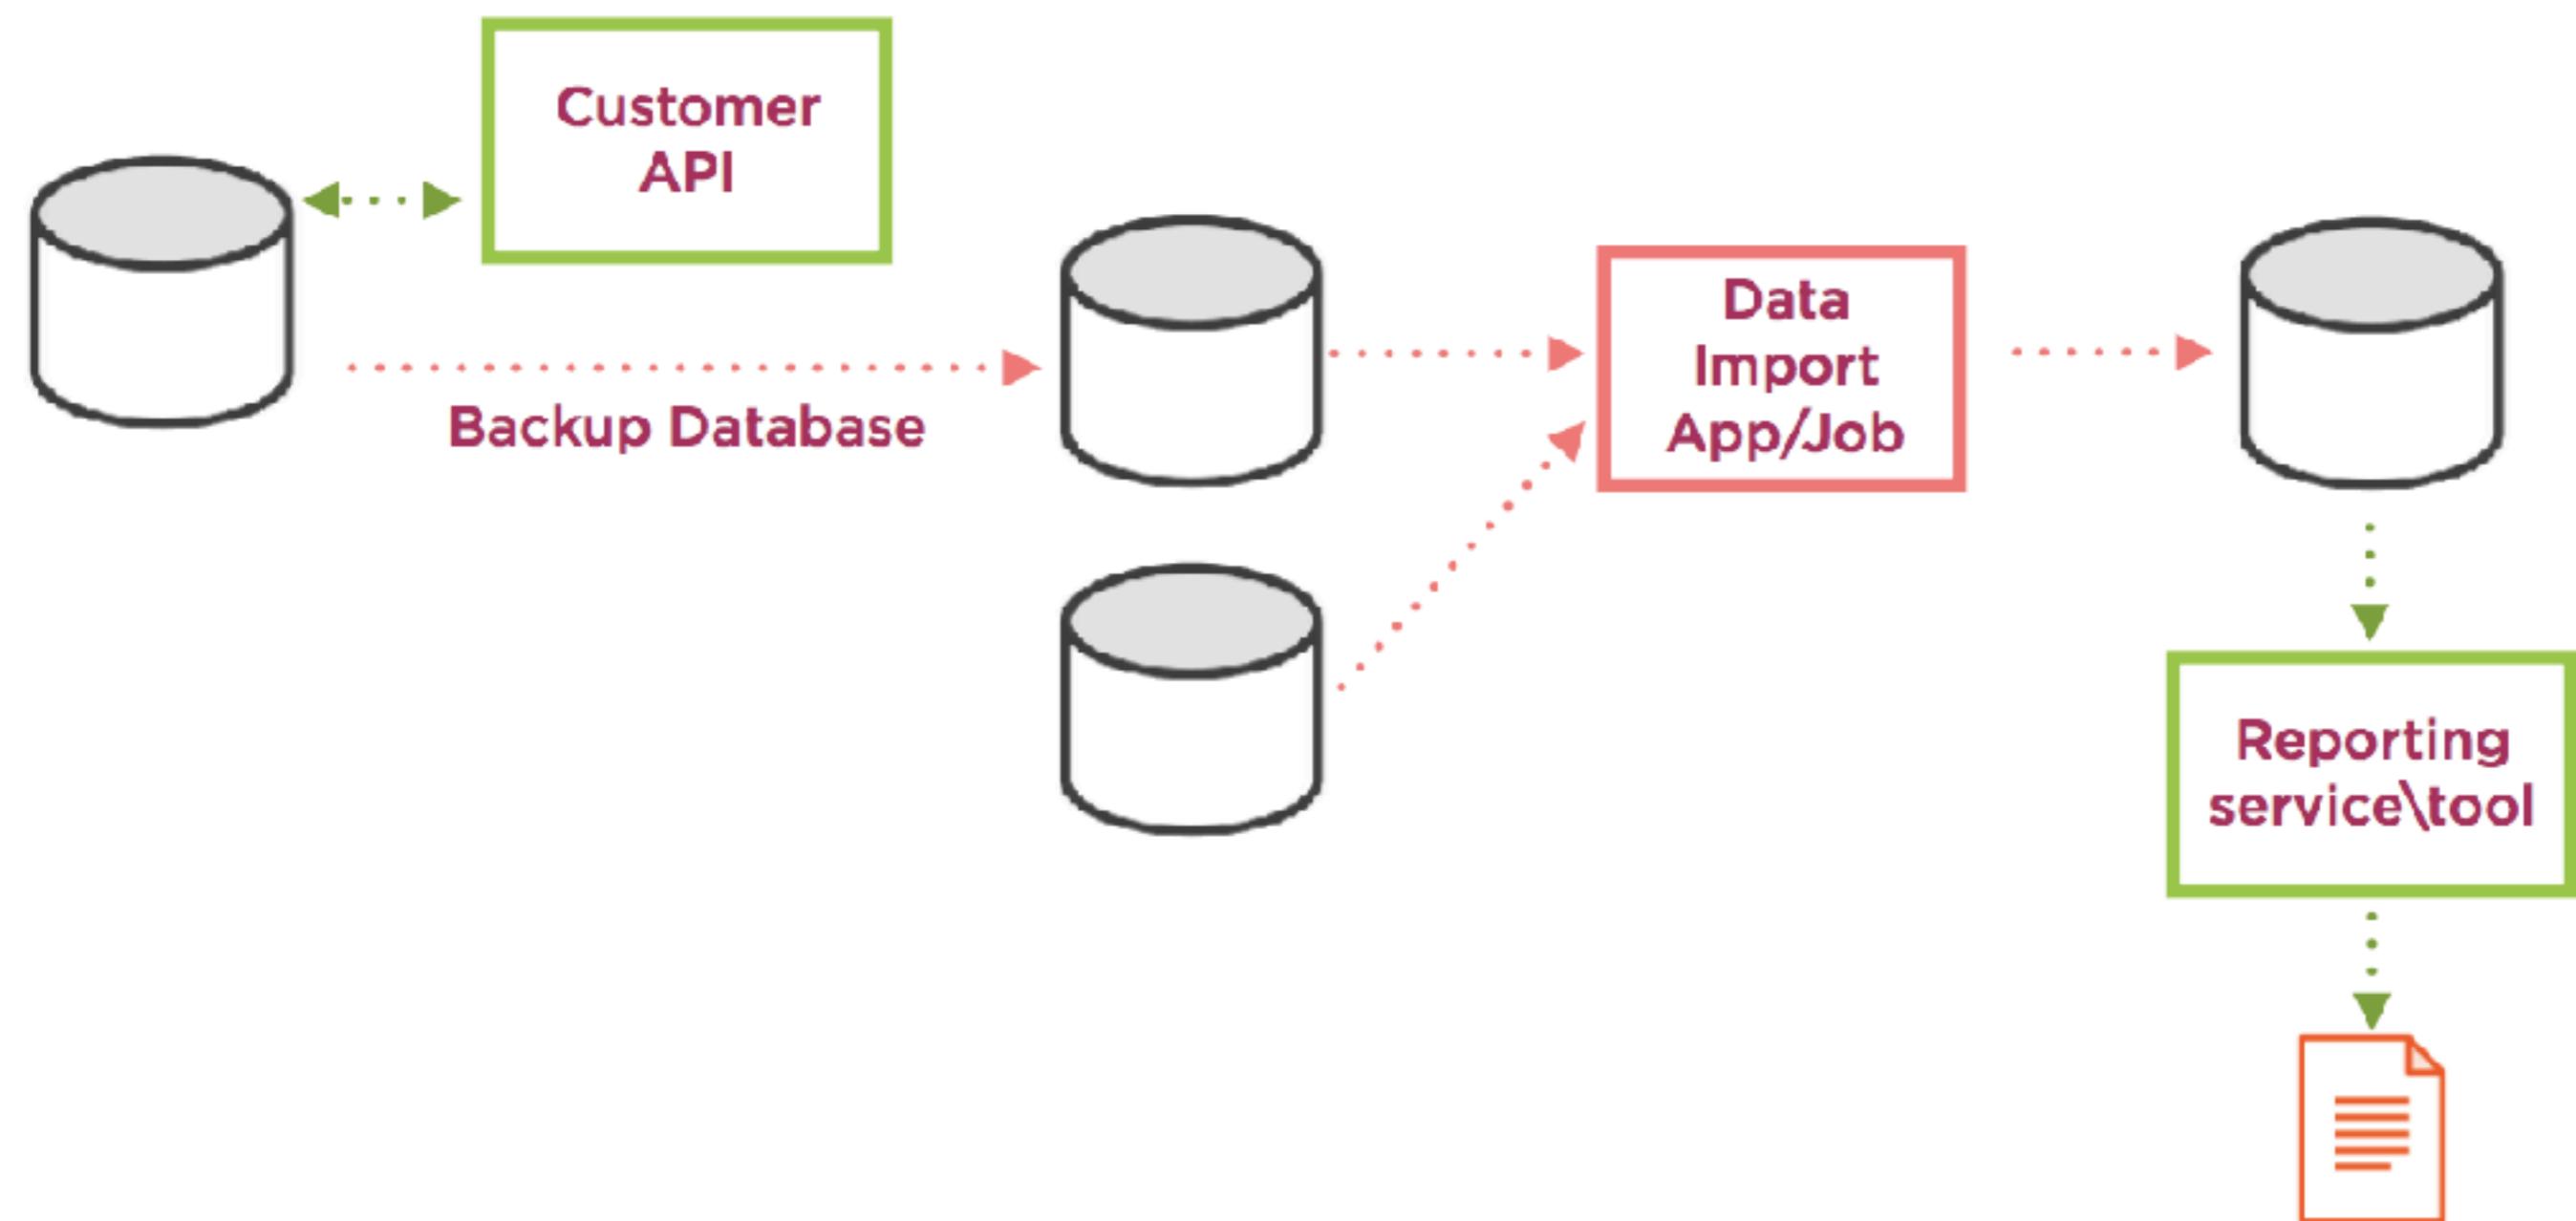
Reporting Event Subscribers



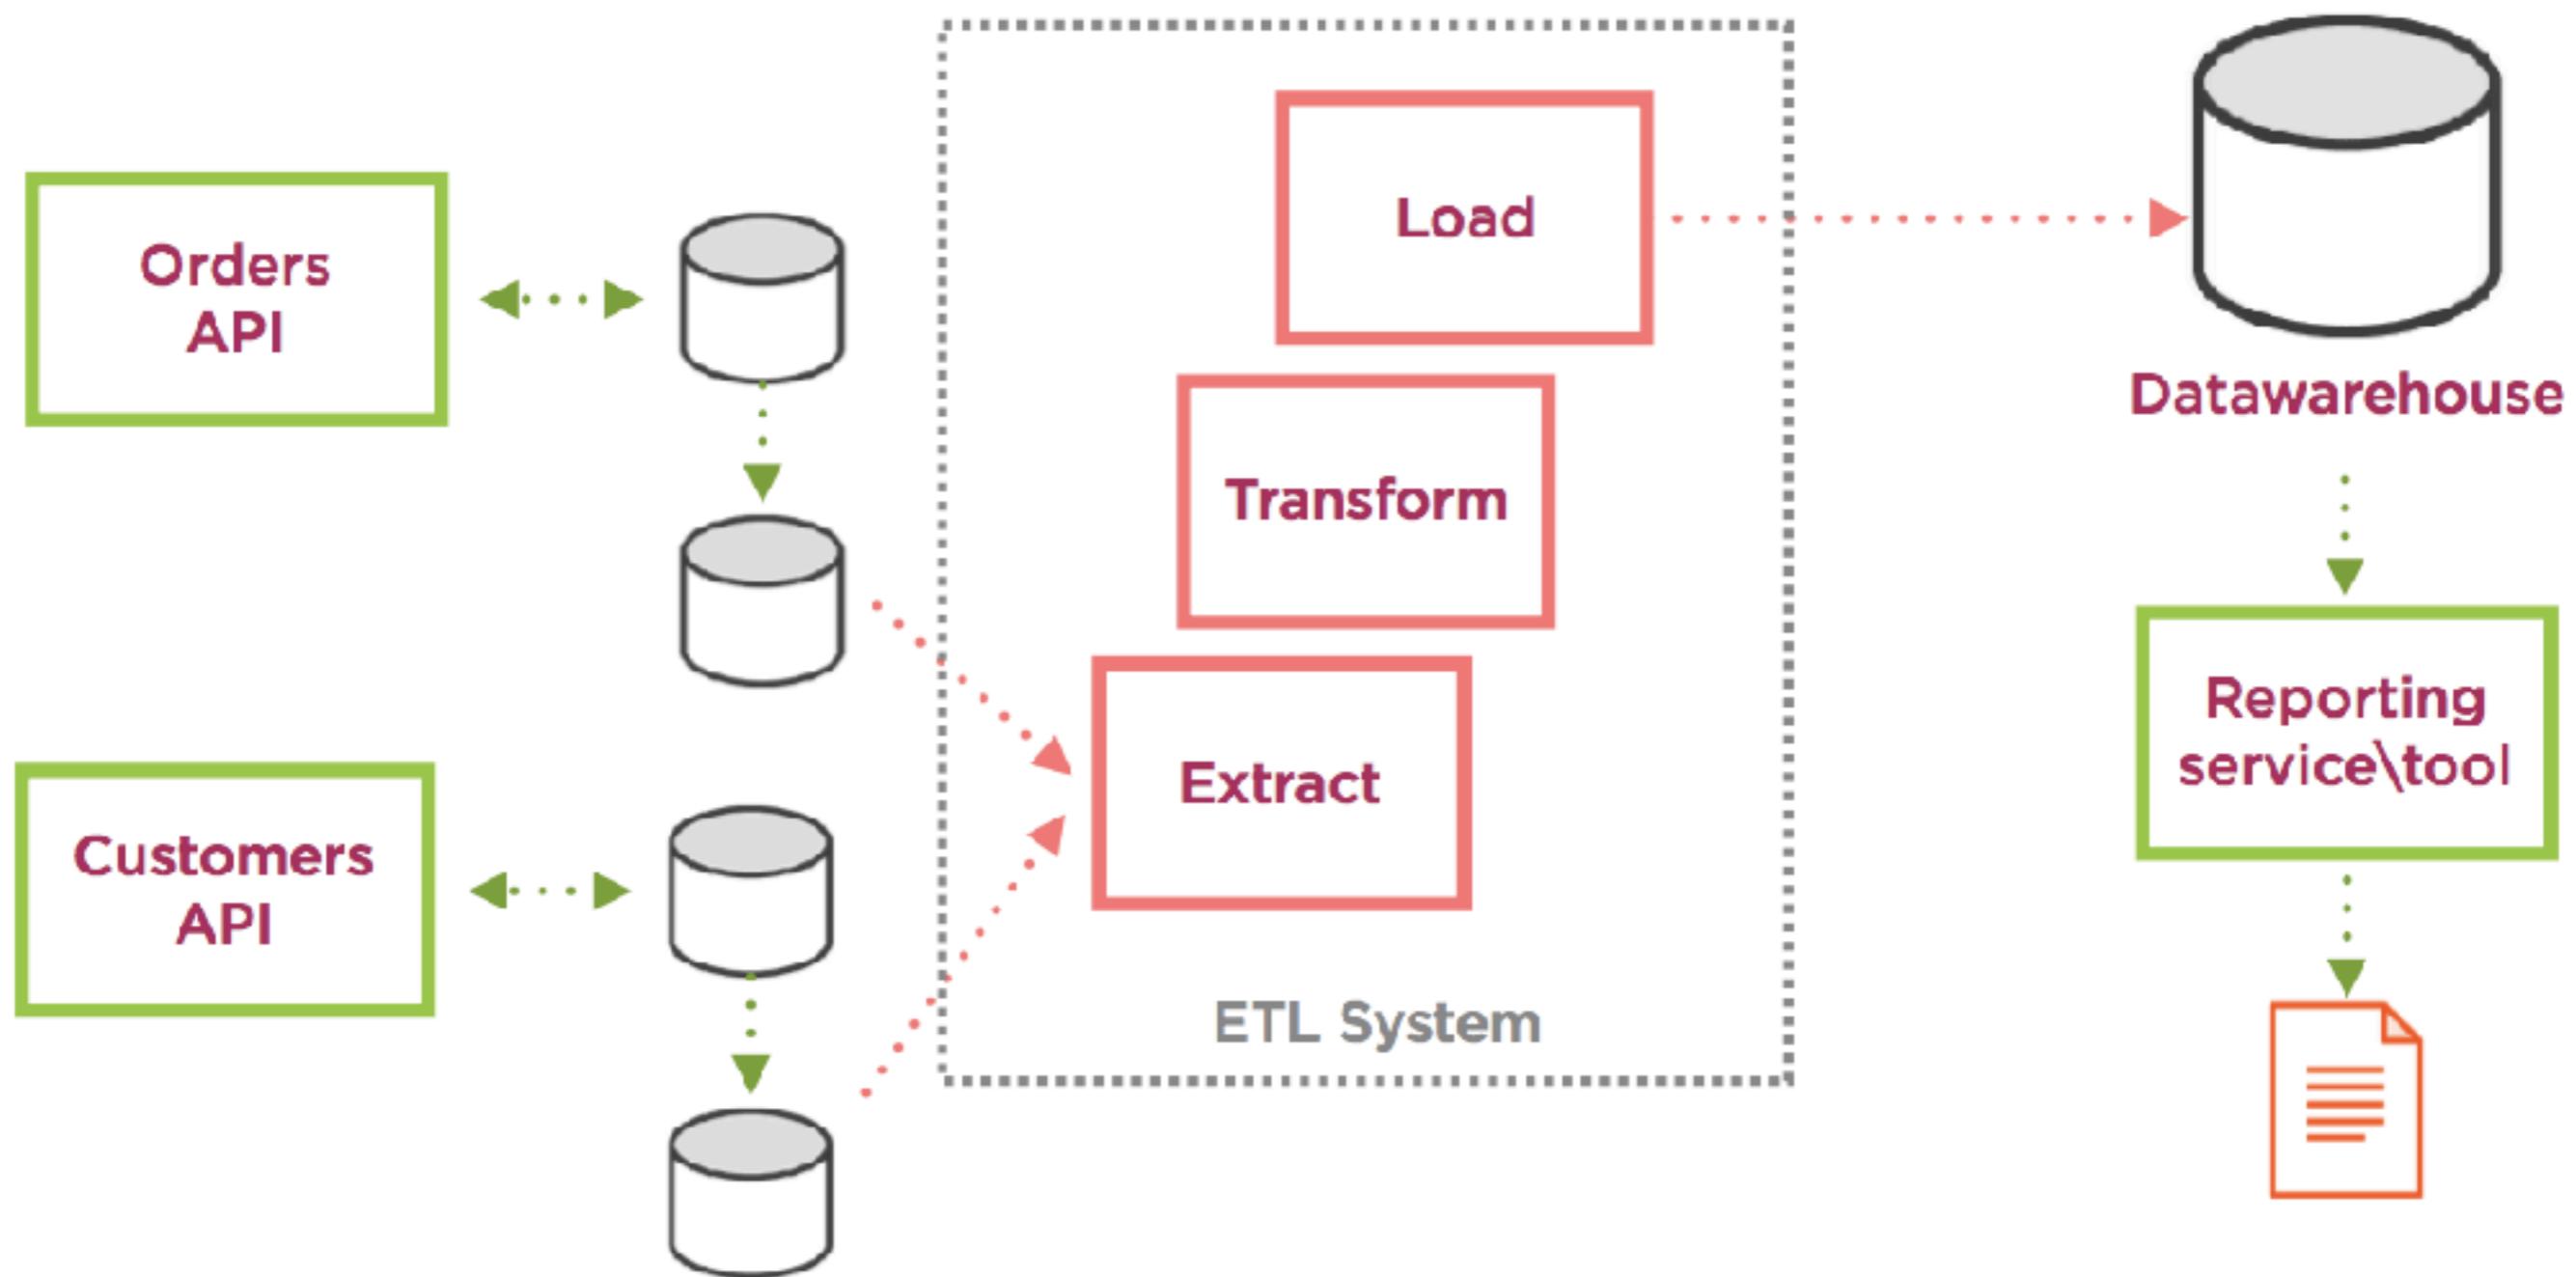
Reporting Event via Gateway



Using Backup Imports



ETL and Data-warehouse



MicroServices Provisos

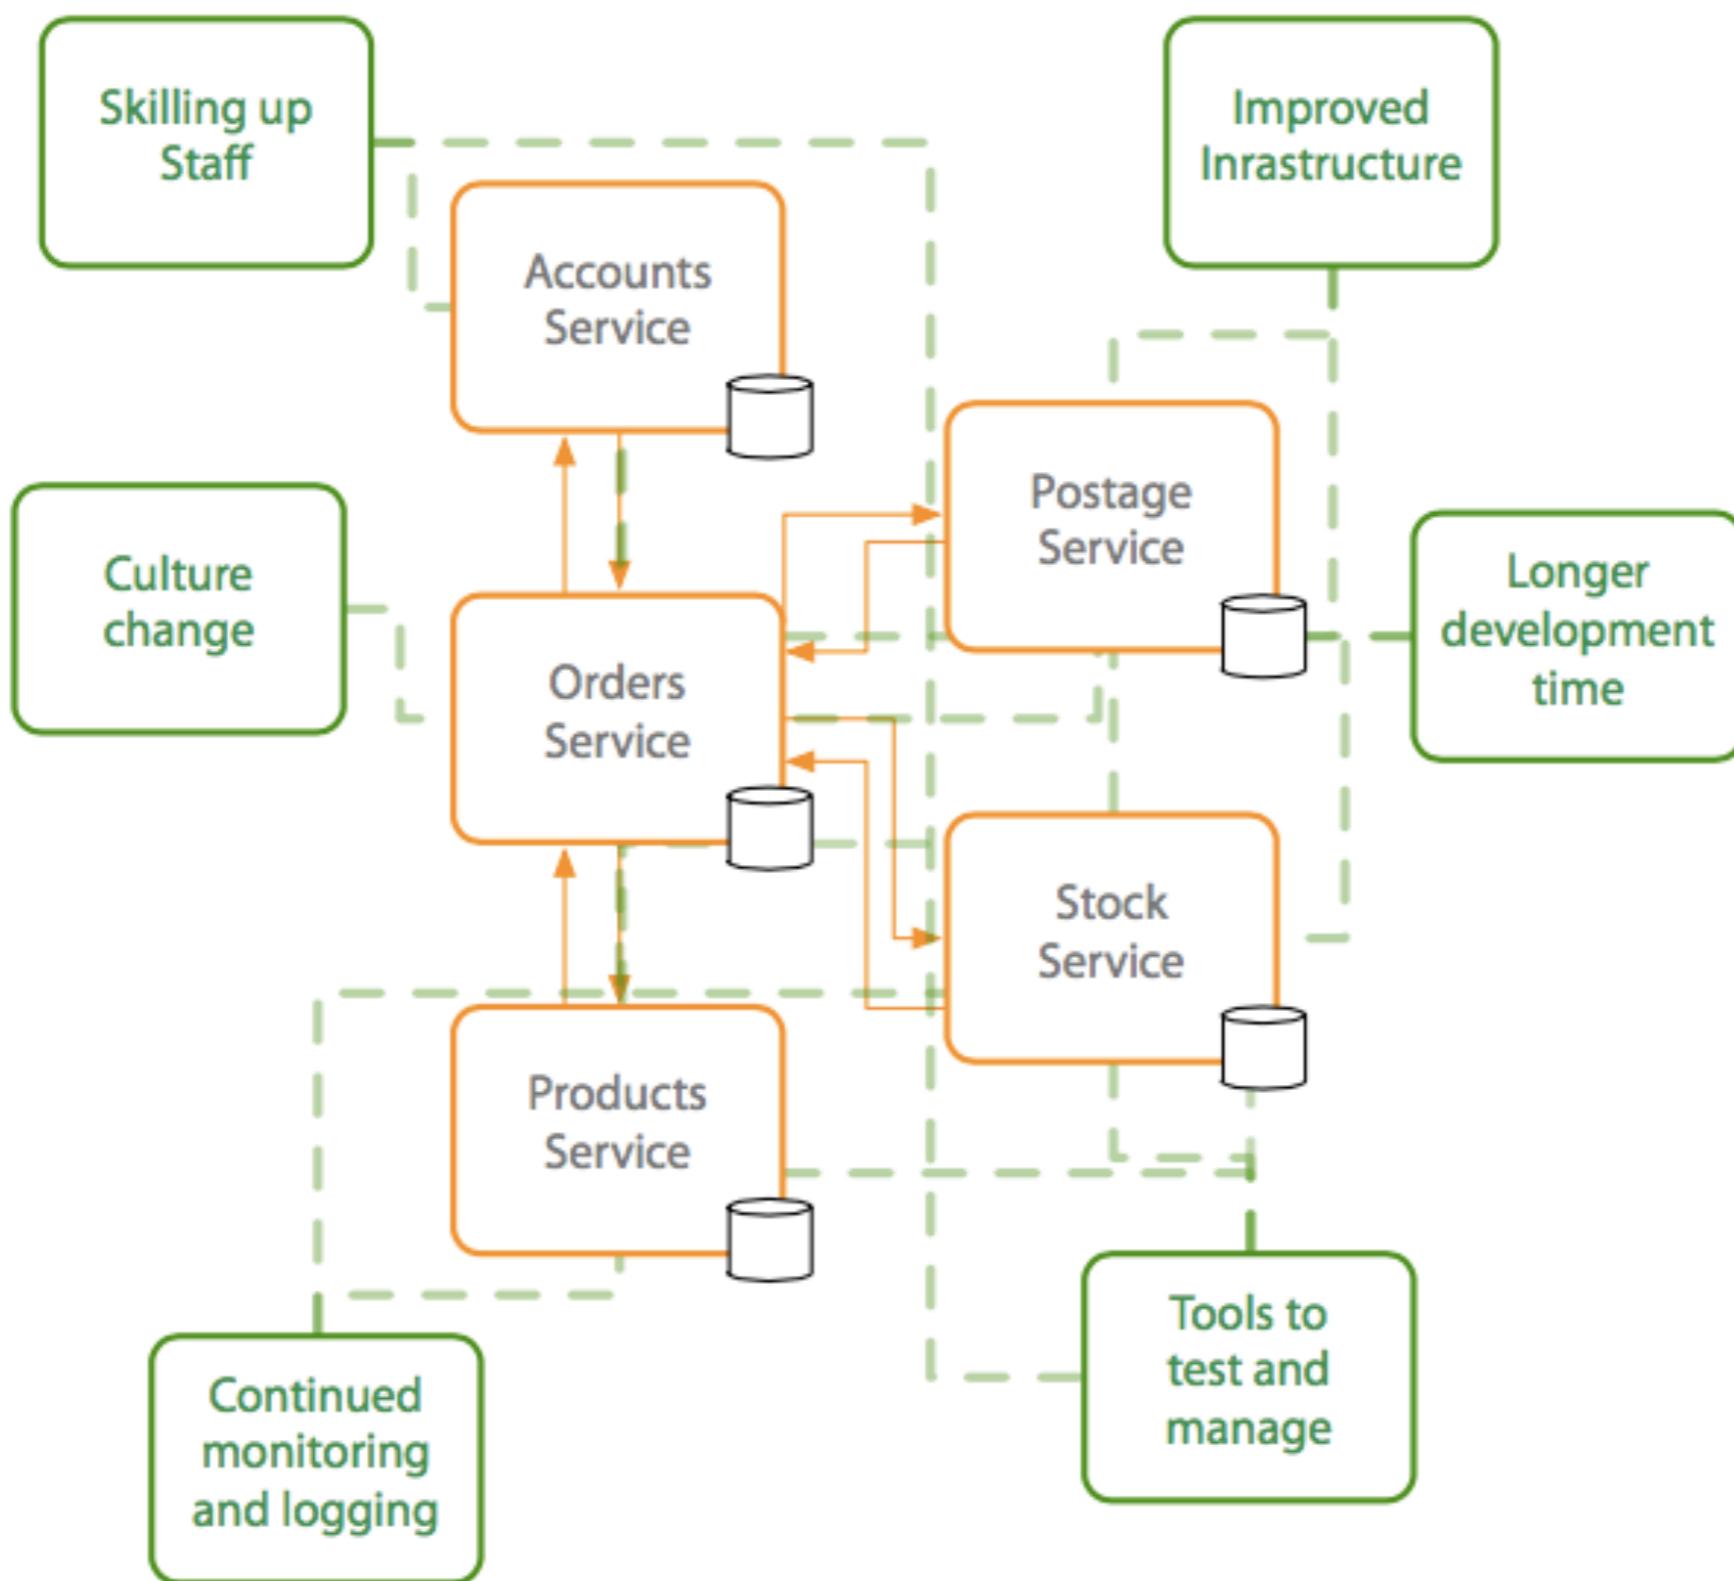
Provisos

- Accepting initial expense
 - Longer development times
 - Cost and training for tools and new skills
- Skilling up for distributed systems
 - Handling distributed transactions
 - Handling reporting
- Additional testing resource
 - Latency and performance testing
 - Testing for resilience

Provisos

- Improving infrastructure
 - Security
 - Performance
 - Reliance
- Overhead to manage microservices
- Cloud technologies
- Culture change

Provisos



Provisos

Business

Technical

Production

Business

Agile and DevOps

Cost

Time to market

Open source

Single support

Standard

Technical

Technical diversity

Performance

Maintainability

Extensibility

Testing

Design

Data store

Distribution

Production

Infrastructure

Portability

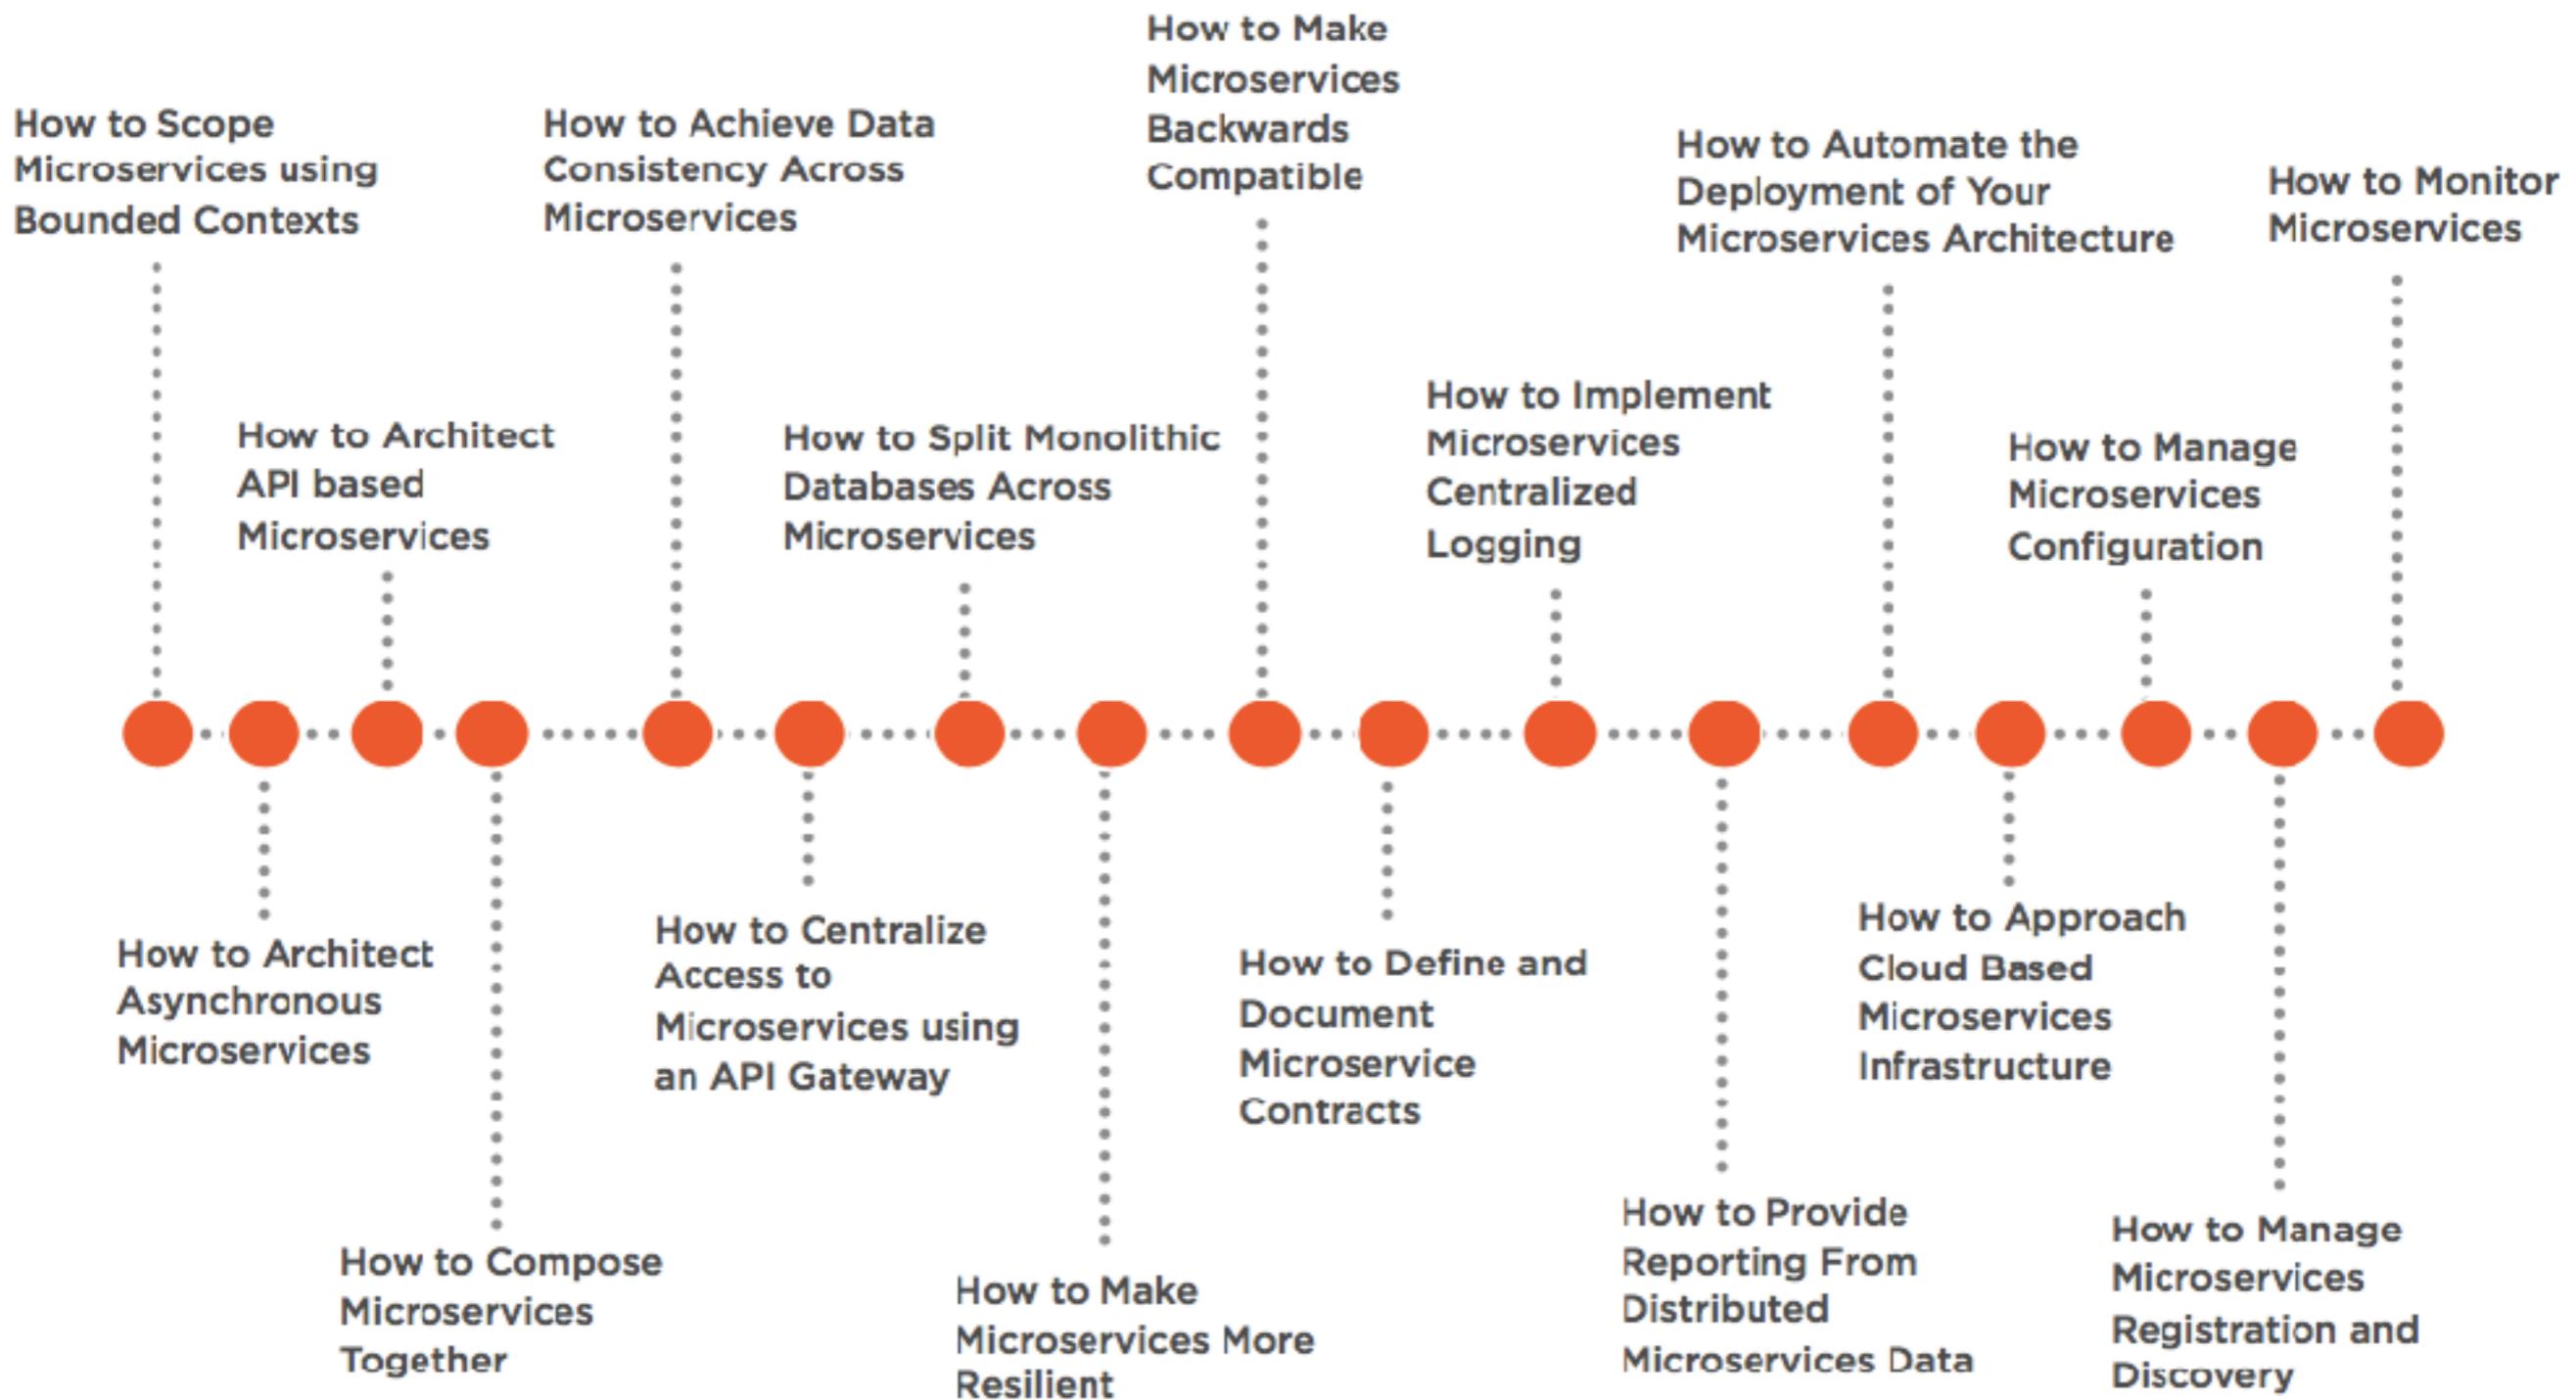
Scalability

Availability

Continuous Integration

Monitoring

Environments



Q&A